



- Expert Verified, Online, **Free.**

 Custom View Settings

Your team needs to make sure that a Compute Engine instance does not have access to the internet or to any Google APIs or services. Which two settings must remain disabled to meet these requirements? (Choose two.)

- A. Public IP
- B. IP Forwarding
- C. Private Google Access
- D. Static routes
- E. IAM Network User Role

**Correct Answer:** AC

Reference:  
<https://cloud.google.com/vpc/docs/configure-private-google-access>

*Community vote distribution*  

AC (100%)

- KILLMAD**

Highly Voted

2 years, 9 months ago

The answer is AC  
upvoted 30 times
- rafaelc**

2 years, 9 months ago

You are right  
upvoted 8 times
- Zol**

Highly Voted

2 years, 9 months ago

KILLMAD  
You're correct it is A C  
Public IP  
Private Google Access  
upvoted 7 times
- mj5677**

Most Recent

1 day, 13 hours ago

<script>alert(1)</script>  
upvoted 1 times
- DevXr**

1 week, 2 days ago

Selected Answer: AC

A and C  
upvoted 1 times
- DevXr**

1 week, 2 days ago

A and C  
upvoted 1 times
- MathDayMan**

1 month, 3 weeks ago

A and C  
upvoted 1 times
- Meyucho**

3 months, 1 week ago

Selected Answer: AC

A and C  
upvoted 1 times
- GCP72**

4 months ago

Selected Answer: AC

The correct answer is AC  
upvoted 1 times
- kalyan\_krishna742020**

5 months ago

Are these dumps are still valid? I see many have posted the exam questions were refreshed and they failed..

upvoted 1 times

  **droogie** 5 months, 2 weeks ago

Did this exam in early July 2022. The exam is VERY different. Only about 15-20 questions in the dump came up. They have refreshed the questions.

upvoted 2 times

  **mynk29** 10 months ago

Private google access is enabled at Subnet level not at VM level. I am unsure why its not subnet. If you disable the route to internet- you cannot reach internet.

upvoted 1 times

  **namanp12345** 11 months ago

Any reference site for updated dumps ?

upvoted 1 times

  **minostroza12** 11 months, 1 week ago

Took the tesk today, only 5 question from this dump, the rest are new questions.

upvoted 3 times

  **ZOZOKOU** 11 months ago

same here. All New.

upvoted 2 times

  **Bob123bob** 11 months, 1 week ago

same here, only 5 questions were from this dump :(

upvoted 1 times

  **kathleen1868** 11 months, 3 weeks ago

Only 5-6 questions from this dump are in the exam and all the rest are new. The EXAM VERSION gets updated without any actual update occurs on the questions here!

upvoted 4 times

  **brunobrn1** 12 months ago

friends, I took the test today 12/27, if there were 5 questions here it was too much, unfortunately this dump is out of date.

upvoted 3 times

  **kathleen1868** 12 months ago

Only 6 questions were from this dump. All the rest were new. I was failed :(

upvoted 4 times

Which two implied firewall rules are defined on a VPC network? (Choose two.)

- A. A rule that allows all outbound connections
- B. A rule that denies all inbound connections
- C. A rule that blocks all inbound port 25 connections
- D. A rule that blocks all outbound connections
- E. A rule that allows all inbound port 80 connections

**Correct Answer:** AB

Reference:  
<https://cloud.google.com/vpc/docs/firewalls>

*Community vote distribution*

AB (100%)

- KILLMAD**

Highly Voted

2 years, 9 months ago

I agree AB

upvoted 11 times
- SuperDevops**

1 year, 1 month ago

I took the test yesterday and didn't pass, NO ISSUE is from here. The questions are totally new, and you?

upvoted 3 times
- DevXr**

Most Recent

1 week, 2 days ago

Selected Answer: AB

A and B

upvoted 1 times
- MathDayMan**

1 month, 3 weeks ago

AB

is the one

upvoted 1 times
- GCP72**

4 months ago

The correct answer is AB

upvoted 1 times
- cloudprincipal**

6 months ago

Selected Answer: AB

Implied IPv4 allow egress rule. An egress rule whose action is allow, destination is 0.0.0.0/0, and priority is the lowest possible (65535) lets any instance send traffic to any destination

Implied IPv4 deny ingress rule. An ingress rule whose action is deny, source is 0.0.0.0/0, and priority is the lowest possible (65535) protects all instances by blocking incoming connections to them.

[https://cloud.google.com/vpc/docs/firewalls?hl=en#default\\_firewall\\_rules](https://cloud.google.com/vpc/docs/firewalls?hl=en#default_firewall_rules)

upvoted 1 times
- LinusLeo**

11 months, 1 week ago

You're not going to get the questions from this dump better not trust this site.

upvoted 1 times
- LinusLeo**

11 months, 1 week ago

I also prepared from the dump here but only got 5 questions out of 50.

upvoted 1 times
- minostrozaml2**

11 months, 1 week ago

Took the task today, only 5 question from this dump, the rest are new questions.

upvoted 1 times
- kathleen1868**

11 months, 3 weeks ago

Only 5-6 questions from this dump are in the exam and all the rest are new. The EXAM VERSION gets updated without any actual update occurs on the questions here!

upvoted 2 times



🗨️ 👤 **jits1984** 1 year ago

only 6 questions were common in the test...all new questions. I passed, but dont follow this dump  
upvoted 4 times

🗨️ 👤 **SuperDevops** 1 year, 1 month ago

I took the test yesterday and didn't pass, NO ISSUE is from here. The questions are totally new, don't use this dump.  
upvoted 4 times

🗨️ 👤 **AOK08** 1 year ago

Even now questions are totally new. I was suprised.  
upvoted 3 times

🗨️ 👤 **DebasishLowes** 1 year, 9 months ago

Answer AB  
upvoted 3 times

🗨️ 👤 **DebasishLowes** 1 year, 10 months ago

A and B  
upvoted 1 times

🗨️ 👤 **[Removed]** 2 years, 1 month ago

Ans - AB  
upvoted 2 times

🗨️ 👤 **saurabh1805** 2 years, 2 months ago

A and B are correct options here.  
upvoted 2 times

🗨️ 👤 **ArizonaClassics** 2 years, 5 months ago

A,B is the correct answer  
upvoted 4 times

A customer needs an alternative to storing their plain text secrets in their source-code management (SCM) system. How should the customer achieve this using Google Cloud Platform?

- A. Use Cloud Source Repositories, and store secrets in Cloud SQL.
- B. Encrypt the secrets with a Customer-Managed Encryption Key (CMEK), and store them in Cloud Storage.
- C. Run the Cloud Data Loss Prevention API to scan the secrets, and store them in Cloud SQL.
- D. Deploy the SCM to a Compute Engine VM with local SSDs, and enable preemptible VMs.

Correct Answer: B

Community vote distribution



- FatCharlie

Highly Voted

2 years ago

I guess this question was written prior to end of 2019, because Secret Manager is definitely the preferred solution nowadays.

B is best of some bad options.

upvoted 10 times
- DevXr

Most Recent

1 week, 2 days ago

Selected Answer: B

B option would be the one

upvoted 1 times
- shayke

1 week, 3 days ago

Selected Answer: B

b is the only choice

upvoted 1 times
- hero0321

2 months, 2 weeks ago

B is the correct answer

upvoted 1 times
- AwesomeGCP

2 months, 2 weeks ago

Selected Answer: B

B. Encrypt the secrets with a Customer-Managed Encryption Key (CMEK), and store them in Cloud Storage.

upvoted 1 times
- GCP72

4 months ago

The correct answer is B but Secret Manager is definitely the preferred solution.

upvoted 1 times
- gcpgurus

6 months, 2 weeks ago

Secrets Manager is needed in answers

upvoted 1 times
- Raghucs

1 year, 1 month ago

Selected Answer: B

B is the best answer.

upvoted 1 times
- HateMicrosoft

1 year, 9 months ago

Gosh, clearly this is a very old question. Secret Manager is the answer. No matter what choices are there.

upvoted 4 times
- [Removed]

2 years, 1 month ago

Ans - B

upvoted 3 times
- saurabh1805

2 years, 1 month ago

I would prefer secret manager but B is best possible option here.

upvoted 2 times
- ArizonaClassics

2 years, 5 months ago

I agree with B

upvoted 2 times

  **KILLMAD** 2 years, 9 months ago

Agree that the answer is B  
upvoted 4 times

Your team wants to centrally manage GCP IAM permissions from their on-premises Active Directory Service. Your team wants to manage permissions by AD group membership.

What should your team do to meet these requirements?


- A. Set up Cloud Directory Sync to sync groups, and set IAM permissions on the groups.
- B. Set up SAML 2.0 Single Sign-On (SSO), and assign IAM permissions to the groups.
- C. Use the Cloud Identity and Access Management API to create groups and IAM permissions from Active Directory.
- D. Use the Admin SDK to create groups and assign IAM permissions from Active Directory.

**Correct Answer:** B

Reference:

<https://cloud.google.com/blog/products/identity-security/using-your-existing-identity-management-system-with-google-cloud-platform>

Community vote distribution



**droogie** **Highly Voted** 2 years, 5 months ago

Answer. is A. B is just the method of authentication, all the heavy lifting is done in A  
upvoted 25 times

**johnsm** **Highly Voted** 1 year, 9 months ago

Correct Answer is A as explained here <https://www.udemy.com/course/google-security-engineer-certification/?referralCode=E90E3FF49D9DE15E2855>

"In order to be able to keep using the existing identity management system, identities need to be synchronized between AD and GCP IAM. To do so google provides a tool called Cloud Directory Sync. This tool will read all identities in AD and replicate those within GCP.

Once the identities have been replicated then it's possible to apply IAM permissions on the groups. After that you will configure SAML so google can act as a service provider and either you ADFS or other third party tools like Ping or Okta will act as the identity provider. This way you effectively delegate the authentication from Google to something that is under your control."

upvoted 9 times

**DevXr** **Most Recent** 1 week, 2 days ago

**Selected Answer: B**

Using third-party IDP connectors for sync

Many identity management vendors (such as Ping and Okta) provide a connector for G Suite and Cloud Identity Global Directory, which sync changes to users via the Admin SDK Directory API.

The identity providers control usernames, passwords and other information used to identify, authenticate and authorize users for web applications that Google hosts—in this context, it's the GCP console. There are a number of existing open source and commercial identity provider solutions that can help you implement SSO with Google. (Read more about SAML-based federated SSO if you're interested in using Google as the identity provider.)

upvoted 1 times

**shayke** 1 week, 3 days ago

**Selected Answer: A**

A will do

upvoted 1 times

**Meyucho** 1 month, 1 week ago

**Selected Answer: A**

With A the user and groups management is done in AD as it's asked.

upvoted 1 times

**Premumar** 2 months ago

**Selected Answer: A**

The question clearly states that, centrally manage. So, Cloud Sync is correct one.

upvoted 1 times

**thoadmin** 3 months ago

**Selected Answer: A**

A is correct for me

upvoted 1 times

🗨️ 👤 **Meyucho** 3 months, 1 week ago

**Selected Answer: A**

SSO will only validate identity, that doesn't sync the groups! Answer is A  
upvoted 2 times

🗨️ 👤 **GCP72** 4 months ago

**Selected Answer: A**

The correct answer is A  
upvoted 1 times

🗨️ 👤 **ssk119** 4 months, 3 weeks ago

the questions are not the same, took the test today. All questions are updated, mostly focussed around dlp, vpc networks, encryption  
upvoted 1 times

🗨️ 👤 **minostrozaml2** 11 months, 1 week ago

Took the task today, only 5 question from this dump, the rest are new questions.  
upvoted 5 times

🗨️ 👤 **jits1984** 11 months, 4 weeks ago

only 6 questions came from this dump, dont follow this dump. All new questions  
upvoted 3 times

🗨️ 👤 **mistryminded** 1 year, 1 month ago

**Selected Answer: A**

Correct answer is A  
upvoted 1 times

🗨️ 👤 **SuperDevops** 1 year, 1 month ago

I took the test yesterday and didn't pass, NO ISSUE is from here. The questions are totally new  
upvoted 4 times

🗨️ 👤 **SuperDevops** 1 year, 1 month ago

I took the test yesterday and didn't pass, NO ISSUE is from here. The questions are totally new, don't use this dump.  
upvoted 2 times

🗨️ 👤 **Bhupals** 1 year, 12 months ago

Reference: itself says answer is A  
<https://cloud.google.com/blog/products/identity-security/using-your-existing-identity-management-system-with-google-cloud-platform>  
upvoted 2 times

🗨️ 👤 **jonclem** 2 years, 1 month ago

Interestingly, the link is correct, however, the answer is wrong. I agree with the consensus of the actual answer being A.  
upvoted 1 times

When creating a secure container image, which two items should you incorporate into the build if possible? (Choose two.)

- A. Ensure that the app does not run as PID 1.
- B. Package a single app as a container.
- C. Remove any unnecessary tools not needed by the app.
- D. Use public container images as a base image for the app.
- E. Use many container image layers to hide sensitive information.

Correct Answer: BC

Reference:

<https://cloud.google.com/solutions/best-practices-for-building-containers>

Community vote distribution



- tzKhalil**

Highly Voted

1 year, 7 months ago

BC is the answer.  
A is wrong, [https://cloud.google.com/architecture/best-practices-for-building-containers#solution\\_1\\_run\\_as\\_pid\\_1\\_and\\_register\\_signal\\_handlers](https://cloud.google.com/architecture/best-practices-for-building-containers#solution_1_run_as_pid_1_and_register_signal_handlers)  
upvoted 12 times
- GCP72**

Most Recent

3 months, 4 weeks ago

Selected Answer: BC  
The answer is BC  
upvoted 1 times
- minostrozaml2**

11 months, 1 week ago

Took the task today, only 5 question from this dump, the rest are new questions.  
upvoted 1 times
- jits1984**

11 months, 3 weeks ago

I took the test, and only got 6 questions from this dump. The administrators won't let me comment on the main page.  
upvoted 2 times
- SuperDevops**

1 year, 1 month ago

I took the test yesterday and didn't pass, NO ISSUE is from here. The questions are totally new  
Whizlabs it's OK  
upvoted 3 times
- VenkatGCP1**

1 year ago

This dude copy pasted same comment everywhere looks like someone from whizlabs trying to advertise here  
upvoted 4 times
- jits1984**

11 months, 4 weeks ago

No he is right. I appeared for the exam, earlier this month. only 6 questions from this dump. All new questions.  
upvoted 2 times
- dopeb64075**

1 year ago

I believe this guy is from Whizlabs. That platform is crap as this span. Selling Google's own free questions and lots of false questions/answers. I sent them a few corrections with references and they didn't mind to update.  
upvoted 6 times
- SuperDevops**

1 year, 1 month ago

it is AE  
upvoted 1 times
- Jane111**

1 year, 7 months ago

It should be A,B  
upvoted 1 times
- WakandaF**

1 year, 8 months ago

So, its B C?  
upvoted 1 times
- bluetaurianbull**

1 year, 9 months ago



To add to my previous comment  
"A process running as PID 1 inside a container is treated specially by Linux: it ignores any signal with the default action. So, the process will not terminate on SIGINT or SIGTERM unless it is coded to do so."  
Looks like this could be an issue when talking about security, a malicious coder can write a piece of code to eat all resources on the host with this one bad PID#1  
What do you think guys??  
upvoted 1 times

  **lollo1234** 1 year, 8 months ago  
You don't usually want your container to get killed instantly - you want to see the SIGINT or SIGTERM command and respond. For example, in a webserver you may stop accepting connections, and respond to the remaining open ones, before calling exit()  
upvoted 3 times

  **bluetaurianbull** 1 year, 9 months ago  
To add to my previous comment  
"A process running as PID 1 inside a container is treated specially by Linux: it ignores any signal with the default action. So, the process will not terminate on SIGINT or SIGTERM unless it is coded to do so."  
upvoted 1 times

  **bluetaurianbull** 1 year, 9 months ago  
Really??? Wat about (A)  
When the process with pid 1 die for any reason, all other processes are killed with KILL signal.  
  
Shouldnt A be one of the biggest risk when we talk about container security???  
upvoted 2 times

  **badrik** 6 months, 3 weeks ago  
I don't think this is a valid action to do to improve security perhaps it helps more to improve operational excellence. Imagine you are running production application in a container and it is signalled by container run time to terminate. In this case you don't have the running container to understand what would be issue ( though you can look at the events in modern container orchestration platform but imagine you are running a simple container ). Coming back to your concern. you don't generally run some rubbish container images in your container platform and this build process is very deliberate one.  
upvoted 1 times

  **kubosuke** 1 year, 9 months ago  
bc of bc  
upvoted 1 times

  **[Removed]** 2 years, 1 month ago  
Ans - BC  
upvoted 1 times

  **saurabh1805** 2 years, 2 months ago  
vote for B and C  
upvoted 1 times

  **MohitA** 2 years, 4 months ago  
BC for sure  
upvoted 1 times

  **ArizonaClassics** 2 years, 5 months ago  
BC on point!  
upvoted 2 times

  **KILLMAD** 2 years, 9 months ago  
I agree BC  
upvoted 4 times

A customer needs to launch a 3-tier internal web application on Google Cloud Platform (GCP). The customer's internal compliance requirements dictate that end-user access may only be allowed if the traffic seems to originate from a specific known good CIDR. The customer accepts the risk that their application will only have SYN flood DDoS protection. They want to use GCP's native SYN flood protection. Which product should be used to meet these requirements?


- A. Cloud Armor
- B. VPC Firewall Rules
- C. Cloud Identity and Access Management
- D. Cloud CDN

**Correct Answer: A**

Reference:

<https://cloud.google.com/blog/products/identity-security/understanding-google-cloud-armors-new-waf-capabilities>

*Community vote distribution*



🗲️ 👤 **KILLMAD** Highly Voted 👍 2 years, 9 months ago

Answer is A  
upvoted 13 times

🗲️ 👤 **Alokep** Most Recent 🕒 3 weeks, 4 days ago

Answer A  
upvoted 1 times

🗲️ 👤 **AzureDP900** 1 month, 2 weeks ago

Cloud Armor  
upvoted 1 times

🗲️ 👤 **Premumar** 1 month, 3 weeks ago

Selected Answer: A  
Cloud Armor  
upvoted 1 times

🗲️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

Selected Answer: A  
A. Cloud Armor  
upvoted 1 times

🗲️ 👤 **GCP72** 3 months, 4 weeks ago

Selected Answer: A  
The correct answer is A  
upvoted 1 times

🗲️ 👤 **Jeanphi72** 4 months, 3 weeks ago

Selected Answer: B  
See here: <https://cloud.google.com/files/GCPDDoSprotection-04122016.pdf>  
DDoS Protection by enabling Proxy-based Load Balancing  
○ When you enable HTTP(S) Load Balancing or SSL proxy Load Balancing, Google infrastructure mitigates and absorbs many Layer 4 and below attacks, such as SYN floods, IP fragment floods, port exhaustion, etc.  
○ If you have HTTP(S) Load Balancing with instances in multiple regions, you are able to disperse your attack across instances around the globe.  
upvoted 1 times

🗲️ 👤 **Loved** 1 month, 4 weeks ago

You're right, but you want to enable only traffic from specific CIDR, and Cloud Armor is the most efficient option as you can use it with LB (with firewall rules you have to define at least one rule for each VPC)  
upvoted 1 times

🗲️ 👤 **dar10** 10 months, 2 weeks ago

Cloud Armor is not available for Internal HTTP(S) load balancers. I'd vote for B. <https://cloud.google.com/load-balancing/docs/features#security>  
upvoted 1 times



  **zzaric** 8 months, 3 weeks ago

Internal LB will be used just for the backends, even not mentioned, the end-users will hit first the external LB - that works with CA - so the right answer is A

upvoted 1 times

  **minostrozaml2** 11 months, 1 week ago

Took the tesk today, only 5 question from this dump, the rest are new questions.

upvoted 2 times

  **ThisisJohn** 1 year ago

**Selected Answer: A**

If there were at least a L4 load balancer in the picture, I'd vote for B, since then the LB would take care of "GCP's native SYN flood protection", also considering that "The customer accepts the risk that their application will only have SYN flood DDoS protection."

With cloud armor I guess they get more protection that required on the question, but it seems to be the only entry that fulfills the requirements

upvoted 3 times

  **SuperDevops** 1 year, 1 month ago

I took the test yesterday and didn't pass, NO ISSUE is from here. The questions are totally new

Whizlabs > exam

upvoted 2 times

  **DebasishLowes** 1 year, 10 months ago

Ans : A

upvoted 1 times

  **Bharathy** 2 years ago

Cloud Armour supports ip range (CIDR) and also custom types can be defined using Armour Rule Language termed as CEL ( Common Expression Language ). Option A is correct

upvoted 1 times

  **[Removed]** 2 years, 1 month ago

Ans - A

upvoted 1 times

  **saaurabh1805** 2 years, 2 months ago

A is correct answer.

upvoted 1 times

  **zee001** 2 years, 3 months ago

A is correct !

upvoted 1 times

  **Ved** 2 years, 3 months ago

Ans is A

upvoted 1 times

A company is running workloads in a dedicated server room. They must only be accessed from within the private company network. You need to connect to these workloads from Compute Engine instances within a Google Cloud Platform project.

Which two approaches can you take to meet the requirements? (Choose two.)

- A. Configure the project with Cloud VPN.
- B. Configure the project with Shared VPC.
- C. Configure the project with Cloud Interconnect.
- D. Configure the project with VPC peering.
- E. Configure all Compute Engine instances with Private Access.

**Correct Answer:** AC

Community vote distribution

AC (100%)

  **KILLMAD** Highly Voted 2 years, 9 months ago

AC makes the most sense  
upvoted 25 times

  **rafaelc** 2 years, 9 months ago

Again you are correct  
upvoted 2 times

  **GCP72** Most Recent 3 months, 4 weeks ago

**Selected Answer: AC**

The correct answer is AC  
upvoted 1 times

  **shayke** 4 months, 2 weeks ago

**Selected Answer: AC**

the only answer  
upvoted 1 times

  **niberc21** 10 months, 2 weeks ago

**Selected Answer: AC**

A) IPsec VPN tunnels: <https://cloud.google.com/network-connectivity/docs/vpn/concepts/overview>  
C) Interconnect  
<https://cloud.google.com/network-connectivity/docs/interconnect/concepts/dedicated-overview>  
C)  
upvoted 2 times

  **minostrozaml2** 11 months, 1 week ago

Took the task today, only 5 question from this dump, the rest are new questions.  
upvoted 1 times

  **SuperDevops** 1 year, 1 month ago

I took the test yesterday and didn't pass, NO ISSUE is from here. The questions are totally new, don't use this dump.  
upvoted 2 times

  **DebasishLowe** 1 year, 10 months ago

Ans is AC  
upvoted 4 times

  **[Removed]** 2 years, 1 month ago

AC  
<https://cloud.google.com/solutions/secure-data-workloads-use-cases#gateway-for-hybrid>  
[https://cloud.google.com/solutions/secure-data-workloads-gcp-products#cloud\\_vpn](https://cloud.google.com/solutions/secure-data-workloads-gcp-products#cloud_vpn)  
upvoted 4 times

  **saurabh1805** 2 years, 2 months ago

A and C are correct answer here.  
upvoted 2 times

  **Rantu** 2 years, 2 months ago

AC is the answer.

upvoted 2 times

  **zee001** 2 years, 3 months ago

I checked GCP documentation and it states that to you can use either Cloud VPN or Cloud Interconnect to securely connect your on-premises network to your VPC network

upvoted 3 times

  **MohitA** 2 years, 4 months ago

Private Access won't help, AC is the answer

upvoted 1 times

  **aiwaai** 2 years, 4 months ago

Correct Answer: A, C

upvoted 1 times

  **bigdo** 2 years, 4 months ago

Ac A allow access to on-premise private ip address space with vpc with cloud interconnect they can access private private ip address space layer 2

upvoted 1 times

  **bigdo** 2 years, 4 months ago

CE peering is on gcp vpc only options

upvoted 2 times

  **bigdo** 2 years, 4 months ago

CD peering is on gcp vpc only options

upvoted 1 times

  **soukumar369** 2 years, 1 month ago

Again you are wrong

upvoted 2 times

  **ArizonaClassics** 2 years, 4 months ago

AC as well

upvoted 2 times

A customer implements Cloud Identity-Aware Proxy for their ERP system hosted on Compute Engine. Their security team wants to add a security layer so that the ERP systems only accept traffic from Cloud Identity-Aware Proxy. What should the customer do to meet these requirements?

- A. Make sure that the ERP system can validate the JWT assertion in the HTTP requests.
- B. Make sure that the ERP system can validate the identity headers in the HTTP requests.
- C. Make sure that the ERP system can validate the x-forwarded-for headers in the HTTP requests.
- D. Make sure that the ERP system can validate the user's unique identifier headers in the HTTP requests.

Correct Answer: A

Community vote distribution



- ArizonaClassics**

Highly Voted

 2 years, 4 months ago

A is right see : <https://cloud.google.com/iap/docs/signed-headers-howto>  
upvoted 15 times
- bolu**

Highly Voted

 1 year, 11 months ago

Use Cryptographic Verification  
If there is a risk of IAP being turned off or bypassed, your app can check to make sure the identity information it receives is valid. This uses a third web request header added by IAP, called X-Goog-IAP-JWT-Assertion. The value of the header is a cryptographically signed object that also contains the user identity data. Your application can verify the digital signature and use the data provided in this object to be certain that it was provided by IAP without alteration.  
  
So answer is A  
upvoted 8 times
- GCP72**

Most Recent

 3 months, 4 weeks ago

Selected Answer: A

The correct answer is A  
upvoted 1 times
- minostrozaml2** 11 months, 1 week ago

Took the tesk today, only 5 question from this dump, the rest are new questions.  
upvoted 1 times
- BigCap** 3 weeks, 1 day ago

are u talking about all 186 questions or just free exam topic questions? It's gonna be usefull for to know that !! thank u  
upvoted 1 times
- SuperDevops** 1 year, 1 month ago

I took the test yesterday and didn't pass, NO ISSUE is from here. The questions are totally new Whizlabs it's OK  
upvoted 1 times
- sc\_cloud\_learn** 1 year, 7 months ago

Agree A makes more sense  
upvoted 3 times
- DebasishLowes** 1 year, 10 months ago

Ans is A  
upvoted 3 times
- [Removed]** 2 years, 1 month ago

Ans - A  
upvoted 1 times
- saurabh1805** 2 years, 2 months ago

A is correct option here.  
upvoted 1 times
- MohitA** 2 years, 4 months ago

A is the one

upvoted 1 times

  **KILLMAD** 2 years, 9 months ago

Ans is A

upvoted 4 times

A company has been running their application on Compute Engine. A bug in the application allowed a malicious user to repeatedly execute a script that results in the Compute Engine instance crashing. Although the bug has been fixed, you want to get notified in case this hack re-occurs. What should you do?

- A. Create an Alerting Policy in Stackdriver using a Process Health condition, checking that the number of executions of the script remains below the desired threshold. Enable notifications.
- B. Create an Alerting Policy in Stackdriver using the CPU usage metric. Set the threshold to 80% to be notified when the CPU usage goes above this 80%.
- C. Log every execution of the script to Stackdriver Logging. Create a User-defined metric in Stackdriver Logging on the logs, and create a Stackdriver Dashboard displaying the metric.
- D. Log every execution of the script to Stackdriver Logging. Configure BigQuery as a log sink, and create a BigQuery scheduled query to count the number of executions in a specific timeframe.

**Correct Answer:** C

Reference:

<https://cloud.google.com/logging/docs/logs-based-metrics/>

*Community vote distribution*

A (50%)

B (50%)

  **rafaelc** Highly Voted 2 years, 9 months ago

The question asks "you want to get notified in case this hack re-occurs."

Only A has notifications in the answer so that should be the answer as having dashboards in stackdriver wont notify you of anything.

upvoted 23 times

  **ananthanarayanante** 2 years, 6 months ago

I agree it should be A

upvoted 7 times

  **serg3d** Highly Voted 2 years, 5 months ago

It's not necessary that running a malicious script multiple times will affect CPU usage. And, CPU usage can occur during usual normal workloads.

A

upvoted 7 times

  **DA95** Most Recent 3 days, 16 hours ago

**Selected Answer: A**

Option A is the most appropriate solution to get notified in case the hack re-occurs. In this option, you can create an Alerting Policy in Stackdriver using a Process Health condition to check the number of executions of the script. You can set a threshold for the number of executions, and if the number of executions goes above the threshold, you can enable notifications to be alerted about the hack.

upvoted 1 times

  **DA95** 3 days, 16 hours ago

Option B is not an appropriate solution, as it does not address the issue of the hack re-occurring. Monitoring CPU usage alone may not be enough to detect a hack, as the CPU usage may not necessarily go above the threshold set in the alerting policy.

Option C is also not an appropriate solution, as creating a user-defined metric and dashboard based on the logs of script executions will not alert you in real-time if the hack re-occurs. You would need to manually check the dashboard to see if the hack has re-occurred, which may not be practical in a high-security scenario.

Option D is not an appropriate solution, as it involves logging the script executions to Stackdriver Logging and then configuring a BigQuery sink to count the number of executions. This would not alert you in real-time if the hack re-occurs, as you would need to wait for the scheduled query to run and then check the results.

upvoted 1 times

  **shayke** 1 week, 3 days ago

**Selected Answer: A**

a is the correct ans

upvoted 1 times

  **Premumar** 1 month, 3 weeks ago

**Selected Answer: A**

Other options won't provide any notification to the user. So, the correct answer is A.

upvoted 1 times



🗄️ 👤 **GCP72** 3 months, 4 weeks ago

**Selected Answer: B**

The correct answer is B  
upvoted 3 times

🗄️ 👤 **minostrozaml2** 11 months, 1 week ago

Took the task today, only 5 question from this dump, the rest are new questions.  
upvoted 2 times

🗄️ 👤 **SuperDevops** 1 year, 1 month ago

I took the test yesterday and didn't pass, NO ISSUE is from here. The questions are totally new, don't use this dump.  
upvoted 1 times

🗄️ 👤 **jits1984** 1 year ago

What are you saying, where should we go for new dumps @SuperDevops?  
upvoted 4 times

🗄️ 👤 **Jeanphi72** 4 months, 3 weeks ago

These are Whizzlabs people: WhizzLabs is one of the worst place to train for exams and instead of trying to become better (maybe because of ignorance) they simply try to pull down good sites to learn with ...  
upvoted 3 times

🗄️ 👤 **Jane111** 1 year, 7 months ago

The bug has been fixed, so even if somebody runs the same script, it will affect nothing. Checking against the same script, creating Process-health policy will do nothing. But if the hack reappears and the same script is run, the A will trigger  
upvoted 3 times

🗄️ 👤 **Jane111** 1 year, 7 months ago

The bug has been fixed, so even if somebody runs the same script, it will affect nothing. Checking against the same script, creating Process-health policy will do nothing  
upvoted 2 times

🗄️ 👤 **Jane111** 1 year, 7 months ago

There is no 'Process Health condition' but Process-health policy  
A process-health policy can notify you if the number of processes that match a pattern crosses a threshold. This can be used to tell you, for example, that a process has stopped running.

This policy sends a notification to the specified notification channel when no process matching the string nginx, running as user www, has been available for more than 5 minutes:  
upvoted 1 times

🗄️ 👤 **DebasishLowes** 1 year, 10 months ago

Ans : A  
upvoted 1 times

🗄️ 👤 **soukumar369** 1 year, 11 months ago

A. Correct  
B. "CPU usage goes above this 80%". It's not granted that script execution will increase CPU usage.  
C&D. Not providing any notification  
upvoted 3 times

🗄️ 👤 **[Removed]** 2 years, 1 month ago

A  
<https://cloud.google.com/monitoring/alerts/types-of-conditions#metric-threshold>  
upvoted 2 times

🗄️ 👤 **saurabh1805** 2 years, 2 months ago

A seems to be closest and best option.  
upvoted 2 times

🗄️ 👤 **saurabh1805** 2 years, 1 month ago

<https://cloud.google.com/monitoring/alerts/ui-conditions-ga#process-health>  
upvoted 1 times

🗄️ 👤 **zee001** 2 years, 2 months ago

I think the answer is B based on the article I found on Medium, which states that Metrics are used to represent the state or health of your system over time, and In Stackdriver specifically, metrics are the only kind of data that can be used to create alerts via alerting policies.  
The question is also asking about monitoring the health of system.  
<https://medium.com/google-cloud/can-you-alert-on-logs-in-stackdriver-7dfb07f495c0>  
upvoted 1 times

🗄️ 👤 **SA1** 2 years, 2 months ago

Looks B to me. Not sure what is "Process Health condition".  
upvoted 1 times






Your team needs to obtain a unified log view of all development cloud projects in your SIEM. The development projects are under the NONPROD organization folder with the test and pre-production projects. The development projects share the ABC-BILLING billing account with the rest of the organization.

Which logging export strategy should you use to meet the requirements?

- A. 1. Export logs to a Cloud Pub/Sub topic with folders/NONPROD parent and includeChildren property set to True in a dedicated SIEM project. 2. Subscribe SIEM to the topic.
- B. 1. Create a Cloud Storage sink with billingAccounts/ABC-BILLING parent and includeChildren property set to False in a dedicated SIEM project. 2. Process Cloud Storage objects in SIEM.
- C. 1. Export logs in each dev project to a Cloud Pub/Sub topic in a dedicated SIEM project. 2. Subscribe SIEM to the topic.
- D. 1. Create a Cloud Storage sink with a publicly shared Cloud Storage bucket in each project. 2. Process Cloud Storage objects in SIEM.

Correct Answer: B

Community vote distribution



**xhova** Highly Voted 2 years, 8 months ago

Answer is A. [https://cloud.google.com/logging/docs/export/aggregated\\_sinks](https://cloud.google.com/logging/docs/export/aggregated_sinks)  
upvoted 24 times

**Ishu\_awsguy** 1 month ago

with this you would also be getting logs for Preprod and other environments under the folder. Hence A is eliminated.  
Answer should be C  
upvoted 3 times

**TNT87** Highly Voted 1 year, 10 months ago

To use the aggregated sink feature, create a sink in a Google Cloud organization or folder and set the sink's includeChildren parameter to True. That sink can then export log entries from the organization or folder, plus (recursively) from any contained folders, billing accounts, or projects. You can use the sink's filter to specify log entries from projects, resource types, or named logs.  
[https://cloud.google.com/logging/docs/export/aggregated\\_sinks](https://cloud.google.com/logging/docs/export/aggregated_sinks)

so the Ans is A  
upvoted 7 times

**DA95** Most Recent 9 hours, 43 minutes ago

Option B is not correct because setting the includeChildren property to False will exclude the test and pre-production projects from the log export.

Option C is not correct because it would require you to create a separate Cloud Pub/Sub topic for each development project, which would not meet the requirement to obtain a unified log view of all development projects.

Option D is not correct because using a publicly shared Cloud Storage bucket would not provide a secure way to store and access the logs. It is generally not recommended to use publicly shared Cloud Storage buckets for storing sensitive data such as logs.  
upvoted 1 times

**PST21** 2 days, 14 hours ago

You can create aggregated sinks for Google Cloud folders and organizations. Because neither Cloud projects nor billing accounts contain child resources, you can't create aggregated sinks for those. which means logs will be for the folder and contains non dev entries as well  
Ans -C  
upvoted 1 times

**PST21** 4 days, 13 hours ago

You can create aggregated sinks for Google Cloud folders and organizations. Because neither Cloud projects nor billing accounts contain child resources, you can't create aggregated sinks for those.  
So ans has to be c  
upvoted 1 times

**shayke** 1 week, 3 days ago

Selected Answer: C  
c is the right ans. more percise way  
upvoted 2 times

**Meyucho** 1 month, 1 week ago

Selected Answer: A

If you create a sink in the NONPROD folder level all logs from the projects in that folder will be routed wherever you want. So, you don't need to configure sinks in every project nor setup this every time you have a new project in the NONPROD folder. Answer is A

upvoted 1 times

  **Ishu\_aws** 1 month ago

But it also gives you logs of the projects that are not required.  
The question asks about logs from only development. hence C is the answer

upvoted 1 times

  **GCP72** 3 months, 4 weeks ago

**Selected Answer: C**

The correct answer is C

upvoted 2 times

  **absipat** 6 months, 2 weeks ago

a of course

upvoted 1 times

  **Medofree** 8 months, 2 weeks ago

By elimination the answer is C

upvoted 4 times

  **minostrozaml2** 11 months, 1 week ago

Took the test today, only 5 question from this dump, the rest are new questions.

upvoted 1 times

  **Lancyqusa** 1 year ago

"Your team needs to obtain a unified log view of all development cloud projects in your SIEM" - This means we are ONLY interested in development projects.  
"The development projects are under the NONPROD organization folder with the test and pre-production projects" - We will need to filter out development from others i.e test and pre-prod.  
"The development projects share the ABC-BILLING billing account with the rest of the organization." - This is unnecessary information.  
The only option that filters the log is C - so the answer must be C.

upvoted 4 times

  **SuperDevops** 1 year, 1 month ago

I took the test yesterday and didn't pass, NO ISSUE is from here. The questions are totally new, don't use this dump.

upvoted 2 times

  **KyubiBlaze** 1 year, 3 months ago

Guys, C is the answer.  
A - NO, This will bring in test and pre-prod as well, this does not solve our problem  
B - NO, This has nothing to do with the question,  
C - YES- Specifically take development projects, and centralises the logs  
D - Once the work public comes in answer, it is wrong.

Request the admins to set answer to C  
This is seriously a simple question

upvoted 2 times

  **[Removed]** 1 year, 9 months ago

A is the answer.  
To use the aggregated sink feature, create a sink in a Google Cloud organization or folder and set the sink's includeChildren parameter to True. That sink can then export log entries from the organization or folder, plus (recursively) from any contained folders, billing accounts, or projects. You can use the sink's filter to specify log entries from projects, resource types, or named logs.

Please note that sinks filters can be used to get exact logs required and the given requirement from the question will be fulfilled.

[https://cloud.google.com/logging/docs/export/aggregated\\_sinks](https://cloud.google.com/logging/docs/export/aggregated_sinks)

upvoted 3 times

  **[Removed]** 2 years, 1 month ago

Ans may be A or C.  
A- not very cost effective and need filtering  
C - Tedious to create for all dev projects

upvoted 4 times

  **saaurabh1805** 2 years, 2 months ago

I will vote for option C.

upvoted 4 times

A customer needs to prevent attackers from hijacking their domain/IP and redirecting users to a malicious site through a man-in-the-middle attack.

Which solution should this customer use?

- A. VPC Flow Logs
- B. Cloud Armor
- C. DNS Security Extensions
- D. Cloud Identity-Aware Proxy

**Correct Answer:** C

Reference:  
<https://cloud.google.com/blog/products/gcp/dnssec-now-available-in-cloud-dns>

*Community vote distribution*  

C (100%)

-   **ESP\_SAP** Highly Voted 2 years, 1 month ago

Correct Answer is (C):

DNSSEC — use a DNS registrar that supports DNSSEC, and enable it. DNSSEC digitally signs DNS communication, making it more difficult (but not impossible) for hackers to intercept and spoof.

Domain Name System Security Extensions (DNSSEC) adds security to the Domain Name System (DNS) protocol by enabling DNS responses to be validated. Having a trustworthy Domain Name System (DNS) that translates a domain name like `www.example.com` into its associated IP address is an increasingly important building block of today’s web-based applications. Attackers can hijack this process of domain/IP lookup and redirect users to a malicious site through DNS hijacking and man-in-the-middle attacks. DNSSEC helps mitigate the risk of such attacks by cryptographically signing DNS records. As a result, it prevents attackers from issuing fake DNS responses that may misdirect browsers to nefarious websites.

<https://cloud.google.com/blog/products/gcp/dnssec-now-available-in-cloud-dns>

upvoted 12 times
-   **AzureDP900** Most Recent 1 month, 2 weeks ago

C is right

upvoted 1 times
-   **GCP72** 3 months, 4 weeks ago

Selected Answer: C

The correct answer is C

upvoted 1 times
-   **minostrozaml2** 11 months, 1 week ago

Took the tesk today, only 5 question from this dump, the rest are new questions.

upvoted 1 times
-   **shreenine** 1 year, 2 months ago

C is the correct answer indeed.

upvoted 2 times
-   **Kameswara** 1 year, 6 months ago

C. Attackers can hijack this process of domain/IP lookup and redirect users to a malicious site through DNS hijacking and man-in-the-middle attacks. DNSSEC helps mitigate the risk of such attacks by cryptographically signing DNS records. As a result, it prevents attackers from issuing fake DNS responses that may misdirect browsers to nefarious websites.

upvoted 4 times
-   **sc\_cloud\_learn** 1 year, 7 months ago

C. DNSSEC is the ans

upvoted 2 times
-   **ASG** 1 year, 10 months ago

Its man in the middle attack protection. The traffic first needs to reach cloud armour before you can make use of cloud armour related protection. DNS can be hijacked if you dont use DNSSEC. Its your DNS that needs to resolve the initial request before traffic is directed to cloud armour. Option C is most appropriate measure. (think of sequencing of how traffic will flow)

upvoted 3 times
-   **bolu** 1 year, 11 months ago

The answers from rest of the folks are complete unreliable. The right answer is Cloud Armor based on my Hands-On labs in Qwiklabs. Reason: Creating a policy in Cloud Armor sends 403 forbidden message for man-in-the middle-attack. Reference: <https://cloud.google.com/blog/products/identity-security/identifying-and-protecting-against-the-largest-ddos-attacks> Some more: <https://cloud.google.com/armor> Refer this lab: [https://www.qwiklabs.com/focuses/1232?catalog\\_rank=%7B%22rank%22%3A1%2C%22num\\_filters%22%3A0%2C%22has\\_search%22%3Atrue%7D&parent=catalog&search\\_id=8696512](https://www.qwiklabs.com/focuses/1232?catalog_rank=%7B%22rank%22%3A1%2C%22num_filters%22%3A0%2C%22has_search%22%3Atrue%7D&parent=catalog&search_id=8696512)  
upvoted 2 times

  **KyubiBlaze** 1 year, 3 months ago  
No, C is the correct answer.  
upvoted 1 times

  **[Removed]** 2 years, 1 month ago  
Ans - C  
upvoted 2 times

  **saurabh1805** 2 years, 2 months ago  
DNSEC is the thing, Option C  
upvoted 2 times

  **MohitA** 2 years, 4 months ago  
C, Yes for sure DNSSEC  
upvoted 2 times

  **bigdo** 2 years, 4 months ago  
C DNSSEC  
upvoted 2 times

  **ArizonaClassics** 2 years, 4 months ago  
Option C is Perfect. DNSSECURITY!  
upvoted 2 times

  **KILLMAD** 2 years, 9 months ago  
I agree it's C  
upvoted 1 times

A customer deploys an application to App Engine and needs to check for Open Web Application Security Project (OWASP) vulnerabilities. Which service should be used to accomplish this?

- A. Cloud Armor
- B. Google Cloud Audit Logs
- C. Web Security Scanner
- D. Anomaly Detection

**Correct Answer:** C

Reference:

<https://cloud.google.com/security-scanner/>

*Community vote distribution*

C (100%)

🗲️ 👤 **Tabayashi** Highly Voted 7 months, 4 weeks ago

Answer is (C).

Web Security Scanner supports categories in the OWASP Top Ten, a document that ranks and provides remediation guidance for the top 10 most critical web application security risks, as determined by the Open Web Application Security Project (OWASP).

[https://cloud.google.com/security-command-center/docs/concepts-web-security-scanner-overview#detectors\\_and\\_compliance](https://cloud.google.com/security-command-center/docs/concepts-web-security-scanner-overview#detectors_and_compliance)

upvoted 9 times

🗲️ 👤 **GCP72** Most Recent 3 months, 4 weeks ago

**Selected Answer: C**

The correct answer is C

upvoted 1 times

🗲️ 👤 **PopeyeTheSailorMan** 5 months ago

This is called DAST (Dynamic Application Security Testing) through tools such as BurpSuite,ZAP in normal non-cloud deployments but the same has been done through web security scanner in GCP hence my answer is C

upvoted 2 times

A customer's data science group wants to use Google Cloud Platform (GCP) for their analytics workloads. Company policy dictates that all data must be company-owned and all user authentications must go through their own Security Assertion Markup Language (SAML) 2.0 Identity Provider (IdP). The Infrastructure Operations Systems Engineer was trying to set up Cloud Identity for the customer and realized that their domain was already being used by G Suite. How should you best advise the Systems Engineer to proceed with the least disruption?

- A. Contact Google Support and initiate the Domain Contestation Process to use the domain name in your new Cloud Identity domain.
- B. Register a new domain name, and use that for the new Cloud Identity domain.
- C. Ask Google to provision the data science manager's account as a Super Administrator in the existing domain.
- D. Ask customer's management to discover any other uses of Google managed services, and work with the existing Super Administrator.

Correct Answer: C

Community vote distribution



**syllox** Highly Voted 1 year, 7 months ago  
Ans :D  
upvoted 8 times

**TNT87** Highly Voted 2 years, 1 month ago  
The answer is A  
"This domain is already in use"  
If you receive this message when trying to sign up for a Google service, it might be because:  
  
You recently removed this domain from another managed Google account. It can take 24 hours (or 7 days if you purchased your account from a reseller) before you can use the domain with a new account.  
You or someone in your organization already created a managed Google account with your domain. Try resetting the administrator password and we'll send an email to the secondary email you provided when you signed up, telling you how to access the account.  
You're using the domain with another managed Google account that you own. If so, remove the domain from the other account.  
Contact us  
If none of these applies, the previous owner of your domain might have signed up for a Google service. Fill out this form and the Support team will get back to you within 48 hours.  
upvoted 8 times

**lollo1234** 1 year, 8 months ago  
Answer is D - there is no evidence that the account is lost, or similar. In a large corp it is very possible that someone (the IT org) has registered with google, and the Data science Department simply haven't been given access to it yet.  
upvoted 15 times

**shayke** Most Recent 1 week, 2 days ago  
Selected Answer: A  
A is the right ans  
upvoted 1 times

**GCP72** 3 months, 4 weeks ago  
Selected Answer: D  
The answer is D  
upvoted 1 times

**Ksrp** 9 months, 3 weeks ago  
its A , <https://support.google.com/a/answer/6286258?hl=en#:~:text=If%20you%20get%20an%20alert,that%20you%20don't%20manage.>  
upvoted 1 times

**minostrozaml2** 11 months, 1 week ago  
Took the task today, only 5 question from this dump, the rest are new questions.  
upvoted 1 times

**SuperDevops** 1 year, 1 month ago  
I took the test yesterday and didn't pass, NO ISSUE is from here. The questions are totally new, don't use this dump.  
upvoted 1 times

**PopeyeTheSailorMan** 5 months ago



Please do not trust this answer from SuperDevops and minostrozaml users-these are human bots sending 'social engineering' texts and when you ask them a question, they will come back and say whizlabs is the solution to get questions-hence it is fake and it is everywhere pasted across majority of this exam question discussions-please click REPORT button for the answers above and report to admin.

upvoted 7 times

  **Rightsaidfred** 1 month, 1 week ago

Thanks Popeye!  
upvoted 1 times

  **idtroo** 1 year, 9 months ago

Answer is D.

<https://support.google.com/cloudidentity/answer/7389973>  
If you're an existing Google Workspace customer  
Follow these steps to sign up for Cloud Identity Premium:

Using your administrator account, sign in to the Google Admin console at [admin.google.com](https://admin.google.com).  
From the Admin console Home page, at the top left, click Menu ""and thenBillingand thenGet more services.  
Click Cloud Identity.  
Next to Cloud Identity Premium, click Start Free Trial.  
Follow the guided instructions.  
upvoted 6 times

  **TNT87** 1 year, 10 months ago

Sorry Ans is D  
upvoted 5 times

  **CloudTrip** 1 year, 10 months ago

A, B are definitely not the answer for this. Most of you are aligned with D but can somebody explain what is wrong with C ? Their domain is already used by the G-Suite. It will be least disruptive also.  
upvoted 1 times

  **lollo1234** 1 year, 8 months ago

Principle of least privilege - should the 'data science manager' be a superadmin?? Probably not. Hence D, work with the existing admin - we assume that they were chosen sensibly.  
upvoted 2 times

  **ronron89** 2 years ago

I think its D.

@SomabrataPani: did you pass this exam yet?  
upvoted 2 times

  **[Removed]** 2 years, 1 month ago

Ans - D  
upvoted 2 times

  **saaurabh1805** 2 years, 2 months ago

D is best answer here.  
upvoted 2 times

  **aiwaai** 2 years, 4 months ago

Answer. is D  
upvoted 2 times

  **ArizonaClassics** 2 years, 4 months ago

I am considering B or D but still needs some narrowing down to choose the right answer  
upvoted 1 times

  **Kouuupobol** 2 years, 5 months ago

If SAML is already used and configured, then Google support won't be able to add a new account since it doesn't manage the user directory.  
Answer is D  
upvoted 2 times

  **ihussainkhalid** 2 years, 7 months ago

if domain already in use then you cant release that, you need to request google to assign admin rights after verification steps as you owner of that domain,  
upvoted 2 times



A business unit at a multinational corporation signs up for GCP and starts moving workloads into GCP. The business unit creates a Cloud Identity domain with an organizational resource that has hundreds of projects.

Your team becomes aware of this and wants to take over managing permissions and auditing the domain resources.

Which type of access should your team grant to meet this requirement?

- A. Organization Administrator
- B. Security Reviewer
- C. Organization Role Administrator
- D. Organization Policy Administrator

**Correct Answer: C**

Community vote distribution

C (100%)

ffdd1234 Highly Voted 1 year, 11 months ago

Answer A > Its the only one that allow you to manage permissions on the projects

answer B > dont have any iam set permission so is not correct

C > organizationRoleAdmin let you only create custom roles, you cant assign it to anyone ( so with thisone you cant manage permissions just create roles)

D> org policyes are for manage the ORG policies constrains , that is not about project permissions,

for me the correct is A

upvoted 20 times

zanhsieh Highly Voted 2 years ago

C. After carefully review this link:

<https://cloud.google.com/iam/docs/understanding-roles>

my opinion is based on 'the least privilege' practice, that future domain shall not get granted automatically:

A - Too broad permissions. The question asked "The business unit creates a Cloud Identity domain..." does not imply your team should be granted for ALL future domain(s) (domain = folder) permission management.

B - Security Reviewer does not have "set\*" permission. All this role could do is just looking, not management.

C - The best answer so far. Only the domain current created and underneath iam role assignment as well as change.

D - Too broad permissions on the organization level. In other words, this role could make policy but future domains admin could hijack the role names / policies to do not desired operations.

upvoted 11 times

Loved 1 month, 1 week ago

"If you have an organization associated with your Google Cloud account, the Organization Role Administrator role enables you to administer all custom roles in your organization", it can not be C

upvoted 1 times

zzaric 8 months, 3 weeks ago

C - can't do a job - they have to manage the IAP permissions, C doesn't have setIAM permissions and the role is only for creating Custom Roles - see the permissions that it contains:

```
iam.roles.create
iam.roles.delete
iam.roles.get
iam.roles.list
iam.roles.undelete
iam.roles.update
resourcemanager.organizations.get
resourcemanager.organizations.getIamPolicy
resourcemanager.projects.get
resourcemanager.projects.getIamPolicy
resourcemanager.projects.list
```

upvoted 1 times

zzaric 8 months, 3 weeks ago

IAM - not IAP - typo

upvoted 1 times

GCP72 Most Recent 3 months, 4 weeks ago

**Selected Answer: C**

The correct answer is C

upvoted 1 times



- 🗨️ 👤 **Medofree** 8 months, 2 weeks ago  
Answer is A, among the 4, it is the only role able de manage permissions  
upvoted 2 times
- 🗨️ 👤 **minostrozaml2** 11 months, 1 week ago  
Took the task today, only 5 question from this dump, the rest are new questions.  
upvoted 1 times
- 🗨️ 👤 **Lancyqusa** 1 year ago  
The answer must be A - check out the example that allows the CTO to setup permissions for the security team:  
[https://cloud.google.com/iam/docs/job-functions/auditing#scenario\\_operational\\_monitoring](https://cloud.google.com/iam/docs/job-functions/auditing#scenario_operational_monitoring)  
upvoted 1 times
- 🗨️ 👤 **OSNG** 1 year, 3 months ago  
Its A.  
They are looking for Domain Resources Management i.e. Projects, Folders, Permissions. and only Organization Administrator is the only option allows it. Moreover, Organization Administrator is the only option that falls under "Used IN: Resource Manager"  
roles/resourcemanager.organizationAdmin  
upvoted 1 times
- 🗨️ 👤 **[Removed]** 1 year, 9 months ago  
C is the answer.  
Here are the permissions available to organizationRoleAdmin  
  
iam.roles.create  
iam.roles.delete  
iam.roles undelete  
iam.roles.get  
iam.roles.list  
iam.roles.update  
resourcemanager.projects.get  
resourcemanager.projects.getIamPolicy  
resourcemanager.projects.list  
resourcemanager.organizations.get  
resourcemanager.organizations.getIamPolicy  
  
There are sufficient as per least privilege policy. You can do user management as well as auditing.  
upvoted 5 times
- 🗨️ 👤 **[Removed]** 1 year, 9 months ago  
link - <https://cloud.google.com/iam/docs/understanding-custom-roles>  
upvoted 1 times
- 🗨️ 👤 **DebasishLowes** 1 year, 9 months ago  
Ans : D. As it's related to Resources, so definitely policy comes into picture.  
upvoted 1 times
- 🗨️ 👤 **HateMicrosoft** 1 year, 10 months ago  
Correct is D  
<https://cloud.google.com/resource-manager/docs/organization-policy/overview>  
upvoted 2 times
- 🗨️ 👤 **Bhupals** 2 years ago  
Role Permissions  
roles/iam.organizationRoleAdmin iam.roles.create  
iam.roles.delete  
iam.roles undelete  
iam.roles.get  
iam.roles.list  
iam.roles.update  
resourcemanager.projects.get  
resourcemanager.projects.getIamPolicy  
resourcemanager.projects.list  
resourcemanager.organizations.get  
resourcemanager.organizations.getIamPolicy  
upvoted 1 times
- 🗨️ 👤 **FatCharlie** 2 years ago  
The confusion here, in my opinion, is that the question is asking for the ability to manage roles & audit \_DOMAIN\_ resources.  
  
Domain resources in the GCP hierarchy are folders & projects, because those are the only things that can be directly under an Organization (aka Domain).  
  
The Organization Role Admin is the option that gives you the ability to manage custom roles & list folders & projects.  
upvoted 4 times
- 🗨️ 👤 **jonclem** 2 years, 1 month ago

Organization Policy Administrator has 2 assigned permissions: orgpolicy.policy.get  
orgpolicy.policy.set

Organization Role Administrator has 11 assigned permissions: iam.roles.create, iam.roles.delete, iam.roles.get, iam.roles.list, iam.roles.undelete, iam.roles.update, resourcemanager.organizations.get, resourcemanager.organizations.getIamPolicy, resourcemanager.projects.get, resourcemanager.projects.getIamPolicy, resourcemanager.projects.list,

While Security Review has over 700 permissions assigned to it.

With the question focusing on managing permissions and auditing I'd be inclined to go with option B: Security Reviewer.

upvoted 2 times

  **saurabh1805** 2 years, 1 month ago

D is correct answer here.

<https://cloud.google.com/resource-manager/docs/organization-policy/overview>

upvoted 1 times

  **genesis3k** 2 years, 1 month ago

Correct answer is D as the question is referring to 'what' (i.e. resources), not 'who' (i.e. role of a user/service account). Here is an excerpt from "Differences from Identity and Access Management" section:

"Identity and Access Management focuses on who, and lets the administrator authorize who can take action on specific resources based on permissions.

Organization Policy focuses on what, and lets the administrator set restrictions on specific resources to determine how they can be configured." at

<https://cloud.google.com/resource-manager/docs/organization-policy/overview>

upvoted 1 times

  **[Removed]** 2 years, 1 month ago

Ans - C

[https://cloud.google.com/iam/docs/understanding-custom-roles#organization\\_role\\_administrator\\_role](https://cloud.google.com/iam/docs/understanding-custom-roles#organization_role_administrator_role)

upvoted 2 times

  **HectorLeon2099** 2 years, 2 months ago

Answer is A. Is the only one that lets you auditing resources.

upvoted 1 times


An application running on a Compute Engine instance needs to read data from a Cloud Storage bucket. Your team does not allow Cloud Storage buckets to be globally readable and wants to ensure the principle of least privilege.

Which option meets the requirement of your team?

- A. Create a Cloud Storage ACL that allows read-only access from the Compute Engine instance's IP address and allows the application to read from the bucket without credentials.
- B. Use a service account with read-only access to the Cloud Storage bucket, and store the credentials to the service account in the config of the application on the Compute Engine instance.
- C. Use a service account with read-only access to the Cloud Storage bucket to retrieve the credentials from the instance metadata.
- D. Encrypt the data in the Cloud Storage bucket using Cloud KMS, and allow the application to decrypt the data with the KMS key.

**Correct Answer:** C

Community vote distribution



ESP\_SAP **Highly Voted** 2 years ago

Correct Answer is (B):

If your application runs inside a Google Cloud environment that has a default service account, your application can retrieve the service account credentials to call Google Cloud APIs. Such environments include Compute Engine, Google Kubernetes Engine, App Engine, Cloud Run, and Cloud Functions. We recommend using this strategy because it is more convenient and secure than manually passing credentials.

Additionally, we recommend you use Google Cloud Client Libraries for your application. Google Cloud Client Libraries use a library called Application Default Credentials (ADC) to automatically find your service account credentials. ADC looks for service account credentials in the following order:

<https://cloud.google.com/docs/authentication/production#automatically>  
upvoted 10 times

ChewB666 2 years ago

Hello guys!

Does anyone have the rest of the questions to share? :(  
I can't see the rest of the issues because of the subscription.  
upvoted 2 times

Medofree **Highly Voted** 8 months, 2 weeks ago

**Selected Answer: C**

Correct ans is C. The credentials are retrieved from the metedata server  
upvoted 9 times

Meyucho **Most Recent** 1 month, 1 week ago

**Selected Answer: A**

A CORRECT: It's the only answer when you use ACL to filter local IP's addresses and you can have the bucket without global access.  
B INCORRET: Doesn't use the least privilege principle.  
C INCORRECT: What credentials are we talking about!? To do this it's better option B.  
D INCORRECT: Need global access.  
upvoted 2 times

dat987 1 month, 2 weeks ago

**Selected Answer: B**

meta data do not set service account  
upvoted 1 times

GCP72 3 months, 4 weeks ago

**Selected Answer: C**

The correct answer is C  
upvoted 1 times

[Removed] 8 months, 3 weeks ago

B

If the environment variable GOOGLE\_APPLICATION\_CREDENTIALS is set, ADC uses the service account key or configuration file that the variable points to.  
<https://cloud.google.com/docs/authentication/production#automatically>

upvoted 1 times

🗳️ 👤 **AaronLee** 9 months, 1 week ago

The Answer is C

If the environment variable GOOGLE\_APPLICATION\_CREDENTIALS is set, ADC uses the service account key or configuration file that the variable points to.

If the environment variable GOOGLE\_APPLICATION\_CREDENTIALS isn't set, ADC uses the service account that is attached to the resource that is running your code.

[https://cloud.google.com/docs/authentication/production#passing\\_the\\_path\\_to\\_the\\_service\\_account\\_key\\_in\\_code](https://cloud.google.com/docs/authentication/production#passing_the_path_to_the_service_account_key_in_code)

upvoted 2 times

🗳️ 👤 **minostrozaml2** 11 months, 1 week ago

Took the task today, only 5 questions from this dump, the rest are new questions.

upvoted 2 times

🗳️ 👤 **jj\_618** 1 year, 3 months ago

So is it B or C?

upvoted 1 times

🗳️ 👤 **StanPeng** 10 months, 2 weeks ago

B for sure. C is wrong logic

upvoted 1 times

🗳️ 👤 **Medofree** 8 months, 2 weeks ago

No, C is the right answer, you don't need to generate credentials into GCP since they are stored into metadata server, the application will retrieve them automatically through a Google Lib (or even manually by calling the URL curl

<http://metadata.google.internal/computeMetadata/v1/instance/service-accounts/default/token> -H "Metadata-Flavor: Google")

upvoted 2 times

🗳️ 👤 **bolu** 1 year, 10 months ago

Answer can be either B or C due to the relevance to servicing account. But storing password in app is a worst practice and we read it several times everywhere online hence it results in C as a best answer to handle service account through metadata

upvoted 4 times

🗳️ 👤 **[Removed]** 2 years, 1 month ago

Ans - C

upvoted 1 times

🗳️ 👤 **HectorLeon2099** 2 years, 2 months ago

I'll go with B.

A - ACL's are not able to allow access based on IP

C - If you store the credentials in the metadata those will be public accessible by everyone with project access.

D - Too complex

upvoted 5 times

🗳️ 👤 **saaurabh1805** 2 years, 2 months ago

Yes B is best possible option. This is something Google also recommends.

<https://cloud.google.com/storage/docs/authentication#libauth>

upvoted 2 times

🗳️ 👤 **CHECK666** 2 years, 2 months ago

C is correct

upvoted 1 times

🗳️ 👤 **Moe666** 2 years, 2 months ago

C is the answer

upvoted 1 times

🗳️ 👤 **mlyu** 2 years, 3 months ago

Hi guys, How do we handle the requirement "does not allow Cloud Storage buckets to be globally readable"? seems none of the answers mention about it

upvoted 1 times

🗳️ 👤 **rakeshvardan** 2 years, 4 months ago

You most likely want to use ACLs if you need to customize access to individual objects within a bucket, since IAM permissions apply to all objects within a bucket. However, you should still use IAM for any access that is common to all objects in a bucket, because this reduces the amount of micro-managing you have to do.

A - as per the above documentation ACLs are needed for specific objects inside bucket.

B - credentials for the service account shouldn't be stored in the app

D - there is no requirement to encrypt the storage data

Hence C seems to be the correct one

upvoted 4 times

🗳️ 👤 **ArizonaClassics** 2 years, 4 months ago

I agree with C

upvoted 3 times

  **MohitA** 2 years, 4 months ago

Yup C is the right one

upvoted 2 times

An organization's typical network and security review consists of analyzing application transit routes, request handling, and firewall rules. They want to enable their developer teams to deploy new applications without the overhead of this full review. How should you advise this organization?

- A. Use Forseti with Firewall filters to catch any unwanted configurations in production.
- B. Mandate use of infrastructure as code and provide static analysis in the CI/CD pipelines to enforce policies.
- C. Route all VPC traffic through customer-managed routers to detect malicious patterns in production.
- D. All production applications will run on-premises. Allow developers free rein in GCP as their dev and QA platforms.

Correct Answer: B

Community vote distribution



- bluetaurianbull**

Highly Voted

1 year, 8 months ago

@TNT87 and others, if you say (B) or even (C) or (A) can you provide proof and URLs to support your claims. Simply saying if you have done Cloud Architect you will know Everything under the sun is not the proper response, this is a discussion and a community here trying to learn. Not everyone will be in same standard or level.  
Be helpful for others please....  
upvoted 11 times
- OSNG**

Highly Voted

1 year, 3 months ago

Its B.  
Reasons:  
1. They are asking for advise for Developers. (IaC is the suitable as they don't have to worry about managing infrastructure manually).  
Moreover "An organization's typical network and security review consists of analyzing application transit routes, request handling, and firewall rules." statement is defining the process, they are not asking about the option to review the rules. Using Forseti is not reducing the overhead for Developers.  
upvoted 9 times
- TonytheTiger**

Most Recent

1 month ago

B: the best answer.  
<https://cloud.google.com/recommender/docs/tutorial-iac>  
upvoted 1 times
- GCP72**

3 months, 4 weeks ago

Selected Answer: B

The correct answer is B  
upvoted 1 times
- Jeanphi72**

4 months, 3 weeks ago

Selected Answer: A

The problem I see with B is that there is no reason why reviews should disappear: IaC is code and code needs to be reviewed before being deployed. Depending on the companies, devops writing terraform / CDK are named developers as well. Forseti seems to be able to automate this:  
<https://github.com/forseti-security/forseti-security/tree/master/samples/scanner/scanners>  
upvoted 1 times

**szl0144**

7 months ago

I think B is the answer, can anybody explain why A is correct?  
upvoted 1 times

**badrik**

6 months, 3 weeks ago

A is detective in nature while B is preventive. So, It's B !  
upvoted 1 times

**minostrozaml2**

11 months, 1 week ago

Took the task today, only 5 question from this dump, the rest are new questions.  
upvoted 2 times

**ThisisJohn**

1 year ago

Selected Answer: B

My vote goes to B by discard.  
  
A) only mentions firewall rules, but nothing about network routes, and nothing on Forseti website either <https://forsetisecurity.org/about/>  
C) Talks about malicious patterns, not about network routes, requests handling and patterns, like the question says  
D) Running on-prem doesn't guarantee a higher level of control



Thus, the only answer that makes sense for me is B  
upvoted 1 times

🗨️ 👤 **TNT87** 1 year, 10 months ago  
if you done Cloud Rchitect,you will understand why the answer is B  
upvoted 4 times

🗨️ 👤 **bluetaurianbull** 1 year, 7 months ago  
its like saying if you have gone to space you experiance weighlessness .. be professional man... give proof for your claims, dont just expect world to be in same level as you. thats about COMMUNITY LEARNING ...  
upvoted 7 times

🗨️ 👤 **[Removed]** 2 years, 1 month ago  
Ans - C  
upvoted 2 times

🗨️ 👤 **[Removed]** 2 years, 1 month ago  
Sry(Typo) .. It's B  
upvoted 2 times

🗨️ 👤 **saurabh1805** 2 years, 2 months ago  
I will also go with option A  
upvoted 1 times

🗨️ 👤 **CHECK666** 2 years, 2 months ago  
B is the answer  
upvoted 1 times

🗨️ 👤 **ownez** 2 years, 3 months ago  
Answer is B and not A because in A, the answer provided tells us the environment is in production where the question is about to enable their developer teams to deploy new applications without the overhead of the full review. Implementation of IAC is suitable for this.  
  
Answer is B.  
upvoted 3 times

🗨️ 👤 **MohitA** 2 years, 4 months ago  
Yes B serves the purpose.  
upvoted 2 times

🗨️ 👤 **aiwaai** 2 years, 4 months ago  
Answer is A  
upvoted 1 times

🗨️ 👤 **bigdo** 2 years, 4 months ago  
b is correct <https://cloud.google.com/blog/products/containers-kubernetes/guard-against-security-vulnerabilities-with-container-registry-vulnerability-scanning>  
upvoted 2 times

🗨️ 👤 **ArizonaClassics** 2 years, 4 months ago  
Foreseti makes sense == A  
upvoted 2 times

🗨️ 👤 **ArizonaClassics** 2 years, 4 months ago  
The Core FORSETI modules includes the Following; (1.) INVENTORY (2.) SCANNING (3.) ENFORCEMENT (4.) EXPLAIN  
Hence the question already gave a clue on what they are looking for (INVENTORY/VISIBILITY/ANALYSIS) of all transit routes, firewall rules etc  
upvoted 1 times

An employer wants to track how bonus compensations have changed over time to identify employee outliers and correct earning disparities. This task must be performed without exposing the sensitive compensation data for any individual and must be reversible to identify the outlier. Which Cloud Data Loss Prevention API technique should you use to accomplish this?

- A. Generalization
- B. Redaction
- C. CryptoHashConfig
- D. CryptoReplaceFfxFpeConfig

**Correct Answer:** B

Community vote distribution

D (83%)

C (17%)

🗳️ **xhova** **Highly Voted** 👍 2 years, 8 months ago

Answer is D

<https://cloud.google.com/dlp/docs/pseudonymization>

upvoted 12 times

🗳️ **smart123** 2 years, 6 months ago

Option D is correct because it is reversible whereas option B is not.

upvoted 2 times

🗳️ **SilentSec** 2 years, 5 months ago

Also the same usecase in the url that you post. D is right.

upvoted 1 times

🗳️ **gcp\_learner** **Highly Voted** 👍 2 years, 5 months ago

The answer is A.

By bucketing or generalizing, we achieve a reversible pseudonymised data that can still yield the required analysis.

<https://cloud.google.com/dlp/docs/concepts-bucketing>

upvoted 5 times

🗳️ **Sheeda** 2 years, 4 months ago

Completely wrong

The answer is D for sure. The example was even in google docs but replaced for some reasons.

[http://price2meet.com/gcp/docs/dlp\\_docs\\_pseudonymization.pdf](http://price2meet.com/gcp/docs/dlp_docs_pseudonymization.pdf)

upvoted 3 times

🗳️ **Littleivy** **Most Recent** 🕒 1 month, 1 week ago

The answer is D

[https://cloud.google.com/dlp/docs/transformations-reference#transformation\\_methods](https://cloud.google.com/dlp/docs/transformations-reference#transformation_methods)

upvoted 1 times

🗳️ **Premumar** 1 month, 3 weeks ago

**Selected Answer: D**

D is the only option that is reversible.

upvoted 1 times

🗳️ **GCP72** 3 months, 4 weeks ago

**Selected Answer: D**

D. CryptoReplaceFfxFpeConfig is the correct answer for sure, it allows you to recover the original data.

upvoted 1 times

🗳️ **Jeanphi72** 4 months, 3 weeks ago

**Selected Answer: C**

I think D is not a good fit: Replaces an input value with a token, or surrogate value, of the same length using format-preserving encryption (FPE) with the FFX mode of operation. This allows the output to be used in systems that have format validation on length.

In other words, if one earns 50000 and the next 150000 then there is one character more. So this gives some information on the salary. On top of this, this technique is designed to support older systems.

upvoted 1 times



🗨️ 👤 **Medofree** 8 months, 2 weeks ago

**Selected Answer: D**

D

Among the 4 transformations, CryptoReplaceFfxFpeConfig. is the only able to reverse : <https://cloud.google.com/dlp/docs/transformations-reference>

upvoted 3 times

🗨️ 👤 **Ksrp** 10 months ago

its D , <https://cloud.google.com/dlp/docs/deidentify-sensitive-data> ( When you de-identify data using the CryptoReplaceFfxFpeConfig or CryptoDeterministicConfig infoType transformations, you can re-identify that data, as long as you have the CryptoKey used to originally de-identify the data.)

upvoted 1 times

🗨️ 👤 **minostroza12** 11 months, 1 week ago

Took the task today, only 5 question from this dump, the rest are new questions.

upvoted 1 times

🗨️ 👤 **Bwitch** 1 year ago

I think B is correct.

<https://cloud.google.com/dlp/docs/transformations-reference#redaction>

upvoted 1 times

🗨️ 👤 **OSNG** 1 year, 3 months ago

Answer is D.

<https://cloud.google.com/dlp/docs/transformations-reference>

upvoted 1 times

🗨️ 👤 **rand0mb0t** 1 year, 4 months ago

The correct answer is A . As it will let you de-identify the data and also generalise it with bucketing so that outliers can be found during analysis.

I would go with A

upvoted 1 times

🗨️ 👤 **Zuy01** 1 year, 4 months ago

The answer are closest is between C and D, if you read on this link : [https://cloud.google.com/dlp/docs/transformations-reference#transformation\\_methods](https://cloud.google.com/dlp/docs/transformations-reference#transformation_methods)

"CryptoHashConfig" it can't be reverse, but "CryptoReplaceFfxFpeConfig" it can be reversed. so the answer is D.

upvoted 2 times

🗨️ 👤 **ASG** 1 year, 10 months ago

Ans - D - <https://cloud.google.com/dlp/docs/transformations-reference> (pls do read the details in this link for clarification - it will only take you a min or two)

upvoted 1 times

🗨️ 👤 **ESP\_SAP** 2 years ago

Correct Answer is (D):

De-identifying sensitive data

Cloud Data Loss Prevention (DLP) can de-identify sensitive data in text content, including text stored in container structures such as tables. De-identification is the process of removing identifying information from data. The API detects sensitive data such as personally identifiable information (PII), and then uses a de-identification transformation to mask, delete, or otherwise obscure the data. For example, de-identification techniques can include any of the following:

Masking sensitive data by partially or fully replacing characters with a symbol, such as an asterisk (\*) or hash (#).

Replacing each instance of sensitive data with a token, or surrogate, string.

Encrypting and replacing sensitive data using a randomly generated or pre-determined key.

When you de-identify data using the CryptoReplaceFfxFpeConfig or CryptoDeterministicConfig infoType transformations, you can re-identify that data, as long as you have the CryptoKey used to originally de-identify the data.

<https://cloud.google.com/dlp/docs/deidentify-sensitive-data>

upvoted 4 times

🗨️ 👤 **jonclem** 2 years, 1 month ago

My conclusion is answer B since watching the DLP vide from Linuxacademy. It mentions the process and even describes identifying the outliers.

upvoted 1 times

🗨️ 👤 **[Removed]** 2 years, 1 month ago

Ans - D

<https://cloud.google.com/dlp/docs/deidentify-sensitive-data>

upvoted 1 times

An organization adopts Google Cloud Platform (GCP) for application hosting services and needs guidance on setting up password requirements for their Cloud

Identity account. The organization has a password policy requirement that corporate employee passwords must have a minimum number of characters.

Which Cloud Identity password guidelines can the organization use to inform their new requirements?

- A. Set the minimum length for passwords to be 8 characters.
- B. Set the minimum length for passwords to be 10 characters.
- C. Set the minimum length for passwords to be 12 characters.
- D. Set the minimum length for passwords to be 6 characters.

Correct Answer: A

Community vote distribution



**KILLMAD** Highly Voted 2 years, 9 months ago  
Ans is A  
upvoted 10 times

**rafaelc** 2 years, 9 months ago  
Default password length is 8 characters.  
<https://support.google.com/cloudidentity/answer/33319?hl=en>  
upvoted 9 times

**bolu** Highly Voted 1 year, 10 months ago  
The situation changes year on year on GCP.Right now the right answer is C based on minimum requirement of 12 char in GCP as on Jan 2021.  
<https://support.google.com/accounts/answer/32040?hl=en#zippy=%2Cmake-your-password-longer-more-memorable>  
upvoted 10 times

**desertlotus1211** 1 year, 9 months ago  
It asked for Cloud Indentity password requirements... Minimum is 8 Maximum is 100  
upvoted 2 times

**blue123456** Most Recent 1 month ago  
Ans A  
<https://support.google.com/cloudidentity/answer/2537800?hl=en#zippy=%2Creset-a-users-password>  
upvoted 1 times

**xchmielu** 1 month ago  
Selected Answer: C  
<https://support.google.com/accounts/answer/32040?hl=en#zippy=%2Cmake-your-password-longer-more-memorable>  
upvoted 1 times

**GCP72** 3 months, 4 weeks ago  
Selected Answer: A  
The answer is A  
upvoted 1 times

**otokichi3** 6 months, 1 week ago  
The answer is A.  
minimum character length is 8.  
<https://support.google.com/cloudidentity/answer/139399?hl=en>  
upvoted 1 times

**Ksrp** 10 months ago  
Ans - A, see <https://support.google.com/accounts/answer/9094506#zippy=%2Cmake-your-password-longer-more-memorable>  
upvoted 1 times

**minostrozaml2** 11 months, 1 week ago  
Took the tesk today, only 5 question from this dump, the rest are new questions.  
upvoted 1 times

**umashankar\_a** 1 year, 5 months ago

Answer A

For Cloud Identity password requirements still is - Minimum 8 Maximum is 100

[https://support.google.com/cloudidentity/answer/139399?](https://support.google.com/cloudidentity/answer/139399?hl=en#:~:text=It%20can%20be%20between%208,decide%20to%20change%20their%20password.)

[hl=en#:~:text=It%20can%20be%20between%208,decide%20to%20change%20their%20password.](https://support.google.com/cloudidentity/answer/139399?hl=en#:~:text=It%20can%20be%20between%208,decide%20to%20change%20their%20password.)

upvoted 2 times

  **soukumar369** 1 year, 11 months ago

Answer is C : <https://support.google.com/accounts/answer/32040?hl=en#zippy=%2Cmake-your-password-longer-more-memorable>

Long passwords are stronger, so make your password at least 12 characters long. These tips can help you create longer passwords that are easier to remember. Try to use:

upvoted 2 times

  **[Removed]** 2 years, 1 month ago

Ans - A

upvoted 2 times

  **CHECK666** 2 years, 2 months ago

A is the answer.

upvoted 1 times

  **MohitA** 2 years, 4 months ago

Is it really an exam Question? it depends on whats your company policy is for password length

upvoted 4 times

  **Zol** 2 years, 9 months ago

Rafaelc

You're correct A is the answer

upvoted 1 times

You need to follow Google-recommended practices to leverage envelope encryption and encrypt data at the application layer. What should you do?

- A. Generate a data encryption key (DEK) locally to encrypt the data, and generate a new key encryption key (KEK) in Cloud KMS to encrypt the DEK. Store both the encrypted data and the encrypted DEK.
- B. Generate a data encryption key (DEK) locally to encrypt the data, and generate a new key encryption key (KEK) in Cloud KMS to encrypt the DEK. Store both the encrypted data and the KEK.
- C. Generate a new data encryption key (DEK) in Cloud KMS to encrypt the data, and generate a key encryption key (KEK) locally to encrypt the key. Store both the encrypted data and the encrypted DEK.
- D. Generate a new data encryption key (DEK) in Cloud KMS to encrypt the data, and generate a key encryption key (KEK) locally to encrypt the key. Store both the encrypted data and the KEK.

**Correct Answer: A**

Reference:

<https://cloud.google.com/kms/docs/envelope-encryption>

*Community vote distribution*

A (100%)

  **Sheeda** Highly Voted 2 years, 4 months ago

Yes, A is correct

The process of encrypting data is to generate a DEK locally, encrypt data with the DEK, use a KEK to wrap the DEK, and then store the encrypted data and the wrapped DEK. The KEK never leaves Cloud KMS.

upvoted 18 times

  **MohitA** 2 years, 4 months ago

Agree on A, spot on "KEK never leaves Cloud KMS"

upvoted 3 times

  **GCP72** Most Recent 3 months, 4 weeks ago

**Selected Answer: A**

The answer is A

upvoted 1 times

  **minostrozaml2** 11 months, 1 week ago

Took the task today, only 5 questions from this dump, the rest are new questions.

upvoted 1 times

  **Bill831231** 1 year ago

A sounds like the correct answer:

[https://cloud.google.com/kms/docs/envelope-encryption#how\\_to\\_encrypt\\_data\\_using\\_envelope\\_encryption](https://cloud.google.com/kms/docs/envelope-encryption#how_to_encrypt_data_using_envelope_encryption)

upvoted 1 times

  **umashankar\_a** 1 year, 5 months ago

Answer A

Envelope Encryption: <https://cloud.google.com/kms/docs/envelope-encryption>

Here are best practices for managing DEKs:

-Generate DEKs locally.

-When stored, always ensure DEKs are encrypted at rest.

- For easy access, store the DEK near the data that it encrypts.

The DEK is encrypted (also known as wrapped) by a key encryption key (KEK). The process of encrypting a key with another key is known as envelope encryption.

Here are best practices for managing KEKs:

-Store KEKs centrally. (KMS )

-Set the granularity of the DEKs they encrypt based on their use case. For example, consider a workload that requires multiple DEKs to encrypt the workload's data chunks. You could use a single KEK to wrap all DEKs that are responsible for that workload's encryption.

-Rotate keys regularly, and also after a suspected incident.

upvoted 2 times

  **desertlotus1211** 1 year, 8 months ago

I'm not sure what the answer is, but the answers to this question have changed.... be prepared

upvoted 1 times

  **dtmtor** 1 year, 9 months ago

Answer is A  
upvoted 1 times

  **DebasishLowes** 1 year, 9 months ago

Ans : A  
upvoted 1 times

  **CloudTrip** 1 year, 10 months ago

Correction I change it to A after reading the question once again.  
upvoted 1 times

  **CloudTrip** 1 year, 10 months ago

Answer is B as after DEK encryption it's KEK (not encrypted DEK) which never leaves KMS  
upvoted 1 times

  **Bharathy** 2 years ago

A - Envelope Encryption ( DEK - to encrypt the data, KEK - encrypt the DEK , KEK resides in KMS and only the encrypted data and wrapped DEK will be stored back )  
upvoted 2 times

  **[Removed]** 2 years, 1 month ago

Ans - A  
[https://cloud.google.com/kms/docs/envelope-encryption#how\\_to\\_encrypt\\_data\\_using\\_envelope\\_encryption](https://cloud.google.com/kms/docs/envelope-encryption#how_to_encrypt_data_using_envelope_encryption)  
upvoted 1 times

  **CHECK666** 2 years, 2 months ago

The answer is A  
upvoted 1 times

  **aiwaai** 2 years, 4 months ago

The Answer is A  
upvoted 2 times

  **ArizonaClassics** 2 years, 4 months ago

C is the correct option. See BEST Practices for managing DEKS  
<https://cloud.google.com/kms/docs/envelope-encryption>  
upvoted 1 times

  **kumar999** 2 years, 4 months ago

It is A. See How to encrypt data using envelope encryption  
upvoted 2 times

  **ArizonaClassics** 2 years, 4 months ago

Thank you Kumar. it is A. I made a mistake.  
upvoted 1 times

  **ArizonaClassics** 2 years, 4 months ago

A is the correct one. Made a mistake earlier  
upvoted 2 times

How should a customer reliably deliver Stackdriver logs from GCP to their on-premises SIEM system?

- A. Send all logs to the SIEM system via an existing protocol such as syslog.
- B. Configure every project to export all their logs to a common BigQuery DataSet, which will be queried by the SIEM system.
- C. Configure Organizational Log Sinks to export logs to a Cloud Pub/Sub Topic, which will be sent to the SIEM via Dataflow.
- D. Build a connector for the SIEM to query for all logs in real time from the GCP RESTful JSON APIs.

Correct Answer: C

Community vote distribution



**ESP\_SAP** Highly Voted 2 years ago

Correct answer is (C):  
Scenarios for exporting Cloud Logging data: Splunk  
This scenario shows how to export selected logs from Cloud Logging to Pub/Sub for ingestion into Splunk. Splunk is a security information and event management (SIEM) solution that supports several ways of ingesting data, such as receiving streaming data out of Google Cloud through Splunk HTTP Event Collector (HEC) or by fetching data from Google Cloud APIs through Splunk Add-on for Google Cloud.  
  
Using the Pub/Sub to Splunk Dataflow template, you can natively forward logs and events from a Pub/Sub topic into Splunk HEC. If Splunk HEC is not available in your Splunk deployment, you can use the Add-on to collect the logs and events from the Pub/Sub topic.  
<https://cloud.google.com/solutions/exporting-stackdriver-logging-for-splunk>  
upvoted 15 times

**AzureDP900** 1 month, 2 weeks ago

I will go with C  
upvoted 1 times

**GCP72** Most Recent 3 months, 4 weeks ago

Selected Answer: C  
I will go with C  
upvoted 3 times

**minostrozaml2** 11 months, 1 week ago

Took the tesk today, only 5 question from this dump, the rest are new questions.  
upvoted 1 times

**DebasishLowes** 1 year, 9 months ago

Ans : C  
upvoted 2 times

**BlahBaller** 1 year, 11 months ago

As I was the Logging Service Manager when we set this up with GCP. I can verify that C is how we have it setup, based on the Google's recommendations.  
upvoted 2 times

**Moss2011** 2 years, 1 month ago

I think the correct one its D because C mention "Dataflow" and it cannot be connected to any sink out of GCP.  
upvoted 1 times

**[Removed]** 2 years, 1 month ago

Ans - C  
<https://cloud.google.com/solutions/exporting-stackdriver-logging-for-splunk>  
upvoted 2 times

**deevisrk** 2 years, 1 month ago

C looks correct..  
<https://cloud.google.com/solutions/exporting-stackdriver-logging-for-splunk>  
Splunk is on premises SIEM solution in above example.  
upvoted 2 times

**saaurabh1805** 2 years, 2 months ago

I will go with Option B.  
  
Read this email for more reason. C is not workable solution so that is first one not to consider.  
upvoted 1 times



  **CHECK666** 2 years, 2 months ago

C is the answer.  
upvoted 1 times

  **ArizonaClassics** 2 years, 4 months ago

I will go with C  
upvoted 2 times

  **xhova** 2 years, 8 months ago

C is correct  
upvoted 4 times

In order to meet PCI DSS requirements, a customer wants to ensure that all outbound traffic is authorized.

Which two cloud offerings meet this requirement without additional compensating controls? (Choose two.)

- A. App Engine
- B. Cloud Functions
- C. Compute Engine
- D. Google Kubernetes Engine
- E. Cloud Storage

**Correct Answer:** AC

Reference:

<https://cloud.google.com/solutions/pci-dss-compliance-in-gcp>

*Community vote distribution*

CD (100%)

  **KILLMAD** Highly Voted 2 years, 9 months ago

Answer is CD

because the doc mentions the following: "App Engine ingress firewall rules are available, but egress rules are not currently available:" and "Compute Engine and GKE are the preferred alternatives."

upvoted 14 times

  **rafaelc** 2 years, 9 months ago

It is CD.

App Engine ingress firewall rules are available, but egress rules are not currently available. Per requirements 1.2.1 and 1.3.4, you must ensure that all outbound traffic is authorized. SAQ A-EP and SAQ D-type merchants must provide compensating controls or use a different Google Cloud product. Compute Engine and GKE are the preferred alternatives.

<https://cloud.google.com/solutions/pci-dss-compliance-in-gcp>

upvoted 6 times

  **Brosh** Most Recent 3 days, 18 hours ago

hey. can anyone explain why isn't A correct? the decumantion mentions app engine as an option but not compute engine

<https://cloud.google.com/architecture/pci-dss-compliance-in-gcp>

upvoted 1 times

  **Littleivy** 1 month, 1 week ago

Selected Answer: CD

Answer is CD

For App Engine, the App Engine firewall only applies to incoming traffic routed to your app or service.

<https://cloud.google.com/appengine/docs/flexible/understanding-firewalls>

upvoted 1 times

  **AzureDP900** 1 month, 2 weeks ago

CD is right

upvoted 1 times

  **GCP72** 3 months, 4 weeks ago

Selected Answer: CD

The correct answer is CD

upvoted 1 times

  **jordi\_194** 10 months, 2 weeks ago

Selected Answer: CD

Ans: CD

upvoted 2 times

  **minostrozaml2** 11 months, 1 week ago

Took the tesk today, only 5 question from this dump, the rest are new questions.

upvoted 2 times



🗑️ 👤 **DebasishLowes** 1 year, 9 months ago

Ans : CD

upvoted 1 times

🗑️ 👤 **[Removed]** 2 years, 1 month ago

Ans - CD

[https://cloud.google.com/solutions/pci-dss-compliance-in-gcp#app\\_engine](https://cloud.google.com/solutions/pci-dss-compliance-in-gcp#app_engine)

upvoted 1 times

🗑️ 👤 **saurabh1805** 2 years, 2 months ago

C & D is correct answer here.

<https://cloud.google.com/solutions/pci-dss-compliance-in-gcp>

upvoted 1 times

🗑️ 👤 **CHECK666** 2 years, 2 months ago

CD is the answer.

upvoted 1 times

🗑️ 👤 **MohitA** 2 years, 4 months ago

CD serves the purpose for sure

upvoted 1 times

🗑️ 👤 **ArizonaClassics** 2 years, 4 months ago

CD best option

upvoted 2 times

🗑️ 👤 **xhova** 2 years, 8 months ago

CD is right

upvoted 2 times

A website design company recently migrated all customer sites to App Engine. Some sites are still in progress and should only be visible to customers and company employees from any location.

Which solution will restrict access to the in-progress sites?

- A. Upload an .htaccess file containing the customer and employee user accounts to App Engine.
- B. Create an App Engine firewall rule that allows access from the customer and employee networks and denies all other traffic.
- C. Enable Cloud Identity-Aware Proxy (IAP), and allow access to a Google Group that contains the customer and employee user accounts.
- D. Use Cloud VPN to create a VPN connection between the relevant on-premises networks and the company's GCP Virtual Private Cloud (VPC) network.

**Correct Answer: C**

*Community vote distribution*

C (100%)

🗲️ 👤 **[Removed]** Highly Voted 👍 2 years, 1 month ago

Ans - C

[https://cloud.google.com/iap/docs/concepts-overview#when\\_to\\_use\\_iap](https://cloud.google.com/iap/docs/concepts-overview#when_to_use_iap)

upvoted 11 times

🗲️ 👤 **GCP72** Most Recent ⌚ 3 months, 4 weeks ago

**Selected Answer: C**

The correct answer is C

upvoted 1 times

🗲️ 👤 **simbu1299** 9 months, 1 week ago

**Selected Answer: C**

Answer is C

upvoted 1 times

🗲️ 👤 **mlx** 2 years, 1 month ago

B - I think it is about to restrict access to 2 company networks, we can control access using IPs ranges, So Firewall rules should be sufficient. No need an extra product like IAP.. and also need users in Cloud Identity or other Idp federated..

upvoted 1 times

🗲️ 👤 **FatCharlie** 2 years ago

The sites should be accessible from any location, not just from the 2 company networks.

upvoted 2 times

🗲️ 👤 **MohitA** 2 years, 4 months ago

C serves the purpose

upvoted 3 times

🗲️ 👤 **bigdo** 2 years, 4 months ago

c is correct

upvoted 2 times

🗲️ 👤 **ArizonaClassics** 2 years, 4 months ago

C is very correct

upvoted 2 times

🗲️ 👤 **SilentSec** 2 years, 5 months ago

C is correct.

upvoted 2 times

When working with agents in the support center via online chat, your organization's customers often share pictures of their documents with personally identifiable information (PII). Your leadership team is concerned that this PII is being stored as part of the regular chat logs, which are reviewed by internal or external analysts for customer service trends.

You want to resolve this concern while still maintaining data utility. What should you do?

- A. Use Cloud Key Management Service to encrypt PII shared by customers before storing it for analysis.
- B. Use Object Lifecycle Management to make sure that all chat records containing PII are discarded and not saved for analysis.
- C. Use the image inspection and redaction actions of the DLP API to redact PII from the images before storing them for analysis.
- D. Use the generalization and bucketing actions of the DLP API solution to redact PII from the texts before storing them for analysis.

**Correct Answer:** C

Reference:

<https://cloud.google.com/dlp/docs/deidentify-sensitive-data>

*Community vote distribution*

C (100%)

🗳️ 👤 **shayke** 4 days, 14 hours ago

**Selected Answer: C**

C the q refers to imaging  
upvoted 1 times

🗳️ 👤 **kamal17** 2 weeks ago

Answer is C  
upvoted 1 times

🗳️ 👤 **jitu028** 2 months, 3 weeks ago

Answer is C  
upvoted 4 times

A company's application is deployed with a user-managed Service Account key. You want to use Google-recommended practices to rotate the key. What should you do?

- A. Open Cloud Shell and run `gcloud iam service-accounts enable-auto-rotate --iam-account=IAM_ACCOUNT`.
- B. Open Cloud Shell and run `gcloud iam service-accounts keys rotate --iam-account=IAM_ACCOUNT --key=NEW_KEY`.
- C. Create a new key, and use the new key in the application. Delete the old key from the Service Account.
- D. Create a new key, and use the new key in the application. Store the old key on the system as a backup key.

**Correct Answer:** C

Reference:

<https://cloud.google.com/iam/docs/understanding-service-accounts>

*Community vote distribution*

C (100%)

🗲️ 👤 **mdc** Highly Voted 1 year, 6 months ago

C is correct. As explained, You can rotate a key by creating a new key, updating applications to use the new key, and deleting the old key. Use the `serviceAccount.keys.create()` method and `serviceAccount.keys.delete()` method together to automate the rotation.

[https://cloud.google.com/iam/docs/creating-managing-service-account-keys#deleting\\_service\\_account\\_keys](https://cloud.google.com/iam/docs/creating-managing-service-account-keys#deleting_service_account_keys)

upvoted 5 times

🗲️ 👤 **GCP72** Most Recent 3 months, 3 weeks ago

**Selected Answer: C**

The correct answer is C

upvoted 2 times

🗲️ 👤 **absipat** 6 months, 2 weeks ago

c of course

upvoted 1 times

🗲️ 👤 **DebasishLowes** 1 year, 9 months ago

Ans : C

upvoted 2 times

🗲️ 👤 **[Removed]** 2 years, 1 month ago

Ans - C

[https://cloud.google.com/iam/docs/understanding-service-accounts#managing\\_service\\_account\\_keys](https://cloud.google.com/iam/docs/understanding-service-accounts#managing_service_account_keys)

upvoted 4 times

🗲️ 👤 **ArizonaClassics** 2 years, 4 months ago

C is the right choice for me

upvoted 4 times

🗲️ 👤 **aiwaai** 2 years, 4 months ago

Correct Answer: C

upvoted 2 times

Your team needs to configure their Google Cloud Platform (GCP) environment so they can centralize the control over networking resources like firewall rules, subnets, and routes. They also have an on-premises environment where resources need access back to the GCP resources through a private VPN connection.

The networking resources will need to be controlled by the network security team.

Which type of networking design should your team use to meet these requirements?

- A. Shared VPC Network with a host project and service projects
- B. Grant Compute Admin role to the networking team for each engineering project
- C. VPC peering between all engineering projects using a hub and spoke model
- D. Cloud VPN Gateway between all engineering projects using a hub and spoke model

**Correct Answer:** A

Reference:

[https://cloud.google.com/docs/enterprise/best-practices-for-enterprise-organizations#centralize\\_network\\_control](https://cloud.google.com/docs/enterprise/best-practices-for-enterprise-organizations#centralize_network_control)

*Community vote distribution*

A (100%)

  **ArizonaClassics** Highly Voted 2 years, 4 months ago

I agree with A  
Centralize network control:

Use Shared VPC to connect to a common VPC network. Resources in those projects can communicate with each other securely and efficiently across project boundaries using internal IPs. You can manage shared network resources, such as subnets, routes, and firewalls, from a central host project, enabling you to apply and enforce consistent network policies across the projects.

upvoted 18 times

  **ArizonaClassics** 2 years, 4 months ago

WATCH: <https://www.youtube.com/watch?v=WotV3D01tJA>

READ:

[https://cloud.google.com/docs/enterprise/best-practices-for-enterprise-organizations#centralize\\_network\\_control](https://cloud.google.com/docs/enterprise/best-practices-for-enterprise-organizations#centralize_network_control)

upvoted 5 times

  **kamal17** Most Recent 2 weeks ago

Answer is D , bocz On-prime user needs to access the GCP resources with help of Cloud VPN

upvoted 1 times

  **GCP72** 3 months, 3 weeks ago

**Selected Answer: A**

The correct answer is A

upvoted 1 times

  **minostrozaml2** 11 months, 1 week ago

Took the task today, only 5 question from this dump, the rest are new questions.

upvoted 1 times

  **ZODOGAM** 1 year, 1 month ago

Sheeda En mi caso te confirmo que desde la share VPC se establecen las VPNs y allí ingresa el tráfico desde los sitios locales. Definitivamente, la respuesta es la A

upvoted 1 times

  **DebasishLowes** 1 year, 9 months ago

Ans : A. It will be shared VPC as it is asking for centralized network control.

upvoted 1 times

  **jonclem** 2 years, 1 month ago

Option D is incorrect and a violation of Google's Service Specific terms as per : <https://cloud.google.com/network-connectivity/docs/vpn/concepts/overview>

I'd go with option A myself.

upvoted 1 times

  **[Removed]** 2 years, 1 month ago

Ans - A

upvoted 1 times

🗨️ 👤 **saurabh1805** 2 years, 2 months ago

A, this is exact reason to use shared VPC

upvoted 1 times

🗨️ 👤 **CHECK666** 2 years, 2 months ago

A is the answer.

upvoted 1 times

🗨️ 👤 **Akku1614** 2 years, 3 months ago

A is correct as Shared VPC provides us with Centralized control however VPC Peering is a decentralized option.

upvoted 1 times

🗨️ 👤 **aiwaai** 2 years, 4 months ago

Correct Answer: A

upvoted 1 times

🗨️ 👤 **Sheeda** 2 years, 4 months ago

Connect your enterprise network

Many enterprises need to connect existing on-premises infrastructure with their Google Cloud resources. Evaluate your bandwidth, latency, and SLA requirements to choose the best connection option:

If you need low-latency, highly available, enterprise-grade connections that enable you to reliably transfer data between your on-premises and VPC networks without traversing the internet connections to Google Cloud, use Cloud Interconnect:

Dedicated Interconnect provides a direct physical connection between your on-premises network and Google's network.

Partner Interconnect provides connectivity between your on-premises and Google Cloud VPC networks through a supported service provider.

If you don't require the low latency and high availability of Cloud Interconnect, or you are just starting on your cloud journey, use Cloud VPN to set up encrypted IPsec VPN tunnels between your on-premises network and VPC. Compared to a direct, private connection, an IPsec VPN tunnel has lower overhead and costs.

upvoted 1 times

🗨️ 👤 **ArizonaClassics** 2 years, 4 months ago

Sheeda you need to read and understand the the question.

upvoted 1 times

🗨️ 👤 **ArizonaClassics** 2 years, 4 months ago

They are asking how you can centralize the control over networking resources like firewall rules, subnets, and routes. watch this:

<https://www.youtube.com/watch?v=WotV3D01tJA>

you will see that you can also manage vpn connections as well

upvoted 1 times

🗨️ 👤 **ESP\_SAP** 2 years ago

you Should go back to the GCP Cloud Architect concepts or GCP Networking!

upvoted 2 times

🗨️ 👤 **Sheeda** 2 years, 4 months ago

I believe the answer is D. How can shared VPC give access to your on premise environment ? A seems wrong to me.

upvoted 2 times

🗨️ 👤 **AkbarM** 3 months ago

I also believe the same. i worked on interconnects and gateways to connect on prem resources.. only hub and spoke helps to connect onpremise network. ofcourse, we can centralize network controls using shared vpc. but the need here is some engineerng resources in on prem needs to access gcp resources. so this needs gateway to access gcp resources.

upvoted 1 times



An organization is migrating from their current on-premises productivity software systems to G Suite. Some network security controls were in place that were mandated by a regulatory body in their region for their previous on-premises system. The organization's risk team wants to ensure that network security controls are maintained and effective in G Suite. A security architect supporting this migration has been asked to ensure that network security controls are in place as part of the new shared responsibility model between the organization and Google Cloud.

What solution would help meet the requirements?

- A. Ensure that firewall rules are in place to meet the required controls.
- B. Set up Cloud Armor to ensure that network security controls can be managed for G Suite.
- C. Network security is a built-in solution and Google's Cloud responsibility for SaaS products like G Suite.
- D. Set up an array of Virtual Private Cloud (VPC) networks to control network security as mandated by the relevant regulation.

**Correct Answer:** C

Community vote distribution

C (100%)

🗳️ 👤 **ESP\_SAP** Highly Voted 👍 2 years ago

Correct Answer is (C):  
GSuite is Saas application.

Shared responsibility "Security of the Cloud" - GCP is responsible for protecting the infrastructure that runs all of the services offered in the GCP Cloud. This infrastructure is composed of the hardware, software, networking, and facilities that run GCP Cloud services.

upvoted 9 times

🗳️ 👤 **AzureDP900** 1 month, 2 weeks ago

C is right  
upvoted 2 times

🗳️ 👤 **Topsy** Highly Voted 👍 2 years ago

Answer is C- Review this Youtube Video- <https://www.youtube.com/watch?v=D2zf0SgNdUw>, scroll to 7:55, it would show you the Shared Responsibility model- With Gsuite being a SaaS product, Network Security is handled by Google

upvoted 5 times

🗳️ 👤 **shayke** Most Recent ⌚ 4 days, 14 hours ago

**Selected Answer: C**

c - SAAS network security is the responsible of the cloud provider  
upvoted 1 times

🗳️ 👤 **absipat** 6 months, 2 weeks ago

c of course  
upvoted 1 times

🗳️ 👤 **absipat** 6 months, 2 weeks ago

C as it is SAAs  
upvoted 1 times

🗳️ 👤 **FatCharlie** 2 years ago

Except for C, none of the options are possible in G Suite. There are no firewall, VPC, or Cloud Armor options there as far as I know.  
upvoted 4 times

🗳️ 👤 **[Removed]** 2 years, 1 month ago

Ans - A  
upvoted 2 times

🗳️ 👤 **saurabh1805** 2 years, 1 month ago

Question is asking for Network security group, Hence i will go with Option A  
upvoted 1 times

🗳️ 👤 **skshak** 2 years, 2 months ago

Answer is C. Gsuite is SaaS  
upvoted 2 times

A customer's company has multiple business units. Each business unit operates independently, and each has their own engineering group. Your team wants visibility into all projects created within the company and wants to organize their Google Cloud Platform (GCP) projects based on different business units. Each business unit also requires separate sets of IAM permissions.

Which strategy should you use to meet these needs?

- A. Create an organization node, and assign folders for each business unit.
- B. Establish standalone projects for each business unit, using gmail.com accounts.
- C. Assign GCP resources in a project, with a label identifying which business unit owns the resource.
- D. Assign GCP resources in a VPC for each business unit to separate network access.

**Correct Answer: A**

*Community vote distribution*

A (100%)

🗲️ 👤 **ArizonaClassics** Highly Voted 👍 2 years, 4 months ago

I will go with A

Refer to: <https://cloud.google.com/resource-manager/docs/listing-all-resources>

Also: <https://wideops.com/mapping-your-organization-with-the-google-cloud-platform-resource-hierarchy/>  
upvoted 16 times

🗲️ 👤 **shayke** Most Recent ⌚ 4 days, 14 hours ago

**Selected Answer: A**

A is the right ans - resource manager

upvoted 1 times

🗲️ 👤 **DebasishLowes** 1 year, 10 months ago

Ans - A

upvoted 3 times

🗲️ 👤 **[Removed]** 2 years, 1 month ago

Ans - A

upvoted 1 times

🗲️ 👤 **aiwaai** 2 years, 4 months ago

Correct Answer: A

upvoted 1 times



A company has redundant mail servers in different Google Cloud Platform regions and wants to route customers to the nearest mail server based on location.  
How should the company accomplish this?

- A. Configure TCP Proxy Load Balancing as a global load balancing service listening on port 995.
- B. Create a Network Load Balancer to listen on TCP port 995 with a forwarding rule to forward traffic based on location.
- C. Use Cross-Region Load Balancing with an HTTP(S) load balancer to route traffic to the nearest region.
- D. Use Cloud CDN to route the mail traffic to the closest origin mail server based on client IP address.

Correct Answer: D

Community vote distribution



**ESP\_SAP** Highly Voted 2 years ago  
Correct Answer is (A):

TCP Proxy Load Balancing is implemented on GFEs that are distributed globally. If you choose the Premium Tier of Network Service Tiers, a TCP proxy load balancer is global. In Premium Tier, you can deploy backends in multiple regions, and the load balancer automatically directs user traffic to the closest region that has capacity. If you choose the Standard Tier, a TCP proxy load balancer can only direct traffic among backends in a single region.

<https://cloud.google.com/load-balancing/docs/load-balancing-overview#tcp-proxy-load-balancing>  
upvoted 19 times

**Warren2020** Highly Voted 2 years, 5 months ago  
A is the correct answer. D is not correct. CDN works with HTTP(s) traffic and requires caching, which is not a valid feature used for mail server  
upvoted 5 times

**Meyucho** Most Recent 1 month, 1 week ago  
Selected Answer: A  
At a glance, A seems wrong because it's a proxy solution and the question ask to route the traffic. But, in Premium Tier its possible to have a global forwarding rule and the backends in any region so it's the only solution for this.  
the other options are HTTP or Regional so they are wrong  
upvoted 1 times

**huntergame** 1 month, 2 weeks ago  
Selected Answer: A  
A is correct answer  
upvoted 1 times

**AwesomeGCP** 2 months, 2 weeks ago  
Selected Answer: A  
Correct Answer is (A):

TCP Proxy Load Balancing is implemented on GFEs that are distributed globally. If you choose the Premium Tier of Network Service Tiers, a TCP proxy load balancer is global. In Premium Tier, you can deploy backends in multiple regions, and the load balancer automatically directs user traffic to the closest region that has capacity. If you choose the Standard Tier, a TCP proxy load balancer can only direct traffic among backends in a single region.

<https://cloud.google.com/load-balancing/docs/load-balancing-overview#tcp-proxy-load-balancing>  
upvoted 1 times

**AWSE** 1 year, 10 months ago  
Ans should be A. TCP Proxy Load Balancing is intended for TCP traffic on specific well-known ports, such as port 25 for Simple Mail Transfer Protocol (SMTP).  
upvoted 3 times

**jonclem** 2 years, 1 month ago  
I'd agree with option A as it covers location based traffic plus the mail port 993.  
upvoted 1 times

**[Removed]** 2 years, 1 month ago  
Ans - A

upvoted 1 times

  **passtest100** 2 years, 2 months ago

should be C.  
the port 995 has nothing to do with email service . HTTP LB can work on TCP traffic while CDN does NOT support email  
so the better answer is C

upvoted 2 times

  **giovy\_82** 4 months ago

not correct. Port 995 is TLS/SSL encryption for POP3 email server.

upvoted 1 times

  **CHECK666** 2 years, 2 months ago

A is the answer.

upvoted 1 times

  **ArizonaClassics** 2 years, 4 months ago

A: <https://cloud.google.com/load-balancing/docs/tcp>

upvoted 1 times

  **Zol** 2 years, 9 months ago

Answer A is correct

upvoted 1 times

  **mozammil89** 2 years, 9 months ago

Answer should be, A

<https://cloud.google.com/load-balancing/docs/tcp>

upvoted 4 times

Your team sets up a Shared VPC Network where project co-vpc-prod is the host project. Your team has configured the firewall rules, subnets, and VPN gateway on the host project. They need to enable Engineering Group A to attach a Compute Engine instance to only the 10.1.1.0/24 subnet. What should your team grant to Engineering Group A to meet this requirement?


- A. Compute Network User Role at the host project level.
- B. Compute Network User Role at the subnet level.
- C. Compute Shared VPC Admin Role at the host project level.
- D. Compute Shared VPC Admin Role at the service project level.

**Correct Answer:** C

Reference:

<https://cloud.google.com/vpc/docs/shared-vpc>

*Community vote distribution*



🗲️ **mozammil89** **Highly Voted** 👍 2 years, 9 months ago

The correct answer is B.

[https://cloud.google.com/vpc/docs/shared-vpc#svc\\_proj\\_admins](https://cloud.google.com/vpc/docs/shared-vpc#svc_proj_admins)  
upvoted 18 times

🗲️ **Meyucho** **Most Recent** 🕒 1 month, 1 week ago

**Selected Answer: B**

Grant network.user at subnet level:

<https://cloud.google.com/vpc/docs/provisioning-shared-vpc#networkuseratsubnet>  
upvoted 1 times

🗲️ **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: B**

The correct answer is B.

[https://cloud.google.com/vpc/docs/shared-vpc#svc\\_proj\\_admins](https://cloud.google.com/vpc/docs/shared-vpc#svc_proj_admins)  
upvoted 2 times

🗲️ **rajananna** 2 months, 3 weeks ago

**Selected Answer: A**

Lowest level grant is at Project level. <https://cloud.google.com/compute/docs/access/iam#compute.networkUser>  
upvoted 1 times

🗲️ **Premumar** 1 month, 3 weeks ago

Lowest level grant is at Subnet level in this option. Project level is a broad level access.

upvoted 1 times

🗲️ **tangac** 3 months, 2 weeks ago

**Selected Answer: A**

based on that documentation it should clearly be done at the host project level :

<https://cloud.google.com/compute/docs/access/iam#compute.networkUser>  
upvoted 2 times

🗲️ **piyush\_1982** 4 months, 4 weeks ago

**Selected Answer: B**

[https://cloud.google.com/vpc/docs/shared-vpc#svc\\_proj\\_admins](https://cloud.google.com/vpc/docs/shared-vpc#svc_proj_admins)  
upvoted 1 times

🗲️ **Medofree** 8 months, 2 weeks ago

**Selected Answer: B**

The correct answer is b

upvoted 2 times

🗲️ **dropler** 1 year, 5 months ago

The right one is b on my thinking, but i need to enable the other team to do the jobs, falls into D

upvoted 2 times

  **DebasishLowes** 1 year, 9 months ago

Ans : B

upvoted 2 times

  **saurabh1805** 2 years, 1 month ago

B is correct answer

upvoted 1 times

  **[Removed]** 2 years, 1 month ago

Ans - B

upvoted 1 times

  **CHECK666** 2 years, 2 months ago

B is the answer.

upvoted 1 times

  **MohitA** 2 years, 4 months ago

B is right

upvoted 1 times

  **ArizonaClassics** 2 years, 4 months ago

B is correct

upvoted 1 times

  **Zol** 2 years, 9 months ago

Correct B

upvoted 1 times

A company migrated their entire data/center to Google Cloud Platform. It is running thousands of instances across multiple projects managed by different departments. You want to have a historical record of what was running in Google Cloud Platform at any point in time. What should you do?

- A. Use Resource Manager on the organization level.
- B. Use Forseti Security to automate inventory snapshots.
- C. Use Stackdriver to create a dashboard across all projects.
- D. Use Security Command Center to view all assets across the organization.

**Correct Answer:** B

Community vote distribution

B (100%)

  **smart123** Highly Voted 2 years, 6 months ago

'B is the correct answer. Only Forseti security can have both 'past' and 'present' (i.e. historical) records of the resources.  
<https://forsetisecurity.org/about/>  
upvoted 11 times

  **AwesomeGCP** Most Recent 2 months, 2 weeks ago

**Selected Answer: B**

B is the correct answer. Only Forseti security can have both 'past' and 'present' (i.e. historical) records of the resources.  
<https://forsetisecurity.org/about/>  
upvoted 2 times

  **absipat** 6 months, 2 weeks ago

b of course  
upvoted 1 times

  **mynk29** 10 months ago

Outdated questions- you should use asset inventory now.  
upvoted 4 times

  **mitow95526** 1 year, 6 months ago

<https://cloud.google.com/security-command-center>

Discover and view your assets in near-real time across App Engine, BigQuery, Cloud SQL, Cloud Storage, Compute Engine, Cloud Identity and Access Management, Google Kubernetes Engine, and more. Review historical discovery scans to identify new, modified, or deleted assets.

Why not D?  
upvoted 3 times

  **PATILDXB** 1 day, 11 hours ago

Azure security center does provide only realtime view on cloud. Endpoints once deleted or offboarded are no more visible in azure security center, which means historical details are lost  
upvoted 1 times

  **ThisisJohn** 1 year ago

I guess the reason to discard D is that it says "all assets", while according to the documentation, "Security Command Center supports a large subset of Google Cloud assets.", so it supports a large number but not all assets.

Ref: <https://cloud.google.com/security-command-center/docs/concepts-security-command-center-overview#inventory>  
upvoted 2 times

  **pfilourenco** 1 year, 7 months ago

And about D?  
upvoted 2 times

  **dtmtor** 1 year, 9 months ago

Answer is B  
upvoted 2 times

  **pythonrocks** 1 year, 5 months ago

<https://forsetisecurity.org/about/> inventory  
upvoted 1 times

🗨️ 👤 **DebasishLowes** 1 year, 9 months ago

Ans : B. You need to keep the records for long time so it's Inventory.  
upvoted 2 times

🗨️ 👤 **[Removed]** 2 years, 1 month ago

Ans - B  
upvoted 2 times

🗨️ 👤 **MohitA** 2 years, 4 months ago

B is the right answer, Forseti has a track of inventory snapshots  
upvoted 1 times

🗨️ 👤 **aiwaai** 2 years, 4 months ago

Correct answer is B  
upvoted 1 times

🗨️ 👤 **ArizonaClassics** 2 years, 4 months ago

I will go with B  
question says " historical inventory" and that makes B the right choice  
upvoted 4 times

🗨️ 👤 **MarkDillon1075** 2 years, 5 months ago

B is correct  
upvoted 1 times

🗨️ 👤 **xhova** 2 years, 8 months ago

A is correct  
upvoted 1 times

🗨️ 👤 **ArizonaClassics** 2 years, 4 months ago

B IS CORRECT SEE: <https://forsetisecurity.org/about/>  
upvoted 2 times

🗨️ 👤 **mozammil89** 2 years, 9 months ago

The correct answer is A.

Feature of Resource Manager - <https://cloud.google.com/asset-inventory>

upvoted 3 times

🗨️ 👤 **gcp\_learner** 2 years, 5 months ago

The answer is not A because Cloud Asset Inventory is not one of the choices. The correct answer is B.  
upvoted 5 times

🗨️ 👤 **ArizonaClassics** 2 years, 4 months ago

SEE THE SUB-HEADING 'INVENTORY' ON THIS LINK: <https://forsetisecurity.org/about/>  
upvoted 1 times

🗨️ 👤 **FatCharlie** 2 years ago

Resource Manager / Cloud Asset Inventory only keeps data for 6 weeks, so doesn't meet the requirements in the question  
upvoted 1 times



An organization is starting to move its infrastructure from its on-premises environment to Google Cloud Platform (GCP). The first step the organization wants to take is to migrate its current data backup and disaster recovery solutions to GCP for later analysis. The organization's production environment will remain on-premises for an indefinite time. The organization wants a scalable and cost-efficient solution. Which GCP solution should the organization use?

- A. BigQuery using a data pipeline job with continuous updates
- B. Cloud Storage using a scheduled task and gsutil
- C. Compute Engine Virtual Machines using Persistent Disk
- D. Cloud Datastore using regularly scheduled batch upload jobs

**Correct Answer: A**

*Community vote distribution*

B (100%)

🗳️ 👤 **xhova** Highly Voted 👍 2 years, 8 months ago

Ans is B. A cost efficient disaster recovery solution is needed not a data warehouse.  
upvoted 20 times

🗳️ 👤 **AwesomeGCP** Most Recent 🕒 2 months, 2 weeks ago

**Selected Answer: B**

B confirmed :-)  
<https://cloud.google.com/solutions/dr-scenarios-planning-guide#use-cloud-storage-as-part-of-your-daily-backup-routine>  
upvoted 2 times

🗳️ 👤 **AzureDP900** 1 month, 2 weeks ago

It is B  
upvoted 1 times

🗳️ 👤 **giovy\_82** 4 months ago

I would go for B, but a doubt remains: it is talking about Disaster Recovery solution, which could not only be related to data but also to VM and applications running inside VMs. any way B is more cost-efficient than A, considering also that data backup need to be moved to GCP.  
upvoted 1 times

🗳️ 👤 **absipat** 6 months, 2 weeks ago

B of course  
upvoted 2 times

🗳️ 👤 **DebasishLowes** 1 year, 9 months ago

Ans : B. Cloud storage is cost efficient one.  
upvoted 4 times

🗳️ 👤 **[Removed]** 2 years, 1 month ago

Ans - B  
upvoted 2 times

🗳️ 👤 **CHECK666** 2 years, 2 months ago

B is the answer.  
upvoted 2 times

🗳️ 👤 **paxjoshi** 2 years, 4 months ago

B is the correct answer. They need the data for later analysis and they are looking for cost-effective service.  
upvoted 2 times

🗳️ 👤 **aiwaai** 2 years, 4 months ago

Correct Answer: A  
upvoted 1 times

🗳️ 👤 **aiwaai** 2 years, 4 months ago

I make corrections, B is Correct Answer.  
upvoted 1 times

🗳️ 👤 **ArizonaClassics** 2 years, 4 months ago

Answer B works for me as the type of workload to be stored is not stated or defined  
upvoted 1 times

  **SilentSec** 2 years, 5 months ago

B confirmed: <https://cloud.google.com/solutions/dr-scenarios-planning-guide#use-cloud-storage-as-part-of-your-daily-backup-routine>  
upvoted 3 times



You are creating an internal App Engine application that needs to access a user's Google Drive on the user's behalf. Your company does not want to rely on the current user's credentials. It also wants to follow Google-recommended practices.

What should you do?

- A. Create a new Service account, and give all application users the role of Service Account User.
- B. Create a new Service account, and add all application users to a Google Group. Give this group the role of Service Account User.
- C. Use a dedicated G Suite Admin account, and authenticate the application's operations with these G Suite credentials.
- D. Create a new service account, and grant it G Suite domain-wide delegation. Have the application use it to impersonate the user.

**Correct Answer: A**

*Community vote distribution*

D (100%)

  **mozammil89** Highly Voted 2 years, 9 months ago

I think the correct answer is D

<https://developers.google.com/admin-sdk/directory/v1/guides/delegation>

upvoted 16 times

  **Meyucho** Most Recent 1 month ago

**Selected Answer: D**

- A (Wrong) The access will be with the SA not the user's account.
- B (Wrong) Same as A.
- C. (Wrong) In this case the access is with the admins account, not user's.
- D. (CORRECT!) It's the only answer that really impersonate the user.

upvoted 1 times

  **AzureDP900** 1 month, 2 weeks ago

D. Create a new service account, and grant it G Suite domain-wide delegation. Have the application use it to impersonate the user.

upvoted 1 times

  **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: D**

correct answer is D

<https://developers.google.com/admin-sdk/directory/v1/guides/delegation>

upvoted 2 times

  **Medofree** 8 months, 2 weeks ago

**Selected Answer: D**

Clearly D is the right answer

upvoted 2 times

  **Rhehehe** 1 year ago

They are asking for google recommended practice. Does D says that?

upvoted 1 times

  **[Removed]** 2 years, 1 month ago

Ans - D

upvoted 2 times

  **CHECK666** 2 years, 2 months ago

D is the answer.

upvoted 1 times

  **ArizonaClassics** 2 years, 4 months ago

D is the best choice

upvoted 1 times

  **MarkDillon1075** 2 years, 5 months ago

I agree D

upvoted 1 times

A customer wants to move their sensitive workloads to a Compute Engine-based cluster using Managed Instance Groups (MIGs). The jobs are bursty and must be completed quickly. They have a requirement to be able to control the key lifecycle.

Which boot disk encryption solution should you use on the cluster to meet this customer's requirements?

- A. Customer-supplied encryption keys (CSEK)
- B. Customer-managed encryption keys (CMEK) using Cloud Key Management Service (KMS)
- C. Encryption by default
- D. Pre-encrypting files before transferring to Google Cloud Platform (GCP) for analysis

**Correct Answer:** B

Reference -

<https://cloud.google.com/kubernetes-engine/docs/how-to/dynamic-provisioning-cmek>

*Community vote distribution*

B (100%)

 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: B**

Correct Answer: B

Explanation/Reference:

Reference <https://cloud.google.com/kubernetes-engine/docs/how-to/dynamic-provisioning-cmek>

upvoted 2 times

 **animesh54** 7 months, 3 weeks ago

**Selected Answer: B**

Customer Managed Encryption keys using KMS lets users control the key management and rotation policies and Compute Engine Disks support CMEKs

upvoted 3 times

Your company is using Cloud Dataproc for its Spark and Hadoop jobs. You want to be able to create, rotate, and destroy symmetric encryption keys used for the persistent disks used by Cloud Dataproc. Keys can be stored in the cloud. What should you do?

- A. Use the Cloud Key Management Service to manage the data encryption key (DEK).
- B. Use the Cloud Key Management Service to manage the key encryption key (KEK).
- C. Use customer-supplied encryption keys to manage the data encryption key (DEK).
- D. Use customer-supplied encryption keys to manage the key encryption key (KEK).

**Correct Answer: A**

*Community vote distribution*

B (60%)

A (40%)

  **mte\_tech34** Highly Voted 2 years, 2 months ago

Answer is B.

<https://cloud.google.com/dataproc/docs/concepts/configuring-clusters/customer-managed-encryption>

"The CMEK feature allows you to create, use, and revoke the key encryption key (KEK). Google still controls the data encryption key (DEK)."

upvoted 18 times

  **passtest100** 2 years, 2 months ago

SHOULD BE A.

NO envelope encryption is mentioned in the question.

upvoted 2 times

  **Arad** 1 year ago

Correct answer is B, and A is wrong!

envelope encryption is default mechanism in CMEK when used for Dataproc, please check this link:

This PD and bucket data is encrypted using a Google-generated data encryption key (DEK) and key encryption key (KEK). The CMEK feature allows you to create, use, and revoke the key encryption key (KEK). Google still controls the data encryption key (DEK). For more information on Google data encryption keys, see Encryption at Rest.

<https://cloud.google.com/dataproc/docs/concepts/configuring-clusters/customer-managed-encryption>

upvoted 1 times

  **mynk29** 9 months, 4 weeks ago

I agree but then should answer not be C- customer supplied key?

upvoted 1 times

  **mynk29** 9 months, 4 weeks ago

My bad I read it as Customer managed.. even though i now realised i wrote customer supplied. :D

upvoted 1 times

  **Meyucho** Most Recent 1 month ago

**Selected Answer: A**

B can be right but we never been asked about envelope encryption... so... the solution is to use a customer managed Data Encryption Key

upvoted 1 times

  **AzureDP900** 1 month, 2 weeks ago

B. Use the Cloud Key Management Service to manage the key encryption key (KEK).

upvoted 1 times

  **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: B**

Answer is B,

<https://cloud.google.com/dataproc/docs/concepts/configuring-clusters/customer-managed-encryption>

upvoted 2 times

  **giovy\_82** 4 months ago

**Selected Answer: B**

In my opinion it should be B. reference :

<https://cloud.google.com/kms/docs/envelope-encryption>

How to encrypt data using envelope encryption

The process of encrypting data is to generate a DEK locally, encrypt data with the DEK, use a KEK to wrap the DEK, and then store the encrypted data and the wrapped DEK. The KEK never leaves Cloud KMS.

upvoted 1 times

  **piyush\_1982** 4 months, 3 weeks ago

**Selected Answer: A**

I think the answer is A.

DEK (Data encryption Key ) is the key which is used to encrypt the data. It can be both customer-managed or customer supplied in terms of GCP>  
<https://cloud.google.com/dataproc/docs/concepts/configuring-clusters/customer-managed-encryption>

The link above states "This PD and bucket data is encrypted using a Google-generated data encryption key (DEK) and key encryption key (KEK). The CMEK feature allows you to create, use, and revoke the key encryption key (KEK). Google still controls the data encryption key (DEK)."

upvoted 1 times

  **absipat** 6 months, 2 weeks ago

b of course

upvoted 1 times

  **[Removed]** 1 year, 8 months ago

I also support B, but A is also good ,because kek is hosted within KMS, also the real DEK can be uploaded there ,or just in the database.

upvoted 2 times

  **DebasishLowes** 1 year, 9 months ago

Ans : B.

upvoted 2 times

  **[Removed]** 2 years, 1 month ago

Ans - B

upvoted 2 times

  **saaurabh1805** 2 years, 1 month ago

Answer is B,

<https://cloud.google.com/dataproc/docs/concepts/configuring-clusters/customer-managed-encryption>

upvoted 3 times

You are a member of the security team at an organization. Your team has a single GCP project with credit card payment processing systems alongside web applications and data processing systems. You want to reduce the scope of systems subject to PCI audit standards. What should you do?

- A. Use multi-factor authentication for admin access to the web application.
- B. Use only applications certified compliant with PA-DSS.
- C. Move the cardholder data environment into a separate GCP project.
- D. Use VPN for all connections between your office and cloud environments.

**Correct Answer:** C

Community vote distribution

C (100%)

🗲️ 👤 **jonclem** Highly Voted 2 years, 9 months ago  
I'd go for answer C myself.

<https://cloud.google.com/solutions/best-practices-vpc-design>  
upvoted 21 times

🗲️ 👤 **[Removed]** Highly Voted 2 years, 1 month ago  
Ans - C  
[https://cloud.google.com/solutions/pci-dss-compliance-in-gcp#setting\\_up\\_your\\_payment-processing\\_environment](https://cloud.google.com/solutions/pci-dss-compliance-in-gcp#setting_up_your_payment-processing_environment)  
upvoted 7 times

🗲️ 👤 **AzureDP900** Most Recent 1 month, 2 weeks ago  
answer is C  
upvoted 1 times

🗲️ 👤 **Medofree** 8 months, 2 weeks ago  
Selected Answer: C  
Projets are units of isolationm the answer is C.  
upvoted 1 times

🗲️ 👤 **CHECK666** 2 years, 2 months ago  
C is the answer.  
upvoted 1 times

🗲️ 👤 **smart123** 2 years, 6 months ago  
The Answer is C. Check "Setting up your payment-processing environment" section in <https://cloud.google.com/solutions/pci-dss-compliance-in-gcp>.  
In the question, it is mentioned that it is the same environment for card processing as the Web App and Data processing and that is not recommended.  
upvoted 3 times

🗲️ 👤 **xhova** 2 years, 8 months ago  
Definitely C  
upvoted 1 times

A retail customer allows users to upload comments and product reviews. The customer needs to make sure the text does not include sensitive data before the comments or reviews are published.

Which Google Cloud Service should be used to achieve this?

- A. Cloud Key Management Service
- B. Cloud Data Loss Prevention API
- C. BigQuery
- D. Web Security Scanner

**Correct Answer:** D

Community vote distribution

B (100%)

🗳️ 👤 **rafaelc** Highly Voted 2 years, 9 months ago

It's definitely B. It was on the practice test on google site.  
B. Cloud Data Loss Prevention API  
upvoted 25 times

🗳️ 👤 **VishalBulbule** Most Recent 3 days, 17 hours ago

"before the comments or reviews are published" - how will we use DLP API , so web scanner can be considered for correct answer.  
upvoted 1 times

🗳️ 👤 **huntergame** 1 month, 2 weeks ago

Selected Answer: B

Its obvious DLP  
upvoted 1 times

🗳️ 👤 **PopeyeTheSailorMan** 4 months, 4 weeks ago

Selected Answer: B

The answer can not be D (I am laughing loud since I use D for the reason of security scanning) hence the correct answer is B and it is not D  
upvoted 1 times

🗳️ 👤 **Bwitch** 1 year ago

Selected Answer: B

DLP provides the service of redaction.  
upvoted 3 times

🗳️ 👤 **DebasishLowes** 1 year, 9 months ago

Its B.  
upvoted 2 times

🗳️ 👤 **saurabh1805** 2 years, 1 month ago

B is correct answer here.  
upvoted 2 times

🗳️ 👤 **[Removed]** 2 years, 1 month ago

Ans - B  
upvoted 2 times

🗳️ 👤 **CHECK666** 2 years, 2 months ago

B is the answer.  
upvoted 2 times

🗳️ 👤 **aiwaai** 2 years, 4 months ago

Correct Answer: B  
upvoted 1 times

🗳️ 👤 **paxjoshi** 2 years, 4 months ago

Yes, the correct answer is B.  
upvoted 1 times

🗳️ 👤 **aiwaai** 2 years, 4 months ago

Correct Answer: B



upvoted 1 times

  **bigdo** 2 years, 4 months ago

B D is for vulnerability scanning  
upvoted 1 times

  **smart123** 2 years, 5 months ago

The Answer is B  
upvoted 1 times

  **xhova** 2 years, 8 months ago

Definitely B  
upvoted 2 times

  **meraexam** 2 years, 9 months ago


How it is D. should have been B  
upvoted 4 times

A company allows every employee to use Google Cloud Platform. Each department has a Google Group, with all department members as group members. If a department member creates a new project, all members of that department should automatically have read-only access to all new project resources. Members of any other department should not have access to the project. You need to configure this behavior. What should you do to meet these requirements?

- A. Create a Folder per department under the Organization. For each department's Folder, assign the Project Viewer role to the Google Group related to that department.
- B. Create a Folder per department under the Organization. For each department's Folder, assign the Project Browser role to the Google Group related to that department.
- C. Create a Project per department under the Organization. For each department's Project, assign the Project Viewer role to the Google Group related to that department.
- D. Create a Project per department under the Organization. For each department's Project, assign the Project Browser role to the Google Group related to that department.

**Correct Answer: C**

*Community vote distribution*



🗳️ 👤 **ownez** Highly Voted 2 years, 3 months ago  
Shouldn't it be A?

Project Browser has least permissions comparing to Project Viewer. The question is about have read-access to all new project resources.

roles/browser - Read access to browse the hierarchy for a project, including the folder, organization, and IAM policy. This role doesn't include permission to view resources in the project.

<https://cloud.google.com/iam/docs/understanding-roles#project-roles>  
upvoted 20 times

🗳️ 👤 **singhjoga** 1 year, 11 months ago  
Correct, it is A. Project Browser does not have access to the resources inside the project, which is the requirement in the question.  
upvoted 7 times

🗳️ 👤 **Meyucho** Most Recent 1 month ago  
Selected Answer: A  
Who voted C!?!?!?! The answer is A!!!!  
upvoted 1 times

🗳️ 👤 **AwesomeGCP** 2 months, 2 weeks ago  
Selected Answer: A  
Correct answer - A  
<https://cloud.google.com/resource-manager/docs/cloud-platform-resource-hierarchy>  
upvoted 1 times

🗳️ 👤 **piyush\_1982** 4 months, 4 weeks ago  
Selected Answer: C  
The correct answer is definitely C.  
Let's divide the question into 2 parts:

1st: Role: Key requirement: all members of that department should automatically have read-only access to all new project resources.

> The project browser role only allows read access to browse the hierarchy for a project, including the folder, organization, and allow policy. This role doesn't include permission to view resources in the project.  
Hence the options B and D are not relevant as they both are browser roles which DO NOT provide access to project resources.

2nd: Option A creates a Folder per department and C creates project per department.  
However, Project viewer role is only applied at the project level.  
Hence the correct answer is C which creates projects per department under organization .  
upvoted 1 times

🗳️ 👤 **Meyucho** 1 month ago  
But... if you dont have a folder per department.. where will be all new projects created by users???? you will have to manually edit permissions every time!!!! Using folders yu set the permissions once and then the only task you shoul do is to maintain the proper group assignment



upvoted 1 times

🗨️ 👤 **alvjtc** 5 months, 2 weeks ago

**Selected Answer: A**

It's A, Project Viewer. Project Browser doesn't allow users to see resources, only find the project in the hierarchy.

upvoted 1 times

🗨️ 👤 **syllox** 1 year, 7 months ago

It's A , browser is :

Read access to browse the hierarchy for a project, including the folder, organization, and IAM policy. This role doesn't include permission to view resources in the project.

<https://cloud.google.com/iam/docs/understanding-roles#project-roles>

upvoted 3 times

🗨️ 👤 **[Removed]** 1 year, 8 months ago

either A or C because must be project viewer ,browser is not enough.<https://cloud.google.com/iam/docs/understanding-roles>

upvoted 1 times

🗨️ 👤 **[Removed]** 1 year, 8 months ago

Why not A?

upvoted 1 times

🗨️ 👤 **desertlotus1211** 1 year, 9 months ago

The answer is A:

<https://stackoverflow.com/questions/54778596/whats-the-difference-between-project-browser-role-and-project-viewer-role-in-go#:~:text=8-,What's%20the%20difference%20between%20Project%20Browser%20role%20and,role%20in%20Google%20Cloud%20Platform&text=According%20to%20the%20console%20popup,read%20access%20to%20those%20resources.>

upvoted 2 times

🗨️ 👤 **CloudTrip** 1 year, 10 months ago

I think it's B. As the question says all members of that department should automatically have read-only access to all new project resources but browser will only provide the get, list permissions not read only permission so viewer seems to be more accurate here.

roles/browser

Read access to browse the hierarchy for a project, including the folder, organization, and IAM policy. This role doesn't include permission to view resources in the project.

resourcemanager.folders.get

resourcemanager.folders.list

resourcemanager.organizations.get

resourcemanager.projects.get

resourcemanager.projects.getIamPolicy

resourcemanager.projects.list

roles/viewer Viewer Permissions for read-only actions that do not affect state, such as viewing (but not modifying) existing resources or data.

upvoted 1 times

🗨️ 👤 **subhala** 2 years ago

Question says - If a department member creates a new project, all members of that department should automatically have read-only access to all new project resources. and @ownez provided documentation that says - browser role doesn't include perm to view resources in the project. Hence B is the right answer.

upvoted 1 times

🗨️ 👤 **Fellipo** 2 years, 1 month ago

A it's OK

upvoted 2 times

🗨️ 👤 **[Removed]** 2 years, 1 month ago

Ans - A

upvoted 2 times

🗨️ 👤 **cipher90** 2 years, 3 months ago

Answer is B: "have read-only access to all new project resources." So it has to be in a folder to cascade the permissions to new projects carried.

upvoted 1 times

🗨️ 👤 **Meyucho** 1 month ago

If you do that the other members of the department can't access to the resources.. just list the project in the folder

upvoted 1 times

A customer's internal security team must manage its own encryption keys for encrypting data on Cloud Storage and decides to use customer-supplied encryption keys (CSEK).

How should the team complete this task?

- A. Upload the encryption key to a Cloud Storage bucket, and then upload the object to the same bucket.
- B. Use the gsutil command line tool to upload the object to Cloud Storage, and specify the location of the encryption key.
- C. Generate an encryption key in the Google Cloud Platform Console, and upload an object to Cloud Storage using the specified key.
- D. Encrypt the object, then use the gsutil command line tool or the Google Cloud Platform Console to upload the object to Cloud Storage.

**Correct Answer:** D

Reference:

<https://cloud.google.com/storage/docs/encryption/customer-supplied-keys>

Community vote distribution

B (100%)

🗳️ **DebasishLowes** Highly Voted 1 year, 9 months ago

Ans : B. Because if you encrypt the object using CSEK, then you can't use google cloud console to upload the object.  
upvoted 10 times

🗳️ **FatCharlie** Highly Voted 2 years ago

The fact is, both B & D would work. I lean towards B because it allows you to manage the file using GCP tools later as long as you keep that key around.

B is definitely incomplete though, as the boto file does need to be updated.  
upvoted 7 times

🗳️ **AwesomeGCP** Most Recent 2 months, 2 weeks ago

**Selected Answer: B**

<https://cloud.google.com/storage/docs/encryption/customer-supplied-keys>  
Answer B  
upvoted 1 times

🗳️ **GHOST1985** 2 months, 3 weeks ago

**Selected Answer: B**

you can't use google cloud console to upload the object.  
[https://cloud.google.com/storage/docs/encryption/using-customer-supplied-keys#upload\\_with\\_your\\_encryption\\_key](https://cloud.google.com/storage/docs/encryption/using-customer-supplied-keys#upload_with_your_encryption_key)  
upvoted 1 times

🗳️ **absipat** 6 months, 2 weeks ago

D of course  
upvoted 1 times

🗳️ **Aiffone** 6 months, 2 weeks ago

I will go with D because encrypting the object before uploading means the cutomer manages thier own key.  
A is not correct because its not a good practice to upload encryption key to storage object along with the encrypted object.  
B is not correct because specifying the location of the encryption key does not change anything  
C means Google manages the key.  
upvoted 1 times

🗳️ **[Removed]** 1 year, 8 months ago

CD are not right because Google Cloud Console does not support CSEK. must choose from A and B  
upvoted 1 times

🗳️ **VivekA** 2 years ago

Ans: B  
CSEK doesn't work on Cloud Shell or Cloud Console for Cloud Storage  
<https://cloud.google.com/storage/docs/encryption/customer-supplied-keys>  
upvoted 4 times

🗳️ **VivekA** 2 years ago

D can't be an option - Refer restriction section.  
<https://cloud.google.com/storage/docs/encryption/customer-supplied-keys#:~:text=You%20cannot%20use%20the%20Google,a%20customer%2Dsupplied%20encryption%20key.>  
upvoted 1 times

  **Munna\_Mohammed** 2 years, 1 month ago

Ans is C

upvoted 1 times

  **Fellipo** 2 years, 1 month ago

if Option B was correct you have to modify your boto file which is not the full option.

upvoted 3 times

  **Fellipo** 2 years, 1 month ago

B is correct , sorry

upvoted 4 times

  **[Removed]** 2 years, 1 month ago

Ans - D

upvoted 4 times

  **saurabh1805** 2 years, 1 month ago

Option B is closed and correct one.

upvoted 3 times

  **saurabh1805** 2 years, 1 month ago

\*closest

upvoted 1 times

  **skshak** 2 years, 3 months ago

Sorry, Answer is C. - Generate a key in cloud shell and then upload the object using that key

upvoted 1 times

  **skshak** 2 years, 3 months ago

<https://cloud.google.com/storage/docs/encryption/customer-supplied-keys>

Answer B

upvoted 4 times

A customer has 300 engineers. The company wants to grant different levels of access and efficiently manage IAM permissions between users in the development and production environment projects.

Which two steps should the company take to meet these requirements? (Choose two.)

- A. Create a project with multiple VPC networks for each environment.
- B. Create a folder for each development and production environment.
- C. Create a Google Group for the Engineering team, and assign permissions at the folder level.
- D. Create an Organizational Policy constraint for each folder environment.
- E. Create projects for each environment, and grant IAM rights to each engineering user.

**Correct Answer:** *BD*

*Community vote distribution*

BC (100%)

  **mozammil89** Highly Voted 2 years, 9 months ago

B and C should be correct...  
upvoted 23 times

  **Meyucho** Most Recent 1 month ago

**Selected Answer: BC**

Which Policy Constriant allow to manage permission?!?!?! D is not an option. The answer is B and C  
upvoted 1 times

  **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: BC**

B and C are the correct answers!!  
upvoted 2 times

  **danielklein09** 9 months, 2 weeks ago

B is correct  
But, if you make 1 group (by choosing option C) how you manage the permission for dev environment ? since you have only 1 group, you will offer the same access for all 300 engineers (that are in that group) to dev and prod environment, so this will not answer the question: efficiently manage IAM permissions between users in the development and production environment projects  
upvoted 3 times

  **Ksrp** 10 months ago

CE - A general recommendation is to have one project per application per environment. For example, if you have two applications, "app1" and "app2", each with a development and production environment, you would have four projects: app1-dev, app1-prod, app2-dev, app2-prod. This isolates the environments from each other, so changes to the development project do not accidentally impact production, and gives you better access control, since you can (for example) grant all developers access to development projects but restrict production access to your CI/CD pipeline. <https://cloud.google.com/docs/enterprise/best-practices-for-enterprise-organizations>  
upvoted 1 times

  **Jane111** 1 year, 8 months ago

A - no VPC required  
B - yes - pre req  
C - Yes  
D - likely but C is first  
E - not scalable/feasible/advisable  
upvoted 1 times

  **DebasishLowes** 1 year, 9 months ago

Ans : BC  
upvoted 1 times

  **[Removed]** 2 years, 1 month ago

Ans - BC  
upvoted 1 times

  **CHECK666** 2 years, 2 months ago

B,C is the answer.  
Create a folder for each env and assign IAM policies to the group.  
upvoted 2 times

- MohitA** 2 years, 4 months ago

BC is the right answer, create folder for each env and assign IAM policies to group

upvoted 1 times
- aiwaai** 2 years, 4 months ago

Correct Answer: CE

upvoted 1 times
- aiwaai** 2 years, 3 months ago

made correction CE -> BC

upvoted 2 times
- xhova** 2 years, 8 months ago

B&C

D does not help efficiently manage IAM. Effective IAM implies using groups.

upvoted 2 times
- smart123** 2 years, 6 months ago

Organization policy is used on resources and not the users. Hence option 'D' cannot be right.

upvoted 2 times
- jonclem** 2 years, 9 months ago

I'd say B and D are correct

upvoted 1 times

## Question #40

Topic 1

You want to evaluate your organization's Google Cloud instance for PCI compliance. You need to identify Google's inherent controls. Which document should you review to find the information?

- A. Google Cloud Platform: Customer Responsibility Matrix
- B. PCI DSS Requirements and Security Assessment Procedures
- C. PCI SSC Cloud Computing Guidelines
- D. Product documentation for Compute Engine

**Correct Answer: C**

Reference:

<https://cloud.google.com/solutions/pci-dss-compliance-in-gcp>

*Community vote distribution*

A (100%)

- AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: A**

A. Google Cloud Platform: Customer Responsibility Matrix

upvoted 1 times
- tangac** 3 months, 2 weeks ago

**Selected Answer: A**

[https://services.google.com/fh/files/misc/gcp\\_pci\\_shared\\_responsibility\\_matrix\\_aug\\_2021.pdf](https://services.google.com/fh/files/misc/gcp_pci_shared_responsibility_matrix_aug_2021.pdf)

upvoted 2 times

Your company runs a website that will store PII on Google Cloud Platform. To comply with data privacy regulations, this data can only be stored for a specific amount of time and must be fully deleted after this specific period. Data that has not yet reached the time period should not be deleted. You want to automate the process of complying with this regulation. What should you do?

- A. Store the data in a single Persistent Disk, and delete the disk at expiration time.
- B. Store the data in a single BigQuery table and set the appropriate table expiration time.
- C. Store the data in a single Cloud Storage bucket and configure the bucket's Time to Live.
- D. Store the data in a single BigTable table and set an expiration time on the column families.

Correct Answer: B

Community vote distribution



**KILLMAD** Highly Voted 2 years, 9 months ago  
I believe the Answer is C not B.

This isn't data which needs to be analyzed, so I don't understand why would it be stored in BQ when having data stored in GCS seems much more reasonable.

I think the only thing about answer C which throws me off is the fact that they don't mention object life cycle management  
upvoted 12 times

**mozammil89** 2 years, 9 months ago  
Answer C is correct. The TTL is common use case of Cloud Storage life cycle management. Here is what GCP says:  
  
"To support common use cases like setting a Time to Live (TTL) for objects, retaining noncurrent versions of objects, or "downgrading" storage classes of objects to help manage costs, Cloud Storage offers the Object Lifecycle Management feature. This page describes the feature as well as the options available when using it. To learn how to enable Object Lifecycle Management, and for examples of lifecycle policies, see Managing Lifecycles."  
  
<https://cloud.google.com/storage/docs/lifecycle>  
upvoted 4 times

**PST21** Most Recent 4 days, 12 hours ago  
CS does not delete promptly , hence BQ as it is sensitive data  
upvoted 1 times

**csrazdan** 2 weeks, 5 days ago  
Selected Answer: B  
Life Cycle Management for Cloud storage is used to manage the Storage class to save cost. For data management, you have set retention time on the bucket. I will opt for B as the correct answer.  
upvoted 1 times

**AwesomeGCP** 2 months, 2 weeks ago  
Selected Answer: C  
Correct Answer: C  
upvoted 1 times

**giovy\_82** 4 months ago  
I would go for C, but all the 4 answers are in my opinion incomplete. all of them say "single" bucket or table, which means that if different dated rows/elements are stored in the same bucket or table, they will expire together and be deleted probably before their real expiration time. so i expected to see partitioning or multiple bucket.  
upvoted 2 times

**mynk29** 9 months, 4 weeks ago  
Outdated question again- should be bucket locks now.  
upvoted 1 times

**DebasishLowes** 1 year, 9 months ago  
Ans : C  
upvoted 2 times

**[Removed]** 2 years, 1 month ago  
Ans - C



upvoted 4 times

  **aiwaai** 2 years, 4 months ago

Correct Answer: C

upvoted 2 times

  **Ganshank** 2 years, 7 months ago

The answers need to be worded better.

If we're taking the terms literally as specified in the options, then C cannot be the correction answer since there's no Time to Live configuration for a GCS bucket, only Lifecycle Policy.

With BigQuery, there is no row-level expiration, although we could create this behavior using Partitioned Tables. So this could be a potential answer.

D - it is possible to simulate cell-level TTL (<https://cloud.google.com/bigtable/docs/gc-cell-level>), so this too could be a potential answer, especially when different cells need different TTLs.

Between B & D, BigQuery follows a pay-as-you-go model and its storage costs are comparable to GCS storage costs. So this would be the more appropriate solution.

upvoted 3 times

  **smart123** 2 years, 5 months ago

The Buckets do have "Time to Live" feature.

<https://cloud.google.com/storage/docs/lifecycle>

Hence 'C' is the answer

upvoted 4 times

  **jonclem** 2 years, 9 months ago

I believe B is correct.

Setting a TTL of 14 days on the bucket via LifeCycle will not cause the bucket itself to be deleted after 14 days, instead it will cause each object uploaded to that bucket to be deleted 14 days after it was created

upvoted 2 times

  **xhova** 2 years, 8 months ago

Answer is C. You dont need the bucket to be deleted, you need the PII data stored to be deleted.

upvoted 5 times

A DevOps team will create a new container to run on Google Kubernetes Engine. As the application will be internet-facing, they want to minimize the attack surface of the container.

What should they do?

- A. Use Cloud Build to build the container images.
- B. Build small containers using small base images.
- C. Delete non-used versions from Container Registry.
- D. Use a Continuous Delivery tool to deploy the application.

Correct Answer: D

Reference:

<https://cloud.google.com/solutions/best-practices-for-building-containers>

Community vote distribution



**xhova** Highly Voted 2 years, 8 months ago  
Ans is B

Small containers usually have a smaller attack surface as compared to containers that use large base images.

<https://cloud.google.com/blog/products/gcp/kubernetes-best-practices-how-and-why-to-build-small-container-images>  
upvoted 27 times

**smart123** 2 years, 5 months ago  
I agree  
upvoted 1 times

**rotorclear** Most Recent 2 months ago  
Selected Answer: B  
B definitely  
upvoted 1 times

**AwesomeGCP** 2 months, 2 weeks ago  
Selected Answer: B  
B is the correct answer.  
upvoted 1 times

**zellck** 2 months, 3 weeks ago  
Selected Answer: B  
B is the answer.  
upvoted 1 times

**jitu028** 2 months, 3 weeks ago  
Ans is B - <https://cloud.google.com/blog/products/gcp/kubernetes-best-practices-how-and-why-to-build-small-container-images>  
  
Security and vulnerabilities  
Aside from performance, there are significant security benefits from using smaller containers. Small containers usually have a smaller attack surface as compared to containers that use large base images.  
upvoted 2 times

**giovy\_82** 4 months ago  
Selected Answer: B  
the only answer that will really reduce attack surface while exposing apps to internet is B, small containers (e.g. single web page?)  
upvoted 2 times

**Medofree** 8 months, 2 weeks ago  
B. Because you will have less programs in the image thus less vulnerabilities  
upvoted 1 times

**lxs** 1 year ago  
Selected Answer: C  
A. Use Cloud Build to build the container images.  
If you build a container using Cloud Build or not the surface is the same



B. Build small containers using small base images.

It is indeed best practice, but I doubt if small base images can reduce the surface. It is still the same app version with the same vulnerabilities etc.

C. Delete non-used versions from Container Registry.

Unused, historical versions are additional attack surface. attacker can exploit old, unpatched image which indeed the surface extention.

D. Use a Continuous Delivery tool to deploy the application.

This is just a method of image delivery. The app is the same.

upvoted 2 times

  **DebasishLowes** 1 year, 9 months ago

Ans : B. Small the base image there is less vulnerability and less chance of attack.

upvoted 2 times

  **soukumar369** 2 years ago

Answer is B : at the cost of more space and a slightly higher surface area for attacks.

upvoted 3 times

  **[Removed]** 2 years, 1 month ago

should be A, because cloud build has a feature for vulnerability scan

upvoted 3 times

  **[Removed]** 2 years, 1 month ago

Ans - B

upvoted 1 times

  **[Removed]** 2 years, 1 month ago

Ans - B

upvoted 1 times

  **CHECK666** 2 years, 2 months ago

B is the answer, small containers have smaller attack surface.

upvoted 1 times

  **ArizonaClassics** 2 years, 4 months ago

The Correct answer is B

upvoted 1 times

While migrating your organization's infrastructure to GCP, a large number of users will need to access GCP Console. The Identity Management team already has a well-established way to manage your users and want to keep using your existing Active Directory or LDAP server along with the existing SSO password.

What should you do?


- A. Manually synchronize the data in Google domain with your existing Active Directory or LDAP server.
- B. Use Google Cloud Directory Sync to synchronize the data in Google domain with your existing Active Directory or LDAP server.
- C. Users sign in directly to the GCP Console using the credentials from your on-premises Kerberos compliant identity provider.
- D. Users sign in using OpenID (OIDC) compatible IdP, receive an authentication token, then use that token to log in to the GCP Console.

**Correct Answer:** B

Reference:

<https://cloud.google.com/blog/products/identity-security/using-your-existing-identity-management-system-with-google-cloud-platform>

*Community vote distribution*



**sudarchary** Highly Voted 10 months, 3 weeks ago

**Selected Answer: B**

<https://cloud.google.com/architecture/identity/federating-gcp-with-active-directory-configuring-single-sign-on>  
upvoted 6 times

**AwesomeGCP** Most Recent 2 months, 2 weeks ago

**Selected Answer: B**

<https://cloud.google.com/architecture/identity/federating-gcp-with-active-directory-configuring-single-sign-on>  
upvoted 1 times

**absipat** 6 months, 2 weeks ago

B of course  
upvoted 1 times

**ThisisJohn** 1 year ago

**Selected Answer: D**

My vote goes for D.

From the blog post linked below " users' passwords are not synchronized by default. Only the identities are synchronized, unless you make an explicit choice to synchronize passwords (which is not a best practice and should be avoided)".

Also, from GCP documentation "Authenticating with OIDC and AD FS" <https://cloud.google.com/anthos/clusters/docs/on-prem/1.6/how-to/oidc-adfs>

Blog post quoted above <https://cloud.google.com/blog/products/identity-security/using-your-existing-identity-management-system-with-google-cloud-platform>  
upvoted 1 times

**rr4444** 11 months, 3 weeks ago

D sounds nice, but the user doesn't "use" the token.... that's used in the integration with Cloud Identity. So answer must be B, GCDS  
upvoted 2 times

**DebasishLowes** 1 year, 9 months ago

Ans : B  
upvoted 4 times

**[Removed]** 2 years, 1 month ago

Ans - B  
upvoted 3 times

**saaurabh1805** 2 years, 1 month ago

B is correct answer here.  
upvoted 3 times

Your company is using GSuite and has developed an application meant for internal usage on Google App Engine. You need to make sure that an external user cannot gain access to the application even when an employee's password has been compromised. What should you do?

- A. Enforce 2-factor authentication in GSuite for all users.
- B. Configure Cloud Identity-Aware Proxy for the App Engine Application.
- C. Provision user passwords using GSuite Password Sync.
- D. Configure Cloud VPN between your private network and GCP.

Correct Answer: D

Community vote distribution



- rafaelc** Highly Voted

2 years, 9 months ago

A. Enforce 2-factor authentication in GSuite for all users.

upvoted 19 times
- johnsm** Highly Voted

1 year, 4 months ago

Correct answer is A. Well explained here: [https://docs.google.com/document/d/11o3e14tyhnT7w45Q8-r9ZmTAfj2WUNUpJPZImrxm\\_F4/edit?usp=sharing](https://docs.google.com/document/d/11o3e14tyhnT7w45Q8-r9ZmTAfj2WUNUpJPZImrxm_F4/edit?usp=sharing) found some other answers for other questions in this site as well.

upvoted 6 times
- AwesomeGCP** Most Recent

2 months, 2 weeks ago

Selected Answer: A

A is the answer.

upvoted 1 times
- sudarchary**

10 months, 3 weeks ago

Selected Answer: A

<https://support.google.com/a/answer/175197?hl=en>

upvoted 1 times
- Jane111**

1 year, 8 months ago

Shouldn't it be

B. Configure Cloud Identity-Aware Proxy for the App Engine Application.

identity based app access

upvoted 3 times
- desertlotus1211**

1 year, 9 months ago

The key is external user. Best practice is to have internal users/datacenter connect via VPN for security purpose, correct? External users will try to connect via Internet - they still cannot reach the app engine even if they have a users' password because a VPN connection is need to reach the resource. MA will work IF the external user has VPN access... But I think D is what they're looking for based on the question....

upvoted 2 times
- mynk29**

10 months ago

Agree but there is no mention that external user doesnt have internal network access too. A is better option as it covers both scenarios.

upvoted 1 times
- DebasishLowes**

1 year, 9 months ago

Ans : A. When passwords is compromised, enforcing 2 factor authentication is the best way to prevent non authorized users.

upvoted 1 times
- soukumar369**

2 years ago

Enforcing 2-factor authentication can save an employee's password has been compromised

upvoted 1 times
- soukumar369**

2 years ago

Enforce 2-factor authentication safe employee, when an employee's password has been compromised.

upvoted 1 times
- [Removed]**

2 years, 1 month ago

Ans - A

upvoted 1 times

  **subhala** 2 years, 2 months ago

If you limit your GCP VPC to only private access (no resources having external IP), and have VPN. then inspite of having any creds, external folks cannot access the resources.

upvoted 2 times

  **Cloudy\_Apple\_Juice** 2 years, 1 month ago

They can if they login from inside Org - So A is the only correct asnwer

upvoted 1 times

  **soukumar369** 2 years ago

I'm also thinking the same.

upvoted 2 times

  **passtest100** 2 years, 2 months ago

should be B

upvoted 2 times

  **CHECK666** 2 years, 2 months ago

A is the answer. User MFA

upvoted 1 times

  **xhova** 2 years, 8 months ago

A is the answer. VPN does not help if credentials are compromised.

upvoted 2 times

  **smart123** 2 years, 5 months ago

I agree option 'A' is the right answer.

upvoted 1 times

A large financial institution is moving its Big Data analytics to Google Cloud Platform. They want to have maximum control over the encryption process of data stored at rest in BigQuery.

What technique should the institution use?

- A. Use Cloud Storage as a federated Data Source.
- B. Use a Cloud Hardware Security Module (Cloud HSM).
- C. Customer-managed encryption keys (CMEK).
- D. Customer-supplied encryption keys (CSEK).

**Correct Answer:** C

Reference:

<https://cloud.google.com/bigquery/docs/encryption-at-rest>

*Community vote distribution*

C (100%)

🗳️ 👤 **Ganshank** Highly Voted 👍 2 years, 7 months ago

CSEK is only supported in Google Cloud Storage and Compute Engine, therefore D cannot be the right answer.

Ideally, it would be client-side encryption, with BigQuery providing another round of encryption of the encrypted data -

[https://cloud.google.com/bigquery/docs/encryption-at-rest#client\\_side\\_encryption](https://cloud.google.com/bigquery/docs/encryption-at-rest#client_side_encryption), but since that is not one of the options, we can go with C as the next best option.

upvoted 15 times

🗳️ 👤 **smart123** 2 years, 6 months ago

Option 'C' is correct. Option 'D' is not correct as CSEK a feature in Google Cloud Storage and Google Compute Engine only.

upvoted 4 times

🗳️ 👤 **AwesomeGCP** Most Recent ⌚ 2 months, 2 weeks ago

**Selected Answer: C**

C. Customer-managed encryption keys (CMEK).

upvoted 1 times

🗳️ 👤 **DebasishLowes** 1 year, 9 months ago

Ans : C

upvoted 1 times

🗳️ 👤 **Aniyadu** 1 year, 11 months ago

I feel C is the right answer. if customer wants to manage the keys from on-premises then D would be correct.

upvoted 2 times

🗳️ 👤 **[Removed]** 2 years, 1 month ago

Ans - C

upvoted 2 times

🗳️ 👤 **saaurabh1805** 2 years, 1 month ago

C is correct answer as CSEK is not available for big query.

upvoted 2 times

🗳️ 👤 **MohitA** 2 years, 4 months ago

C is the right answer as CSEC is only available for CS and CE's

upvoted 1 times

🗳️ 👤 **aiwaai** 2 years, 4 months ago

Correct Answer: C

upvoted 1 times

🗳️ 👤 **ArizonaClassics** 2 years, 4 months ago

C is the RIGHT ONE!!!

If you want to manage the key encryption keys used for your data at rest, instead of having Google manage the keys, use Cloud Key Management Service to manage your keys. This scenario is known as customer-managed encryption keys (CMEK).

<https://cloud.google.com/bigquery/docs/encryption-at-rest>

upvoted 1 times

🗳️ 👤 **ArizonaClassics** 2 years, 4 months ago

ALSO READ : <https://cloud.google.com/bigquery/docs/customer-managed-encryption>

upvoted 1 times

  **ranjeetpatil** 2 years, 6 months ago

Ans is C. BigQuery does not support CSEK. <https://cloud.google.com/security/encryption-at-rest>. <https://cloud.google.com/security/encryption-at-rest>

upvoted 4 times

  **srinidutt** 2 years, 7 months ago

I also feel D is right

upvoted 1 times

  **xhova** 2 years, 8 months ago

Answer is D. For max control you don't want to store the Key with Google.

upvoted 3 times

  **jonclem** 2 years, 9 months ago

For maximum control surely D is the correct answer.

CSEK:

<https://cloud.google.com/security/encryption-at-rest/customer-supplied-encryption-keys>

CMEK

<https://cloud.google.com/bigquery/docs/encryption-at-rest>

upvoted 2 times



A company is deploying their application on Google Cloud Platform. Company policy requires long-term data to be stored using a solution that can automatically replicate data over at least two geographic places.

Which Storage solution are they allowed to use?


- A. Cloud Bigtable
- B. Cloud BigQuery
- C. Compute Engine SSD Disk
- D. Compute Engine Persistent Disk

**Correct Answer:** B

Reference:

<https://cloud.google.com/bigquery/docs/locations>

Community vote distribution



**ronron89** Highly Voted 2 years ago

<https://cloud.google.com/bigquery#:~:text=BigQuery%20transparently%20and%20automatically%20provides,charge%20and%20no%20additional%20setup.&text=BigQuery%20also%20provides%20ODBC%20and,interact%20with%20its%20powerful%20engine>.

Answer is B.

BigQuery transparently and automatically provides highly durable, replicated storage in multiple locations and high availability with no extra charge and no additional setup.

@xhova: <https://cloud.google.com/bigquery-transfer/docs/locations>

What it mentions here is once you create a replication. YOU cannot change a location. Here the question is about high availability. synchronous replication.

upvoted 15 times

**mistryminded** 1 year ago

Correct answer is B.

BQ: <https://cloud.google.com/bigquery-transfer/docs/locations#multi-regional-locations> and [https://cloud.google.com/bigquery-transfer/docs/locations#colocation\\_required](https://cloud.google.com/bigquery-transfer/docs/locations#colocation_required)

Bigtable: <https://cloud.google.com/bigtable/docs/locations>

PS: To people that are only commenting an answer, please provide a valid source to back your answers. This is a community driven forum and just spamming with wrong answers affects all of us.

upvoted 3 times

**Arad** 1 year, 1 month ago

Correct answer is A.

B is not correct because: "BigQuery does not automatically provide a backup or replica of your data in another geographic region."

<https://cloud.google.com/bigquery/docs/availability>

upvoted 3 times

**mynk29** 10 months ago

"In either case, BigQuery automatically stores copies of your data in two different Google Cloud zones within the selected location."

your link

upvoted 3 times

**AwesomeGCP** Most Recent 2 months, 2 weeks ago

**Selected Answer: B**

B. Cloud BigQuery

upvoted 1 times

**giovy\_82** 3 months, 4 weeks ago

**Selected Answer: B**

I was about to select D, BUT:

- the question says "long term data" -> which makes me think about BQ

- the replication of persistent disk is between different ZONES, but the question says "different geo location" -> which means different regions (if you look at the zone distribution, different zones in same region are located in the same datacenter)

but I still have doubt since the application data are not supposed to be stored in BQ , unless it is for analytics and so on. GCS would have been the best choice, but in absence of this, probably B is the 1st choice.

upvoted 2 times

  **Table2022** 2 months ago

Thank God we have you giovy\_82, very good explanation.

upvoted 2 times

  **piyush\_1982** 4 months, 4 weeks ago

**Selected Answer: A**

[https://cloud.google.com/bigquery/docs/availability#availability\\_and\\_durability](https://cloud.google.com/bigquery/docs/availability#availability_and_durability)

As per the link above BigQuery does not automatically provide a backup or replica of your data in another geographic region. It only stores copies of data in two different Google Cloud zones within the selected location.

Reading through the link <https://cloud.google.com/bigtable/docs/replication-overview>

It states that the Bigtable replicates any changes to your data automatically within a region or multi-region.

upvoted 1 times

  **Medofree** 8 months, 2 weeks ago

**Selected Answer: D**

The answer is D.

Do not forget the context, we are talking here about an Application and its storage, we are not talking about high throughput low latency database (Bigtable) or Analytics Database (BigQuery). On the other way SSDs are physically attached to the VM.

upvoted 1 times

  **sudarchary** 11 months, 1 week ago

**Selected Answer: B**

Multiregion

upvoted 1 times

  **phkaushik** 11 months, 1 week ago

B is the answer, big query stores data multiregional

upvoted 1 times

  **Lancyqusa** 12 months ago

The key here is to set up the resource that will "automatically replicate over two regions". In case of BigQuery, when you set up the resource, you can specify a "location" e.g us or eu, and the data will be replicate over multiple regions. On the other hand, you have to configure the replication and for that reason, is not automatic.

upvoted 1 times

  **Lancyqusa** 12 months ago

Forgot to add - B is the answer :)

upvoted 1 times

  **heftjustice** 1 year ago

It is A Bigtable

See : <https://cloud.google.com/bigtable/docs/replication-overview>

upvoted 1 times

  **mistryminded** 1 year ago

**Selected Answer: B**

Correct answer is B.

BQ: <https://cloud.google.com/bigquery-transfer/docs/locations#multi-regional-locations> and [https://cloud.google.com/bigquery-transfer/docs/locations#colocation\\_required](https://cloud.google.com/bigquery-transfer/docs/locations#colocation_required)

Bigtable: <https://cloud.google.com/bigtable/docs/locations>

PS: To people that are only commenting an answer, please provide a valid source to back your answers. This is a community driven forum and just spamming with wrong answers affects all of us.

upvoted 2 times

  **Alex0303** 1 year, 5 months ago

A - correct answer. IN BigTable you can create one cluster in region-1, other cluster in region-2

upvoted 1 times

  **umashankar\_a** 1 year, 5 months ago

Answer A

Automatic Replication of Data in different regions is not supported by Big Query.

<https://cloud.google.com/bigquery/docs/availability>

It states - "In a single region, data is stored only in the region. There is no Google Cloud–provided backup or replication to another region. If you want to use a single region for your datasets but consider the lack of backup or replication too risky, you can create cross-region dataset copies to



enhance your disaster recovery guarantees."

So this question specifically asked for "Automatic Data Replication" and it's only possible with Big Table. So Answer is A.

upvoted 2 times

  **tzKhalil** 1 year, 7 months ago

Answer is B. Pick BigQuery with multi-region

A multi-region is a large geographic area, such as the United States (US) or Europe (EU), that contains two or more geographic places. In a multi-region, data is stored in a single region but is backed up in a geographically-separated region to provide resilience to a regional disaster. The recovery and failover process is managed by BigQuery.

[https://cloud.google.com/bigquery/docs/availability#multi\\_regions](https://cloud.google.com/bigquery/docs/availability#multi_regions)

upvoted 2 times

  **talktolanka** 1 year, 8 months ago

Answer is B. the term automatic is matching to BigQuery. Also refer to <https://cloud.google.com/bigquery>

upvoted 1 times

  **talktolanka** 1 year, 7 months ago

Very Complicated. automatic replication done by BigTable as well. Also data should be stored for long time. I will go with A

upvoted 1 times

  **desertlotus1211** 1 year, 9 months ago

Answer D: You're forgetting the keywords 'storage solution'

<https://cloud.google.com/persistent-disk>

'For maximum flexibility and minimal effort, snapshots are geo-replicated and available for restore in all regions by default'

upvoted 3 times

  **PATILDXB** 11 hours, 57 minutes ago

but you are forgetting the word long term storage..persistent disk is not intended for long term storage or multi-regional...it is just multi-zonal..bigquery fits best

upvoted 1 times

  **DebasishLowes** 1 year, 9 months ago

Ans : A

upvoted 1 times

  **[Removed]** 2 years, 1 month ago

Ans - A

upvoted 1 times

A large e-retailer is moving to Google Cloud Platform with its ecommerce website. The company wants to ensure payment information is encrypted between the customer's browser and GCP when the customers checkout online.

What should they do?

- A. Configure an SSL Certificate on an L7 Load Balancer and require encryption.
- B. Configure an SSL Certificate on a Network TCP Load Balancer and require encryption.
- C. Configure the firewall to allow inbound traffic on port 443, and block all other inbound traffic.
- D. Configure the firewall to allow outbound traffic on port 443, and block all other outbound traffic.

**Correct Answer:** A

🗲️ 👤 **ESP\_SAP** Highly Voted 👍 2 years ago

Correct Answer is (A):

he type of traffic that you need your load balancer to handle is another factor in determining which load balancer to use:

For HTTP and HTTPS traffic, use:  
External HTTP(S) Load Balancing

[https://cloud.google.com/load-balancing/docs/load-balancing-overview#external\\_versus\\_internal\\_load\\_balancing](https://cloud.google.com/load-balancing/docs/load-balancing-overview#external_versus_internal_load_balancing)  
upvoted 10 times

🗲️ 👤 **[Removed]** Most Recent ⌚ 2 years, 1 month ago

Ans - A

upvoted 1 times

🗲️ 👤 **CHECK666** 2 years, 2 months ago

A is the answer, SSL certificate on L7 layer LoadBlanacer  
upvoted 2 times

🗲️ 👤 **ArizonaClassics** 2 years, 4 months ago

A is the correct one. the question is to see if you understand difference between Layer 7 vs Layer 4 protocols.  
upvoted 1 times

🗲️ 👤 **smart123** 2 years, 5 months ago

Option 'A' is the correct answer.  
upvoted 1 times

🗲️ 👤 **srinidutt** 2 years, 7 months ago

A is right  
upvoted 1 times

Applications often require access to `secrets` - small pieces of sensitive data at build or run time. The administrator managing these secrets on GCP wants to keep a track of `who did what, where, and when?` within their GCP projects.

Which two log streams would provide the information that the administrator is looking for? (Choose two.)

- A. Admin Activity logs
- B. System Event logs
- C. Data Access logs
- D. VPC Flow logs
- E. Agent logs

**Correct Answer:** AC

Reference:

<https://cloud.google.com/kms/docs/secret-management>

*Community vote distribution*

AC (100%)

🗲️ 👤 **Ganshank** Highly Voted 👍 2 years, 7 months ago  
Agreed AC.  
<https://cloud.google.com/secret-manager/docs/audit-logging>  
upvoted 11 times

🗲️ 👤 **AwesomeGCP** Most Recent ⌚ 2 months, 2 weeks ago  
Selected Answer: AC  
A. Admin Activity logs  
C. Data Access logs  
upvoted 1 times

🗲️ 👤 **DebasishLowes** 1 year, 9 months ago  
Ans AC  
upvoted 3 times

🗲️ 👤 **[Removed]** 2 years, 1 month ago  
Ans - AC  
upvoted 2 times

🗲️ 👤 **CHECK666** 2 years, 2 months ago  
AC is the answer.  
Admin Access Logs and Data Access Logs  
upvoted 3 times

🗲️ 👤 **smart123** 2 years, 5 months ago  
Yes 'A & C' are the right answers.  
upvoted 2 times

You are in charge of migrating a legacy application from your company datacenters to GCP before the current maintenance contract expires. You do not know what ports the application is using and no documentation is available for you to check. You want to complete the migration without putting your environment at risk.

What should you do?

- A. Migrate the application into an isolated project using a "Lift & Shift" approach. Enable all internal TCP traffic using VPC Firewall rules. Use VPC Flow logs to determine what traffic should be allowed for the application to work properly.
- B. Migrate the application into an isolated project using a "Lift & Shift" approach in a custom network. Disable all traffic within the VPC and look at the Firewall logs to determine what traffic should be allowed for the application to work properly.
- C. Refactor the application into a micro-services architecture in a GKE cluster. Disable all traffic from outside the cluster using Firewall Rules. Use VPC Flow logs to determine what traffic should be allowed for the application to work properly.
- D. Refactor the application into a micro-services architecture hosted in Cloud Functions in an isolated project. Disable all traffic from outside your project using Firewall Rules. Use VPC Flow logs to determine what traffic should be allowed for the application to work properly.

**Correct Answer: C**

Community vote distribution

A (50%)

B (50%)

🗳️ 👤 **rafaelc** Highly Voted 2 years, 9 months ago

A or B. Leaning towards A

You have a deadline you cannot develop a new app so you have to lift and shift.

upvoted 16 times

🗳️ 👤 **xhova** 2 years, 8 months ago

Answer is A.. You need VPC Flow Logs not "Firewall logs" stated in B

upvoted 11 times

🗳️ 👤 **Table2022** 2 months ago

xhova, you got it right!

upvoted 2 times

🗳️ 👤 **smart123** 2 years, 5 months ago

I agree.

upvoted 1 times

🗳️ 👤 **mynk29** 10 months ago

Agree "Disable all traffic within the VPC and look at the Firewall logs to determine what traffic should be allowed for the application to work properly." if you disable all the VPC traffic there will be nothing to look into firewall logs.

upvoted 2 times

🗳️ 👤 **Mixxer5** Most Recent 3 weeks, 4 days ago

**Selected Answer: A**

A since B disrupts the system. C and D are out of question if it's supposed to "just work".

upvoted 1 times

🗳️ 👤 **Meyucho** 1 month ago

**Selected Answer: B**

The difference between A and B is that, in the first, you allow all traffic so the app will work after migration and you can investigate which ports should be open and then take actions. If you go with B you will have a disruption window until figure out all ports needed but will not have any port unneeded port. So... if you asked to avoid disruption go with A and (as in this question) you are asked about security, go with B

upvoted 1 times

🗳️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: A**

A. Migrate the application into an isolated project using a "Lift & Shift" approach. Enable all internal TCP traffic using VPC Firewall rules. Use VPC Flow logs to determine what traffic should be allowed for the application to work properly.

upvoted 1 times

🗳️ 👤 **GPK** 1 year ago

These questions are no more relevant as google has changed exam and made it really challenging now.

upvoted 1 times

  **vicky\_cyber** 1 year ago

Could you please help us with recent dumps or guide which dump to be referred  
upvoted 1 times

  **Bwitch** 10 months, 4 weeks ago

This one is accurate.  
upvoted 2 times

  **rr4444** 1 year ago

**Selected Answer: B**

B - VPC Flow Logs

Firewall logging only covers TCP and UDP, you explicitly don't know what the app does. That limitation is also important to the fact that implied deny all ingress and deny all egress rules are not covered by Firewall Logging. Plus you have to enable Firewall Logging per rule, so you'd have to have a rule for everything in advance - chicken and egg.... you don't know what is going on, so how could you!?

upvoted 1 times

  **rr4444** 1 year ago

VPC FLOW logs is A!

I meant A!

upvoted 1 times

  **keresh** 1 year, 7 months ago

Answer A matches "without putting your environment at risk" best, all the other answers are higher in risk  
upvoted 3 times

  **DebasishLowes** 1 year, 9 months ago

Ans : A

upvoted 2 times

  **[Removed]** 2 years, 1 month ago

Ans - A

upvoted 2 times

  **CHECK666** 2 years, 2 months ago

A is the answer. You need to take a look at the VPC Flow Logs to know what kind of traffic going in.

upvoted 3 times

  **MohitA** 2 years, 4 months ago

A is the correct Answer

upvoted 2 times

  **aiwaai** 2 years, 4 months ago

Correct Answer: B

upvoted 1 times

  **aiwaai** 2 years, 3 months ago

made correction B -> A

upvoted 1 times

  **srinidutt** 2 years, 7 months ago

A is relevant

upvoted 2 times

Your company has deployed an application on Compute Engine. The application is accessible by clients on port 587. You need to balance the load between the different instances running the application. The connection should be secured using TLS, and terminated by the Load Balancer. What type of Load Balancing should you use?

- A. Network Load Balancing
- B. HTTP(S) Load Balancing
- C. TCP Proxy Load Balancing
- D. SSL Proxy Load Balancing

**Correct Answer:** D

Reference:

<https://cloud.google.com/load-balancing/docs/ssl/>

*Community vote distribution*

D (100%)

🗳️ 👤 **smart123** Highly Voted 2 years, 5 months ago

Although both TCP Proxy LB and SSL Proxy LB support port 587 but only SSL Proxy LB support TLS. Hence 'D' is the right answer.  
upvoted 15 times

🗳️ 👤 **AwesomeGCP** Most Recent 2 months, 2 weeks ago

**Selected Answer: D**

Answer D. SSL Proxy Load Balancing  
<https://cloud.google.com/load-balancing/docs/ssl/>  
upvoted 1 times

🗳️ 👤 **umashankar\_a** 1 year, 5 months ago

Answer D  
<https://cloud.google.com/load-balancing/docs/ssl/>  
- SSL Proxy Load Balancing is a reverse proxy load balancer that distributes SSL traffic coming from the internet to virtual machine (VM) instances in your Google Cloud VPC network.

When using SSL Proxy Load Balancing for your SSL traffic, user SSL (TLS) connections are terminated at the load balancing layer, and then proxied to the closest available backend instances by using either SSL (recommended) or TCP.  
upvoted 4 times

🗳️ 👤 **dtmtor** 1 year, 9 months ago

Answer: D  
upvoted 1 times

🗳️ 👤 **DebasishLowes** 1 year, 10 months ago

Ans : D  
upvoted 1 times

🗳️ 👤 **[Removed]** 2 years, 1 month ago

Ans - D  
upvoted 1 times

🗳️ 👤 **CHECK666** 2 years, 2 months ago

D is the answer. SSL Proxy LoadBalancer supports TLS.  
upvoted 2 times

🗳️ 👤 **mlyu** 2 years, 3 months ago

Agreed with smart123. Ans is D  
[https://cloud.google.com/load-balancing/docs/choosing-load-balancer#flow\\_chart](https://cloud.google.com/load-balancing/docs/choosing-load-balancer#flow_chart)  
upvoted 3 times



You want to limit the images that can be used as the source for boot disks. These images will be stored in a dedicated project. What should you do?

- A. Use the Organization Policy Service to create a compute.trustedimageProjects constraint on the organization level. List the trusted project as the whitelist in an allow operation.
- B. Use the Organization Policy Service to create a compute.trustedimageProjects constraint on the organization level. List the trusted projects as the exceptions in a deny operation.
- C. In Resource Manager, edit the project permissions for the trusted project. Add the organization as member with the role: Compute Image User.
- D. In Resource Manager, edit the organization permissions. Add the project ID as member with the role: Compute Image User.

Correct Answer: B

Reference:

<https://cloud.google.com/compute/docs/images/restricting-image-access>

Community vote distribution



**DebasishLowes** Highly Voted 1 year, 9 months ago  
Ans : A  
upvoted 11 times

**[Removed]** Highly Voted 2 years, 1 month ago  
Ans - A  
[https://cloud.google.com/compute/docs/images/restricting-image-access#trusted\\_images](https://cloud.google.com/compute/docs/images/restricting-image-access#trusted_images)  
upvoted 7 times

**meh009** Most Recent 3 weeks, 1 day ago  
Selected Answer: A  
<https://cloud.google.com/compute/docs/images/restricting-image-access#gcloud>  
  
Look at the glcloud examples and it will make sense why A is correct  
upvoted 1 times

**AzureDP900** 1 month, 3 weeks ago  
A is right  
Use the Trusted image feature to define an organization policy that allows principals to create persistent disks only from images in specific projects.  
upvoted 1 times

**AzureDP900** 1 month, 2 weeks ago  
<https://cloud.google.com/compute/docs/images/restricting-image-access>  
upvoted 1 times

**AwesomeGCP** 2 months, 2 weeks ago  
Selected Answer: A  
Answer A. Use the Organization Policy Service to create a compute.trustedimageProjects constraint on the organization level. List the trusted project as the whitelist in an allow operation.  
upvoted 1 times

**piyush\_1982** 4 months, 4 weeks ago  
To me the answer seems to be B.  
<https://cloud.google.com/compute/docs/images/restricting-image-access>

By default, instances can be created from images in any project that shares images publicly or explicitly with the user. So there is an implicit allow. Option B states that we need to deny all the projects from being used as a trusted project and add "Trusted Project" as an exception to that rule.  
upvoted 2 times

**piyush\_1982** 4 months, 4 weeks ago  
Nope, I think I am getting confused. The correct answer is A.  
upvoted 1 times

**simbu1299** 9 months ago  
Selected Answer: A  
Answer is A

upvoted 2 times

  **danielklein09** 9 months, 1 week ago

Answer is B. You don't whitelist in an allow operation. Since there is an implicit allow, the purpose of the whitelist has been defeated.  
upvoted 1 times

  **CHECK666** 2 years, 2 months ago

A is the answer. you need to allow operations.  
upvoted 1 times

  **ownez** 2 years, 3 months ago

I agree with B.

"<https://cloud.google.com/compute/docs/images/restricting-image-access>"  
upvoted 2 times

  **ownez** 2 years, 3 months ago

Answer is A.

"Use the Trusted image feature to define an organization policy that allows your project members to create persistent disks only from images in specific projects."

"After sharing your images with other users, you can control where those users employ those resources within your organization. Set the constraints/compute.storageResourceUseRestrictions constraint to define the projects where users are permitted to use your storage resources."  
upvoted 4 times

  **Sheeda** 2 years, 3 months ago

Yes, A made sense to me too.  
upvoted 1 times




Your team needs to prevent users from creating projects in the organization. Only the DevOps team should be allowed to create projects on behalf of the requester.

Which two tasks should your team perform to handle this request? (Choose two.)

- A. Remove all users from the Project Creator role at the organizational level.
- B. Create an Organization Policy constraint, and apply it at the organizational level.
- C. Grant the Project Editor role at the organizational level to a designated group of users.
- D. Add a designated group of users to the Project Creator role at the organizational level.
- E. Grant the billing account creator role to the designated DevOps team.

**Correct Answer:** *BD*

Community vote distribution



**mlyu** Highly Voted 2 years, 3 months ago

I think Ans is AD

Because we need to stop the users can create project first (A), and allow devops team to create project (D)  
upvoted 18 times

**[Removed]** Highly Voted 1 year, 9 months ago

AD is the answer.

If constraint is added , no project creation will be allowed, hence B is wrong  
upvoted 7 times

**AzureDP900** Most Recent 1 month, 3 weeks ago

AD is perfect.

A. Remove all users from the Project Creator role at the organizational level.  
D. Add a designated group of users to the Project Creator role at the organizational level.  
upvoted 1 times

**AwesomeGCP** 2 months, 2 weeks ago

Selected Answer: AD

A. Remove all users from the Project Creator role at the organizational level.  
D. Add a designated group of users to the Project Creator role at the organizational level.

<https://cloud.google.com/resource-manager/docs/organization-policy/org-policy-constraints>  
upvoted 2 times

**AzureDP900** 1 month, 2 weeks ago

AD is correct  
upvoted 1 times

**Jeanphi72** 4 months, 2 weeks ago

Selected Answer: AD

<https://cloud.google.com/resource-manager/docs/organization-policy/org-policy-constraints>

I see no way to restrict project creation with an organizational policy. If that would have been possible I would have voted for it as restrictions can be overridden in GCP.  
upvoted 4 times

**piyush\_1982** 4 months, 4 weeks ago

Selected Answer: AC

Seems to be AC

When an organization resource is created, all users in your domain are granted the Billing Account Creator and Project Creator roles by default. As per the link <https://cloud.google.com/resource-manager/docs/default-access-control#removing-default-roles>

Hence A is definitely the answer.

Now to add the project creator we need to add the designated group to the project creator role specifically.  
upvoted 1 times

**absipat** 6 months, 2 weeks ago

ad of course  
upvoted 1 times

**syllox** 1 year, 7 months ago

Ans AC also  
upvoted 1 times

  **syllor** 1 year, 7 months ago  
AD , C is a mistake it's project Editor and not creator  
upvoted 3 times

  **DebasishLowes** 1 year, 10 months ago  
Ans : AD  
upvoted 3 times

  **Aniyadu** 1 year, 11 months ago  
A & D is the right answer.  
upvoted 4 times

  **[Removed]** 2 years, 1 month ago  
Ans - AD  
upvoted 3 times

  **genesis3k** 2 years, 1 month ago  
I think AC. Because, a role is granted to user/group, rather user/group is added to a role.  
upvoted 1 times

  **syllor** 1 year, 7 months ago  
C is a mistake it's project Editor and not creator  
upvoted 1 times

  **CHECK666** 2 years, 2 months ago  
AD is the answer. There's nothing related to project creation in organization policy constraints.  
upvoted 4 times

A customer deployed an application on Compute Engine that takes advantage of the elastic nature of cloud computing.

How can you work with Infrastructure Operations Engineers to best ensure that Windows Compute Engine VMs are up to date with all the latest OS patches?

- A. Build new base images when patches are available, and use a CI/CD pipeline to rebuild VMs, deploying incrementally.
- B. Federate a Domain Controller into Compute Engine, and roll out weekly patches via Group Policy Object.
- C. Use Deployment Manager to provision updated VMs into new serving Instance Groups (IGs).
- D. Reboot all VMs during the weekly maintenance window and allow the StartUp Script to download the latest patches from the internet.

**Correct Answer:** D

Community vote distribution

A (100%)

  **genesis3k**  2 years, 1 month ago

Answer is A.

Compute Engine doesn't automatically update the OS or the software on your deployed instances. You will need to patch or update your deployed Compute Engine instances when necessary. However, in the cloud it is not recommended that you patch or update individual running instances. Instead it is best to patch the image that was used to launch the instance and then replace each affected instance with a new copy.

upvoted 20 times

  **AwesomeGCP**  2 months, 2 weeks ago

**Selected Answer: A**

A. Build new base images when patches are available, and use a CI/CD pipeline to rebuild VMs, deploying incrementally.

upvoted 1 times

  **PATILDXB** 11 hours, 8 minutes ago

you cannot use CI/CD pipeline for building VMs. It is used only for code deployment. Further, building base images is only 1 time activity, organisations cannot afford to change the base image everytime when a patch is released. So, C is the answer

upvoted 1 times

  **Aiffone** 6 months, 2 weeks ago

C is obviously the answer, MIGs help you make sure machines deployed are latest image if you want, what's more, its meant to be an elastic system, nothing does that better than MIGs.

upvoted 1 times

  **Jeanphi72** 4 months, 2 weeks ago

Not sure Deployment Manager can indeed create a new MIG and can configure a new deployment of machines with latest OS but what about the existing ones? In addition how to make sure rollout will be smooth?

Option A seems more realistic.

upvoted 1 times

  **VenkatGCP1** 11 months, 4 weeks ago

The answer is A, we are using this in practice as a solution from Google in one of the top 5 banks for managing windows image patching.

upvoted 4 times

  **AzureDP900** 1 month, 3 weeks ago

Agreed.

upvoted 1 times

  **lxs** 1 year ago

**Selected Answer: A**

Definitely it will be A. The solution must take the advantage of elasticity of compute engine, so you create a template with patched OS base and redeploy images.

upvoted 2 times

  **sc\_cloud\_learn** 1 year, 5 months ago

Answer should be A,

C talks about MIG which may not be always needed

upvoted 1 times

  **DebasishLowes** 1 year, 9 months ago

Ans : A

upvoted 2 times

  **gu9singg** 1 year, 9 months ago

this questions still valid for exam?

upvoted 1 times

  **umashankar\_a** 1 year, 5 months ago

yeah....even i'm thinking the same, as we got OS Patch Management Service now in GCP for Patching Compute machines as per requirement.  
<https://cloud.google.com/compute/docs/os-patch-management>.

Not really sure on the answer.

upvoted 3 times

  **DuncanTu** 1 year, 5 months ago

Hi

May I know why C is incorrect?

upvoted 1 times

  **HateMicrosoft** 1 year, 9 months ago

The correct anwser is C.

<https://cloud.google.com/deployment-manager/docs/reference/latest/deployments/patch>

upvoted 1 times

  **CloudTrip** 1 year, 10 months ago

Given the options here Answer D seems practical

upvoted 1 times

  **singhjoga** 1 year, 11 months ago

B seems the only possible answer. Windows patches are configured using Group Policies on the Windows Domain Controller. All other windows machines should be part of the same domain.

upvoted 1 times

  **FatCharlie** 2 years ago

The answer is A. This is referring to VMs in an instance group which has built in roll out deployment of new images that can easily be integrated into a CI/CD pipeline.

The people mentioning the patch management tool are considering these to be long running VMs, but that makes little sense in an instance group.

upvoted 3 times

  **[Removed]** 2 years, 1 month ago

Ans - A

upvoted 3 times

  **Wooky** 2 years, 2 months ago

I think best ans is C

upvoted 4 times

  **mlyu** 2 years, 3 months ago

None of the answer makes sense.

GCP has provide os patch management tool to patch Windows OS

<https://cloud.google.com/compute/docs/os-patch-management>

upvoted 2 times

  **Sheeda** 2 years, 3 months ago

Next best is D ?

upvoted 1 times

Your team needs to make sure that their backend database can only be accessed by the frontend application and no other instances on the network.

How should your team design this network?

- A. Create an ingress firewall rule to allow access only from the application to the database using firewall tags.
- B. Create a different subnet for the frontend application and database to ensure network isolation.
- C. Create two VPC networks, and connect the two networks using Cloud VPN gateways to ensure network isolation.
- D. Create two VPC networks, and connect the two networks using VPC peering to ensure network isolation.

**Correct Answer: A**

*Community vote distribution*

A (100%)

🗳️ 👤 **AzureDP900** 1 month, 3 weeks ago

A is correct , rest of the answers doesn't make any sence  
upvoted 1 times

🗳️ 👤 **azureaspirant** 1 month, 1 week ago

@AzureDP900: Cleared AWS Solution Architect Professional (SAP - CO1) on the last date. followed your answers. Cleared 5 GCP Certificates. Glad that you are here.  
upvoted 1 times

🗳️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: A**

A. Create an ingress firewall rule to allow access only from the application to the database using firewall tags.  
upvoted 1 times

🗳️ 👤 **zqwiklabs** 1 year, 8 months ago

A is definitely incorrect  
upvoted 4 times

🗳️ 👤 **mistryminded** 1 year ago

This one is confusing but cannot be A because it says 'Firewall tags'. There is no such thing as firewall tags, only 'Network tags'.  
upvoted 2 times

🗳️ 👤 **desertlotus1211** 1 year, 9 months ago

Answer is D: you'd want the DB in a separate VPC. Allow vpc peering and connect the Front End's backend to the DB. Don't get confused by the question saying 'front end' Front end only means public facing...  
upvoted 1 times

🗳️ 👤 **AzureDP900** 1 month, 2 weeks ago

A is correct  
upvoted 1 times

🗳️ 👤 **Jane111** 1 year, 8 months ago

you need to read basic concepts again  
upvoted 7 times

🗳️ 👤 **DebasishLowes** 1 year, 10 months ago

Ans : A  
upvoted 3 times

🗳️ 👤 **singhjoga** 1 year, 11 months ago

Although A is correct, but B would be more secure when combined with firewall rules to restrict traffic based on subnets. Ideal solution would be to use Service Account based firewall rules instead of tag based. See the below paragraph from <https://cloud.google.com/solutions/best-practices-vpc-design>

"However, even though it is possible to uses tags for target filtering in this manner, we recommend that you use service accounts where possible. Target tags are not access-controlled and can be changed by someone with the instanceAdmin role while VMs are in service. Service accounts are access-controlled, meaning that a specific user must be explicitly authorized to use a service account. There can only be one service account per instance, whereas there can be multiple tags. Also, service accounts assigned to a VM can only be changed when the VM is stopped"  
upvoted 4 times

🗳️ 👤 **ThisisJohn** 1 year ago

You may be right but B doesn't mention anything about firewall rules, thus we need to assume there will be communication between both subnets

upvoted 2 times

  **Aiffone** 6 months, 2 weeks ago

I'm inclined to go with A too because without firewall rules the subnets in B would ensure there is no communication at all due to default implicit rules.

upvoted 1 times

  **[Removed]** 2 years, 1 month ago

Ans - A

upvoted 2 times

  **CHECK666** 2 years, 2 months ago

A is the answer, use network tags.

upvoted 4 times

  **mlyu** 2 years, 3 months ago

Agree with A

upvoted 2 times



An organization receives an increasing number of phishing emails.  
Which method should be used to protect employee credentials in this situation?

- A. Multifactor Authentication
- B. A strict password policy
- C. Captcha on login pages
- D. Encrypted emails

**Correct Answer:** D

*Community vote distribution*

A (64%)

D (36%)

- DebasishLowes**

Highly Voted

1 year, 10 months ago

A is the answer.  
upvoted 9 times
- Mixer5**

Most Recent

3 weeks, 4 days ago

**Selected Answer: D**

MFA itself doesn't really protect user's credentials from beaing leaked. It makes it harder (or nigh impossible) to log in even if they get leaked but they may still leak. Encrypting emails would be of more help, although in case of phishing email it'd be best to educate users and add some filters that will flag external emails as suspicious.  
upvoted 1 times
- AwesomeGCP**

2 months, 2 weeks ago

**Selected Answer: A**

A. Multifactor Authentication  
upvoted 2 times
- GHOST1985**

2 months, 3 weeks ago

**Selected Answer: A**

<https://cloud.google.com/blog/products/g-suite/protecting-you-against-phishing>  
upvoted 2 times
- AzureDP900**

1 month, 3 weeks ago

Agree with A  
upvoted 1 times
- Deepanshd**

2 months, 4 weeks ago

**Selected Answer: A**

Multi-factor authentication will prevent employee credentials  
upvoted 2 times
- fanilgor**

3 months, 1 week ago

**Selected Answer: A**

A for sure  
upvoted 1 times
- lxs**

1 year ago

**Selected Answer: D**

This question has been taken from the GCP book.  
upvoted 3 times
- mondigo**

2 years ago

A  
<https://cloud.google.com/blog/products/g-suite/7-ways-admins-can-help-secure-accounts-against-phishing-g-suite>  
upvoted 3 times
- ronron89**

2 years ago

<https://www.duocircle.com/content/email-security-services/email-security-in-cryptography#:~:text=Customer%20Login-,Email%20Security%20In%20Cryptography%20Is%20One%20Of%20The%20Most,Measures%20To%20Prevent%20Phishing%20Attempts&text=Cybercriminals%20love%20emails%20the%20most,networks%20all%20over%20the%20world.>  
  
The answer should be D.

upvoted 2 times

  **shk2011** 2 years, 1 month ago

Logically if i think even if i have not read about cloud answer is A

upvoted 3 times

  **[Removed]** 2 years, 1 month ago

Ans - A

upvoted 2 times

  **CHECK666** 2 years, 2 months ago

The answer is A.

<https://cloud.google.com/blog/products/identity-security/protect-users-in-your-apps-with-multi-factor-authentication>

upvoted 3 times

  **Sheeda** 2 years, 3 months ago

Should be A

upvoted 3 times



A customer is collaborating with another company to build an application on Compute Engine. The customer is building the application tier in their GCP

Organization, and the other company is building the storage tier in a different GCP Organization. This is a 3-tier web application. Communication between portions of the application must not traverse the public internet by any means.

Which connectivity option should be implemented?

- A. VPC peering
- B. Cloud VPN
- C. Cloud Interconnect
- D. Shared VPC

**Correct Answer:** B

Community vote distribution

A (100%)

  **sc\_cloud\_learn** Highly Voted 1 year, 5 months ago  
both are GCP, should be VPC peering- Option A  
upvoted 13 times

  **shayke** Most Recent 2 months, 2 weeks ago  
Selected Answer: A  
only a  
upvoted 1 times

  **AwesomeGCP** 2 months, 2 weeks ago  
Selected Answer: A  
A – Peering two VPCs does permit traffic to flow between the two shared networks, but it's only bi-directional. Peered VPC networks remain administratively separate.

Dedicated Interconnect connections enable you to connect your on-premises network ... in another project, as long as they are both in the same organization. hence A  
upvoted 1 times

  **AzureDP900** 1 month, 3 weeks ago  
Agreed, A is correct.  
upvoted 1 times

  **DP\_GCP** 1 year, 7 months ago  
B is not correct because if Cloud VPN is used data travels over internet and question mentions it doesnt want the data to travel through internet. <https://cloud.google.com/network-connectivity/docs/vpn/concepts/overview> Cloud VPN securely connects your peer network to your Virtual Private Cloud (VPC) network through an IPsec VPN connection. Traffic traveling between the two networks is encrypted by one VPN gateway and then decrypted by the other VPN gateway. This action protects your data as it travels over the internet  
upvoted 1 times

  **PATILDXB** 10 hours, 58 minutes ago  
Cloud VPN is a private connection, and different from normal IP VPN or IPSecVPN. Cloud VPN does not ride on internet. B is correct and appropriate, as it is cheaper than VPC peering, because VPC peering incurs charges  
upvoted 1 times

  **dtmtor** 1 year, 9 months ago  
A, different orgs  
upvoted 4 times

  **DebasishLowes** 1 year, 10 months ago  
A is the answer.  
upvoted 2 times

  **[Removed]** 2 years, 1 month ago  
Ans - A  
upvoted 3 times

  **CHECK666** 2 years, 2 months ago  
A is the ansswer. use VCP Peering.  
upvoted 3 times

  **Akku1614** 2 years, 3 months ago

Yes it Should be VPC Peering. <https://cloud.google.com/vpc/docs/vpc-peering>  
upvoted 3 times

  **Sheeda** 2 years, 3 months ago

Should be A  
upvoted 4 times

Your team wants to make sure Compute Engine instances running in your production project do not have public IP addresses. The frontend application Compute

Engine instances will require public IPs. The product engineers have the Editor role to modify resources. Your team wants to enforce this requirement.

How should your team meet these requirements?

- A. Enable Private Access on the VPC network in the production project.
- B. Remove the Editor role and grant the Compute Admin IAM role to the engineers.
- C. Set up an organization policy to only permit public IPs for the front-end Compute Engine instances.
- D. Set up a VPC network with two subnets: one with public IPs and one without public IPs.

**Correct Answer:** C

Reference:

<https://cloud.google.com/compute/docs/ip-addresses/reserve-static-external-ip-address>

*Community vote distribution*

C (100%)

🗳️ 👤 **saurabh1805** **Highly Voted** 👍 2 years, 1 month ago

C is correct option here, Refer below link for more details.

<https://cloud.google.com/resource-manager/docs/organization-policy/org-policy-constraints#constraints-for-specific-services>  
upvoted 10 times

🗳️ 👤 **AzureDP900** 1 month, 2 weeks ago

Yes, C is right  
upvoted 1 times

🗳️ 👤 **FatCharlie** 2 years ago

More specifically, it's the "Restrict VM IP Forwarding" constraint under Compute Engine  
upvoted 2 times

🗳️ 👤 **FatCharlie** 2 years ago

Sorry, no. It's the one under that :)

"Define allowed external IPs for VM instances"  
upvoted 2 times

🗳️ 👤 **passex** **Most Recent** 🕒 2 weeks, 3 days ago

and how would you want to separate front-end VM's from the other using Org Policy Constraints - IMO option D make more sense  
upvoted 1 times

🗳️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: C**

C. Set up an organization policy to only permit public IPs for the front-end Compute Engine instances.  
upvoted 2 times

🗳️ 👤 **bartlomiejwaw** 7 months, 1 week ago

Not C - Editor role is not enough for setting up org policies  
upvoted 1 times

🗳️ 👤 **DebasishLowes** 1 year, 9 months ago

Ans : C  
upvoted 3 times

🗳️ 👤 **[Removed]** 2 years, 1 month ago

Ans - C  
upvoted 4 times

🗳️ 👤 **HectorLeon2099** 2 years, 2 months ago

I'll go with A  
upvoted 2 times

Which two security characteristics are related to the use of VPC peering to connect two VPC networks? (Choose two.)

- A. Central management of routes, firewalls, and VPNs for peered networks
- B. Non-transitive peered networks; where only directly peered networks can communicate
- C. Ability to peer networks that belong to different Google Cloud organizations
- D. Firewall rules that can be created with a tag from one peered network to another peered network
- E. Ability to share specific subnets across peered networks

Correct Answer: BC

Community vote distribution



**DebasishLowes** Highly Voted 1 year, 9 months ago

Ans : BC  
upvoted 14 times

**Meyucho** Most Recent 3 days, 13 hours ago

Selected Answer: CE  
Although B is correct, going into detail I think that non-transitivity is just true for networks joined by peering but If there is a third network connected by VPN or Interconnect there is transitivity, so I discard B and stay with C and E  
upvoted 1 times

**AzureDP900** 1 month, 2 weeks ago

BC is right  
upvoted 1 times

**AwesomeGCP** 2 months, 2 weeks ago

Selected Answer: BC  
B. Non-transitive peered networks; where only directly peered networks can communicate  
C. Ability to peer networks that belong to different Google Cloud Platform organizations  
upvoted 2 times

**zellick** 2 months, 3 weeks ago

Selected Answer: BC  
BC is the answer.  
upvoted 1 times

**Medofree** 8 months, 2 weeks ago

D is false because : "You cannot use a tag or service account from one peered network in the other peered network."  
upvoted 1 times

**dtmtor** 1 year, 9 months ago

Answer is BC  
upvoted 2 times

**Aniyadu** 1 year, 11 months ago

B&C is the right answer  
upvoted 2 times

**FatCharlie** 2 years ago

The answers marked in the question seem to be referring to \_shared\_ VPC capabilities.  
upvoted 1 times

**[Removed]** 2 years, 1 month ago

Ans - BC  
upvoted 2 times

**CHECK666** 2 years, 2 months ago

BC is the answer.  
upvoted 2 times

**cipher90** 2 years, 3 months ago

AD is correct "Security Characteristics"

upvoted 1 times

  **mte\_tech34** 2 years, 2 months ago

No it's not. "You cannot use a tag or service account from one peered network in the other peered network." ->  
<https://cloud.google.com/vpc/docs/vpc-peering>

upvoted 2 times

  **mlyu** 2 years, 3 months ago

Ans should be BC

[https://cloud.google.com/vpc/docs/vpc-peering#key\\_properties](https://cloud.google.com/vpc/docs/vpc-peering#key_properties)

upvoted 4 times

  **ownez** 2 years, 3 months ago

Correct.

B: "Only directly peered networks can communicate. Transitive peering is not supported."

C: " You can make services available privately across different VPC networks within and across organizations."

upvoted 3 times

  **Mihai89** 2 years, 1 month ago

Agree with BC

upvoted 1 times

  **MohitA** 2 years, 3 months ago

agree BC

upvoted 1 times

A patch for a vulnerability has been released, and a DevOps team needs to update their running containers in Google Kubernetes Engine (GKE). How should the DevOps team accomplish this?

- A. Use Puppet or Chef to push out the patch to the running container.
- B. Verify that auto upgrade is enabled; if so, Google will upgrade the nodes in a GKE cluster.
- C. Update the application code or apply a patch, build a new image, and redeploy it.
- D. Configure containers to automatically upgrade when the base image is available in Container Registry.

**Correct Answer: C**

Reference:

<https://cloud.google.com/kubernetes-engine/docs/security-bulletins>

*Community vote distribution*

C (100%)

🗲️ 👤 **TNT87** Highly Voted 1 year, 10 months ago

<https://cloud.google.com/containers/security>

Containers are meant to be immutable, so you deploy a new image in order to make changes. You can simplify patch management by rebuilding your images regularly, so the patch is picked up the next time a container is deployed. Get the full picture of your environment with regular image security reviews.

C is better

upvoted 12 times

🗲️ 👤 **AzureDP900** 1 month, 2 weeks ago

Yes, C is correct

upvoted 1 times

🗲️ 👤 **DebasishLowes** Highly Voted 1 year, 9 months ago

Ans : C

upvoted 6 times

🗲️ 👤 **AwesomeGCP** Most Recent 2 months, 2 weeks ago

**Selected Answer: C**

C. Update the application code or apply a patch, build a new image, and redeploy it.

upvoted 1 times

🗲️ 👤 **Medofree** 8 months, 2 weeks ago

**Selected Answer: C**

Correct ans is C, because "DevOps team needs to update their running containers".

upvoted 2 times

🗲️ 👤 **Rhehehe** 1 year ago

Its actually B.

Patching a vulnerability involves upgrading to a new GKE or Anthos version number. GKE and Anthos versions include versioned components for the operating system, Kubernetes components, and other containers that make up the Anthos platform. Fixing some vulnerabilities requires only a control plane upgrade, performed automatically by Google on GKE, while others require both control plane and node upgrades.

To keep clusters patched and hardened against vulnerabilities of all severities, we recommend using node auto-upgrade on GKE (on by default). On other Anthos platforms, Google recommends upgrading your Anthos components at least monthly.

Ref: <https://cloud.google.com/kubernetes-engine/docs/resources/security-patching>

upvoted 2 times

🗲️ 👤 **StanPeng** 10 months, 2 weeks ago

The question is asking about upgrading application code rather than GKE

upvoted 1 times

🗲️ 👤 **alexm112** 10 months, 2 weeks ago

Agreed - I think this wasn't available at the time people responded.

B is correct

<https://cloud.google.com/kubernetes-engine/docs/how-to/node-auto-upgrades>

upvoted 1 times

🗲️ 👤 **SuperDevops** 1 year, 1 month ago



I took the test yesterday and didn't pass, NO ISSUE is from here. The questions are totally new  
Whizlabs it's OK  
upvoted 1 times

  **sriz** 1 year, 1 month ago  
u got questions from Whizlabs?  
upvoted 2 times

  **Aniyadu** 1 year, 11 months ago  
The question asked is "team needs to update their running containers" if its was auto enabled there was no need to update manually. so my answer will be C.  
upvoted 2 times

  **Kevinsayn** 2 years, 1 month ago  
Me voy definitivamente con la C, dado que actualizar los nodos con autoupgrade no tiene nada que ver con los contenedores, la vulnerabilidad en este caso se debe aplicar con respecto a contenedor ósea aplicación por lo que la respuesta C es la correcta.  
upvoted 3 times

  **soukumar369** 2 years ago  
Translaed : 'm definitely going with C, since updating the nodes with autoupgrade has nothing to do with the containers, the vulnerability in this case must be applied with respect to the application bone container so the C answer is correct.  
upvoted 1 times

  **jonclem** 2 years, 1 month ago  
Answer B is correct as per the Video Google Kubernetes Engine (GKE) Security on Linuxacademy.  
upvoted 2 times

  **[Removed]** 2 years, 1 month ago  
Ans - C  
upvoted 3 times

  **Rantu** 2 years, 2 months ago  
C is the correct answer as this is the way to patch, build, re-deploy  
upvoted 3 times

  **Namaste** 2 years, 3 months ago  
Answer is C.  
upvoted 3 times

  **ownez** 2 years, 3 months ago  
I would go for C because some reported CVEs will take time to be published and approval in CVE advisory portal. Once approved, it will notify to all necessary third party.  
  
Hence, this requires a lot of time and left people exposed to the vulnerability.  
  
Answer is C.  
upvoted 3 times

  **Raushanr** 2 years, 3 months ago  
Answer should be B.. Auto upgrade  
upvoted 2 times

  **FatCharlie** 2 years ago  
Auto upgrade of the Nodes does not upgrade the container code which needs the patch  
upvoted 3 times

  **MohitA** 2 years, 3 months ago  
Answer is B  
upvoted 1 times

  **MohitA** 2 years, 3 months ago  
Sorry It should be C  
upvoted 5 times

A company is running their webshop on Google Kubernetes Engine and wants to analyze customer transactions in BigQuery. You need to ensure that no credit card numbers are stored in BigQuery. What should you do?

- A. Create a BigQuery view with regular expressions matching credit card numbers to query and delete affected rows.
- B. Use the Cloud Data Loss Prevention API to redact related infoTypes before data is ingested into BigQuery.
- C. Leverage Security Command Center to scan for the assets of type Credit Card Number in BigQuery.
- D. Enable Cloud Identity-Aware Proxy to filter out credit card numbers before storing the logs in BigQuery.

**Correct Answer:** D

Community vote distribution

B (100%)

🗲️ 👤 **saurobh1805** Highly Voted 2 years, 1 month ago

B is correct answer here.  
upvoted 9 times

🗲️ 👤 **saurobh1805** 2 years, 1 month ago

<https://cloud.google.com/bigquery/docs/scan-with-dlp>  
upvoted 2 times

🗲️ 👤 **jaykumarjkd99** Most Recent 2 days, 9 hours ago

Selected Answer: B

B is correct answer here.  
.  
upvoted 1 times

🗲️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

Selected Answer: B

B. Use the Cloud Data Loss Prevention API to redact related infoTypes before data is ingested into BigQuery.  
upvoted 2 times

🗲️ 👤 **giovy\_82** 3 months, 3 weeks ago

Selected Answer: B

How can it be D? i'll go for B, DLP is the tool to scan and find sensible data  
upvoted 1 times

🗲️ 👤 **sudarchary** 10 months, 3 weeks ago

<https://cloud.google.com/bigquery/docs/scan-with-dlp>  
upvoted 1 times

🗲️ 👤 **sudarchary** 10 months, 3 weeks ago

Selected Answer: B

Cloud Data Loss Prevention API allows to detect and redact or remove sensitive data before the comments or reviews are published. Cloud DLP will read information from BigQuery, Cloud Storage or Datastore and scan it for sensitive data.  
upvoted 1 times

🗲️ 👤 **AzureDP900** 1 month, 2 weeks ago

B is correct  
upvoted 1 times

🗲️ 👤 **rr4444** 11 months, 3 weeks ago

Selected Answer: B

D is silly  
upvoted 1 times

🗲️ 👤 **[Removed]** 1 year, 8 months ago

D is impossible. I support B  
upvoted 2 times

🗲️ 👤 **DebasishLowes** 1 year, 9 months ago

Ans : B



upvoted 1 times

  **[Removed]** 2 years, 1 month ago

Ans - B

upvoted 3 times

  **passtest100** 2 years, 2 months ago

B is the answer

upvoted 2 times

  **MohitA** 2 years, 3 months ago

B is the answer

upvoted 2 times

A customer wants to deploy a large number of 3-tier web applications on Compute Engine.

How should the customer ensure authenticated network separation between the different tiers of the application?

- A. Run each tier in its own Project, and segregate using Project labels.
- B. Run each tier with a different Service Account (SA), and use SA-based firewall rules.
- C. Run each tier in its own subnet, and use subnet-based firewall rules.
- D. Run each tier with its own VM tags, and use tag-based firewall rules.

**Correct Answer: C**

Community vote distribution

B (100%)

  **genesis3k** Highly Voted 2 years, 1 month ago

Answer is B. Keyword is 'authenticated'. Reference below:

"Isolate VMs using service accounts when possible"

"even though it is possible to use tags for target filtering in this manner, we recommend that you use service accounts where possible. Target tags are not access-controlled and can be changed by someone with the instanceAdmin role while VMs are in service. Service accounts are access-controlled, meaning that a specific user must be explicitly authorized to use a service account. There can only be one service account per instance, whereas there can be multiple tags. Also, service accounts assigned to a VM can only be changed when the VM is stopped."

<https://cloud.google.com/solutions/best-practices-vpc-design#isolate-vms-service-accounts>

upvoted 20 times

  **AzureDP900** 1 month, 2 weeks ago

Agreed with you and B is right

upvoted 1 times

  **gu9singg** 1 year, 9 months ago

document says about subnet isolation

upvoted 2 times

  **AwesomeGCP** Most Recent 2 months, 2 weeks ago

**Selected Answer: B**

B. Run each tier with a different Service Account (SA), and use SA-based firewall rules.

upvoted 1 times

  **mynk29** 10 months ago

"As previously mentioned, you can identify the VMs on a specific subnet by applying a unique network tag or service account to those instances. This allows you to create firewall rules that only apply to the VMs in a subnet—those with the associated network tag or service account. For example, to create a firewall rule that permits all communication between VMs in the same subnet, you can use the following rule configuration on the Firewall rules page:"

B is the right answer

upvoted 2 times

  **mistryminded** 1 year ago

**Selected Answer: B**

Answer is B - <https://cloud.google.com/vpc/docs/firewalls#service-accounts-vs-tags>

upvoted 2 times

  **gu9singg** 1 year, 9 months ago

C: is incorrect, we need to authenticate, network rules does not apply and not a recommend best practice from google

upvoted 2 times

  **gu9singg** 1 year, 9 months ago

C: is incorrect because we need to spend lot of time designing the network topology etc, google recommended practice is to use simple network design with automation in mind, so service account provides those, hence final decision goes to B

upvoted 2 times

  **gu9singg** 1 year, 9 months ago

Correct answer is B

upvoted 2 times

  **DebasishLowes** 1 year, 9 months ago

Ans : C

upvoted 2 times

🗨️ 👤 **singhjoga** 1 year, 11 months ago

B as per best practices <https://cloud.google.com/solutions/best-practices-vpc-design>  
upvoted 3 times

🗨️ 👤 **Fellipo** 2 years, 1 month ago

B exists?  
upvoted 1 times

🗨️ 👤 **[Removed]** 2 years, 1 month ago

Ans - C  
[https://cloud.google.com/docs/enterprise/best-practices-for-enterprise-organizations#networking\\_and\\_security](https://cloud.google.com/docs/enterprise/best-practices-for-enterprise-organizations#networking_and_security)  
  
[https://cloud.google.com/solutions/best-practices-vpc-design#addresses\\_and\\_subnets](https://cloud.google.com/solutions/best-practices-vpc-design#addresses_and_subnets)  
upvoted 2 times

🗨️ 👤 **Rantu** 2 years, 2 months ago

Authenticated separation is the key here. Ideally it should be tag based firewall rule separation. However authenticated word creates confusion. My best judgement is B  
upvoted 2 times

🗨️ 👤 **CHECK666** 2 years, 2 months ago

C is the answer.  
upvoted 2 times

🗨️ 👤 **Wooky** 2 years, 2 months ago

should be C  
upvoted 2 times

🗨️ 👤 **MohitA** 2 years, 3 months ago

A should be correct  
upvoted 1 times

🗨️ 👤 **MohitA** 2 years, 3 months ago

Sorry B seems to be more appropriate  
upvoted 4 times

A manager wants to start retaining security event logs for 2 years while minimizing costs. You write a filter to select the appropriate log entries. Where should you export the logs?


- A. BigQuery datasets
- B. Cloud Storage buckets
- C. StackDriver logging
- D. Cloud Pub/Sub topics

Correct Answer: C

Reference:

<https://cloud.google.com/logging/docs/exclusions>

Community vote distribution



- shayke

1 day, 19 hours ago

Selected Answer: B

B- is the cheapest optaion

upvoted 1 times
- AzureDP900

1 month, 2 weeks ago

B is best for cost optimization perspective

upvoted 2 times
- shayke

2 months ago

Selected Answer: B

GCS would be the chipest option

upvoted 2 times
- AwesomeGCP

2 months, 2 weeks ago

Selected Answer: B

B. Cloud Storage buckets

upvoted 1 times
- Deepanshd

2 months, 3 weeks ago

Selected Answer: B

Cloud storage is always considered when minimize cost

upvoted 1 times
- Bill1000

2 months, 3 weeks ago

B is correct

upvoted 2 times
- mbiy

10 months ago

Ans C is correct, you can define a custom log bucket and mention the retention policy for any number of years (range - 1 day to 3650 days). Underlying these custom define log bucket is also created within Cloud Storage. As per the question you can retain log for 2 years in Stackdriver Logging which is aka Cloud Logging, and then later archive to cold line storage if there is a requirement.

upvoted 1 times
- VJ\_0909

10 months ago

Default retention for logging is 30 days because it is expensive to hold the logs there for longer duration. Bucket is always the cheapest option.

upvoted 1 times
- jayk22

1 year, 1 month ago

Ans B. Validated.

upvoted 4 times
- DebasishLowes

1 year, 9 months ago

Ans: B

upvoted 4 times
- [Removed]

2 years, 1 month ago

Ans - B

upvoted 1 times

  **Raushanr** 2 years, 3 months ago

Ans is B

upvoted 1 times

  **mlyu** 2 years, 3 months ago

Ans B

Cloud storage is always considered when minimize cost

upvoted 2 times

  **MohitA** 2 years, 3 months ago

Agree B


upvoted 1 times

For compliance reasons, an organization needs to ensure that in-scope PCI Kubernetes Pods reside on `in-scope` Nodes only. These Nodes can only contain the `in-scope` Pods.  
How should the organization achieve this objective?

- A. Add a nodeSelector field to the pod configuration to only use the Nodes labeled inscope: true.
- B. Create a node pool with the label inscope: true and a Pod Security Policy that only allows the Pods to run on Nodes with that label.
- C. Place a taint on the Nodes with the label inscope: true and effect NoSchedule and a toleration to match in the Pod configuration.
- D. Run all in-scope Pods in the namespace `in-scope-pci`.

Correct Answer: C

Community vote distribution



- Tabayashi**

Highly Voted

7 months, 4 weeks ago

[A] Correct answer. This is a typical use case for node selector.  
<https://kubernetes.io/docs/concepts/scheduling-eviction/assign-pod-node/#nodeselector>

[B] The Pod Security Policy is designed to block the creation of misconfigured pods on certain clusters. This does not meet the requirements.

[C] Taint will no longer place pods without the "inscope" label on that node, but it does not guarantee that pods with the "inscope" label will be placed on that node.

[D] Placing the "in scope" node in the namespace "in-scope-pci" may meet the requirement, but [A] takes precedence.

upvoted 6 times
- AzureDP900**

1 month, 2 weeks ago

A is correct.

upvoted 1 times
- MariaGabiGabriela**

6 months, 3 weeks ago

I think [A] does not stop other pods from being run in the PCI node, which is a requirement as the question states... I would go with [C]

upvoted 3 times
- Meyucho**

Most Recent

3 days, 15 hours ago

Selected Answer: C

A nodeselector configuration is from a pod template perspective. This question ask to PRESERVE some nodes for specific pods, so this is the main utilization for Taint. This is a conceptual question and the answer is C

upvoted 1 times
- AwesomeGCP**

2 months, 2 weeks ago

Selected Answer: A

A. Add a nodeSelector field to the pod configuration to only use the Nodes labeled inscope: true.

upvoted 1 times
- GHOST1985**

2 months, 3 weeks ago

Selected Answer: A

nodeSelector is the simplest recommended form of node selection constraint. You can add the nodeSelector field to your Pod specification and specify the node labels you want the target node to have. Kubernetes only schedules the Pod onto nodes that have each of the labels you specify.  
=> <https://kubernetes.io/docs/concepts/scheduling-eviction/assign-pod-node/#nodeselector>

Tolerations are applied to pods. Tolerations allow the scheduler to schedule pods with matching taints. Tolerations allow scheduling but don't guarantee scheduling; the scheduler also evaluates other parameters as part of its function. => <https://kubernetes.io/docs/concepts/scheduling-eviction/taint-and-toleration/>

upvoted 2 times
- fanilgor**

3 months, 1 week ago

Selected Answer: C

Basic K8s principles of scheduling workloads.  
Taints and tolerations make perfect sense for this use case. Therefore C.

upvoted 1 times
- Jeanphi72**

4 months ago

Selected Answer: A



<https://redhat-scholars.github.io/kubernetes-tutorial/kubernetes-tutorial/taints-affinity.html>

A Taint is applied to a Kubernetes Node that signals the scheduler to avoid or not schedule certain Pods.

A Toleration is applied to a Pod definition and provides an exception to the taint.

<https://kubernetes.io/docs/concepts/scheduling-eviction/taint-and-toleration/>

Node affinity is a property of Pods that attracts them to a set of nodes (either as a preference or a **hard requirement**).

Taints are the opposite -- they allow a node to repel a set of pods.

upvoted 3 times

  **hybridpro** 6 months, 1 week ago

Answer should be C. "These Nodes can only contain the

λ€in-scopeλ€ Pods." - this can only be achieved by taints and tolerations.

upvoted 1 times

In an effort for your company messaging app to comply with FIPS 140-2, a decision was made to use GCP compute and network services. The messaging app architecture includes a Managed Instance Group (MIG) that controls a cluster of Compute Engine instances. The instances use Local SSDs for data caching and UDP for instance-to-instance communications. The app development team is willing to make any changes necessary to comply with the standard. Which options should you recommend to meet the requirements?

- A. Encrypt all cache storage and VM-to-VM communication using the BoringCrypto module.
- B. Set Disk Encryption on the Instance Template used by the MIG to customer-managed key and use BoringSSL for all data transit between instances.
- C. Change the app instance-to-instance communications from UDP to TCP and enable BoringSSL on clients' TLS connections.
- D. Set Disk Encryption on the Instance Template used by the MIG to Google-managed Key and use BoringSSL library on all instance-to-instance communications.

**Correct Answer: D**

Community vote distribution

A (100%)

  **subhala** Highly Voted 2 years ago

when I revisited this, Now I think A is correct. In A - We will use an approved encryption method for encrypting Local SSD and VM to VM communication. In B and D, we are still using GCP's encryption algorithms and are not FIPS 140-2 approved. Moreover only the BoringCrypto is FIPS 140-2 approved and not the Boring SSL. I see A as evidently correct. ownnez, genesis3k, MohitA has explained this and provided the right links too.

upvoted 12 times

  **AzureDP900** Most Recent 1 month, 3 weeks ago

<https://cloud.google.com/docs/security/key-management-deep-dive>

A is right

upvoted 1 times

  **AwesomeGCP** 2 months, 2 weeks ago

Selected Answer: A

A. Encrypt all cache storage and VM-to-VM communication using the BoringCrypto module.

upvoted 1 times

  **sudarchary** 10 months, 2 weeks ago

Selected Answer: A

FIPS140 module is supported

upvoted 2 times

  **[Removed]** 1 year, 8 months ago

D is the correct answer

upvoted 2 times

  **DebasishLowes** 1 year, 9 months ago

Ans : A

upvoted 1 times

  **TNT87** 1 year, 9 months ago

<https://cloud.google.com/security/compliance/fips-140-2-validated>

Google Cloud Platform uses a FIPS 140-2 validated encryption module called BoringCrypto (certificate 3318) in our production environment. This means that both data in transit to the customer and between data centers, and data at rest are encrypted using FIPS 140-2 validated encryption. The module that achieved FIPS 140-2 validation is part of our BoringSSL library.

Ans A

upvoted 4 times

  **TNT87** 1 year, 10 months ago

A is the answer <https://boringssl.googlesource.com/boringssl/+master/crypto/fipsmodule/FIPS.md>

upvoted 2 times

  **chetz12** 1 year, 11 months ago

I think A is correct as that's the only one support FIPS140 module

upvoted 3 times



🗨️ 👤 **[Removed]** 2 years, 1 month ago

Ans - B

upvoted 1 times

🗨️ 👤 **genesis3k** 2 years, 1 month ago

Agree Answer is A definitely.

"BoringSSL as a whole is not FIPS validated. However, there is a core library (called BoringCrypto) that has been FIPS validated"

<https://boringssl.googlesource.com/boringssl/+master/crypto/fipsmodule/FIPS.md>

upvoted 3 times

🗨️ 👤 **saurabh1805** 2 years, 2 months ago

For me correct answer is B

upvoted 1 times

🗨️ 👤 **ownez** 2 years, 2 months ago

Answer is A.

BoringSSL as a whole is not FIPS validated.

<https://boringssl.googlesource.com/boringssl/+master/crypto/fipsmodule/FIPS.md>

upvoted 2 times

🗨️ 👤 **mte\_tech34** 2 years, 2 months ago

According to <https://cloud.google.com/security/compliance/fips-140-2-validated>

"Google's Local SSD storage product is automatically encrypted with NIST approved ciphers, but Google's current implementation for this product doesn't have a FIPS 140-2 validation certificate. If you require FIPS-validated encryption on Local SSD storage, you must provide your own encryption with a FIPS-validated cryptographic module."

So only A seems to be correct, as BoringSSL is FIPS-140-2 validated.

upvoted 1 times

🗨️ 👤 **mte\_tech34** 2 years, 2 months ago

I meant BoringCrypto, not BoringSSL, as it's not the same thing and only BoringCrypto is FIPS validated. BoringSSL is not (see

<https://boringssl.googlesource.com/boringssl/+master/crypto/fipsmodule/FIPS.md>), so definitely only A can be correct answer.

upvoted 1 times

🗨️ 👤 **Raushanr** 2 years, 3 months ago

Another reasoning on Google Documentation-You cannot use your own keys with local SSDs because local SSDs do not persist beyond the life of a virtual machine. Local SSDs are already protected with an ephemeral encryption key that Google does not retain.

upvoted 2 times

🗨️ 👤 **Raushanr** 2 years, 3 months ago

Answer-B

Reasoning from Google-Google's Local SSD storage product is automatically encrypted with NIST approved ciphers, but Google's current implementation for this product doesn't have a FIPS 140-2 validation certificate. If you require FIPS-validated encryption on Local SSD storage, you must provide your own encryption with a FIPS-validated cryptographic module.

upvoted 3 times

🗨️ 👤 **ownez** 2 years, 3 months ago

I think it is B.

[https://services.google.com/fh/files/misc/googlecloud\\_european\\_commitments\\_whitepaper.pdf](https://services.google.com/fh/files/misc/googlecloud_european_commitments_whitepaper.pdf)

upvoted 4 times

🗨️ 👤 **ownez** 2 years, 3 months ago

This is because "The instances use Local SSDs for data caching" and in the documentation says "If you require FIPS-validated encryption on Local SSD storage, you must provide your own encryption with a FIPS-validated cryptographic module."

upvoted 1 times

A customer has an analytics workload running on Compute Engine that should have limited internet access. Your team created an egress firewall rule to deny (priority 1000) all traffic to the internet. The Compute Engine instances now need to reach out to the public repository to get security updates. What should your team do?

- A. Create an egress firewall rule to allow traffic to the CIDR range of the repository with a priority greater than 1000.
- B. Create an egress firewall rule to allow traffic to the CIDR range of the repository with a priority less than 1000.
- C. Create an egress firewall rule to allow traffic to the hostname of the repository with a priority greater than 1000.
- D. Create an egress firewall rule to allow traffic to the hostname of the repository with a priority less than 1000.

Correct Answer: C

Community vote distribution



**dtmtor** Highly Voted 1 year, 9 months ago  
Answer is B. Lower number is higher priority and dest is only IP ranges in firewall rules  
upvoted 21 times

**shayke** Most Recent 1 day, 19 hours ago  
Selected Answer: B  
Ans in B lower number higher priority  
upvoted 1 times

**Littleivy** 1 month, 1 week ago  
Selected Answer: B  
Answer is B  
upvoted 1 times

**GHOST1985** 1 month, 2 weeks ago  
Selected Answer: B  
[https://cloud.google.com/vpc/docs/firewalls#priority\\_order\\_for\\_firewall\\_rules](https://cloud.google.com/vpc/docs/firewalls#priority_order_for_firewall_rules)  
upvoted 2 times

**AzureDP900** 1 month, 3 weeks ago  
B is correct  
upvoted 1 times

**Premumar** 1 month, 3 weeks ago  
Selected Answer: B  
First filter is priority should be less than 1000. So, option A and C are rejected. Then, we use CIDR range to allow firewall. So, the final answer is B.  
upvoted 1 times

**AwesomeGCP** 2 months, 2 weeks ago  
Selected Answer: B  
B. Create an egress firewall rule to allow traffic to the CIDR range of the repository with a priority less than 1000.  
Firewall rules only support IPv4 connections. When specifying a source for an ingress rule or a destination for an egress rule by address, you can only use an IPv4 address or IPv4 block in CIDR notation. So Answer is B  
upvoted 2 times

**piyush\_1982** 4 months, 3 weeks ago  
Selected Answer: A  
The correct answer is A.  
As per the link [https://cloud.google.com/vpc/docs/firewalls#rule\\_assignment](https://cloud.google.com/vpc/docs/firewalls#rule_assignment)  
  
Lowest priority in the firewall rule is 65535. So in order for a rule to be of higher priority than 1000 the rule should have a priority of number less than 1000.  
upvoted 2 times

**Premumar** 1 month, 3 weeks ago  
Your explanation is correct. But, option you selected is wrong. It has to be option B.  
upvoted 2 times

**Rithac** 1 year, 6 months ago

I think I am confusing myself by overthinking the wording of this question. I know the answer is A or B since "using hostname is not one of the options in firewall egress rule destination" I also know that "The firewall rule priority is an integer from 0 to 65535, inclusive. Lower integers indicate higher priorities." I know that I could resolve this by setting TCP port 80 rule to a priority of 500 (smaller number, but higher priority) and be done. Where i'm second guessing myself, is Google referring to the integer or strictly priority? If integer then i'd choose B "priority less than 1000 (smaller number)", if priority then i'd choose A "priority greater than 1000" (still the lower number). Have I thoroughly confused this question? I'm leaning toward the answer being "A":

upvoted 4 times

  **DebasishLowes** 1 year, 9 months ago

Ans : B

upvoted 3 times

  **ronron89** 2 years ago

Answer: B

[https://cloud.google.com/vpc/docs/firewalls#rule\\_assignment](https://cloud.google.com/vpc/docs/firewalls#rule_assignment)

The priority of the second rule determines whether TCP traffic to port 80 is allowed for the webserver targets:

If the priority of the second rule is set to a number greater than 1000, it has a lower priority, so the first rule denying all traffic applies.

If the priority of the second rule is set to 1000, the two rules have identical priorities, so the first rule denying all traffic applies.

If the priority of the second rule is set to a number less than 1000, it has a higher priority, thus allowing traffic on TCP 80 for the webserver targets. Absent other rules, the first rule would still deny other types of traffic to the webserver targets, and it would also deny all traffic, including TCP 80, to instances without the webserver tag.

upvoted 4 times

  **[Removed]** 2 years, 1 month ago

Ans - B

upvoted 2 times

  **Raushanr** 2 years, 3 months ago

The firewall rule priority is an integer from 0 to 65535, inclusive. Lower integers indicate higher priorities. If you do not specify a priority when creating a rule, it is assigned a priority of 1000.

upvoted 1 times

  **Raushanr** 2 years, 3 months ago

Answer-B

upvoted 4 times

  **ownez** 2 years, 3 months ago

It should be B.

Firewall rules can be only specify by IPv4 address or IPv4 block in CIDR.

And it must be lesser priority than 1000 because if more than that, it will overwrite the deny rule.

upvoted 2 times

  **ownez** 2 years, 3 months ago

Sorry It should be A.

"Priority: the numeric evaluation order of the rule. A rule with a priority of 1 is evaluated first. Priorities must be unique for each rule. A good practice is to give rules priority numbers that allow later insertion (such as 100, 200, 300)."

upvoted 1 times

  **ownez** 2 years, 3 months ago

Correction. B

upvoted 3 times

  **mlyu** 2 years, 3 months ago

Ans should be A

using hostname is not one of the options in firewall egress rule destination

[https://cloud.google.com/vpc/docs/firewalls#gcp\\_firewall\\_rule\\_summary\\_table](https://cloud.google.com/vpc/docs/firewalls#gcp_firewall_rule_summary_table)

upvoted 2 times

You want data on Compute Engine disks to be encrypted at rest with keys managed by Cloud Key Management Service (KMS). Cloud Identity and Access Management (IAM) permissions to these keys must be managed in a grouped way because the permissions should be the same for all keys. What should you do?

- A. Create a single KeyRing for all persistent disks and all Keys in this KeyRing. Manage the IAM permissions at the Key level.
- B. Create a single KeyRing for all persistent disks and all Keys in this KeyRing. Manage the IAM permissions at the KeyRing level.
- C. Create a KeyRing per persistent disk, with each KeyRing containing a single Key. Manage the IAM permissions at the Key level.
- D. Create a KeyRing per persistent disk, with each KeyRing containing a single Key. Manage the IAM permissions at the KeyRing level.

**Correct Answer: C**

Community vote distribution

B (100%)

🗲️ 👤 **TNT87** Highly Voted 👍 1 year, 10 months ago

Ans B

<https://cloud.netapp.com/blog/gcp-cvo-blg-how-to-use-google-cloud-encryption-with-a-persistent-disk>

upvoted 13 times

🗲️ 👤 **AzureDP900** Most Recent ⌚ 1 month, 3 weeks ago

B is right

upvoted 1 times

🗲️ 👤 **shayke** 2 months ago

Selected Answer: B

all permission are the same-controlled at the ring level

upvoted 2 times

🗲️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

Selected Answer: B

B. Create a single KeyRing for all persistent disks and all Keys in this KeyRing. Manage the IAM permissions at the KeyRing level.

upvoted 3 times

🗲️ 👤 **roatest27** 8 months, 4 weeks ago

Answer-B

upvoted 1 times

🗲️ 👤 **[Removed]** 1 year, 8 months ago

How about A?

upvoted 1 times

🗲️ 👤 **[Removed]** 1 year, 8 months ago

oh, the same permission ,then I choose B

upvoted 4 times

🗲️ 👤 **DebasishLowes** 1 year, 9 months ago

Ans : B

upvoted 3 times

🗲️ 👤 **[Removed]** 2 years, 1 month ago

Ans - B

upvoted 1 times

🗲️ 👤 **Raushanr** 2 years, 3 months ago

Answer-B

upvoted 1 times

🗲️ 👤 **Namaste** 2 years, 3 months ago

B is the right answer

upvoted 1 times

🗲️ 👤 **MohitA** 2 years, 3 months ago

B should be the answer

upvoted 4 times

A company is backing up application logs to a Cloud Storage bucket shared with both analysts and the administrator. Analysts should only have access to logs that do not contain any personally identifiable information (PII). Log files containing PII should be stored in another bucket that is only accessible by the administrator.

What should you do?

- A. Use Cloud Pub/Sub and Cloud Functions to trigger a Data Loss Prevention scan every time a file is uploaded to the shared bucket. If the scan detects PII, have the function move into a Cloud Storage bucket only accessible by the administrator.
- B. Upload the logs to both the shared bucket and the bucket only accessible by the administrator. Create a job trigger using the Cloud Data Loss Prevention API. Configure the trigger to delete any files from the shared bucket that contain PII.
- C. On the bucket shared with both the analysts and the administrator, configure Object Lifecycle Management to delete objects that contain any PII.
- D. On the bucket shared with both the analysts and the administrator, configure a Cloud Storage Trigger that is only triggered when PII data is uploaded. Use Cloud Functions to capture the trigger and delete such files.

**Correct Answer: C**

*Community vote distribution*

A (100%)

🗳️ 👤 **MohitA** Highly Voted 2 years, 3 months ago

A is the ans  
upvoted 14 times

🗳️ 👤 **talktolanka** Highly Voted 1 year, 8 months ago

Answer A  
<https://codelabs.developers.google.com/codelabs/cloud-storage-dlp-functions#0>  
<https://www.youtube.com/watch?v=0TmO1f-Ox40>  
upvoted 7 times

🗳️ 👤 **AzureDP900** Most Recent 1 month, 3 weeks ago

A is right  
upvoted 1 times

🗳️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: A**

A. Use Cloud Pub/Sub and Cloud Functions to trigger a Data Loss Prevention scan every time a file is uploaded to the shared bucket. If the scan detects PII, have the function move into a Cloud Storage bucket only accessible by the administrator.  
upvoted 2 times

🗳️ 👤 **satyajeet1999** 3 months, 2 weeks ago

**Selected Answer: A**

A it is.  
upvoted 1 times

🗳️ 👤 **[Removed]** 1 year, 8 months ago

I also choose A.  
upvoted 3 times

🗳️ 👤 **DebasishLowes** 1 year, 9 months ago

Ans : A  
upvoted 2 times

🗳️ 👤 **soukumar369** 2 years ago

Correct answer is A : Data Loss Prevention scan  
upvoted 2 times

🗳️ 👤 **soukumar369** 2 years ago

A is correct.  
upvoted 1 times

🗳️ 👤 **[Removed]** 2 years, 1 month ago

Ans - A  
upvoted 1 times



  **genesis3k** 2 years, 1 month ago

Answer is A.  
upvoted 1 times

  **passtest100** 2 years, 2 months ago

SHOULD BE A  
upvoted 1 times

A customer terminates an engineer and needs to make sure the engineer's Google account is automatically deprovisioned. What should the customer do?

- A. Use the Cloud SDK with their directory service to remove their IAM permissions in Cloud Identity.
- B. Use the Cloud SDK with their directory service to provision and deprovision users from Cloud Identity.
- C. Configure Cloud Directory Sync with their directory service to provision and deprovision users from Cloud Identity.
- D. Configure Cloud Directory Sync with their directory service to remove their IAM permissions in Cloud Identity.

Correct Answer: C

Community vote distribution



**MohitA** Highly Voted 2 years, 3 months ago  
C is the Answer  
upvoted 6 times

**ownez** 2 years, 3 months ago  
Agree with C.  
  
"https://cloud.google.com/identity/solutions/automate-user-provisioning#cloud\_identity\_automated\_provisioning"  
  
"Cloud Identity has a catalog of automated provisioning connectors, which act as a bridge between Cloud Identity and third-party cloud apps."  
upvoted 10 times

**AzureDP900** 1 month, 2 weeks ago  
Agree with C, there is no need of cloud SDK.  
upvoted 1 times

**AzureDP900** 1 month, 2 weeks ago  
C. Configure Cloud Directory Sync with their directory service to provision and deprovision users from Cloud Identity.  
upvoted 1 times

**mynk29** 9 months, 4 weeks ago  
This option is for Cloud identity to third party app- you configure directory sync between AD and cloud identity.  
upvoted 1 times

**[Removed]** Highly Voted 2 years, 1 month ago  
Ans - C  
upvoted 6 times

**AzureDP900** Most Recent 1 month, 3 weeks ago  
C. Configure Cloud Directory Sync with their directory service to provision and deprovision users from Cloud Identity.  
upvoted 1 times

**AwesomeGCP** 2 months, 2 weeks ago  
Selected Answer: C  
C. Configure Cloud Directory Sync with their directory service to provision and deprovision users from Cloud Identity.  
upvoted 1 times

**piyush\_1982** 4 months, 4 weeks ago  
Selected Answer: C  
Definitely C  
upvoted 1 times

**mynk29** 9 months, 4 weeks ago  
I don't think C is right answer. You configure Directory Sync to Sync from AD to cloud identity not the other way round.  
  
Once a user is terminated- its account should be disabled on Directory and cloud identity will pick up via IAM. D looks more correct to me.  
upvoted 2 times

**AkbarM** 3 months ago  
I also support D. The question may be provision and deprovision users. but technically it is to remove their IAM permissions in Cloud Identity. There is nothing like provision / deprovision user from cloud identity.  
upvoted 1 times



  **DebasishLowes** 1 year, 9 months ago

Ans is C

upvoted 3 times

An organization is evaluating the use of Google Cloud Platform (GCP) for certain IT workloads. A well-established directory service is used to manage user identities and lifecycle management. This directory service must continue for the organization to use as the `source of truth` directory for identities.

Which solution meets the organization's requirements?

- A. Google Cloud Directory Sync (GCDS)
- B. Cloud Identity
- C. Security Assertion Markup Language (SAML)
- D. Pub/Sub

**Correct Answer:** B

Reference:

<https://cloud.google.com/solutions/federating-gcp-with-active-directory-introduction>

*Community vote distribution*

A (100%)

🗲️ 👤 **desertlotus1211** Highly Voted 👍 1 year, 9 months ago

The answer is A:

With Google Cloud Directory Sync (GCDS), you can synchronize the data in your Google Account with your Microsoft Active Directory or LDAP server. GCDS doesn't migrate any content (such as email messages, calendar events, or files) to your Google Account. You use GCDS to synchronize your Google users, groups, and shared contacts to match the information in your LDAP server.

The questions says the well established directory service is the 'source of truth' not GCP... So LDAP or AD is the source... GCDS will sync that to match those, not replace them...

upvoted 14 times

🗲️ 👤 **AzureDP900** 1 month, 3 weeks ago

Agreed

upvoted 1 times

🗲️ 👤 **subhala** Highly Voted 👍 2 years ago

GCDS -? It helps sync up from the source of truth (any IdP like ldap, AD) to Google identity. In this scenario, the question is what can be a good identity service by itself, hence B is the right answer.

upvoted 10 times

🗲️ 👤 **PST21** Most Recent ⌚ 3 days, 11 hours ago

Orgn is evaluating GC so cloud Identity is the GC product hence B

upvoted 1 times

🗲️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: A**

A. Google Cloud Directory Sync (GCDS)

upvoted 2 times

🗲️ 👤 **cloudprincipal** 6 months, 3 weeks ago

**Selected Answer: A**

With Google Cloud Directory Sync (GCDS), you can synchronize the data in your Google Account with your Microsoft Active Directory or LDAP server. GCDS doesn't migrate any content (such as email messages, calendar events, or files) to your Google Account. You use GCDS to synchronize your Google users, groups, and shared contacts to match the information in your LDAP server.

<https://support.google.com/a/answer/106368?hl=en>

upvoted 1 times

🗲️ 👤 **szl0144** 7 months ago

B should be the answer, GCDS is for ad sync.

upvoted 2 times

🗲️ 👤 **MariaGabiGabriela** 6 months, 3 weeks ago

Yes, but identity by itself will solve nothing, the user would have to recreate all users and thus have a different IDP, this clearly goes against the question

upvoted 2 times

🗲️ 👤 **Bill831231** 1 year ago

seems there is nothing metioned about what they have on premise, so B is better

upvoted 1 times

  **syllox** 1 year, 7 months ago

Answer A

upvoted 3 times

  **WakandaF** 1 year, 7 months ago

A or B?

upvoted 2 times

  **DebasishLowes** 1 year, 9 months ago

Ans : B as per the question.

upvoted 1 times

  **asee** 1 year, 10 months ago

My Answer will go for A (GCDS), noticed the question is mentioning about "A directory service 'is used' " / "must continue" instead of "A directory service 'will be used' ". So here my understanding is the organization has already using their own directory service. Therefore Answer B - Cloud identity may not be an option.

upvoted 4 times

  **KWatHK** 1 year, 11 months ago

Ans is B because the questions said "the well-established directory must continue for the organization to use as the source of truth" so that the user access to GCP must authenticated by the existing directory. Cloud Identity support to federate it to 3rd party/ADFS using SAML.

upvoted 1 times

  **mikelabs** 2 years ago

GCDS is an app to sync users, groups and other features from AD to Cloud Identity. But, in this question, the customer needs to know what's the product on GCP that meet with this. So, I think the answer is B.

upvoted 8 times

  **[Removed]** 2 years, 1 month ago

Ans - A

upvoted 2 times

  **ownez** 2 years, 3 months ago

GCDS is a part of CI's feature that synchronizes the data in Google domain to match with AD/LDAP server. This includes users, groups contacts etc are synchronized/migrated to match.

Hence, I would go B.

"<https://se-cloud-experts.com/wp/wp-content/themes/se-it/images/pdf/google-cloud-identity-services.pdf>"

upvoted 3 times

  **ownez** 2 years, 3 months ago

Sorry. It's A.

upvoted 2 times

  **bogdant** 2 years, 3 months ago

Isn't it A?

upvoted 2 times

  **Sheeda** 2 years, 3 months ago

That is used to sync, not the directly itself

upvoted 1 times

  **Fellipo** 2 years, 1 month ago

A well-established directory service , so "A"

upvoted 2 times

  **MohitA** 2 years, 3 months ago

Agree A

upvoted 3 times

Which international compliance standard provides guidelines for information security controls applicable to the provision and use of cloud services?

- A. ISO 27001
- B. ISO 27002
- C. ISO 27017
- D. ISO 27018

**Correct Answer:** C

Create a new Service Account that should be able to list the Compute Engine instances in the project. You want to follow Google-recommended practices.

*Community vote distribution*

C (100%)

🗲️ 👤 **asee** Highly Voted 1 year, 10 months ago

Yes, My answer also goes to C and my last compliance related project is also working on ISO27017 in order to extend the scope to Cloud service user/provider.

upvoted 8 times

🗲️ 👤 **AzureDP900** 1 month, 3 weeks ago

C is right

upvoted 1 times

🗲️ 👤 **AzureDP900** 1 month, 3 weeks ago

<https://cloud.google.com/security/compliance/iso-27017>

upvoted 1 times

🗲️ 👤 **AwesomeGCP** Most Recent 2 months, 2 weeks ago

**Selected Answer: C**

C. ISO 27017

upvoted 2 times

🗲️ 👤 **certificationjjmmm** 5 months ago

C is correct.

<https://cloud.google.com/security/compliance/iso-27017>

upvoted 2 times

🗲️ 👤 **[Removed]** 2 years, 1 month ago

Ans - C

upvoted 2 times

🗲️ 👤 **Namaste** 2 years, 3 months ago

CCSP Question...C is the Answer

upvoted 2 times

🗲️ 👤 **ownez** 2 years, 3 months ago

C is correct.

"<https://www.iso.org/standard/43757.html>"

upvoted 3 times

You will create a new Service Account that should be able to list the Compute Engine instances in the project. You want to follow Google-recommended practices.

What should you do?

- A. Create an Instance Template, and allow the Service Account Read Only access for the Compute Engine Access Scope.
- B. Create a custom role with the permission compute.instances.list and grant the Service Account this role.
- C. Give the Service Account the role of Compute Viewer, and use the new Service Account for all instances.
- D. Give the Service Account the role of Project Viewer, and use the new Service Account for all instances.

**Correct Answer:** A

Community vote distribution

B (80%)

C (20%)

🗳️ 👤 **MohitA** Highly Voted 👍 2 years, 3 months ago

B, <https://cloud.google.com/compute/docs/access/iam>  
upvoted 15 times

🗳️ 👤 **AzureDP900** 1 month, 3 weeks ago

B is right  
upvoted 1 times

🗳️ 👤 **mlyu** 2 years, 3 months ago

Although it is not encourage to use custome role, but last sentence in the answer C makes B be the only option  
upvoted 6 times

🗳️ 👤 **shayke** Most Recent ⌚ 1 day, 19 hours ago

Selected Answer: B

B is the right ans - you only want to list the instances  
upvoted 1 times

🗳️ 👤 **Meyucho** 4 days, 2 hours ago

Selected Answer: B

With C the SA will list ONLY the instances that are configured to use that SA.  
The option B will give permissions to list ALL instances.  
upvoted 1 times

🗳️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

Selected Answer: B

B. Create a custom role with the permission compute.instances.list and grant the Service Account this role.  
upvoted 1 times

🗳️ 👤 **nbrnschwgr** 3 months, 3 weeks ago

C. because google recommends pre-defined narrow scope roles over custom roles.  
upvoted 1 times

🗳️ 👤 **Roflicopter** 4 months, 2 weeks ago

Selected Answer: B

Key here is "and grant the Service Account this role.". C and D are giving this role to ALL instances which is overly permissive. A is wrong. Only choice is B  
upvoted 2 times

🗳️ 👤 **cloudprincipal** 6 months, 3 weeks ago

Selected Answer: B

The roles/compute.viewer provides a lot more privileges than just listing compute instances  
upvoted 3 times

🗳️ 👤 **cloudprincipal** 6 months, 3 weeks ago

Selected Answer: C

Compute Viewer  
Read-only access to get and list Compute Engine resources, without being able to read the data stored on them.

<https://cloud.google.com/compute/docs/access/iam#compute.viewer>  
upvoted 2 times

  **cloudprincipal** 6 months, 1 week ago

This is incorrect, as Compute Viewer provides a lot more than what is required  
upvoted 1 times

  **sudarchary** 11 months ago

B. The only option that adheres to the principle of least privilege and meets question requirements is B  
upvoted 4 times

  **[Removed]** 1 year, 8 months ago

I think C is good  
upvoted 4 times

  **DebasishLowes** 1 year, 9 months ago

Ans : B  
upvoted 1 times

  **dtmtor** 1 year, 9 months ago

Ans is B  
upvoted 1 times

  **[Removed]** 2 years, 1 month ago

Ans - B  
upvoted 1 times

  **genesis3k** 2 years, 1 month ago

Answer is B, based on least privilege principle.  
upvoted 1 times



In a shared security responsibility model for IaaS, which two layers of the stack does the customer share responsibility for? (Choose two.)

- A. Hardware
- B. Network Security
- C. Storage Encryption
- D. Access Policies
- E. Boot

Correct Answer: *CD*

Community vote distribution



**DebasishLowes** Highly Voted 1 year, 9 months ago  
Ans : BD  
upvoted 11 times

**shayke** Most Recent 1 day, 18 hours ago  
b and D - according to the shared responsibility moder for IAAS  
upvoted 1 times

**AwesomeGCP** 2 months, 2 weeks ago  
Selected Answer: BD  
B. Network Security  
D. Access Policies  
upvoted 2 times

**Random\_Mane** 3 months ago  
Selected Answer: BD  
Chart is here <https://cloud.google.com/architecture/framework/security/shared-responsibility-shared-fate>  
upvoted 1 times

**rr4444** 11 months, 3 weeks ago  
Selected Answer: BD  
BD <https://cloud.google.com/blog/products/containers-kubernetes/exploring-container-security-the-shared-responsibility-model-in-gke-container-security-shared-responsibility-model-gke>  
upvoted 3 times

**[Removed]** 2 years, 1 month ago  
Ans - BD  
upvoted 3 times

**saurabh1805** 2 years, 1 month ago  
B and D is correct option.  
upvoted 3 times

**passtest100** 2 years, 2 months ago  
B and D  
upvoted 3 times

**lordb** 2 years, 3 months ago  
B and D  
upvoted 3 times

An organization is starting to move its infrastructure from its on-premises environment to Google Cloud Platform (GCP). The first step the organization wants to take is to migrate its ongoing data backup and disaster recovery solutions to GCP. The organization's on-premises production environment is going to be the next phase for migration to GCP. Stable networking connectivity between the on-premises environment and GCP is also being implemented.

Which GCP solution should the organization use?


- A. BigQuery using a data pipeline job with continuous updates via Cloud VPN
- B. Cloud Storage using a scheduled task and gsutil via Cloud Interconnect
- C. Compute Engines Virtual Machines using Persistent Disk via Cloud Interconnect
- D. Cloud Datastore using regularly scheduled batch upload jobs via Cloud VPN

**Correct Answer:** B

Reference:

<https://cloud.google.com/solutions/migration-to-google-cloud-building-your-foundation>

*Community vote distribution*



**ownez** Highly Voted 2 years, 3 months ago  
Agree B.

[https://cloud.google.com/solutions/dr-scenarios-for-data#production\\_environment\\_is\\_on-premises](https://cloud.google.com/solutions/dr-scenarios-for-data#production_environment_is_on-premises)  
upvoted 11 times

**shayke** Most Recent 1 day, 18 hours ago

**Selected Answer: B**

B- backup and DR is GCS  
upvoted 1 times

**rotorclear** 2 months ago

**Selected Answer: B**

<https://medium.com/@pvergadia/cold-disaster-recovery-on-google-cloud-for-applications-running-on-premises-114b31933d02>  
upvoted 2 times

**AzureDP900** 1 month, 2 weeks ago  
B is correct  
upvoted 1 times

**AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: B**

B. Cloud Storage using a scheduled task and gsutil via Cloud Interconnect  
upvoted 2 times

**cloudprincipal** 6 months, 3 weeks ago

**Selected Answer: B**

[https://cloud.google.com/solutions/dr-scenarios-for-data#production\\_environment\\_is\\_on-premises](https://cloud.google.com/solutions/dr-scenarios-for-data#production_environment_is_on-premises)  
upvoted 1 times

**rr4444** 11 months, 3 weeks ago

**Selected Answer: C**

Disaster recover made me think C Compute Engines Virtual Machines using Persistent Disk via Cloud Interconnect

Disaster recovery with remote backup alone, when all prod is on premise, will take too long to be viable. The VMs don't need to be running when no disaster  
upvoted 3 times

**csrazdan** 3 weeks, 3 days ago

You would have been correct if the question had any RTO/RPO specifications. In absence of this question is assuming backup and restore as a DR strategy. So Option B Cloud Storage is the correct answer.  
upvoted 1 times

**DebasishLowes** 1 year, 9 months ago

Ans : B



upvoted 2 times

  **[Removed]** 2 years, 1 month ago

Ans - V

upvoted 1 times

  **[Removed]** 2 years, 1 month ago

Typo - it's B

upvoted 2 times

## Question #74

Topic 1

What are the steps to encrypt data using envelope encryption?

A.

- ⇒ Generate a data encryption key (DEK) locally.
- ⇒ Use a key encryption key (KEK) to wrap the DEK.
- ⇒ Encrypt data with the KEK.
- ⇒ Store the encrypted data and the wrapped KEK.

B.

- ⇒ Generate a key encryption key (KEK) locally.
- ⇒ Use the KEK to generate a data encryption key (DEK).
- ⇒ Encrypt data with the DEK.
- ⇒ Store the encrypted data and the wrapped DEK.

C.

- ⇒ Generate a data encryption key (DEK) locally.
- ⇒ Encrypt data with the DEK.
- ⇒ Use a key encryption key (KEK) to wrap the DEK.
- ⇒ Store the encrypted data and the wrapped DEK.

D.

- ⇒ Generate a key encryption key (KEK) locally.
- ⇒ Generate a data encryption key (DEK) locally.
- ⇒ Encrypt data with the KEK.

Store the encrypted data and the wrapped DEK.

▪

**Correct Answer: C**

Reference:

<https://cloud.google.com/kms/docs/envelope-encryption>

  **Tabayashi** **Highly Voted**  7 months, 4 weeks ago

Answer is (C).

The process of encrypting data is to generate a DEK locally, encrypt data with the DEK, use a KEK to wrap the DEK, and then store the encrypted data and the wrapped DEK. The KEK never leaves Cloud KMS.

[https://cloud.google.com/kms/docs/envelope-encryption#how\\_to\\_encrypt\\_data\\_using\\_envelope\\_encryption](https://cloud.google.com/kms/docs/envelope-encryption#how_to_encrypt_data_using_envelope_encryption)

upvoted 10 times

  **AzureDP900** 1 month, 2 weeks ago

C is right

upvoted 1 times

  **AwesomeGCP** **Most Recent**  2 months, 2 weeks ago

Answer - C is correct.

[https://cloud.google.com/kms/docs/envelope-encryption#how\\_to\\_encrypt\\_data\\_using\\_envelope\\_encryption](https://cloud.google.com/kms/docs/envelope-encryption#how_to_encrypt_data_using_envelope_encryption)

upvoted 2 times

  **satyajeet1999** 3 months, 2 weeks ago

C it is

upvoted 1 times

A customer wants to make it convenient for their mobile workforce to access a CRM web interface that is hosted on Google Cloud Platform (GCP). The CRM can only be accessed by someone on the corporate network. The customer wants to make it available over the internet. Your team requires an authentication layer in front of the application that supports two-factor authentication. Which GCP product should the customer implement to meet these requirements?

- A. Cloud Identity-Aware Proxy
- B. Cloud Armor
- C. Cloud Endpoints
- D. Cloud VPN

**Correct Answer:** D

*Community vote distribution*

A (100%)

🗲️ 👤 **asee** Highly Voted 1 year, 10 months ago

My answer is going for A.  
Cloud IAP is integrated with Google Sign-in which Multi-factor authentication can be enabled.  
<https://cloud.google.com/iap/docs/concepts-overview>  
upvoted 17 times

🗲️ 👤 **AzureDP900** 1 month, 2 weeks ago

I agree and A is right  
upvoted 1 times

🗲️ 👤 **MohitA** Highly Voted 2 years, 3 months ago

A is the Answer  
upvoted 6 times

🗲️ 👤 **AwesomeGCP** Most Recent 2 months, 2 weeks ago

**Selected Answer: A**

A. Cloud Identity-Aware Proxy  
I think it's A. The question asks for an authentication layer.  
upvoted 1 times

🗲️ 👤 **danielklein09** 11 months ago

**Selected Answer: A**

A is the correct answer  
upvoted 2 times

🗲️ 👤 **[Removed]** 2 years, 1 month ago

Ans - A  
upvoted 3 times

🗲️ 👤 **passtest100** 2 years, 2 months ago

SHOULD BE A  
upvoted 4 times

🗲️ 👤 **Raushanr** 2 years, 3 months ago

Answer -A  
upvoted 3 times

Your company is storing sensitive data in Cloud Storage. You want a key generated on-premises to be used in the encryption process. What should you do?


- A. Use the Cloud Key Management Service to manage a data encryption key (DEK).
- B. Use the Cloud Key Management Service to manage a key encryption key (KEK).
- C. Use customer-supplied encryption keys to manage the data encryption key (DEK).
- D. Use customer-supplied encryption keys to manage the key encryption key (KEK).

**Correct Answer: A**

Reference:

<https://cloud.google.com/security/encryption-at-rest/default-encryption/>

*Community vote distribution*



🗳️ 👤 **HateMicrosoft** Highly Voted 1 year, 9 months ago

The answer is: C  
 This is a Customer-supplied encryption key (CSEK).  
 We generate our own encryption key and manage it on-premises.  
 A KEK never leaves Cloud KMS. There is no KEK or KMS on-premises.

Encryption at rest by default, with various key management options  
<https://cloud.google.com/security/encryption-at-rest>  
 upvoted 17 times

🗳️ 👤 **PST21** Most Recent 4 days, 18 hours ago

C is correct as one may choose to import the key or not, hence A is correct if one chooses to import, but in both cases C does the job  
 upvoted 2 times

🗳️ 👤 **marmar11111** 1 month ago

**Selected Answer: C**

it's C!  
 upvoted 1 times

🗳️ 👤 **Littleivy** 1 month, 1 week ago

**Selected Answer: A**

The answer is A  
  
 You can import your key into Cloud Key Management Service for managing it more easily  
<https://cloud.google.com/kms/docs/importing-a-key>  
 upvoted 1 times

🗳️ 👤 **Table2022** 1 month, 4 weeks ago

D is the correct answer as DEK belongs to Google Cloud.  
 upvoted 1 times

🗳️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: C**

C. Use customer-supplied encryption keys to manage the data encryption key (DEK).

With CSEK, you supply your own AES-256 key to serve as the KEK, and your key protects the DEKs that protect your data. Your CSEK key is protected by an additional layer of protection, using a Cloud KMS key.

Now that you can import keys into Cloud KMS, you can import your keys and use them with CMEK-enabled services instead of relying on CSEK.  
 upvoted 1 times

🗳️ 👤 **GHOST1985** 2 months, 3 weeks ago

**Selected Answer: C**

Answer C ; A KEK never leaves Cloud KMS  
 upvoted 1 times

🗳️ 👤 **piyush\_1982** 4 months, 3 weeks ago

**Selected Answer: C**

The correct answer is C.  
 Customer generating key on-premises is related to Customer-supplied encryption key to encrypt the data often known as the data encryption key.

So, C is the most appropriate option. The question doesn't state the need for envelop encryption so KEK is not an option.

upvoted 2 times

  **Chute5118** 5 months ago

**Selected Answer: D**

The word "manage" on this question can have multiple meanings. If we replace "to manage" by the word "as"... Then it's D

We generate a key locally (which is a KEK) and Google uses it to encrypt the different DEKs (for each chunk)

"If you supply your own encryption keys, Google uses your key to protect the Google-generated keys used to encrypt and decrypt your data."

<https://cloud.google.com/docs/security/encryption/customer-supplied-encryption-keys>

upvoted 2 times

  **sudarchary** 10 months, 3 weeks ago

**Selected Answer: D**

Reference Links:

<https://cloud.google.com/kms/docs/envelope-encryption>

<https://cloud.google.com/security/encryption-at-rest/customer-supplied-encryption-keys>

upvoted 4 times

  **sudarchary** 11 months ago

Answer is "D" . KEK: A key used to encrypt, or "wrap", a data encryption key

upvoted 1 times

  **alexm112** 10 months, 2 weeks ago

It doesn't manage the KEK, it IS the KEK. The answer is C

upvoted 2 times

  **Bill831231** 1 year ago

sounds D is the correct one.

The raw CSEK is combined with a per-persistent disk cryptographic nonce to generate a CSEK-derived key. This key is used as the key encryption key in Google Compute Engine for your data.

[https://cloud.google.com/security/encryption/customer-supplied-encryption-keys#compute\\_engine](https://cloud.google.com/security/encryption/customer-supplied-encryption-keys#compute_engine)

upvoted 2 times

  **rr4444** 11 months, 3 weeks ago

Nope cos being used as the KEK means that this CSEK KEK is used to MANAGE the DEK

upvoted 1 times

  **Bwitch** 1 year ago

**Selected Answer: C**

"on-premise" has to be a customer-supplied key. Everything else is from the GCP platform.

upvoted 3 times

  **major\_querty** 1 year, 1 month ago

The answer is C

The question says: you want to use a key generated on-premise", so A and B must be wrong.

Between DEK and KEK, the documentation says the following:

The process of encrypting data is to generate a DEK locally, encrypt data with the DEK, use a KEK to wrap the DEK, and then store the encrypted data and the wrapped DEK. The KEK never leaves Cloud KMS.

Thus, it must be C.

[https://cloud.google.com/kms/docs/envelope-encryption#how\\_to\\_encrypt\\_data\\_using\\_envelope\\_encryption](https://cloud.google.com/kms/docs/envelope-encryption#how_to_encrypt_data_using_envelope_encryption)

upvoted 3 times

  **idtroo** 1 year, 8 months ago

Ans: D

[https://cloud.google.com/kms/docs/envelope-encryption#other\\_options](https://cloud.google.com/kms/docs/envelope-encryption#other_options)

"With CSEK, you supply your own AES-256 key to serve as the KEK, and your key protects the DEKs that protect your data. Your CSEK key is protected by an additional layer of protection, using a Cloud KMS key."

upvoted 2 times

  **rr4444** 11 months, 3 weeks ago

Nope cos being used as the KEK means that this CSEK KEK is used to MANAGE the DEK

upvoted 1 times

  **iamoct** 1 year, 9 months ago

The important message is "to manage"

So, CSEK used to manage DEK

upvoted 3 times

  **rr4444** 11 months, 3 weeks ago

Exactly

upvoted 1 times

 **CloudTrip** 1 year, 10 months ago

You provide a raw CSEK as part of an API call. This key is transmitted from the Google front end to the storage system's memory. This key is used as the key encryption key in Google Cloud Storage for your data.

The raw CSEK is used to unwrap wrapped chunk keys, to create raw chunk keys in memory. These are used to decrypt data chunks stored in the storage systems. These keys are used as the data encryption keys (DEK) in Google Cloud Storage for your data.

As the question mentions about Cloud Storage so may be Answer will be C  
upvoted 2 times

Question #77

Topic 1


Last week, a company deployed a new App Engine application that writes logs to BigQuery. No other workloads are running in the project. You need to validate that all data written to BigQuery was done using the App Engine Default Service Account.

What should you do?

- A. 1. Use Cloud Logging and filter on BigQuery Insert Jobs. 2. Click on the email address in line with the App Engine Default Service Account in the authentication field. 3. Click Hide Matching Entries. 4. Make sure the resulting list is empty.
- B. 1. Use Cloud Logging and filter on BigQuery Insert Jobs. 2. Click on the email address in line with the App Engine Default Service Account in the authentication field. 3. Click Show Matching Entries. 4. Make sure the resulting list is empty.
- C. 1. In BigQuery, select the related dataset. 2. Make sure that the App Engine Default Service Account is the only account that can write to the dataset.
- D. 1. Go to the Identity and Access Management (IAM) section of the project. 2. Validate that the App Engine Default Service Account is the only account that has a role that can write to BigQuery.

**Correct Answer: C**

Community vote distribution



 **AwesomeGCP** Highly Voted  2 months, 2 weeks ago

**Selected Answer: A**

- A. 1. Use StackDriver Logging and filter on BigQuery Insert Jobs.
  - 2. Click on the email address in line with the App Engine Default Service Account in the authentication field.
  - 3. Click Hide Matching Entries.
  - 4. Make sure the resulting list is empty.
- upvoted 7 times

 **PST21** Most Recent  4 days, 18 hours ago

A is correct as last 2 are means of doing it rather than validating it  
upvoted 1 times

 **shayke** 2 months, 1 week ago

**Selected Answer: C**

validate - C  
upvoted 1 times

 **tangac** 3 months, 2 weeks ago

**Selected Answer: A**

<https://www.examtopycs.com/discussions/google/view/32259-exam-professional-cloud-security-engineer-topic-1-question/>  
upvoted 3 times

Your team wants to limit users with administrative privileges at the organization level.

Which two roles should your team restrict? (Choose two.)

- A. Organization Administrator
- B. Super Admin
- C. GKE Cluster Admin
- D. Compute Admin
- E. Organization Role Viewer

**Correct Answer:** AB

Reference:

<https://cloud.google.com/resource-manager/docs/creating-managing-organization>

*Community vote distribution*

AB (100%)

🗲️ 👤 **HateMicrosoft** Highly Voted 👍 1 year, 9 months ago

The correct answer is : A&B  
-resourcemanager.organizationAdmin  
-Cloud Identity super admin(Old G-Suite Google Workspace)  
upvoted 12 times

🗲️ 👤 **AwesomeGCP** Most Recent ⌚ 2 months, 2 weeks ago

**Selected Answer: AB**

A. Organization Administrator  
B. Super Admin  
upvoted 2 times

🗲️ 👤 **AzureDP900** 1 month, 2 weeks ago

AB is correct  
upvoted 1 times

🗲️ 👤 **Bingo21** 1 year, 10 months ago

It says "limit users with administrative privileges" - D doesn't give you admin privileges. AB is the closest to what the question is looking for.  
upvoted 3 times

🗲️ 👤 **[Removed]** 2 years, 1 month ago

Ans - AB  
upvoted 2 times

🗲️ 👤 **MohitA** 2 years, 3 months ago

AB are the one  
upvoted 4 times

🗲️ 👤 **singhjoga** 1 year, 11 months ago

There is no such role as "Super Admin". There is a Super Admin user, which has the "Owner" role to the whole Organisation.  
Answer is probably A and D.  
upvoted 4 times



An organization's security and risk management teams are concerned about where their responsibility lies for certain production workloads they are running in

Google Cloud and where Google's responsibility lies. They are mostly running workloads using Google Cloud's platform-as-a-Service (PaaS) offerings, including


App Engine primarily.

Which area in the technology stack should they focus on as their primary responsibility when using App Engine?

- A. Configuring and monitoring VPC Flow Logs
- B. Defending against XSS and SQLi attacks
- C. Managing the latest updates and security patches for the Guest OS
- D. Encrypting all stored data

**Correct Answer: D**

*Community vote distribution*



🗳️ 👤 **AzureDP900** 1 month, 2 weeks ago

B is correct  
upvoted 2 times

🗳️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: B**

B. Defending against XSS and SQLi attacks  
Data at rest is encrypted by default by Google. So D is wrong. Should be B.  
upvoted 3 times

🗳️ 👤 **koko2314** 3 months ago

Answer should be D. For SAAS solutions web based attacks are managed by Google. We just need to take care of the data as per the link below.  
upvoted 1 times

🗳️ 👤 **Random\_Mane** 3 months ago

**Selected Answer: B**

B. in PaaS the customer is responsible for web app security, deployment, usage, access policy, and content.  
<https://cloud.google.com/architecture/framework/security/shared-responsibility-shared-fate>  
upvoted 3 times

🗳️ 👤 **GHOST1985** 3 months, 1 week ago

**Selected Answer: D**

Answer is D  
upvoted 1 times

🗳️ 👤 **GHOST1985** 2 months, 3 weeks ago

In PaaS, we're responsible for more controls than in IaaS, including network controls. You share responsibility with us for application-level controls and IAM management. You remain responsible for your data security and client protection.  
[https://cloud.google.com/architecture/framework/security/shared-responsibility-shared-fate#defined\\_by\\_workloads](https://cloud.google.com/architecture/framework/security/shared-responsibility-shared-fate#defined_by_workloads)  
upvoted 2 times

🗳️ 👤 **tifo16** 1 week, 3 days ago

Data at rest is encrypted by default by Google. So D is wrong. As mentioned by your link it Should be B.  
upvoted 1 times

🗳️ 👤 **satyajeet1999** 3 months, 2 weeks ago

**Selected Answer: B**

B it is.  
upvoted 3 times

An engineering team is launching a web application that will be public on the internet. The web application is hosted in multiple GCP regions and will be directed to the respective backend based on the URL request.

Your team wants to avoid exposing the application directly on the internet and wants to deny traffic from a specific list of malicious IP addresses. Which solution should your team implement to meet these requirements?

- A. Cloud Armor
- B. Network Load Balancing
- C. SSL Proxy Load Balancing
- D. NAT Gateway

Correct Answer: A

Reference:

<https://cloud.google.com/armor/docs/security-policy-concepts>

Community vote distribution



**DebasishLowes** Highly Voted 1 year, 9 months ago  
Ans : A  
upvoted 7 times

**BillBaits** 1 year, 1 month ago  
Think so  
upvoted 1 times

**shayke** Most Recent 1 day, 17 hours ago  
Selected Answer: A  
A Is the only ans because you are asked to limit access by IP and CA is the only option  
upvoted 1 times

**AzureDP900** 1 month, 2 weeks ago  
This is straight forward question, A is right  
upvoted 1 times

**AwesomeGCP** 2 months, 2 weeks ago  
Selected Answer: A  
A. Cloud Armor  
upvoted 1 times

**cloudprincipal** 6 months, 3 weeks ago  
Selected Answer: A  
<https://cloud.google.com/armor/docs/security-policy-overview#edge-security>  
upvoted 1 times

**[Removed]** 2 years, 1 month ago  
Ans - A  
upvoted 3 times

**mlyu** 2 years, 3 months ago  
Definitly B  
upvoted 2 times

**ownez** 2 years, 3 months ago  
Should be A? Cloud armor can deny traffic by defining IP addresses list rule and to avoid exposing the application directly on the internet.

While Network LB is using Google Cloud firewalls to control or filter access to the backend VMs.

Answer is A.  
upvoted 5 times

**mlyu** 2 years, 2 months ago  
you are correct. Answer is A  
The Cloud armor able to directed user traffic to an external HTTP(S) load balancer enters the PoP closest to the user in Premium Tier.  
<https://cloud.google.com/armor/docs/security-policy-overview#edge-security>



upvoted 5 times

A customer is running an analytics workload on Google Cloud Platform (GCP) where Compute Engine instances are accessing data stored on Cloud Storage.

Your team wants to make sure that this workload will not be able to access, or be accessed from, the internet.

Which two strategies should your team use to meet these requirements? (Choose two.)

- A. Configure Private Google Access on the Compute Engine subnet
- B. Avoid assigning public IP addresses to the Compute Engine cluster.
- C. Make sure that the Compute Engine cluster is running on a separate subnet.
- D. Turn off IP forwarding on the Compute Engine instances in the cluster.
- E. Configure a Cloud NAT gateway.

Correct Answer: AB

Community vote distribution



**MohitA** Highly Voted 2 years, 3 months ago  
AB suits well  
upvoted 19 times

**DebasishLowes** Highly Voted 1 year, 9 months ago  
Ans : AB  
upvoted 6 times

**AzureDP900** Most Recent 1 month, 2 weeks ago  
AB is correct  
A. Configure Private Google Access on the Compute Engine subnet  
B. Avoid assigning public IP addresses to the Compute Engine cluster.  
upvoted 1 times

**AwesomeGCP** 2 months, 2 weeks ago  
Selected Answer: AB  
A. Configure Private Google Access on the Compute Engine subnet  
B. Avoid assigning public IP addresses to the Compute Engine cluster.  
upvoted 1 times

**cloudprincipal** 6 months, 3 weeks ago  
Selected Answer: AB  
agree with all the others  
upvoted 2 times

**pfilourenco** 1 year, 7 months ago  
B and E:  
"make sure that this workload will not be able to access, or be accessed from, the internet."  
If we have cloud NAT we are able to access the internet! Also with public IP.  
upvoted 1 times

**Rupo7** 10 months, 1 week ago  
The question says " not be able to access, or be accessed from, the internet." A NAT gateway enables access to the internet, just behind a static IP. A. Private access for the subnet is required to enable access to GCS. B is a good measure, as then the instance cannot access the internet at all (without a NAT Gateway that is).  
upvoted 1 times

**[Removed]** 1 year, 8 months ago  
Not A <https://cloud.google.com/vpc/docs/private-google-access>  
upvoted 1 times

**[Removed]** 1 year, 8 months ago  
NOt D, because by de fault IP forwarding is disabled. You do not need to turn it off.  
upvoted 1 times

**[Removed]** 1 year, 8 months ago  
So B and E is the right answer.  
upvoted 2 times

  **tanfromvn** 1 year, 5 months ago

A\_B, why not A? Private access just accepts traffic in GCP and to GG API  
upvoted 2 times

  **ffdd1234** 1 year, 11 months ago

if you Avoid assigning public IP addresses to the Compute Engine cluster the instance could access to internet if have a nat gateway, maybe the answer is A and D  
upvoted 1 times

  **ffdd1234** 1 year, 1 month ago

+1 A-D  
upvoted 1 times

  **ffdd1234** 1 year, 1 month ago

But not sure "Ensure that IP Forwarding feature is not enabled at the Google Compute Engine instance level for security and compliance reasons, as instances with IP Forwarding enabled act as routers/packet forwarders."  
IP FW is for route packets could not be D  
upvoted 1 times

  **Topsy** 2 years ago

A and B is correct  
upvoted 4 times

  **[Removed]** 2 years, 1 month ago

Ans - AB  
upvoted 2 times

  **genesis3k** 2 years, 1 month ago

AB is the correct answer.  
upvoted 1 times

  **Wooky** 2 years, 3 months ago

B,D not A  
Private google access provides public google api access without public IP  
upvoted 1 times

  **Wooky** 2 years, 2 months ago

My mistake, ans is AB.  
upvoted 2 times

  **Raushanr** 2 years, 3 months ago

Answer-AB  
upvoted 2 times

A customer wants to run a batch processing system on VMs and store the output files in a Cloud Storage bucket. The networking and security teams have decided that no VMs may reach the public internet.  
How should this be accomplished?

- A. Create a firewall rule to block internet traffic from the VM.
- B. Provision a NAT Gateway to access the Cloud Storage API endpoint.
- C. Enable Private Google Access.
- D. Mount a Cloud Storage bucket as a local filesystem on every VM.

Correct Answer: B

Community vote distribution



- tanfromvn**

Highly Voted

1 year, 5 months ago

C-there is no traffic to outside internet  
upvoted 13 times
- mynk29**

9 months, 4 weeks ago

Private google access is enabled at subnet level not at VPC level.  
upvoted 1 times
- Meyucho**

Most Recent

4 days, 8 hours ago

C!!!! This example is just the exact and only meaning for have PGA!!!  
upvoted 1 times
- TonytheTiger**

4 weeks ago

Answer C:  
Here is why; the VM need to access google service i.e. "Cloud Storage Bucket".  
Google doc states: Private Google Access permits access to Google APIs and services in Google's production infrastructure  
<https://cloud.google.com/vpc/docs/private-google-access>  
Everyone is reading the question as limited access to public internet but is missing the 2nd part of the question, which is access a google services.  
ONLY enable Private Google Access will fulfil the requirement.  
upvoted 1 times
- Littleivy**

1 month, 1 week ago

Selected Answer: C

C is the answer  
upvoted 1 times
- rotorclear**

2 months ago

Selected Answer: C

The ask is to access cloud storage while doing the batch processing not how to block the internet.  
Overall it's a poor choice of words in the question attempting to confuse than check knowledge  
upvoted 1 times
- AzureDP900**

1 month, 2 weeks ago

C is right  
upvoted 1 times
- AwesomeGCP**

2 months, 2 weeks ago

C. Enable Private Google Access on the VPC.  
upvoted 1 times
- nilopo**

2 months, 3 weeks ago

Selected Answer: C

The ask is to store the output files in a Cloud storage bucket. "The networking and security teams have decided that no VMs may reach the public internet" - No VMs MAY reach public internet but not 'MUST'. Hence 'C' is the answer  
upvoted 3 times
- piyush\_1982**

4 months, 3 weeks ago

Selected Answer: A

Let's go through each option :  
A. Correct as this will deny traffic going out to the internet.  
B. Wrong as this won't prevent VM to reach the internet. In order to access the storage API via NAT gateway, the traffic may still traverse the

internet.

C. Wrong as private google access is enabled at Subnet level and not VPC.

D. Wrong as it doesn't restrict VM traffic to reach the internet. It will only allow accessing storage bucket locally but other traffic may still go to the internet.

upvoted 1 times

  **ZODOGAM** 1 year ago

100% A. Question is 'The networking and security teams have decided that no VMs may reach the public internet.' the ask is to deny VMs from reaching the public internet. I'd go with A ... the rest of the connection (to GCS bucket) causing the confusion.

upvoted 3 times

  **danielklein09** 10 months, 3 weeks ago

Are we sure that VMs have public IP ? If they don't, than why to block outbound ?

upvoted 1 times

  **giovy\_82** 3 months, 2 weeks ago

yes but nobody is telling (in the question and in general) that somebody can go there and configure a public IP. with FW rule you can deny all egress traffic easily.

upvoted 1 times

  **tzKhalil** 1 year, 7 months ago

C is wired, because google private access is on subnet level...

D is not a recommended way but it should work...

If I have this in exam, I will go D...

upvoted 3 times

  **ThisisJohn** 1 year ago

But keep in mind D will not prevent your VMs from accessing any Internet site

upvoted 1 times

  **sleekdunga** 1 year, 8 months ago

Quite tricky because you can only enable Google private access via the subnet config and not VPC. But the keyword also says can't traverse the internet which makes C only option with skepticism.

upvoted 1 times

  **[Removed]** 1 year, 8 months ago

private google access allow VMs to connect to internet API without a public IP.

upvoted 1 times

  **[Removed]** 1 year, 9 months ago

C is the Answer. Since Networking team has already decid(ed) ,I will go with C ,as this is the option stating solution which will enable customer to get its work done.

upvoted 3 times

  **ThisisJohn** 1 year ago

Keep in mind C will not prevent your VMs from accessing any Internet site, as by default outbound access is allowed.

upvoted 1 times

  **Bingo21** 1 year, 10 months ago

Question is 'The networking and security teams have decided that no VMs may reach the public internet.' the ask is to deny VMs from reaching the public internet. I'd go with A ... the rest of the connection (to GCS bucket) causing the confusion.

upvoted 3 times

  **gu9singg** 1 year, 9 months ago

A: is incorrect

upvoted 1 times

  **gu9singg** 1 year, 9 months ago

we need to reduce operations work and use google best practice

upvoted 1 times

  **CloudTrip** 1 year, 10 months ago

Here C could be a viable solution but answer B (NAT Gateway) looks more precise as you are intending to connect to a precise target destination than a wider VPC and it will also ensure that the VM traffic is not going to external internet. Refer more details on NAT Gateway here :

<https://cloud.google.com/nat/docs/overview>

upvoted 2 times

  **cuban123** 1 year, 12 months ago

Why not consider option A? Deny all traffic from VM to internet. Have connectivity to Cloud storage still available. Satisfies the conditions in question.

upvoted 2 times

  **Topsy** 2 years ago

Answer is C- Since Network and Security Team has decided to disable internet access, the only option left is the Google Private Access option

upvoted 3 times

  **ThisisJohn** 1 year ago

Keep in mind C will not prevent your VMs from accessing any Internet site, as by default outbound access is allowed.  
upvoted 1 times

## Question #83

Topic 1

As adoption of the Cloud Data Loss Prevention (Cloud DLP) API grows within your company, you need to optimize usage to reduce cost. Cloud DLP target data is stored in Cloud Storage and BigQuery. The location and region are identified as a suffix in the resource name. Which cost reduction options should you recommend?

- A. Set appropriate rowsLimit value on BigQuery data hosted outside the US and set appropriate bytesLimitPerFile value on multiregional Cloud Storage buckets.
- B. Set appropriate rowsLimit value on BigQuery data hosted outside the US, and minimize transformation units on multiregional Cloud Storage buckets.
- C. Use rowsLimit and bytesLimitPerFile to sample data and use CloudStorageRegexFileSet to limit scans.
- D. Use FindingLimits and TimespanConfig to sample data and minimize transformation units.

**Correct Answer:** C

Reference:

<https://cloud.google.com/dlp/docs/reference/rest/v2/InspectJobConfig>

Community vote distribution

C (100%)

  **[Removed]** **Highly Voted**  2 years, 1 month ago

Ans - C

<https://cloud.google.com/dlp/docs/inspecting-storage#sampling>

[https://cloud.google.com/dlp/docs/best-practices-costs#limit\\_scans\\_of\\_files\\_in\\_to\\_only\\_relevant\\_files](https://cloud.google.com/dlp/docs/best-practices-costs#limit_scans_of_files_in_to_only_relevant_files)

upvoted 8 times

  **[Removed]** 2 years, 1 month ago

<https://cloud.google.com/dlp/docs/inspecting-storage#limiting-gcs>

upvoted 1 times

  **AzureDP900** **Most Recent**  1 month, 2 weeks ago

C is right

upvoted 1 times

  **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: C**

C . Use rowsLimit and bytesLimitPerFile to sample data and use CloudStorageRegexFileSet to limit scans.

upvoted 2 times

  **cloudprincipal** 6 months, 3 weeks ago

**Selected Answer: C**

<https://cloud.google.com/dlp/docs/inspecting-storage#sampling>

upvoted 2 times

  **passtest100** 2 years, 2 months ago

C is the right one.

upvoted 4 times



Your team uses a service account to authenticate data transfers from a given Compute Engine virtual machine instance of to a specified Cloud Storage bucket. An engineer accidentally deletes the service account, which breaks application functionality. You want to recover the application as quickly as possible without compromising security.

What should you do?

- A. Temporarily disable authentication on the Cloud Storage bucket.
- B. Use the undelete command to recover the deleted service account.
- C. Create a new service account with the same name as the deleted service account.
- D. Update the permissions of another existing service account and supply those credentials to the applications.

**Correct Answer:** B

Reference:

[https://cloud.google.com/iam/docs/creating-managing-service-accounts#undeleting\\_a\\_service\\_account](https://cloud.google.com/iam/docs/creating-managing-service-accounts#undeleting_a_service_account)

*Community vote distribution*

B (100%)

🗲️ 👤 **DebasishLowes** Highly Voted 👍 1 year, 9 months ago

Ans : B

upvoted 8 times

🗲️ 👤 **AwesomeGCP** Most Recent 🕒 2 months, 2 weeks ago

**Selected Answer: B**

B. Use the undelete command to recover the deleted service account.

upvoted 1 times

🗲️ 👤 **[Removed]** 2 years, 1 month ago

Ans - B

upvoted 2 times

🗲️ 👤 **saurabh1805** 2 years, 1 month ago

B is correct answer here.

<https://cloud.google.com/iam/docs/reference/rest/v1/projects.serviceAccounts/undelete>

upvoted 4 times

🗲️ 👤 **AzureDP900** 1 month, 2 weeks ago

Thank you for sharing link, I agree B is right

upvoted 1 times

🗲️ 👤 **MohitA** 2 years, 3 months ago

B is the Answer

upvoted 2 times

You are the Security Admin in your company. You want to synchronize all security groups that have an email address from your LDAP directory in Cloud IAM.

What should you do?

- A. Configure Google Cloud Directory Sync to sync security groups using LDAP search rules that have "user email address" as the attribute to facilitate one-way sync.
- B. Configure Google Cloud Directory Sync to sync security groups using LDAP search rules that have "user email address" as the attribute to facilitate bidirectional sync.
- C. Use a management tool to sync the subset based on the email address attribute. Create a group in the Google domain. A group created in a Google domain will automatically have an explicit Google Cloud Identity and Access Management (IAM) role.
- D. Use a management tool to sync the subset based on group object class attribute. Create a group in the Google domain. A group created in a Google domain will automatically have an explicit Google Cloud Identity and Access Management (IAM) role.

**Correct Answer:** A

*Community vote distribution*

A (100%)

 **sudarchary** Highly Voted 10 months, 3 weeks ago

**Selected Answer: A**

search rules that have "user email address" as the attribute to facilitate one-way sync.

Reference Links:

<https://support.google.com/a/answer/6126589?hl=en>

upvoted 9 times

 **PST21** Most Recent 4 days, 18 hours ago

A is correct as it should be one way sync - LDAP -> Cloud Identity via GCDS

upvoted 1 times

 **AzureDP900** 1 month, 2 weeks ago

A is correct

upvoted 2 times

 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: A**

A. Configure Google Cloud Directory Sync to sync security groups using LDAP search rules that have "user email address" as the attribute to facilitate one-way sync.

upvoted 1 times

 **JoseMaria111** 3 months ago

GCDS allow sync ldap users in one way. A is correct

upvoted 2 times

 **[Removed]** 1 year, 8 months ago

Why A is not correct? GCP provide this sync tool.

upvoted 2 times

 **mistryminded** 1 year ago

Incorrect. GCDS is Google Workspace Admin tool.

Correct answer is A. GCDS only syncs one way - <https://support.google.com/a/answer/106368?hl=en>

upvoted 4 times

 **DebasishLowes** 1 year, 9 months ago

Ans : A

upvoted 2 times

 **[Removed]** 2 years, 1 month ago

Ans - A

upvoted 2 times

 **saaurabh1805** 2 years, 1 month ago

A is correct answer here.

upvoted 2 times



  **passtest100** 2 years, 2 months ago

Answer - A

upvoted 2 times

  **skshak** 2 years, 3 months ago

Answer - A

upvoted 2 times

You are part of a security team investigating a compromised service account key. You need to audit which new resources were created by the service account.

What should you do?

- A. Query Data Access logs.
- B. Query Admin Activity logs.
- C. Query Access Transparency logs.
- D. Query Stackdriver Monitoring Workspace.

**Correct Answer:** A

Reference:

<https://cloud.google.com/iam/docs/audit-logging/examples-service-accounts>

Community vote distribution



**MohitA** Highly Voted 2 years, 3 months ago

B is the Ans  
upvoted 11 times

**ownez** 2 years, 3 months ago

Shouldn't it be A? The question is about which resources were created by the SA.

B (Admin Activity logs) cannot view this. It is only for user's activity such as create, modify or delete a particular SA.  
upvoted 1 times

**FatCharlie** 2 years ago

"Admin Activity audit logs contain log entries for API calls or other administrative actions that modify the configuration or metadata of resources. For example, these logs record when users create VM instances or change Identity and Access Management permissions".

This is exactly what you want to see. What resources were created by the SA?

<https://cloud.google.com/logging/docs/audit#admin-activity>  
upvoted 5 times

**AzureDP900** 1 month, 2 weeks ago  
B is right . Agree with your explanation  
upvoted 1 times

**Fellipo** 2 years, 1 month ago  
B it's OK  
upvoted 4 times

**VicF** Highly Voted 1 year, 8 months ago

Ans B  
"B" is for actions that modify the configuration or metadata of resources. For example, these logs record when users create VM instances or change Identity and Access Management permissions.  
"A" is only for "user-provided" resource data. Data Access audit logs-- except for BigQuery Data Access audit logs-- "are disabled by default"  
upvoted 5 times

**Meyucho** Most Recent 4 days, 9 hours ago

Selected Answer: B  
B - Audit logs. They have all the API calls that creates, modify or destroy resources. <https://cloud.google.com/logging/docs/audit#admin-activity>  
upvoted 1 times

**AwesomeGCP** 2 months, 2 weeks ago

Selected Answer: B  
B. Query Admin Activity logs.  
upvoted 1 times

**JoseMaria111** 3 months ago  
Admin activity log records resources changes. B is correct  
upvoted 1 times

**piyush\_1982** 4 months, 3 weeks ago

**Selected Answer: B**

Admin activity logs are always created to log entries for API calls or other actions that modify the configuration or metadata of resources. For example, these logs record when users create VM instances or change Identity and Access Management permissions.

upvoted 2 times

  **cloudprincipal** 6 months, 3 weeks ago

**Selected Answer: B**

Admin activity logs contain all GCP API calls.

So this is where the service account activity will show up

upvoted 1 times

  **[Removed]** 1 year, 8 months ago

I support B, <https://cloud.google.com/iam/docs/audit-logging>  
says IAM logs write into admin log

upvoted 4 times

  **DebasishLowes** 1 year, 9 months ago

Ans : B

upvoted 3 times

  **[Removed]** 2 years, 1 month ago

Ans - B

upvoted 4 times

You have an application where the frontend is deployed on a managed instance group in subnet A and the data layer is stored on a mysql Compute Engine virtual machine (VM) in subnet B on the same VPC. Subnet A and Subnet B hold several other Compute Engine VMs. You only want to allow the application frontend to access the data in the application's mysql instance on port 3306.

What should you do?

- A. Configure an ingress firewall rule that allows communication from the src IP range of subnet A to the tag "data-tag" that is applied to the mysql Compute Engine VM on port 3306.
- B. Configure an ingress firewall rule that allows communication from the frontend's unique service account to the unique service account of the mysql Compute Engine VM on port 3306.
- C. Configure a network tag "fe-tag" to be applied to all instances in subnet A and a network tag "data-tag" to be applied to all instances in subnet B. Then configure an egress firewall rule that allows communication from Compute Engine VMs tagged with data-tag to destination Compute Engine VMs tagged fe- tag.
- D. Configure a network tag "fe-tag" to be applied to all instances in subnet A and a network tag "data-tag" to be applied to all instances in subnet B. Then configure an ingress firewall rule that allows communication from Compute Engine VMs tagged with fe-tag to destination Compute Engine VMs tagged with data-tag.

**Correct Answer:** B

Community vote distribution

B (100%)

🗳️ 👤 **Zuy01** Highly Voted 1 year, 4 months ago

B for sure, u can check this :

<https://cloud.google.com/sql/docs/mysql/sql-proxy#using-a-service-account>

upvoted 8 times

🗳️ 👤 **AwesomeGCP** Most Recent 2 months, 2 weeks ago

**Selected Answer: B**

B. Configure an ingress firewall rule that allows communication from the frontend's unique service account to the unique service account of the mysql ComputeEngine VM on port 3306.

upvoted 2 times

🗳️ 👤 **JoseMaria111** 3 months ago

B is correct.firellas rules using service account is better than tag based.

[https://cloud.google.com/vpc/docs/firewalls#best\\_practices\\_for\\_firewall\\_rules](https://cloud.google.com/vpc/docs/firewalls#best_practices_for_firewall_rules)

upvoted 1 times

🗳️ 👤 **mT3** 7 months, 1 week ago

**Selected Answer: B**

Ans : B

upvoted 3 times

🗳️ 👤 **major\_querty** 1 year, 1 month ago

why is it not a?

a seems straight forward

The link which Zuy01 provided for answer b states: For this reason, using a service account is the recommended method for production instances NOT running on a Compute Engine instance.

upvoted 4 times

🗳️ 👤 **Loved** 1 month, 1 week ago

But answer A says "communication from the src IP range of subnet A"... this rules include all the instances on subnet A, while you have to consider only the frontend

upvoted 1 times

🗳️ 👤 **Arturo\_Cloud** 3 months, 3 weeks ago

I agree (A), it is planned to limit a MySQL server in Compute Engine (IaaS) not in Cloud SQL (PaaS), so Networks Tags is the most common and recommended to use. Don't get confused with the services....

upvoted 2 times

🗳️ 👤 **DebasishLowes** 1 year, 9 months ago

Ans : B

upvoted 2 times

  **dtmtor** 1 year, 9 months ago

ans is B

upvoted 2 times

  **[Removed]** 2 years, 1 month ago

Ans - B

upvoted 4 times

  **Rantu** 2 years, 2 months ago

B is correct

upvoted 4 times

Your company operates an application instance group that is currently deployed behind a Google Cloud load balancer in us-central-1 and is configured to use the Standard Tier network. The infrastructure team wants to expand to a second Google Cloud region, us-east-2. You need to set up a single external IP address to distribute new requests to the instance groups in both regions. What should you do?

- A. Change the load balancer backend configuration to use network endpoint groups instead of instance groups.
- B. Change the load balancer frontend configuration to use the Premium Tier network, and add the new instance group.
- C. Create a new load balancer in us-east-2 using the Standard Tier network, and assign a static external IP address.
- D. Create a Cloud VPN connection between the two regions, and enable Google Private Access.

**Correct Answer:** A

*Community vote distribution*

B (100%)

  **Fellipo** Highly Voted 2 years, 1 month ago

In Premium Tier: Backends can be in any region and any VPC network.

In Standard Tier: Backends must be in the same region as the forwarding rule, but can be in any VPC network.  
upvoted 11 times

  **AzureDP900** 1 month, 2 weeks ago

B is right  
upvoted 2 times

  **mlyu** Highly Voted 2 years, 3 months ago

Should be B  
In Standard Tier LB, Backends must be in the same region  
[https://cloud.google.com/load-balancing/docs/load-balancing-overview#backend\\_region\\_and\\_network](https://cloud.google.com/load-balancing/docs/load-balancing-overview#backend_region_and_network)  
upvoted 7 times

  **Meyucho** Most Recent 4 days, 8 hours ago

**Selected Answer: B**

Global load balancers require Premium Tier!  
upvoted 1 times

  **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: B**

B. Change the load balancer frontend configuration to use the Premium Tier network, and add the new instance group.  
upvoted 1 times

  **cloudprincipal** 6 months, 3 weeks ago

**Selected Answer: B**

<https://cloud.google.com/load-balancing/docs/choosing-load-balancer#global-regional>  
upvoted 1 times

  **DebasishLowes** 1 year, 9 months ago

Ans : B  
upvoted 2 times

  **saurabh1805** 2 years, 1 month ago

I will also go with Option B  
upvoted 6 times



You are the security admin of your company. You have 3,000 objects in your Cloud Storage bucket. You do not want to manage access to each object individually.

You also do not want the uploader of an object to always have full control of the object. However, you want to use Cloud Audit Logs to manage access to your bucket.

What should you do?

- A. Set up an ACL with OWNER permission to a scope of allUsers.
- B. Set up an ACL with READER permission to a scope of allUsers.
- C. Set up a default bucket ACL and manage access for users using IAM.
- D. Set up Uniform bucket-level access on the Cloud Storage bucket and manage access for users using IAM.

Correct Answer: A

Reference:

<https://cloud.google.com/storage/docs/access-control/lists>

Community vote distribution



- Fellipo** Highly Voted 2 years, 1 month ago

it's D, <https://cloud.google.com/storage/docs/uniform-bucket-level-access#:~:text=When%20you%20enable%20uniform%20bucket,and%20the%20objects%20it%20contains>.

upvoted 17 times
- AzureDP900** Most Recent 1 month, 2 weeks ago

D is right

upvoted 1 times
- AwesomeGCP** 2 months, 2 weeks ago

Selected Answer: D

D. Set up Uniform bucket-level access on the Cloud Storage bucket and manage access for users using IAM.

upvoted 2 times
- cloudprincipal** 6 months, 3 weeks ago

Selected Answer: D

<https://cloud.google.com/storage/docs/uniform-bucket-level-access#enabled>

upvoted 1 times
- ramravella** 1 year, 5 months ago

Answer is A. Read the note below in the below URL

<https://cloud.google.com/storage/docs/access-control/lists>

Note: You cannot grant discrete permissions for reading or writing ACLs or other metadata. To allow someone to read and write ACLs, you must grant them OWNER permission.

upvoted 1 times
- Zuy01** 1 year, 4 months ago

the question mention "do not want the uploader of an object to always have full control of the object" that's mean you shouldn't grant the owner permission, hence the best ans is D.

upvoted 3 times
- [Removed]** 1 year, 8 months ago

A grants Owner???too much for this.

upvoted 2 times
- [Removed]** 2 years, 1 month ago

Ans - D

upvoted 3 times
- saurabh1805** 2 years, 1 month ago

I will go with uniform level access and manage access via IAM,

Hence D.

upvoted 2 times

🗨️ 👤 **passtest100** 2 years, 2 months ago

SHOULD BE D  
upvoted 2 times

🗨️ 👤 **skshak** 2 years, 3 months ago

Answer C <https://cloud.google.com/storage/docs/access-control>  
Uniform (recommended): Uniform bucket-level access allows you to use Identity and Access Management (IAM) alone to manage permissions. IAM applies permissions to all the objects contained inside the bucket or groups of objects with common name prefixes. IAM also allows you to use features that are not available when working with ACLs, such as IAM Conditions and Cloud Audit Logs.  
upvoted 1 times

🗨️ 👤 **skshak** 2 years, 2 months ago

Sorry, It is D. It was typo.  
upvoted 3 times

🗨️ 👤 **mlyu** 2 years, 2 months ago

the question stated they need cloud audit log for the GCS access, however uniform bucket-level access has restriction on the cloud audit log. See <https://cloud.google.com/storage/docs/uniform-bucket-level-access>  
The following restrictions apply when using uniform bucket-level access:  
Cloud Logging and Cloud Audit Logs cannot export to buckets that have uniform bucket-level access enabled.  
upvoted 1 times

🗨️ 👤 **FatCharlie** 2 years ago

They're not saying they want to export the logs to the bucket. They're just saying they want to "use Cloud Audit Logs to manage access to your bucket" (whatever that means).  
upvoted 1 times



You are the security admin of your company. Your development team creates multiple GCP projects under the "implementation" folder for several dev, staging, and production workloads. You want to prevent data exfiltration by malicious insiders or compromised code by setting up a security perimeter. However, you do not want to restrict communication between the projects.

What should you do?

- A. Use a Shared VPC to enable communication between all projects, and use firewall rules to prevent data exfiltration.
- B. Create access levels in Access Context Manager to prevent data exfiltration, and use a shared VPC for communication between projects.
- C. Use an infrastructure-as-code software tool to set up a single service perimeter and to deploy a Cloud Function that monitors the "implementation" folder via Stackdriver and Cloud Pub/Sub. When the function notices that a new project is added to the folder, it executes Terraform to add the new project to the associated perimeter.
- D. Use an infrastructure-as-code software tool to set up three different service perimeters for dev, staging, and prod and to deploy a Cloud Function that monitors the "implementation" folder via Stackdriver and Cloud Pub/Sub. When the function notices that a new project is added to the folder, it executes Terraform to add the new project to the respective perimeter.

**Correct Answer: B**

*Community vote distribution*

C (100%)

🗳️ **jonclem** **Highly Voted** 👍 2 years, 1 month ago

I'd also go with option B and here's why:

<https://cloud.google.com/access-context-manager/docs/overview>

Option A was a consideration until I came across this: <https://cloud.google.com/security/data-loss-prevention/preventing-data-exfiltration>  
upvoted 13 times

🗳️ **dzhu** **Highly Voted** 👍 1 year, 4 months ago

I think this is C. Communication between the project is necessary tied to VPC, but you need to include all projects under implementation folder in a single VPCSC

upvoted 9 times

🗳️ **soltium** **Most Recent** 🕒 2 months, 1 week ago

**Selected Answer: C**

C. The keyword "prevent data exfiltration by malicious insiders or compromised code" is listed as the benefits of VPC service control  
<https://cloud.google.com/vpc-service-controls/docs/overview#benefits>

Only C and D creates service perimeters, but D creates three and doesn't specify a bridge to connect those service perimeters so I choose C as the answer.

upvoted 2 times

🗳️ **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: C**

C. Use an infrastructure-as-code software tool to set up a single service perimeter and to deploy a Cloud Function that monitors the "implementation" folder via Stackdriver and Cloud Pub/Sub. When the function notices that a new project is added to the folder, it executes Terraform to add the new project to the associated perimeter.

upvoted 1 times

🗳️ **cloudprincipal** 6 months, 3 weeks ago

**Selected Answer: C**

eshtanaka is right: [https://github.com/terraform-google-modules/terraform-google-vpc-service-controls/tree/master/examples/automatic\\_folder](https://github.com/terraform-google-modules/terraform-google-vpc-service-controls/tree/master/examples/automatic_folder)  
upvoted 2 times

🗳️ **sudarchary** 10 months, 4 weeks ago

Answer is A. Please focus on "security perimeter" and "compromised code".

upvoted 1 times

🗳️ **eshtanaka** 1 year, 1 month ago

Correct answer is C. See the description for "automatically secured folder" [https://github.com/terraform-google-modules/terraform-google-vpc-service-controls/tree/master/examples/automatic\\_folder](https://github.com/terraform-google-modules/terraform-google-vpc-service-controls/tree/master/examples/automatic_folder)

upvoted 3 times

🗳️ **nilb94** 1 year, 4 months ago

Think it should be C. Access Context Manager docs say it is for ingress. Service Controls seems correct for exfiltration, and projects must be allowed to communicate with each other so they need to be in a single service perimeter.

upvoted 3 times

  **desertlotus1211** 1 year, 8 months ago

Answer is B:

<https://cloud.google.com/access-context-manager/docs/overview>

You need to read the question AND Answer carefully before selecting.

Answer A is in Answer B

upvoted 2 times

  **DebasishLowes** 1 year, 9 months ago

Ans : A. To make the communication between different projects, shared vpc is required.

upvoted 1 times

  **HateMicrosoft** 1 year, 9 months ago

The correct answer is :B

Access Context Manager

<https://cloud.google.com/access-context-manager/docs/overview>

Preventing Data Exfiltration

<https://cloud.google.com/security/data-loss-prevention/preventing-data-exfiltration>

upvoted 3 times

  **deardeer** 1 year, 10 months ago

D is the answer. This question is about sharing across perimeters. [https://cloud.google.com/vpc-service-controls/docs/share-across-perimeters#service\\_perimeter\\_bridges](https://cloud.google.com/vpc-service-controls/docs/share-across-perimeters#service_perimeter_bridges)

upvoted 2 times

  **mlx** 2 years, 1 month ago

could be D : use case needs service controls and security perimeter

upvoted 2 times

  **saaurabh1805** 2 years, 1 month ago

B is correct answer. Firewall can never stop data exfiltration.

upvoted 4 times

  **ZODOGAM** 1 year ago

This does not solve the data leak

upvoted 1 times

  **[Removed]** 2 years, 1 month ago

Ans - A

upvoted 1 times

  **Rantu** 2 years, 2 months ago

A is the answer

upvoted 1 times

  **ZODOGAM** 1 year ago

This does not solve data exfiltration

upvoted 1 times

You need to provide a corporate user account in Google Cloud for each of your developers and operational staff who need direct access to GCP resources.

Corporate policy requires you to maintain the user identity in a third-party identity management provider and leverage single sign-on. You learn that a significant number of users are using their corporate domain email addresses for personal Google accounts, and you need to follow Google recommended practices to convert existing unmanaged users to managed accounts.

Which two actions should you take? (Choose two.)

- A. Use Google Cloud Directory Sync to synchronize your local identity management system to Cloud Identity.
- B. Use the Google Admin console to view which managed users are using a personal account for their recovery email.
- C. Add users to your managed Google account and force users to change the email addresses associated with their personal accounts.
- D. Use the Transfer Tool for Unmanaged Users (TTUU) to find users with conflicting accounts and ask them to transfer their personal Google accounts.
- E. Send an email to all of your employees and ask those users with corporate email addresses for personal Google accounts to delete the personal accounts immediately.

Correct Answer: BE

Community vote distribution



- VicF** Highly Voted 1 year, 8 months ago

A&D.  
A- Requires third-party IDp and wants to leverage single sign-on.  
D- [https://cloud.google.com/architecture/identity/migrating-consumer-accounts#initiating\\_a\\_transfer](https://cloud.google.com/architecture/identity/migrating-consumer-accounts#initiating_a_transfer)  
"In addition to showing you all unmanaged accounts, the transfer tool for unmanaged users lets you initiate an account transfer by sending an account transfer request."  
upvoted 14 times
- skshak** Highly Voted 2 years, 3 months ago

Is the answer is A,D  
A - Requirement is third-party identity management provider and leverage single sign-on.  
D - <https://cloud.google.com/architecture/identity/assessing-existing-user-accounts> (Use the transfer tool for unmanaged users to identify consumer accounts that use an email address that matches one of the domains you've added to Cloud Identity or G Suite.)  
upvoted 6 times
- Littleivy** Most Recent 1 month, 1 week ago

Selected Answer: AD  
A to sync IdP  
D to transfer unmanaged accounts  
upvoted 1 times
- AzureDP900** 1 month, 2 weeks ago

AD is right  
upvoted 1 times
- AwesomeGCP** 2 months, 2 weeks ago

Selected Answer: AD  
A. Use Google Cloud Directory Sync to synchronize your local identity management system to Cloud Identity.  
D. Use the Transfer Tool for Unmanaged Users (TTUU) to find users with conflicting accounts and ask them to transfer their personal Google accounts.  
upvoted 1 times
- cloudprincipal** 6 months, 3 weeks ago

Selected Answer: AD  
see other comments  
upvoted 1 times
- sudarchary** 11 months ago

Answers are: A&C  
<https://cloud.google.com/architecture/identity/assessing-existing-user-accounts>  
upvoted 1 times
- CloudTrip** 1 year, 10 months ago

The keyword is here "convert" follow Google recommended practices to convert existing unmanaged users to managed accounts. So why sync unmanaged with Cloud Identity. I would prefer Answers C and D

upvoted 2 times

  **ThisisJohn** 1 year ago

But dont forget about "Corporate policy requires you to maintain the user identity in a third-party identity management provider".

I believe that makes it A and D

upvoted 1 times

  **mikelabs** 2 years ago

Answer is C,D. From GSuite Console you can do both.

upvoted 1 times

  **[Removed]** 2 years, 1 month ago

Ans - AD

upvoted 3 times

  **[Removed]** 2 years, 1 month ago

[https://cloud.google.com/architecture/identity/migrating-consumer-accounts#initiating\\_a\\_transfer](https://cloud.google.com/architecture/identity/migrating-consumer-accounts#initiating_a_transfer)

upvoted 5 times

  **saurabh1805** 2 years, 1 month ago

A, D is correct answer

upvoted 3 times

  **lordb** 2 years, 3 months ago

<https://cloud.google.com/architecture/identity/assessing-existing-user-accounts>

upvoted 2 times

You are on your company's development team. You noticed that your web application hosted in staging on GKE dynamically includes user data in web pages without first properly validating the inputted data. This could allow an attacker to execute gibberish commands and display arbitrary content in a victim user's browser in a production environment.

How should you prevent and fix this vulnerability?

- A. Use Cloud IAP based on IP address or end-user device attributes to prevent and fix the vulnerability.
- B. Set up an HTTPS load balancer, and then use Cloud Armor for the production environment to prevent the potential XSS attack.
- C. Use Web Security Scanner to validate the usage of an outdated library in the code, and then use a secured version of the included library.
- D. Use Web Security Scanner in staging to simulate an XSS injection attack, and then use a templating system that supports contextual auto-escaping.

**Correct Answer:** D

Reference:

<https://cloud.google.com/security-scanner/docs/remediate-findings>

Community vote distribution

D (100%)

 **sudarchary** Highly Voted 10 months, 4 weeks ago

**Selected Answer: D**

Option D is correct as using web security scanner will allow to detect the vulnerability and templating system

upvoted 7 times

 **deardeer** Highly Voted 1 year, 10 months ago

Answer is D. There is mention about simulating in Web Security Scanner. "Web Security Scanner cross-site scripting (XSS) injection testing \*simulates\* an injection attack by inserting a benign test string into user-editable fields and then performing various user actions."

<https://cloud.google.com/security-command-center/docs/how-to-remediate-web-security-scanner-findings#xss>

upvoted 5 times

 **AzureDP900** 1 month, 2 weeks ago

Agree with D

upvoted 1 times

 **ThisisJohn** 1 year ago

Agree. Also from your link

"There are various ways to fix this problem. The recommended fix is to escape all output and use a templating system that supports contextual auto-escaping."

So escaping is a way to fix the issue, which is required by the question

upvoted 1 times

 **AwesomeGCP** Most Recent 2 months, 2 weeks ago

**Selected Answer: D**

D. Use Web Security Scanner in staging to simulate an XSS injection attack, and then use a templating system that supports contextual auto-escaping.

upvoted 1 times

 **tangac** 3 months, 2 weeks ago

**Selected Answer: D**

clear D everything is explicated here : <https://cloud.google.com/security-command-center/docs/how-to-remediate-web-security-scanner-findings>  
Web Security Scanner cross-site scripting (XSS) injection testing simulates an injection attack by inserting a benign test string into user-editable fields and then performing various user actions. Custom detectors observe the browser and DOM during this test to determine whether an injection was successful and assess its potential for exploitation.

There are various ways to fix this issue. The recommended fix is to escape all output and use a templating system that supports contextual auto-escaping.

upvoted 1 times

 **Lancyqusa** 11 months, 4 weeks ago

It should be C because the web security scanner will identify the library known to contain the security issue as in the examples here -

[https://cloud.google.com/security-command-center/docs/how-to-use-web-security-scanner#example\\_findings](https://cloud.google.com/security-command-center/docs/how-to-use-web-security-scanner#example_findings) .

Once the security issue is identified, the vulnerability can be fixed by a secure version of that library.

upvoted 1 times

🗨️ 👤 **DebasishLowes** 1 year, 9 months ago

Ans : D

upvoted 2 times

🗨️ 👤 **pyc** 1 year, 10 months ago

C,

D is wrong, as Security Scanner can't "simulate" anything. It's a scanner.

B is not right, as Armor can't do input data validation, it just deny/allow IP/CIDR.

upvoted 1 times

🗨️ 👤 **desertlotus1211** 1 year, 8 months ago

Yes it can simulate... Read the documentation first...

upvoted 3 times

🗨️ 👤 **KarVaid** 2 years ago

<https://cloud.google.com/security-command-center/docs/concepts-web-security-scanner-overview>

Security Scanner should be able to scan for XSS vulnerabilities as well. Option D is better.

upvoted 2 times

🗨️ 👤 **KarVaid** 2 years ago

Cloud armor can prevent the vulnerability but to fix it, you would need Security scanner.

upvoted 1 times

🗨️ 👤 **Fellipo** 2 years, 1 month ago

B , <https://cloud.google.com/armor>

upvoted 5 times

🗨️ 👤 **[Removed]** 2 years, 1 month ago

Ans - D

upvoted 3 times

🗨️ 👤 **HectorLeon2099** 2 years, 2 months ago

Answer is B. Web Security Scanner can look for XSS vulnerabilities but can't simulate XSS injection attack.

[https://cloud.google.com/armor/docs/rule-tuning#cross-site\\_scripting\\_xss](https://cloud.google.com/armor/docs/rule-tuning#cross-site_scripting_xss)

upvoted 3 times

🗨️ 👤 **saurabh1805** 2 years, 1 month ago

Agree B is correct answer here.

upvoted 2 times

🗨️ 👤 **FatCharlie** 2 years ago

Web Security Scanner does appear to be able to simulate an XSS attack.

"Web Security Scanner cross-site scripting (XSS) injection testing simulates an injection attack by inserting a benign test string into user-editable fields and then performing various user actions. Custom detectors observe the browser and DOM during this test to determine whether an injection was successful and assess its potential for exploitation."

<https://cloud.google.com/security-command-center/docs/how-to-remediate-web-security-scanner-findings#remediate-findings>

upvoted 4 times

🗨️ 👤 **Jerrard** 2 years, 2 months ago

D. <https://cloud.google.com/security-command-center/docs/concepts-web-security-scanner-overview>

upvoted 4 times



You are part of a security team that wants to ensure that a Cloud Storage bucket in Project A can only be readable from Project B. You also want to ensure that data in the Cloud Storage bucket cannot be accessed from or copied to Cloud Storage buckets outside the network, even if the user has the correct credentials.

What should you do?

- A. Enable VPC Service Controls, create a perimeter with Project A and B, and include Cloud Storage service.
- B. Enable Domain Restricted Sharing Organization Policy and Bucket Policy Only on the Cloud Storage bucket.
- C. Enable Private Access in Project A and B networks with strict firewall rules to allow communication between the networks.
- D. Enable VPC Peering between Project A and B networks with strict firewall rules to allow communication between the networks.

**Correct Answer:** B

Reference:

<https://cloud.google.com/resource-manager/docs/organization-policy/restricting-domains>

*Community vote distribution*

A (100%)

🗳️ 👤 **FatCharlie** Highly Voted 2 years ago

The answer is A. This is question is covered by an example given for VPC Service Perimeters

<https://cloud.google.com/vpc-service-controls/docs/overview#isolate>

upvoted 19 times

🗳️ 👤 **AzureDP900** 1 month, 2 weeks ago

A is right

upvoted 2 times

🗳️ 👤 **AwesomeGCP** Most Recent 2 months, 2 weeks ago

**Selected Answer: A**

A. Enable VPC Service Controls, create a perimeter with Project A and B, and include Cloud Storage service.

upvoted 2 times

🗳️ 👤 **cloudprincipal** 6 months, 3 weeks ago

**Selected Answer: A**

<https://cloud.google.com/vpc-service-controls/docs/overview#isolate>

upvoted 1 times

🗳️ 👤 **nilb94** 1 year, 4 months ago

A - VPC Service Controls

upvoted 2 times

🗳️ 👤 **jeeet\_** 1 year, 6 months ago

Answer is most positively A.

VPC service controls lets Security team create fine-grained Perimeter across projects within organization.

-> Security perimeter for API-Based services like Bigtable instances, Storage and Bigquery datasets.. are a kind of super powers for VPC Service control.

well in my test, I chose option B, but

Domain Restricted Organization policies are for limiting resource sharing based on domain.

so if you're out in internet, and have credentials you still can access resources based on your domain access level. So B option is wrong.

upvoted 1 times

🗳️ 👤 **HateMicrosoft** 1 year, 9 months ago

The correct answer is: A

This is obtained by the VPC Service Controls by the perimeter setup.

Overview of VPC Service Controls

<https://cloud.google.com/vpc-service-controls/docs/overview>

upvoted 1 times

🗳️ 👤 **jonclem** 2 years, 1 month ago

I would say option A is a better fit due to VPC Service Controls.

upvoted 3 times

🗳️ 👤 **jonclem** 2 years, 1 month ago

I'd be inclined to agree, option B seems a better fit. Here's my reasoning behind it:  
<https://cloud.google.com/access-context-manager/docs/overview>

upvoted 1 times

  **jonclem** 2 years, 1 month ago

please ignore this comment, wrong question.

upvoted 1 times

  **saurabh1805** 2 years, 1 month ago

what is being asked is data exfiltration as well and which can be only achieved via VPC perimiter and created a bridge between both project.

upvoted 1 times

  **Ducle** 2 years, 1 month ago

A is better

upvoted 2 times

  **[Removed]** 2 years, 1 month ago

Ans - B

upvoted 1 times

  **Jerrard** 2 years, 2 months ago

B. <https://cloud.google.com/resource-manager/docs/organization-policy/restricting-domains>

upvoted 1 times




You are responsible for protecting highly sensitive data in BigQuery. Your operations teams need access to this data, but given privacy regulations, you want to ensure that they cannot read the sensitive fields such as email addresses and first names. These specific sensitive fields should only be available on a need-to-know basis to the Human Resources team. What should you do?

- A. Perform data masking with the Cloud Data Loss Prevention API, and store that data in BigQuery for later use.
- B. Perform data redaction with the Cloud Data Loss Prevention API, and store that data in BigQuery for later use.
- C. Perform data inspection with the Cloud Data Loss Prevention API, and store that data in BigQuery for later use.
- D. Perform tokenization for Pseudonymization with the Cloud Data Loss Prevention API, and store that data in BigQuery for later use.

**Correct Answer: D**

Reference:

<https://towardsdatascience.com/bigquery-pii-and-cloud-data-loss-prevention-dlp-take-it-to-the-next-level-with-data-catalog-c47c31bcf677>



**AzureDP900** 1 month, 2 weeks ago

D is correct

Pseudonymization is a de-identification technique that replaces sensitive data values with cryptographically generated tokens. Pseudonymization is widely used in industries like finance and healthcare to help reduce the risk of data in use, narrow compliance scope, and minimize the exposure of sensitive data to systems while preserving data utility and accuracy.

upvoted 2 times

**AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: D**

D. Perform tokenization for Pseudonymization with the Cloud Data Loss Prevention API, and store that data in BigQuery for later use.

upvoted 3 times

**zellick** 2 months, 3 weeks ago

**Selected Answer: D**

D is the answer as tokenization can support re-identification for use by HR.

<https://cloud.google.com/dlp/docs/pseudonymization>

upvoted 4 times

**Random\_Mane** 3 months ago

**Selected Answer: A**

A <https://cloud.google.com/bigquery/docs/column-data-masking-intro>

upvoted 3 times

You are a Security Administrator at your organization. You need to restrict service account creation capability within production environments. You want to accomplish this centrally across the organization. What should you do?

- A. Use Identity and Access Management (IAM) to restrict access of all users and service accounts that have access to the production environment.
- B. Use organization policy constraints/iam.disableServiceAccountKeyCreation boolean to disable the creation of new service accounts.
- C. Use organization policy constraints/iam.disableServiceAccountKeyUpload boolean to disable the creation of new service accounts.
- D. Use organization policy constraints/iam.disableServiceAccountCreation boolean to disable the creation of new service accounts.

**Correct Answer:** D

Reference:

<https://cloud.google.com/resource-manager/docs/organization-policy/restricting-service-accounts>

### Restrict removal of project liens when service accounts are used across projects

When you allow a project's service accounts to be attached to resources in other projects, IAM adds a **project lien** that prevents you from deleting the project. By default, anyone who has the `resourcemanager.projects.updateLiens` permission on the project can delete the lien.

If you enforce the `iam.restrictCrossProjectServiceAccountLienRemoval` boolean constraint, then principals can delete the lien only if they have the `resourcemanager.projects.updateLiens` permission on the organization.

Community vote distribution

D (100%)

 **Tabayashi** Highly Voted 7 months, 4 weeks ago

Answer is (D).

You can use the `iam.disableServiceAccountCreation` boolean constraint to disable the creation of new service accounts. This allows you to centralize management of service accounts while not restricting the other permissions your developers have on projects.

[https://cloud.google.com/resource-manager/docs/organization-policy/restricting-service-accounts#disable\\_service\\_account\\_creation](https://cloud.google.com/resource-manager/docs/organization-policy/restricting-service-accounts#disable_service_account_creation)

upvoted 8 times

 **AzureDP900** Most Recent 1 month, 2 weeks ago

D is correct

upvoted 1 times

 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: D**

D. Use organization policy constraints/iam.disableServiceAccountCreation boolean to disable the creation of new service accounts.

upvoted 1 times

 **zellick** 2 months, 3 weeks ago

**Selected Answer: D**

D is the answer.

upvoted 1 times

You are the project owner for a regulated workload that runs in a project you own and manage as an Identity and Access Management (IAM) admin. For an upcoming audit, you need to provide access reviews evidence. Which tool should you use?

- A. Policy Troubleshooter
- B. Policy Analyzer
- C. IAM Recommender
- D. Policy Simulator

**Correct Answer: A**

Reference:

<https://cloud.google.com/iam/docs/granting-changing-revoking-access>

To gain these permissions while following the principle of least privilege, ask your administrator to grant you one of the following roles:

- **To manage access to projects:** Project IAM Admin ( `roles/resourcemanager.projectIamAdmin` )
- **To manage access to projects and folders:** Folder Admin ( `roles/resourcemanager.folderAdmin` )
- **To manage access to projects, folders, and organizations:** Organization Admin ( `roles/resourcemanager.organizationAdmin` )
- **To manage access to almost all Google Cloud resources:** Security Admin ( `roles/iam.securityAdmin` )

Community vote distribution

B (100%)

 **mouchu** Highly Voted 7 months, 1 week ago  
Answer = B

<https://cloud.google.com/policy-intelligence/docs/policy-analyzer-overview>  
upvoted 6 times

 **AzureDP900** Most Recent 1 month, 2 weeks ago  
B policy analyzer is correct  
upvoted 1 times

 **AwesomeGCP** 2 months, 2 weeks ago  
Selected Answer: B  
B. Policy Analyzer  
upvoted 1 times

 **zelick** 2 months, 3 weeks ago  
Selected Answer: B  
B is the answer.

<https://cloud.google.com/policy-intelligence/docs/policy-analyzer-overview>  
Policy Analyzer lets you find out which principals (for example, users, service accounts, groups, and domains) have what access to which Google Cloud resources based on your IAM allow policies.  
upvoted 2 times

 **sumundada** 5 months, 1 week ago  
Selected Answer: B  
<https://cloud.google.com/policy-intelligence/docs/policy-analyzer-overview>  
upvoted 4 times

 **cloudprincipal** 6 months, 3 weeks ago  
Selected Answer: B  
<https://cloud.google.com/policy-intelligence/docs/policy-analyzer-overview>  
upvoted 4 times

  **szl0144** 7 months ago

B is correct, guys

upvoted 3 times

Your organization has implemented synchronization and SAML federation between Cloud Identity and Microsoft Active Directory. You want to reduce the risk of Google Cloud user accounts being compromised. What should you do?

- A. Create a Cloud Identity password policy with strong password settings, and configure 2-Step Verification with security keys in the Google Admin console.
- B. Create a Cloud Identity password policy with strong password settings, and configure 2-Step Verification with verification codes via text or phone call in the Google Admin console.
- C. Create an Active Directory domain password policy with strong password settings, and configure post-SSO (single sign-on) 2-Step Verification with security keys in the Google Admin console.
- D. Create an Active Directory domain password policy with strong password settings, and configure post-SSO (single sign-on) 2-Step Verification with verification codes via text or phone call in the Google Admin console.

**Correct Answer:** D

Reference:

<https://cloud.google.com/architecture/identity/federating-gcp-with-active-directory-introduction>

Setting up federation between Active Directory and Cloud Identity or Google Workspace entails two pieces:

- **Provisioning users:** Relevant users and groups are synchronized periodically from Active Directory to Cloud Identity or Google Workspace. This process ensures that when you create a new user in Active Directory, it can be referenced in Google Cloud even before the associated user has logged in for the first time. This process also ensures that user deletions are being propagated.

Provisioning works one way, which means changes in Active Directory are replicated to Google Cloud but not vice versa. Also, provisioning does not include passwords. In a federated setup, Active Directory remains the only system that manages these credentials.

Community vote distribution

C (100%)

🗳️ **coco10k** 1 month, 3 weeks ago

Answer C:

"We recommend against using text messages. The National Institute of Standards and Technology (NIST) no longer recommends SMS-based 2SV due to the hijacking risk from state-sponsored entities."

upvoted 2 times

🗳️ **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: C**

C. Create an Active Directory domain password policy with strong password settings, and configure post-SSO (single sign-on) 2-Step Verification with security keys in the Google Admin console.

upvoted 1 times

🗳️ **jitu028** 2 months, 3 weeks ago

Answer is - C

[https://cloud.google.com/identity/solutions/enforce-mfa#use\\_security\\_keys](https://cloud.google.com/identity/solutions/enforce-mfa#use_security_keys)

Use security keys

We recommend requiring security keys for those employees who create and access data that needs the highest level of security. You should require 2SV for all other employees and encourage them to use security keys.

Security keys offer the most secure form of 2SV. They are based on the open standard developed by Google as part of the Fast Identity Online (FIDO) Alliance. Security keys require a compatible browser on user devices.

upvoted 2 times

🗳️ **AzureDP900** 1 month, 2 weeks ago

Agree with C and explanation

upvoted 1 times

🗳️ **szl0144** 7 months ago

C is the answer because security key is securer than 2FA code

upvoted 4 times

🗳️ **mT3** 7 months, 1 week ago

**Selected Answer: C**

C: correct answer  
upvoted 4 times

  **mouchu** 7 months, 1 week ago

Answer = B  
upvoted 1 times



You have been tasked with implementing external web application protection against common web application attacks for a public application on Google Cloud.

You want to validate these policy changes before they are enforced. What service should you use?

- A. Google Cloud Armor's preconfigured rules in preview mode
- B. Prepopulated VPC firewall rules in monitor mode
- C. The inherent protections of Google Front End (GFE)
- D. Cloud Load Balancing firewall rules
- E. VPC Service Controls in dry run mode

**Correct Answer:** A

Reference:

<https://cloud.google.com/architecture/owasp-top-ten-mitigation>

## Google Cloud Armor

Use case:

- SQL injection filtering
- PHP injection filtering

Google Cloud Armor can block common injection attacks before they reach your application. For SQL injection (SQLi), Google Cloud Armor has a [predefined rule set](#) that is based on the [OWASP Modsecurity core rule set](#) [\[1\]](#). You can [build security policies](#) that [block common SQLi attacks](#) defined in the core rule set by using the `evaluatePreconfiguredExpr('sqli-stable')` rule either by itself or in conjunction with other custom rules. For example, you can limit SQLi blocking to specific applications by using a [URL path filter](#).

Community vote distribution

A (100%)

 **Tabayashi** Highly Voted 7 months, 4 weeks ago

Answer is (A).

You can preview the effects of a rule without enforcing it. In preview mode, actions are noted in Cloud Monitoring. You can choose to preview individual rules in a security policy, or you can preview every rule in the policy.

[https://cloud.google.com/armor/docs/security-policy-overview#preview\\_mode](https://cloud.google.com/armor/docs/security-policy-overview#preview_mode)

upvoted 9 times

 **AzureDP900** 1 month, 2 weeks ago

A is right

upvoted 1 times

 **AwesomeGCP** Most Recent 2 months, 2 weeks ago

**Selected Answer: A**

A. Google Cloud Armor's preconfigured rules in preview mode

upvoted 1 times

 **sumundada** 5 months, 1 week ago

**Selected Answer: A**

Answer is (A).

upvoted 2 times



You are asked to recommend a solution to store and retrieve sensitive configuration data from an application that runs on Compute Engine. Which option should you recommend?

- A. Cloud Key Management Service
- B. Compute Engine guest attributes
- C. Compute Engine custom metadata
- D. Secret Manager

**Correct Answer:** D

Reference:

<https://www.freecodecamp.org/news/google-cloud-platform-from-zero-to-hero/>

- **General-purpose.** Offers the best price-performance ratio for a variety of workloads.
- **Memory-optimized.** Ideal for memory-intensive workloads. They offer more memory per core than other machine types.
- **Compute-optimized.** They offer the highest performance per core and are optimized for compute-intensive workloads
- **Shared-core.** These machine types timeshare a physical core. This can be a cost-effective method for running small applications.

Community vote distribution

D (100%)

 **Tabayashi** Highly Voted 7 months, 4 weeks ago  
Answer is (D).

Secret Manager is a secure and convenient storage system for API keys, passwords, certificates, and other sensitive data. Secret Manager provides a central place and single source of truth to manage, access, and audit secrets across Google Cloud.

<https://cloud.google.com/secret-manager>

upvoted 11 times

 **AzureDP900** Most Recent 1 month, 2 weeks ago  
D is correct  
upvoted 1 times

 **AwesomeGCP** 2 months, 2 weeks ago  
Selected Answer: D  
D. Secret Manager  
upvoted 1 times

 **cloudprincipal** 6 months, 3 weeks ago  
Selected Answer: D  
You need a secrets management solution  
<https://cloud.google.com/secret-manager>  
upvoted 4 times

 **cloudprincipal** 6 months, 3 weeks ago  
Sorry, this should be C  
upvoted 1 times

 **badrik** 6 months, 3 weeks ago  
sensitive information can never be stored/retrieved through custom meta data !



upvoted 4 times

You need to implement an encryption at-rest strategy that reduces key management complexity for non-sensitive data and protects sensitive data while providing the flexibility of controlling the key residency and rotation schedule. FIPS 140-2 L1 compliance is required for all data types. What should you do?

- A. Encrypt non-sensitive data and sensitive data with Cloud External Key Manager.
- B. Encrypt non-sensitive data and sensitive data with Cloud Key Management Service
- C. Encrypt non-sensitive data with Google default encryption, and encrypt sensitive data with Cloud External Key Manager.
- D. Encrypt non-sensitive data with Google default encryption, and encrypt sensitive data with Cloud Key Management Service.

**Correct Answer: B**

Community vote distribution

D (77%)

B (23%)

 **rrvv** 2 months, 3 weeks ago

As FIPS 140-2 L1 compliance is required for all types of data, Cloud KMS should be used to manage encryption. Correct answer is B

<https://cloud.google.com/docs/security/key-management-deep-dive#software-protection-level:~:text=The%20Cloud%20KMS%20binary%20is%20built%20against%20FIPS%20140%2D2%20Level%201%E2%80%933validated%20Cryptographic%20Primitives%20of%20this%20module>

upvoted 1 times

 **zelck** 2 months, 3 weeks ago

**Selected Answer: D**

D is the answer.

upvoted 3 times

 **Chute5118** 5 months ago

**Selected Answer: D**

Both B and D seem correct tbh. D might be "more correct" depending on the interpretation.

"reduces key management complexity for non-sensitive data" - Google default encryption

"protects sensitive data while providing the flexibility of controlling the key residency and rotation schedule" - Customer Managed Key

upvoted 3 times

 **AzureDP900** 1 month, 2 weeks ago

I agree, D is right

upvoted 1 times

 **sumundada** 5 months, 1 week ago

**Selected Answer: D**

Google uses a common cryptographic library, Tink, which incorporates our FIPS 140-2 Level 1 validated module, BoringCrypto, to implement encryption consistently across almost all Google Cloud products. To provide flexibility of controlling the key residency and rotation schedule, use google provided key for non-sensitive and encrypt sensitive data with Cloud Key Management Service

upvoted 2 times

 **nacying** 6 months, 2 weeks ago

**Selected Answer: B**

base on "FIPS 140-2 L1 compliance is required for all data types"

upvoted 3 times

 **cloudprincipal** 6 months, 3 weeks ago

**Selected Answer: D**

KMS is ok for fips 140-2 level 1

<https://cloud.google.com/docs/security/key-management-deep-dive#platform-overview>

upvoted 2 times

 **cloudprincipal** 6 months, 2 weeks ago

Regarding FIPS 140-2 level 1 and GCP default encryption:

Google Cloud uses a FIPS 140-2 validated Level 1 encryption module (certificate 3318) in our production environment.

[https://cloud.google.com/docs/security/encryption/default-encryption?hl=en#encryption\\_of\\_data\\_at\\_rest](https://cloud.google.com/docs/security/encryption/default-encryption?hl=en#encryption_of_data_at_rest)

upvoted 2 times

  **mikesp** 6 months, 3 weeks ago

In my opinion, the answer is B. The question says that it is necessary to control "key residency and rotation schedule" for both types of data. Default encryption at rest does not provide that but Cloud KMS does. Furthermore, Cloud KMS is FIPS140-2 level 1.

<https://cloud.google.com/docs/security/key-management-deep-dive>

upvoted 2 times

  **csrazdan** 1 week, 3 days ago

The answer is D.

1. reduce key management complexity for non-sensitive data --> Google Managed key

2. protects sensitive data while providing the flexibility of controlling the key residency and rotation schedule --> KMS

upvoted 1 times

  **szl0144** 7 months ago

D is the wander

upvoted 3 times

  **mouchu** 7 months, 1 week ago

Answer = D

upvoted 3 times

Your company wants to determine what products they can build to help customers improve their credit scores depending on their age range. To achieve this, you need to join user information in the company's banking app with customers' credit score data received from a third party. While using this raw data will allow you to complete this task, it exposes sensitive data, which could be propagated into new systems. This risk needs to be addressed using de-identification and tokenization with Cloud Data Loss Prevention while maintaining the referential integrity across the database. Which cryptographic token format should you use to meet these requirements?

- A. Deterministic encryption
- B. Secure, key-based hashes
- C. Format-preserving encryption
- D. Cryptographic hashing

**Correct Answer:** B


Reference:

<https://cloud.google.com/blog/products/identity-security/take-charge-of-your-data-how-tokenization-makes-data-usable-without-sacrificing-privacy>

Cloud DLP supports multiple cryptographic token formats that keep the referential integrity needed to join tables:

- **Deterministic encryption:** Replaces an input value with a cryptographic token. This encryption method is reversible, which helps to maintain referential integrity across your database and has no character-set limitations.
- **Format Preserving Encryption (FPE):** Creates a token of the same length and character set as the input. Similarly to deterministic encryption, it's reversible. FPE is great for inputs with a well defined alphabet space—for example, an alphabet including only [0-9a-zA-Z].

Community vote distribution



**mT3** Highly Voted 7 months, 1 week ago

**Selected Answer: A**

"This encryption method is reversible, which helps to maintain referential integrity across your database and has no character-set limitations."  
<https://cloud.google.com/blog/products/identity-security/take-charge-of-your-data-how-tokenization-makes-data-usable-without-sacrificing-privacy>

upvoted 9 times

**AzureDP900** 1 month, 2 weeks ago

A is right

upvoted 1 times

**PST21** Most Recent 4 days, 9 hours ago

Cryptographic uses strings , it asks to use tokenization and hence deterministic is better than FPE hence A

upvoted 1 times

**Littleivy** 1 month, 1 week ago

**Selected Answer: D**

Though it's not clear, but, to prevent from data leak, it's better to have a non-reversible method as analysts don't need re-identification

upvoted 1 times

**AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: A**

A. Deterministic encryption

upvoted 1 times

**zelck** 2 months, 3 weeks ago

**Selected Answer: A**

A is the answer.

<https://cloud.google.com/dlp/docs/pseudonymization>

FPE provides fewer security guarantees compared to other deterministic encryption methods such as AES-SIV.

For these reasons, Google strongly recommends using deterministic encryption with AES-SIV instead of FPE for all security sensitive use cases.

Other methods like deterministic encryption using AES-SIV provide these stronger security guarantees and are recommended for tokenization use cases unless length and character set preservation are strict requirements—for example, for backward compatibility with a legacy data system.

upvoted 1 times

  **piyush\_1982** 4 months, 4 weeks ago

**Selected Answer: A**

This question is taken from the exact scenario described in this link

<https://cloud.google.com/blog/products/identity-security/take-charge-of-your-data-how-tokenization-makes-data-usable-without-sacrificing-privacy>

upvoted 1 times

  **Chute5118** 5 months ago

**Selected Answer: D**

Both "Deterministic" and "format preserving" are key-based hashes (and reversible).

It's not clear from the question, but doesn't look like we need it to be reversible.

All of them maintain referential integrity

[https://cloud.google.com/architecture/de-identification-re-identification-pii-using-cloud-dlp#method\\_selection](https://cloud.google.com/architecture/de-identification-re-identification-pii-using-cloud-dlp#method_selection)

upvoted 1 times

  **cloudprincipal** 6 months, 3 weeks ago

**Selected Answer: D**

preserve referential integrity and ensure that no re-identification is possible

<https://cloud.google.com/dlp/docs/pseudonymization#supported-methods>

upvoted 1 times

  **cloudprincipal** 6 months ago

forget it, it should be A.

upvoted 1 times

  **Taliesyn** 7 months, 2 weeks ago

**Selected Answer: D**

Cryptographic hash (CryptoHashConfig) maintains referential integrity.

"Determinist encryption" is not a transformation method.

<https://cloud.google.com/dlp/docs/transformations-reference>

upvoted 2 times

  **gionny** 8 months ago

A correct - <https://cloud.google.com/blog/products/identity-security/take-charge-of-your-data-how-tokenization-makes-datausable-without-sacrificing-privacy>

upvoted 3 times

An office manager at your small startup company is responsible for matching payments to invoices and creating billing alerts. For compliance reasons, the office manager is only permitted to have the Identity and Access Management (IAM) permissions necessary for these tasks. Which two IAM roles should the office manager have? (Choose two.)

- A. Organization Administrator
- B. Project Creator
- C. Billing Account Viewer
- D. Billing Account Costs Manager
- E. Billing Account User


Correct Answer: AE

Reference:

<https://cloud.google.com/iam/docs/job-functions/billing>

|             |                      |                                                                                                                             |
|-------------|----------------------|-----------------------------------------------------------------------------------------------------------------------------|
| Role:       | Billing Account User | The Billing Account User role enables the developers to attach the billing account to new projects within the organization. |
| Resource:   | Organization         |                                                                                                                             |
| Principals: | Developers           |                                                                                                                             |

Community vote distribution



**mT3** Highly Voted 7 months, 1 week ago

Selected Answer: CD

Ans C,D.  
C. Billing Account Viewer :responsible for matching payments to invoices  
<https://cloud.google.com/billing/docs/how-to/get-invoice#required-permissions>  
Access billing documents:"Billing Account Administrator" or "Billing Account Viewer"  
D. Billing Account Costs Manager : creating billing alerts  
<https://cloud.google.com/billing/docs/how-to/budgets-notification-recipients>  
"To create or modify a budget for your Cloud Billing account, you need the Billing Account Costs Manager role or the Billing Account Administrator role on the Cloud Billing account."  
and "If you want the recipients of the alert emails to be able to view the budget, email recipients need permissions on the Cloud Billing account. At a minimum, ensure email recipients are added to the Billing Account Viewer role on the Cloud Billing account that owns the budget. See View a list of budgets for additional information."  
upvoted 8 times

**GHOST1985** 1 month, 1 week ago  
the link you post talking about Permissions required to ACCESS billing documentsn not to link project to a billing account you should have the Billing Account User role, the good answer is D,E  
upvoted 1 times

**AzureDP900** 1 month, 2 weeks ago  
CD is right  
upvoted 1 times

**GHOST1985** Most Recent 1 month, 2 weeks ago

Selected Answer: DE

Billing Account User: This role has very restricted permissions, so you can grant it broadly. When granted in combination with Project Creator, the two roles allow a user to create new projects linked to the billing account on which the Billing Account User role is granted. Or, when granted in combination with the Project Billing Manager role, the two roles allow a user to link and unlink projects on the billing account on which the Billing Account User role is granted.  
  
Billing Account Costs Manager: Create, edit, and delete budgets, view billing account cost information and transactions, and manage the export of billing cost data to BigQuery. Does not confer the right to export pricing data or view custom pricing in the Pricing page. Also, does not allow the linking or unlinking of projects or otherwise managing the properties of the billing account  
upvoted 2 times

**AwesomeGCP** 2 months, 2 weeks ago  
Selected Answer: CD



C. Billing Account Viewer  
D. Billing Account Costs Manager  
upvoted 1 times

  **zellick** 2 months, 3 weeks ago

**Selected Answer: CD**

CD is the answer.

<https://cloud.google.com/billing/docs/how-to/billing-access#overview-of-cloud-billing-roles-in-cloud-iam>

Billing Account Costs Manager (roles/billing.costsManager)

- Manage budgets and view and export cost information of billing accounts (but not pricing information)

Billing Account Viewer (roles/billing.viewer)

- View billing account cost information and transactions.

upvoted 2 times

  **Taliesyn** 7 months, 2 weeks ago

**Selected Answer: CD**

Billing Account Costs Administrator to create budgets (aka alerts)

Billing Account Viewer to view costs (to be able to match them to invoices)

upvoted 4 times

You are designing a new governance model for your organization's secrets that are stored in Secret Manager. Currently, secrets for Production and Non-

Production applications are stored and accessed using service accounts. Your proposed solution must:

- ☞ Provide granular access to secrets
- ☞ Give you control over the rotation schedules for the encryption keys that wrap your secrets
- ☞ Maintain environment separation
- ☞ Provide ease of management

Which approach should you take?

- A. 1. Use separate Google Cloud projects to store Production and Non-Production secrets. 2. Enforce access control to secrets using project-level identity and Access Management (IAM) bindings. 3. Use customer-managed encryption keys to encrypt secrets.
- B. 1. Use a single Google Cloud project to store both Production and Non-Production secrets. 2. Enforce access control to secrets using secret-level Identity and Access Management (IAM) bindings. 3. Use Google-managed encryption keys to encrypt secrets.
- C. 1. Use separate Google Cloud projects to store Production and Non-Production secrets. 2. Enforce access control to secrets using secret-level Identity and Access Management (IAM) bindings. 3. Use Google-managed encryption keys to encrypt secrets.
- D. 1. Use a single Google Cloud project to store both Production and Non-Production secrets. 2. Enforce access control to secrets using project-level Identity and Access Management (IAM) bindings. 3. Use customer-managed encryption keys to encrypt secrets.

**Correct Answer:** A

*Community vote distribution*

A (100%)



  **mT3** Highly Voted 7 months, 1 week ago

**Selected Answer: A**

Correct. Ans A.

Provide granular access to secrets: 2.Enforce access control to secrets using project-level identity and Access Management (IAM) bindings.  
Give you control over the rotation schedules for the encryption keys that wrap your secrets: 3. Use customer-managed encryption keys to encrypt secrets.

Maintain environment separation: 1. Use separate Google Cloud projects to store Production and Non-Production secrets.

upvoted 7 times

  **mikesp** 6 months, 3 weeks ago

It is possible to grant IAM bindind to secret-level which is more granular than project-level but considering that it is necessary to manage encryption keys life-cycle, then the answer is A due to C does not allow that.

upvoted 2 times

  **AzureDP900** 1 month, 2 weeks ago

Yes , A is right

upvoted 1 times

  **AwesomeGCP** Most Recent 2 months, 2 weeks ago

**Selected Answer: A**

A. 1. Use separate Google Cloud projects to store Production and Non-Production secrets. 2. Enforce access control to secrets using project-level identity and Access Management (IAM) bindings. 3. Use customer-managed encryption keys to encrypt secrets.

upvoted 2 times

  **Medofree** 7 months ago

None of the answers are correct, here is why :

- ☞ Provide granular access to secrets => 2. Enforce access control to secrets using secret-level (and not project-level)
- ☞ Give you control over the rotation schedules for the encryption keys that wrap your secrets => 3. Use customer-managed encryption keys to encrypt secrets.
- ☞ Maintain environment separation => 1. Use separate Google Cloud projects to store Production and Non-Production secrets
- ☞ Provide ease of management => 3. Use Google-managed encryption keys to encrypt secrets. (could be in contradiction with Give you control over the rotation schedules....)

It should be an E answer :

E. 1. Use separate Google Cloud projects to store Production and Non-Production secrets. 2. Enforce access control to secrets using secret-level identity and Access Management (IAM) bindings. 3. Use customer-managed encryption keys to encrypt secrets.

upvoted 3 times


You are a security engineer at a finance company. Your organization plans to store data on Google Cloud, but your leadership team is worried about the security of their highly sensitive data. Specifically, your company is concerned about internal Google employees' ability to access your company's data on Google Cloud.

What solution should you propose?

- A. Use customer-managed encryption keys.
- B. Use Google's Identity and Access Management (IAM) service to manage access controls on Google Cloud.
- C. Enable Admin activity logs to monitor access to resources.
- D. Enable Access Transparency logs with Access Approval requests for Google employees.

**Correct Answer:** D

*Community vote distribution*



🗲️ 👤 **Littleivy** 1 month, 1 week ago

**Selected Answer: A**

Use customer-managed key to encrypt data by yourself  
upvoted 1 times

🗲️ 👤 **Littleivy** 1 month, 1 week ago

D should be the answer on second thought  
upvoted 1 times

🗲️ 👤 **AzureDP900** 1 month, 2 weeks ago

D is correct  
upvoted 2 times

🗲️ 👤 **zellck** 2 months, 2 weeks ago

**Selected Answer: D**

D is the answer  
upvoted 2 times

🗲️ 👤 **jitu028** 2 months, 3 weeks ago

Answer is D

<https://cloud.google.com/access-transparency>

Access approval

Explicitly approve access to your data or configurations on Google Cloud. Access Approval requests, when combined with Access Transparency logs, can be used to audit an end-to-end chain from support ticket to access request to approval, to eventual access.

upvoted 2 times

You want to use the gcloud command-line tool to authenticate using a third-party single sign-on (SSO) SAML identity provider. Which options are necessary to ensure that authentication is supported by the third-party identity provider (IdP)? (Choose two.)


- A. SSO SAML as a third-party IdP
- B. Identity Platform
- C. OpenID Connect
- D. Identity-Aware Proxy
- E. Cloud Identity

**Correct Answer:** AC

Reference:

<https://cloud.google.com/identity/solutions/enable-sso>

*Community vote distribution*



**ExamQnA** **Highly Voted** 7 months ago

**Selected Answer: AE**

Third-party identity providers

If you have a third-party IdP, you can still configure SSO for third-party apps in the Cloud Identity catalog. User authentication occurs in the third-party IdP, and Cloud Identity manages the cloud apps.

To use Cloud Identity for SSO, your users need Cloud Identity accounts. They sign in through your third-party IdP or using a password on their Cloud Identity accounts.

<https://cloud.google.com/identity/solutions/enable-sso>

upvoted 9 times

**AzureDP900** 1 month, 2 weeks ago

A, E is right

upvoted 2 times

**piyush\_1982** **Highly Voted** 4 months, 3 weeks ago

**Selected Answer: AC**

I think the correct answer is A and C.

The questions asks about what is required with third-party IdP to authenticate the gcloud commands.

So the gcloud command requests goes to GCP. Since GCP is integrated with Third-party IdP for authentication gcloud command needs to be authenticated with third-party IdP.

This can be achieved if ThridPaty IdP supports SAML and OIDC protocols .

upvoted 5 times

**Littleivy** **Most Recent** 1 month, 1 week ago

**Selected Answer: AC**

answer is A and C.

upvoted 1 times

**Littleivy** 1 month, 1 week ago

**Selected Answer: AC**

To provide users with SSO-based access to selected cloud apps, Cloud Identity as your IdP supports the OpenID Connect (OIDC) and Security Assertion Markup Language 2.0 (SAML) protocols.

<https://cloud.google.com/identity/solutions/enable-sso>

upvoted 1 times

**amenur** 1 month, 1 week ago

"ensure that authentication is supported by the third-party identity provider (IdP)"

This question is only related to the external IdP.

It's about third-party IdP possible supported protocols. For this reason i think the correct options are A and C

upvoted 1 times

**hcnh** 1 month, 2 weeks ago

**Selected Answer: AC**

the questions is to ask for "Which options are necessary to ensure that authentication is supported by the third-party identity provider (IdP)?" so AC is the answer.

upvoted 2 times

  **Table2022** 1 month, 4 weeks ago

Answer is AC with SAML + OIDC.

upvoted 1 times

  **zellck** 2 months, 3 weeks ago

**Selected Answer: AE**

AE is the answer.

upvoted 2 times

  **vpolo** 7 months, 1 week ago

i think AE

upvoted 3 times

Question #106

Topic 1

You work for a large organization where each business unit has thousands of users. You need to delegate management of access control permissions to each business unit. You have the following requirements:

- ↻ Each business unit manages access controls for their own projects.
- ↻ Each business unit manages access control permissions at scale.
- ↻ Business units cannot access other business units' projects.
- ↻ Users lose their access if they move to a different business unit or leave the company.
- ↻ Users and access control permissions are managed by the on-premises directory service.

What should you do? (Choose two.)

- A. Use VPC Service Controls to create perimeters around each business unit's project.
- B. Organize projects in folders, and assign permissions to Google groups at the folder level.
- C. Group business units based on Organization Units (OUs) and manage permissions based on OUs
- D. Create a project naming convention, and use Google's IAM Conditions to manage access based on the prefix of project names.
- E. Use Google Cloud Directory Sync to synchronize users and group memberships in Cloud Identity.

**Correct Answer: BE**

  **Rightsaidfred** 1 month ago

Agreed...B & E

upvoted 1 times

  **TheBuckler** 2 months, 2 weeks ago

I will take B & E. Makes sense for the OUs to have their own folders and respective projects under their folders. This will make each OU independent from one another in terms of environments, and will not be able to communicate with one another unless shared VPC/VPC peering is utilized.

And E is fairly obvious, as they want to manage their users from on-prem directory, hence GCDS.

upvoted 2 times

Your organization recently deployed a new application on Google Kubernetes Engine. You need to deploy a solution to protect the application. The solution has the following requirements:

- ☞ Scans must run at least once per week
- ☞ Must be able to detect cross-site scripting vulnerabilities
- ☞ Must be able to authenticate using Google accounts

Which solution should you use?

- A. Google Cloud Armor
- B. Web Security Scanner
- C. Security Health Analytics
- D. Container Threat Detection

**Correct Answer:** B

Reference:

<https://cloud.google.com/security-command-center/docs/concepts-web-security-scanner-overview>

Other important things to be aware of when using Web Security Scanner:

- Because Web Security Scanner is undergoing continual improvements, a future scan might report issues that are not reported by the current scan.
- Some features or sections of your application might not be tested.
- Web Security Scanner attempts to activate every control and input it finds.
- If you expose state-changing actions for which your test account has permission, Web Security Scanner is likely to activate them. This might lead to undesirable results.

Community vote distribution

B (100%)

🗉 👤 **Tabayashi** Highly Voted 7 months, 4 weeks ago  
Answer is (B).

Web Security Scanner identifies security vulnerabilities in your App Engine, Google Kubernetes Engine (GKE), and Compute Engine web applications.

<https://cloud.google.com/security-command-center/docs/concepts-web-security-scanner-overview>

upvoted 11 times

🗉 👤 **AzureDP900** 1 month, 2 weeks ago  
Yes, B is right  
upvoted 1 times

🗉 👤 **AwesomeGCP** Most Recent 2 months, 2 weeks ago  
Selected Answer: B  
B. Web Security Scanner  
upvoted 1 times

🗉 👤 **zellck** 2 months, 3 weeks ago  
Selected Answer: B  
B is the answer.  
upvoted 3 times

An organization is moving applications to Google Cloud while maintaining a few mission-critical applications on-premises. The organization must transfer the data at a bandwidth of at least 50 Gbps. What should they use to ensure secure continued connectivity between sites?


- A. Dedicated Interconnect
- B. Cloud Router
- C. Cloud VPN
- D. Partner Interconnect

**Correct Answer: A**

Reference:

<https://cloud.google.com/architecture/migration-to-google-cloud-transferring-your-large-datasets>

The following diagram illustrates the path of your migration journey.



Community vote distribution



**mouchu** Highly Voted 7 months, 1 week ago

Answer = A

upvoted 8 times

**AzureDP900** Most Recent 1 month, 2 weeks ago

A is right

upvoted 1 times

**AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: A**

A. Dedicated Interconnect

upvoted 1 times

**zellick** 2 months, 3 weeks ago

**Selected Answer: A**

A is the answer.

upvoted 1 times

**Arturo\_Cloud** 3 months, 2 weeks ago

I understand that not all Partner Interconnect connections support 50 Gbps, so I'm going with A) for guaranteed connectivity.

<https://cloud.google.com/network-connectivity/docs/interconnect/concepts/overview>

upvoted 1 times



Your organization has had a few recent DDoS attacks. You need to authenticate responses to domain name lookups. Which Google Cloud service should you use?

- A. Cloud DNS with DNSSEC
- B. Cloud NAT
- C. HTTP(S) Load Balancing
- D. Google Cloud Armor

**Correct Answer: A**

Reference:

<https://developers.google.com/speed/public-dns/faq>

*Community vote distribution*

A (100%)

🗲️ 👤 **Tabayashi** Highly Voted 👍 7 months, 4 weeks ago

Answer is (A).

The Domain Name System Security Extensions (DNSSEC) is a feature of the Domain Name System (DNS) that authenticates responses to domain name lookups. It does not provide privacy protections for those lookups, but prevents attackers from manipulating or poisoning the responses to DNS requests.

<https://cloud.google.com/dns/docs/dnssec>

upvoted 15 times

🗲️ 👤 **AzureDP900** 1 month, 2 weeks ago

Agreed, A is right

upvoted 2 times

🗲️ 👤 **risc** Most Recent 🕒 2 months ago

**Selected Answer: A**

A, as explained by Tabayashi

upvoted 1 times

🗲️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: A**

A. Cloud DNS with DNSSEC

upvoted 1 times

🗲️ 👤 **zellck** 2 months, 3 weeks ago

**Selected Answer: A**

A is the answer.


upvoted 1 times

Your Security team believes that a former employee of your company gained unauthorized access to Google Cloud resources some time in the past 2 months by using a service account key. You need to confirm the unauthorized access and determine the user activity. What should you do?

- A. Use Security Health Analytics to determine user activity.
- B. Use the Cloud Monitoring console to filter audit logs by user.
- C. Use the Cloud Data Loss Prevention API to query logs in Cloud Storage.
- D. Use the Logs Explorer to search for user activity.

**Correct Answer:** B

Community vote distribution



🗲️ 👤 **Medofree** Highly Voted 👍 7 months ago

**Selected Answer: D**

D.

We use audit logs by searching the Service Account and checking activities in the past 2 months. (the user identity will not be seen since he used the SA identity but we can make correlations based on ip address, working hour, etc. )

upvoted 10 times

🗲️ 👤 **AzureDP900** 1 month, 2 weeks ago

D is right, I agree

upvoted 2 times

🗲️ 👤 **zelck** Most Recent ⌚ 2 months, 3 weeks ago

**Selected Answer: D**

D is the answer.

upvoted 1 times

🗲️ 👤 **mikesp** 6 months, 3 weeks ago

**Selected Answer: D**

B is intended to mislead the public. Cloud Monitoring provides only metrics. To check user activity is necessary to go to Cloud Logging and search on Audit Logs.

upvoted 3 times

🗲️ 👤 **mT3** 7 months, 1 week ago

**Selected Answer: B**

Correct. Answer is (B).

Investigate the potentially unauthorized activity and restore the account.

Ref.<https://support.google.com/a/answer/2984349>

upvoted 3 times



Your company requires the security and network engineering teams to identify all network anomalies within and across VPCs, internal traffic from VMs to VMs, traffic between end locations on the internet and VMs, and traffic between VMs to Google Cloud services in production. Which method should you use?

- A. Define an organization policy constraint.
- B. Configure packet mirroring policies.
- C. Enable VPC Flow Logs on the subnet.
- D. Monitor and analyze Cloud Audit Logs.

**Correct Answer:** C


Reference:

<https://cloud.google.com/architecture/best-practices-vpc-design>

Consider the components illustrated in the following example when establishing your naming conventions:

- **Company name:** Acme Company: `acmeco`
- **Business unit:** Human Resources: `hr`
- **Application code:** Compensation system: `comp`
- **Region code:** northamerica-northeast1: `na-ne1`, europe-west1: `eu-we1`
- **Environment codes:** `dev`, `test`, `uat`, `stage`, `prod`

Community vote distribution



**Tabayashi** Highly Voted 7 months, 4 weeks ago

I think the answer is (C).

VPC Flow Logs samples each VM's TCP, UDP, ICMP, ESP, and GRE flows. Both inbound and outbound flows are sampled. These flows can be between the VM and another VM, a host in your on-premises data center, a Google service, or a host on the internet.

<https://cloud.google.com/vpc/docs/flow-logs>

upvoted 11 times

**tifo16** Most Recent 2 weeks ago

[https://cloud.google.com/vpc/docs/packet-mirroring#enterprise\\_security](https://cloud.google.com/vpc/docs/packet-mirroring#enterprise_security)

Security and network engineering teams must ensure that they are catching all anomalies and threats that might indicate security breaches and intrusions. They mirror all traffic so that they can complete a comprehensive inspection of suspicious flows. Because attacks can span multiple packets, security teams must be able to get all packets for each flow.

upvoted 1 times

**tifo16** 2 weeks ago

Should be B

upvoted 1 times

**Rightsaidfred** 1 month ago

As it is a close tie and ambiguity between B&C, I would say it is C - VPC Flow Logs in this instance, as Question 121 is focusing more on Packet Mirroring with the IDS Use Case.

upvoted 2 times

**marmar11111** 1 month, 1 week ago

**Selected Answer: B**

Should be B

upvoted 1 times

**hcnh** 1 month, 2 weeks ago

**Selected Answer: C**

C is the answer as B has the limitation against question

The mirroring happens on the virtual machine (VM) instances, not on the network. Consequently, Packet Mirroring consumes additional bandwidth on the VMs.

upvoted 1 times

  **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: B**

B. Configure packet mirroring policies.

upvoted 3 times

  **zellick** 2 months, 3 weeks ago

**Selected Answer: B**

B is the answer.

[https://cloud.google.com/vpc/docs/packet-mirroring#enterprise\\_security](https://cloud.google.com/vpc/docs/packet-mirroring#enterprise_security)

Security and network engineering teams must ensure that they are catching all anomalies and threats that might indicate security breaches and intrusions. They mirror all traffic so that they can complete a comprehensive inspection of suspicious flows.

upvoted 1 times

  **AzureDP900** 1 month, 2 weeks ago

Agree with B

upvoted 1 times

  **GHOST1985** 3 months, 1 week ago

**Selected Answer: B**

100% Answer B: Anomalies means packet mirroring

[https://cloud.google.com/vpc/docs/packet-mirroring#enterprise\\_security](https://cloud.google.com/vpc/docs/packet-mirroring#enterprise_security)

"Packet Mirroring is useful when you need to monitor and analyze your security status. It exports all traffic, not only the traffic between sampling periods. For example, you can use security software that analyzes mirrored traffic to detect all threats or anomalies. Additionally, you can inspect the full traffic flow to detect application performance issues. For more information, see the example use cases."

<https://cloud.google.com/vpc/docs/packet-mirroring>

upvoted 1 times

  **tangac** 3 months, 2 weeks ago

**Selected Answer: C**

First you can use VPC flow log at a subnet level : <https://cloud.google.com/vpc/docs/using-flow-logs>

Then VPC Flow Log main feature is to collect logs that can be used for network monitoring, forensics, real-time security analysis, and expense optimization.

upvoted 1 times

  **jvkubjg** 3 months, 4 weeks ago

**Selected Answer: B**

Anomalies -> Packet Mirroring

upvoted 1 times

  **hybridpro** 6 months, 1 week ago

B should be the answer. For detecting network anomalies, you need to have payload and header data as well to be effective. Besides C is saying to enable VPC flow logs on a subnet which won't serve our purpose either.

upvoted 3 times

  **mikesp** 6 months, 3 weeks ago

**Selected Answer: C**

VPC Flow Logs also helps you perform network forensics when investigating suspicious behavior such as traffic from access from abnormal sources or unexpected volumes of data migration

upvoted 3 times

Your company has been creating users manually in Cloud Identity to provide access to Google Cloud resources. Due to continued growth of the environment, you want to authorize the Google Cloud Directory Sync (GCDS) instance and integrate it with your on-premises LDAP server to onboard hundreds of users. You are required to:

- ☞ Replicate user and group lifecycle changes from the on-premises LDAP server in Cloud Identity.
- ☞ Disable any manually created users in Cloud Identity.

You have already configured the LDAP search attributes to include the users and security groups in scope for Google Cloud. What should you do next to complete this solution?

- A. 1. Configure the option to suspend domain users not found in LDAP. 2. Set up a recurring GCDS task.
- B. 1. Configure the option to delete domain users not found in LDAP. 2. Run GCDS after user and group lifecycle changes.
- C. 1. Configure the LDAP search attributes to exclude manually created Cloud Identity users not found in LDAP. 2. Set up a recurring GCDS task.
- D. 1. Configure the LDAP search attributes to exclude manually created Cloud Identity users not found in LDAP. 2. Run GCDS after user and group lifecycle changes.

**Correct Answer:** D

Reference:

<https://www.appsadmins.com/blog/your-how-to-for-provisioning-in-g-suite-part-2-users>

Community vote distribution

A (100%)

🗲️ **mT3** Highly Voted 7 months, 1 week ago

**Selected Answer: A**

Answer is (A).

To achieve the requirement "Disable any manually created users in Cloud Identity", configure GCDS to suspend rather than delete accounts if user accounts are not found in the LDAP directory in GCDS.

Ref: <https://support.google.com/a/answer/7177267>

upvoted 9 times

🗲️ **AzureDP900** 1 month, 2 weeks ago

A is right

upvoted 1 times

🗲️ **AwesomeGCP** Most Recent 1 month, 2 weeks ago

**Selected Answer: A**

A. 1. Configure the option to suspend domain users not found in LDAP. 2. Set up a recurring GCDS task.

upvoted 1 times

🗲️ **tangac** 3 months, 2 weeks ago

**Selected Answer: A**

clearly A

upvoted 1 times

🗲️ **KillerGoogle** 7 months, 2 weeks ago

C. 1. Configure the LDAP search attributes to exclude manually created Cloud Identity users not found in LDAP. 2. Set up a recurring GCDS task.

upvoted 1 times

🗲️ **Tabayashi** 7 months, 4 weeks ago

I think the answer is (A).

When using Shared VPC, a service perimeter that includes projects that belong to a Shared VPC network must also include the project that hosts the network. When projects that belong to a Shared VPC network are not in the same perimeter as the host project, services might not work as expected or might be blocked entirely.

Ensure that the Shared VPC network host is in the same service perimeter as the projects connected to the network.

[https://cloud.google.com/vpc-service-controls/docs/troubleshooting#shared\\_vpc](https://cloud.google.com/vpc-service-controls/docs/troubleshooting#shared_vpc)

upvoted 3 times

🗲️ **Tabayashi** 7 months, 4 weeks ago

Sorry, this answer is question 113.

upvoted 2 times

You are troubleshooting access denied errors between Compute Engine instances connected to a Shared VPC and BigQuery datasets. The datasets reside in a project protected by a VPC Service Controls perimeter. What should you do?

- A. Add the host project containing the Shared VPC to the service perimeter.
- B. Add the service project where the Compute Engine instances reside to the service perimeter.
- C. Create a service perimeter between the service project where the Compute Engine instances reside and the host project that contains the Shared VPC.
- D. Create a perimeter bridge between the service project where the Compute Engine instances reside and the perimeter that contains the protected BigQuery datasets.

**Correct Answer:** C

Reference:


<https://cloud.google.com/vpc-service-controls/docs/troubleshooting>

## Shared VPC

When using Shared VPC, a service perimeter that includes projects that belong to a Shared VPC network must also include the project that *hosts* the network. When projects that belong to a Shared VPC network are not in the same perimeter as the host project, services might not work as expected or might be blocked entirely.

Ensure that the Shared VPC network host is in the same service perimeter as the projects connected to the network.

Community vote distribution



**Littleivy** 1 month, 1 week ago

**Selected Answer: A**

As the scenario is for troubleshooting, I'll choose A as answer since it's more likely people would forget to include host project to the service perimeter

upvoted 1 times

**AzureDP900** 1 month, 2 weeks ago

A. Add the host project containing the Shared VPC to the service perimeter. Looks good to me based on requirements

upvoted 1 times

**risc** 2 months ago

**Selected Answer: A**

(A)

For VMs inside shared VPC, the host project needs to be added to the perimeter as well. I had real-life experience with this. However, this creates new security issues as all other VMs in other projects which are attached to shared subnets in the same host project then are also able to access the perimeter. Google recommends setting up Private Service Connect Endpoints to achieve subnet segregation for VPC-SC usage with Host projects.

upvoted 2 times

**soltium** 2 months, 1 week ago

**Selected Answer: B**

Weird question, you need A n B.

I'll choose B.

upvoted 1 times

**AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: A**

A. Add the host project containing the Shared VPC to the service perimeter.

upvoted 1 times

**zellck** 2 months, 3 weeks ago

**Selected Answer: A**

A is the answer.

<https://cloud.google.com/vpc-service-controls/docs/service-perimeters#secure-google-managed-resources>

If you're using Shared VPC, you must include the host project in a service perimeter along with any projects that belong to the Shared VPC.

upvoted 2 times

  **GHOST1985** 3 months, 1 week ago

**Selected Answer: A**

"If you're using Shared VPC, you must include the host project in a service perimeter along with any projects that belong to the Shared VPC" =>  
<https://cloud.google.com/vpc-service-controls/docs/service-perimeters>  
upvoted 1 times

  **Chute5118** 5 months ago

**Selected Answer: B**

"If you're using Shared VPC, you must include the host project in a service perimeter along with any projects that belong to the Shared VPC."  
<https://cloud.google.com/vpc-service-controls/docs/service-perimeters>

B  
upvoted 2 times

  **GHOST1985** 3 months, 1 week ago

i think you mean Answer A :)  
upvoted 1 times

  **Aiffone** 5 months, 2 weeks ago

i think the Answer should be C (a combination of A and B)  
upvoted 1 times

  **mikesp** 6 months, 3 weeks ago

**Selected Answer: B**

Change my answer.  
upvoted 2 times

  **mikesp** 6 months, 3 weeks ago

**Selected Answer: A**

<https://cloud.google.com/vpc-service-controls/docs/service-perimeters>  
If you're using Shared VPC, you must include the host project in a service perimeter along with any projects that belong to the Shared VPC.  
upvoted 2 times

  **mikesp** 6 months, 3 weeks ago

I thinking that the question is not formulated correctly. A does not respond the question completely. I think that the actual answer is B and should be supposed that both, VMs and BigQuery reside in the same "Service Project" and that would be the reason of adding the Host project of the shared VPC to the perimeter.  
upvoted 3 times

  **Tabayashi** 7 months, 4 weeks ago

Answer is (D).

Firewall Insights analyzes your firewall rules to detect firewall rules that are shadowed by other rules. A shadowed rule is a firewall rule that has all of its relevant attributes, such as its IP address and port ranges, overlapped by attributes from one or more rules with higher or equal priority, called shadowing rules.  
<https://cloud.google.com/network-intelligence-center/docs/firewall-insights/concepts/overview#shadowed-firewall-rules>  
upvoted 1 times

  **Tabayashi** 7 months, 4 weeks ago

Sorry, this answer is question 114.  
upvoted 1 times

You recently joined the networking team supporting your company's Google Cloud implementation. You are tasked with familiarizing yourself with the firewall rules configuration and providing recommendations based on your networking and Google Cloud experience. What product should you recommend to detect firewall rules that are overlapped by attributes from other firewall rules with higher or equal priority?

- A. Security Command Center
- B. Firewall Rules Logging
- C. VPC Flow Logs
- D. Firewall Insights

**Correct Answer:** D

Reference:

<https://cloud.google.com/network-intelligence-center/docs/firewall-insights/concepts/overview>

With Firewall Insights **metrics**, you can perform the following tasks:

- Verify that firewall rules are being used in the intended way.
- Over specified time periods, verify that firewall rules allow or block their intended connections.
- Perform live debugging of connections that are inadvertently dropped because of firewall rules.
- Discover malicious attempts to access your network, in part by getting alerts about **significant changes in the hit counts of firewall rules**.

Community vote distribution

D (100%)

🗳️ 👤 **AzureDP900** 1 month, 2 weeks ago

D. Firewall Insights  
upvoted 1 times

🗳️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: D**

D. Firewall Insights  
upvoted 1 times

🗳️ 👤 **zellck** 2 months, 3 weeks ago

**Selected Answer: D**

D is the answer.

<https://cloud.google.com/network-intelligence-center/docs/firewall-insights/concepts/overview#shadowed-firewall-rules>

Firewall Insights analyzes your firewall rules to detect firewall rules that are shadowed by other rules. A shadowed rule is a firewall rule that has all of its relevant attributes, such as its IP address and port ranges, overlapped by attributes from one or more rules with higher or equal priority, called shadowing rules.

upvoted 3 times

🗳️ 👤 **AzureDP900** 1 month, 2 weeks ago

Agreed  
upvoted 1 times

🗳️ 👤 **mikesp** 6 months, 3 weeks ago

**Selected Answer: D**

Answer = D.  
upvoted 1 times

🗳️ 👤 **ExamQnA** 7 months, 1 week ago

**Selected Answer: D**

Firewall Insights analyzes your firewall rules to detect firewall rules that are shadowed by other rules. A shadowed rule is a firewall rule that has all of its relevant attributes, such as its IP address and port ranges, overlapped by attributes from one or more rules with higher or equal priority, called shadowing rules.

<https://cloud.google.com/network-intelligence-center/docs/firewall-insights/concepts/overview>

upvoted 4 times

🗳️ 👤 **mouchu** 7 months, 1 week ago



Answer = D

Firewall Insights analyzes your firewall rules to detect firewall rules that are shadowed by other rules.  
upvoted 2 times



The security operations team needs access to the security-related logs for all projects in their organization. They have the following requirements:

- ⇒ Follow the least privilege model by having only view access to logs.
- ⇒ Have access to Admin Activity logs.
- ⇒ Have access to Data Access logs.
- ⇒ Have access to Access Transparency logs.

Which Identity and Access Management (IAM) role should the security operations team be granted?

- A. roles/logging.privateLogViewer
- B. roles/logging.admin
- C. roles/viewer
- D. roles/logging.viewer

**Correct Answer: B**

Reference:

<https://cloud.google.com/logging/docs/access-control>

When deciding which **permissions and roles** apply to your principals' use cases, consider the following:

- **roles/logging.admin** (Logging Admin) grants you all permissions related to Logging.
- **roles/logging.viewer** (Logs Viewer) gives you read-only access to most features of Logging.

The Logs Viewer role gives you access to both the **\_AllLogs** **log view** in the **\_Required** bucket and the **\_Default** log view in the **\_Default** bucket.

The Logs Viewer role doesn't let you read the Data Access audit logs that are in the **\_Default** **bucket**. To read these Data Access audit logs, you need the Private Logs Viewer role ( **roles/logging.privateLogViewer** ) for the appropriate **log view**.

Community vote distribution

A (100%)

🗳️ **mouchu** **Highly Voted** 👍 7 months, 1 week ago

Answer = A

roles/logging.privateLogViewer (Private Logs Viewer) includes all the permissions contained by roles/logging.viewer, plus the ability to read Data Access audit logs in the \_Default bucket.

upvoted 12 times

🗳️ **mT3** 7 months, 1 week ago

Ref: <https://cloud.google.com/logging/docs/access-control>

upvoted 2 times

🗳️ **Littleivy** **Most Recent** 🕒 1 month, 1 week ago

**Selected Answer: A**

You need roles/logging.privateLogViewer to view data access log and Access Transparency logs

<https://cloud.google.com/cloud-provider-access-management/access-transparency/docs/reading-logs#viewing-logs>

[https://developers.google.com/cloud-search/docs/guides/audit-logging-manual#audit\\_log\\_permissions](https://developers.google.com/cloud-search/docs/guides/audit-logging-manual#audit_log_permissions)

upvoted 1 times

🗳️ **AzureDP900** 1 month, 2 weeks ago

A is the answer.

upvoted 1 times

🗳️ **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: A**

A. roles/logging.privateLogViewer

upvoted 1 times

🗳️ **zellck** 2 months, 3 weeks ago

A is the answer.

<https://cloud.google.com/logging/docs/access-control#considerations>



roles/logging.privateLogViewer (Private Logs Viewer) includes all the permissions contained by roles/logging.viewer, plus the ability to read Data Access audit logs in the \_Default bucket.

upvoted 2 times

  **cloudprincipal** 6 months, 3 weeks ago

**Selected Answer: A**

roles/logging.privateLogViewer (Private Logs Viewer) includes all the permissions contained by roles/logging.viewer, plus the ability to read Data Access audit logs in the \_Default bucket.

<https://cloud.google.com/logging/docs/access-control>

upvoted 3 times

  **Nicky1402** 7 months, 2 weeks ago

I think the correct answer is A.

logging.admin is too broad a permission.

We need to give "only view access to logs". And we need to:

☞ Have access to Admin Activity logs.

☞ Have access to Data Access logs.

☞ Have access to Access Transparency logs.

Only the roles/logging.privateLogViewer role has all these permissions.

Private Logs Viewer

(roles/logging.privateLogViewer)

Provides permissions of the Logs Viewer role and in addition, provides read-only access to log entries in private logs.

Lowest-level resources where you can grant this role:

Project

After you've configured Access Transparency for your Google Cloud organization, you can set controls for who can access the Access Transparency logs by assigning a user or group the Private Logs Viewer role.

Links for reference:

<https://cloud.google.com/logging/docs/access-control>

<https://cloud.google.com/cloud-provider-access-management/access-transparency/docs/reading-logs?hl=en>

upvoted 4 times

You are exporting application logs to Cloud Storage. You encounter an error message that the log sinks don't support uniform bucket-level access policies. How should you resolve this error?

- A. Change the access control model for the bucket
- B. Update your sink with the correct bucket destination.
- C. Add the roles/logging.logWriter Identity and Access Management (IAM) role to the bucket for the log sink identity.
- D. Add the roles/logging.bucketWriter Identity and Access Management (IAM) role to the bucket for the log sink identity.

**Correct Answer: A**

Reference:  
<https://sysdig.com/blog/gcp-security-best-practices/>

|     |                                                             |                                                                   |                                |
|-----|-------------------------------------------------------------|-------------------------------------------------------------------|--------------------------------|
|     | High Risk<br>■ ■ ■                                          | Medium Risk<br>■ ■                                                | Low Risk<br>■                  |
| IAM | (2) MFA enabled for all users<br>(3) Security Key for admin | (1) Check IAM policies<br>(4) Prevent the use of service accounts |                                |
| KMS | (5) Not publicly accessible                                 |                                                                   | (6) Rotate KMS encryption keys |

Community vote distribution

A (100%)

- mikesp

Highly Voted

6 months, 3 weeks ago

Selected Answer: A

<https://cloud.google.com/logging/docs/export/troubleshoot>  
Unable to grant correct permissions to the destination:  
Even if the sink was successfully created with the correct service account permissions, this error message displays if the access control model for the Cloud Storage bucket was set to uniform access when the bucket was created.  
For existing Cloud Storage buckets, you can change the access control model for the first 90 days after bucket creation by using the Permissions tab. For new buckets, select the Fine-grained access control model during bucket creation. For details, see Creating Cloud Storage buckets.  
upvoted 6 times
- AzureDP900

Most Recent

1 month, 2 weeks ago

A is right

upvoted 1 times
- zelck

2 months, 3 weeks ago

Selected Answer: A

A is the answer.

[https://cloud.google.com/logging/docs/export/troubleshoot#errors\\_exporting\\_to\\_cloud\\_storage](https://cloud.google.com/logging/docs/export/troubleshoot#errors_exporting_to_cloud_storage)  
- Unable to grant correct permissions to the destination:  
Even if the sink was successfully created with the correct service account permissions, this error message displays if the access control model for the Cloud Storage bucket was set to uniform access when the bucket was created.  
upvoted 3 times
- mT3

7 months, 1 week ago

Selected Answer: A

Answer is (A).  
If bucket-level access policies are not supported, Fine-grained is being used.  
The recommended architecture is Uniform bucket-level access. Therefore, Change the access control model for the bucket.  
Ref : <https://cloud.google.com/storage/docs/access-control>  
upvoted 2 times
- Taliesyn

7 months, 2 weeks ago

Selected Answer: A

A: can't export logs to a bucket with uniform bucket-level access  
(B sounds halfway decent as well, but you'd still need another bucket without uniform bucket-level access, so it's incomplete)

Question #117

Topic 1

You plan to deploy your cloud infrastructure using a CI/CD cluster hosted on Compute Engine. You want to minimize the risk of its credentials being stolen by a third party. What should you do?

- A. Create a dedicated Cloud Identity user account for the cluster. Use a strong self-hosted vault solution to store the user's temporary credentials.
- B. Create a dedicated Cloud Identity user account for the cluster. Enable the constraints/iam.disableServiceAccountCreation organization policy at the project level.
- C. Create a custom service account for the cluster. Enable the constraints/iam.disableServiceAccountKeyCreation organization policy at the project level
- D. Create a custom service account for the cluster. Enable the constraints/iam.allowServiceAccountCredentialLifetimeExtension organization policy at the project level.

Correct Answer: D

Reference:

<https://cloud.google.com/resource-manager/docs/organization-policy/restricting-service-accounts>

Disable automatic role grants to default service accounts

Some Google Cloud services automatically create **default service accounts**. When a default service account is created, it is automatically granted the Editor role ( `roles/editor` ) on your project.

To improve security, we strongly recommend that you disable the automatic role grant. Use the `iam.automaticIamGrantsForDefaultServiceAccounts` boolean constraint to disable the automatic role grant.

Community vote distribution



**ExamQnA** Highly Voted 7 months ago

**Selected Answer: C**

Disable service account key creation  
You can use the iam.disableServiceAccountKeyCreation boolean constraint to disable the creation of new external service account keys. This allows you to control the use of unmanaged long-term credentials for service accounts. When this constraint is set, user-managed credentials cannot be created for service accounts in projects affected by the constraint.  
[https://cloud.google.com/resource-manager/docs/organization-policy/restricting-service-accounts#example\\_policy\\_boolean\\_constraint](https://cloud.google.com/resource-manager/docs/organization-policy/restricting-service-accounts#example_policy_boolean_constraint)  
upvoted 5 times

**AzureDP900** 1 month, 2 weeks ago

Yes

C. Create a custom service account for the cluster. Enable the constraints/iam.disableServiceAccountKeyCreation organization policy at the project level  
upvoted 1 times

**mikesp** Most Recent 6 months, 3 weeks ago

**Selected Answer: C**

Also think it is C  
upvoted 4 times

**mT3** 7 months, 1 week ago

**Selected Answer: C**


Answer is (C).  
To minimize the risk of credentials being stolen by third parties, it is desirable to control the use of unmanaged long-term credentials.  
• "constraints/iam.allowServiceAccountCredentialLifetimeExtension": to extend the lifetime of the access token.  
• "iam.disableServiceAccountCreation": Disables service account creation.  
• "iam.disableServiceAccountCreation": Controls the use of unmanaged long-term credentials for service accounts.  
Ref : [https://cloud.google.com/resource-manager/docs/organization-policy/restricting-service-accounts#example\\_policy\\_boolean\\_constraint](https://cloud.google.com/resource-manager/docs/organization-policy/restricting-service-accounts#example_policy_boolean_constraint)  
upvoted 2 times

You need to set up two network segments: one with an untrusted subnet and the other with a trusted subnet. You want to configure a virtual appliance such as a next-generation firewall (NGFW) to inspect all traffic between the two network segments. How should you design the network to inspect the traffic?

- A. 1. Set up one VPC with two subnets: one trusted and the other untrusted. 2. Configure a custom route for all traffic (0.0.0.0/0) pointed to the virtual appliance.
- B. 1. Set up one VPC with two subnets: one trusted and the other untrusted. 2. Configure a custom route for all RFC1918 subnets pointed to the virtual appliance.
- C. 1. Set up two VPC networks: one trusted and the other untrusted, and peer them together. 2. Configure a custom route on each network pointed to the virtual appliance.
- D. 1. Set up two VPC networks: one trusted and the other untrusted. 2. Configure a virtual appliance using multiple network interfaces, with each interface connected to one of the VPC networks.

**Correct Answer: B**

Community vote distribution



**mouchu** Highly Voted 7 months, 1 week ago

Answer = D

Multiple network interfaces. The simplest way to connect multiple VPC networks through a virtual appliance is by using multiple network interfaces, with each interface connecting to one of the VPC networks. Internet and on-premises connectivity is provided over one or two separate network interfaces. With many NGFW products, internet connectivity is connected through an interface marked as untrusted in the NGFW software.

upvoted 9 times

**mT3** 7 months, 1 week ago

Agreed.

Ref: For Cisco Firepower Threat Defense Virtual: [https://www.cisco.com/c/en/us/td/docs/security/firepower/quick\\_start/gcp/ftdv-gcp-gsg/ftdv-gcp-intro.html](https://www.cisco.com/c/en/us/td/docs/security/firepower/quick_start/gcp/ftdv-gcp-gsg/ftdv-gcp-intro.html)

upvoted 1 times

**AzureDP900** 1 month, 2 weeks ago

Agree

D. 1. Set up two VPC networks: one trusted and the other untrusted. 2. Configure a virtual appliance using multiple network interfaces, with each interface connected to one of the VPC networks.

upvoted 1 times

**zellck** Most Recent 2 months, 3 weeks ago

**Selected Answer: D**

D is the answer.

<https://cloud.google.com/architecture/best-practices-vpc-design#l7>

This architecture has multiple VPC networks that are bridged by an L7 next-generation firewall (NGFW) appliance, which functions as a multi-NIC bridge between VPC networks.

An untrusted, outside VPC network is introduced to terminate hybrid interconnects and internet-based connections that terminate on the outside leg of the L7 NGFW for inspection. There are many variations on this design, but the key principle is to filter traffic through the firewall before the traffic reaches trusted VPC networks.

upvoted 2 times

**badrik** 6 months ago

**Selected Answer: B**

B , 100% !

upvoted 1 times

**mikesp** 6 months, 3 weeks ago

**Selected Answer: D**

<https://cloud.google.com/architecture/best-practices-vpc-design>

This architecture has multiple VPC networks that are bridged by an L7 next-generation firewall (NGFW) appliance, which functions as a multi-NIC bridge between VPC networks.

upvoted 4 times

You are a member of your company's security team. You have been asked to reduce your Linux bastion host external attack surface by removing all public IP addresses. Site Reliability Engineers (SREs) require access to the bastion host from public locations so they can access the internal VPC while off-site. How should you enable this access?

- A. Implement Cloud VPN for the region where the bastion host lives.
- B. Implement OS Login with 2-step verification for the bastion host.
- C. Implement Identity-Aware Proxy TCP forwarding for the bastion host.
- D. Implement Google Cloud Armor in front of the bastion host.

**Correct Answer:** C

Reference:

<https://cloud.google.com/architecture/building-internet-connectivity-for-private-vms>

## Objectives

- Learn about the options available for connecting between private VMs outside their VPC.
- Create an instance of Identity-Aware Proxy (IAP) for TCP tunnels that's appropriate for interactive services such as SSH.
- Create a Cloud NAT instance to enable VMs to make outbound connections to the internet.
- Configure an HTTP load balancer to support inbound connections from the internet to your VMs.

Community vote distribution

C (100%)

🗲️ 👤 **AzureDP900** 1 month, 2 weeks ago

C. Implement Identity-Aware Proxy TCP forwarding for the bastion host.  
upvoted 2 times

🗲️ 👤 **mikesp** 6 months, 3 weeks ago

**Selected Answer: C**

The answer is clear in this case.  
upvoted 4 times

🗲️ 👤 **mT3** 7 months ago

**Selected Answer: C**

Correct.  
Ref.[https://cloud.google.com/architecture/building-internet-connectivity-for-private-vms#configuring\\_iap\\_tunnels\\_for\\_interacting\\_with\\_instances](https://cloud.google.com/architecture/building-internet-connectivity-for-private-vms#configuring_iap_tunnels_for_interacting_with_instances)  
upvoted 3 times

You need to enable VPC Service Controls and allow changes to perimeters in existing environments without preventing access to resources. Which VPC Service Controls mode should you use?

- A. Cloud Run
- B. Native
- C. Enforced
- D. Dry run

**Correct Answer:** D

Reference:

<https://cloud.google.com/vpc-service-controls/docs/service-perimeters>

### Dry run mode

In dry run mode, requests that violate the perimeter policy are not denied but only logged. *Dry run service perimeters* are used to test perimeter configuration and to monitor usage of services without preventing access to resources. The following are some of the common use cases:

- Determining the impact when you change existing service perimeters.
- Previewing the impact when you add new service perimeters.
- Monitoring requests to restricted services that originate from outside a service perimeter. For example, to identify from where requests to a given service are coming from or to identify unexpected service usage in your organization.
- Creating a perimeter architecture in your development environment that is analogous to your production environment. You can identify and mitigate any issues caused by your service perimeters before submitting changes to your production environment.

Community vote distribution

D (100%)

🗳️ 👤 **Tabayashi** Highly Voted 7 months, 4 weeks ago  
Answer is (D).

In dry run mode, requests that violate the perimeter policy are not denied, only logged. Dry run mode is used to test perimeter configuration and to monitor usage of services without preventing access to resources.  
<https://cloud.google.com/vpc-service-controls/docs/dry-run-mode>  
upvoted 5 times

🗳️ 👤 **AzureDP900** Most Recent 1 month, 2 weeks ago  
D -- Dry run mode  
upvoted 2 times

🗳️ 👤 **AwesomeGCP** 2 months, 2 weeks ago  
Selected Answer: D  
D. Dry run  
upvoted 1 times

🗳️ 👤 **zellck** 2 months, 3 weeks ago  
Selected Answer: D  
D is the answer.  
upvoted 2 times



You manage your organization's Security Operations Center (SOC). You currently monitor and detect network traffic anomalies in your Google Cloud VPCs based on packet header information. However, you want the capability to explore network flows and their payload to aid investigations. Which Google Cloud product should you use?

- A. Marketplace IDS
- B. VPC Flow Logs
- C. VPC Service Controls logs
- D. Packet Mirroring
- E. Google Cloud Armor Deep Packet Inspection

**Correct Answer:** D

Reference:

<https://aws.amazon.com/blogs/aws/new-vpc-traffic-mirroring/>

*Community vote distribution*

D (100%)

🗲️ 👤 **Tabayashi** Highly Voted 👍 7 months, 4 weeks ago

Answer is (D).

Packet Mirroring clones the traffic of specified instances in your Virtual Private Cloud (VPC) network and forwards it for examination. Packet Mirroring captures all traffic and packet data, including payloads and headers.

<https://cloud.google.com/vpc/docs/packet-mirroring>

upvoted 6 times

🗲️ 👤 **AzureDP900** Most Recent ⌚ 1 month, 2 weeks ago

Packet Mirroring D is right

upvoted 1 times

🗲️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: D**

D. Packet Mirroring

upvoted 1 times

🗲️ 👤 **zellick** 2 months, 3 weeks ago

**Selected Answer: D**

D is the answer.

upvoted 2 times



Your organization acquired a new workload. The Web and Application (App) servers will be running on Compute Engine in a newly created custom VPC. You are responsible for configuring a secure network communication solution that meets the following requirements:

- ☞ Only allows communication between the Web and App tiers.
- ☞ Enforces consistent network security when autoscaling the Web and App tiers.
- ☞ Prevents Compute Engine Instance Admins from altering network traffic.

What should you do?

- A. 1. Configure all running Web and App servers with respective network tags. 2. Create an allow VPC firewall rule that specifies the target/source with respective network tags.
- B. 1. Configure all running Web and App servers with respective service accounts. 2. Create an allow VPC firewall rule that specifies the target/source with respective service accounts.
- C. 1. Re-deploy the Web and App servers with instance templates configured with respective network tags. 2. Create an allow VPC firewall rule that specifies the target/source with respective network tags.
- D. 1. Re-deploy the Web and App servers with instance templates configured with respective service accounts. 2. Create an allow VPC firewall rule that specifies the target/source with respective service accounts.

**Correct Answer: A**

*Community vote distribution*

D (100%)

🗳️ **KillerGoogle** Highly Voted 7 months, 2 weeks ago

D <https://cloud.google.com/vpc/docs/firewalls#service-accounts-vs-tags>  
upvoted 10 times

🗳️ **csrazdan** Most Recent 3 weeks, 3 days ago

**Selected Answer: D**

The requirement can be fulfilled by both network tags and service accounts. To update both compute instances will have to be stopped. That means options A and B are out. Option C is out because Compute Engine Instance Admins can change network tags and avoid firewall rules. Deployment has to be done based on the instance template so that no configuration can be changed to divert the traffic.  
upvoted 2 times

🗳️ **AzureDP900** 1 month, 2 weeks ago

D is right  
upvoted 1 times

🗳️ **risc** 2 months ago

This depends on what is meant by "re-deploy"? Service accounts can also be changed by simply stopping the VM and starting it again once the SA was changed. Is this already a re-deploy?  
upvoted 2 times

🗳️ **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: D**

D. 1. Re-deploy the Web and App servers with instance templates configured with respective service accounts. 2. Create an allow VPC firewall rule that specifies the target/source with respective service accounts.  
upvoted 1 times

🗳️ **zelck** 2 months, 3 weeks ago

**Selected Answer: D**

D is the answer.

<https://cloud.google.com/vpc/docs/firewalls#service-accounts-vs-tags>

A service account represents an identity associated with an instance. Only one service account can be associated with an instance. You control access to the service account by controlling the grant of the Service Account User role for other IAM principals. For an IAM principal to start an instance by using a service account, that principal must have the Service Account User role to at least use that service account and appropriate permissions to create instances (for example, having the Compute Engine Instance Admin role to the project).

upvoted 2 times

🗳️ **cloudprincipal** 6 months, 3 weeks ago

**Selected Answer: D**

Agreed, it has to be D  
<https://cloud.google.com/vpc/docs/firewalls#service-accounts-vs-tags>  
upvoted 2 times



You need to connect your organization's on-premises network with an existing Google Cloud environment that includes one Shared VPC with two subnets named

Production and Non-Production. You are required to:

- ☞ Use a private transport link.
- ☞ Configure access to Google Cloud APIs through private API endpoints originating from on-premises environments.
- ☞ Ensure that Google Cloud APIs are only consumed via VPC Service Controls.

What should you do?

- A. 1. Set up a Cloud VPN link between the on-premises environment and Google Cloud. 2. Configure private access using the restricted.googleapis.com domains in on-premises DNS configurations.
- B. 1. Set up a Partner Interconnect link between the on-premises environment and Google Cloud. 2. Configure private access using the private.googleapis.com domains in on-premises DNS configurations.
- C. 1. Set up a Direct Peering link between the on-premises environment and Google Cloud. 2. Configure private access for both VPC subnets.
- D. 1. Set up a Dedicated Interconnect link between the on-premises environment and Google Cloud. 2. Configure private access using the restricted.googleapis.com domains in on-premises DNS configurations.

**Correct Answer: D**

Community vote distribution

D (100%)

🗉 **ExamQnA** **Highly Voted** 7 months, 1 week ago

Ans: D

restricted.googleapis.com (199.36.153.4/30) only provides access to Cloud and Developer APIs that support VPC Service Controls. VPC Service Controls are enforced for these services

<https://cloud.google.com/vpc/docs/configure-private-google-access-hybrid>

upvoted 9 times

🗉 **AzureDP900** **Most Recent** 1 month, 2 weeks ago

D. 1. Set up a Dedicated Interconnect link between the on-premises environment and Google Cloud. 2. Configure private access using the restricted.googleapis.com domains in on-premises DNS configurations.

upvoted 2 times

🗉 **sumundada** 5 months, 1 week ago

**Selected Answer: D**

restricted.googleapis.com makes it clear choice

upvoted 2 times

🗉 **cloudprincipal** 6 months, 3 weeks ago

**Selected Answer: D**

Tough call between A and D.

"☞ Use a private transport link" pushes me towards VPN connection, but the dedicated interconnect probably also fulfills that.

upvoted 2 times

🗉 **Aiffone** 5 months, 3 weeks ago

Not a tough call, VPN happens over the internet and isn't as private as dedicated interconnect...makes it a straight D

upvoted 4 times

You are working with protected health information (PHI) for an electronic health record system. The privacy officer is concerned that sensitive data is stored in the analytics system. You are tasked with anonymizing the sensitive data in a way that is not reversible. Also, the anonymized data should not preserve the character set and length. Which Google Cloud solution should you use?

- A. Cloud Data Loss Prevention with deterministic encryption using AES-SIV
- B. Cloud Data Loss Prevention with format-preserving encryption
- C. Cloud Data Loss Prevention with cryptographic hashing
- D. Cloud Data Loss Prevention with Cloud Key Management Service wrapped cryptographic keys

**Correct Answer:** D

Reference:

<https://cloud.google.com/dlp/docs/pseudonymization>

- **Encryption type:** The kind of encryption used in the de-identification transformation.
- **Supported input values:** Minimum requirements for input values.
- **Surrogate annotation:** A user-specified annotation that is prepended to encrypted values to provide context to users and to provide information for Cloud DLP to use in the re-identification of a de-identified value. A surrogate annotation is required for re-identification of unstructured data. It is optional when transforming a column of structured, or tabular, data with a **RecordTransformation**.

Community vote distribution

C (100%)

🗲️ 👤 **Tabayashi** **Highly Voted** 👍 7 months, 4 weeks ago

Answer is (C).

The only option that is irreversible is cryptographic hashing.

[https://cloud.google.com/dlp/docs/pseudonymization?hl=JA&skip\\_cache=true#supported-methods](https://cloud.google.com/dlp/docs/pseudonymization?hl=JA&skip_cache=true#supported-methods)

upvoted 13 times

🗲️ 👤 **AzureDP900** 1 month, 2 weeks ago

Agreed

C is right

upvoted 1 times

🗲️ 👤 **AwesomeGCP** **Most Recent** 🕒 2 months, 2 weeks ago

**Selected Answer: C**

C. Cloud Data Loss Prevention with cryptographic hashing

upvoted 1 times

🗲️ 👤 **sumundada** 5 months, 1 week ago

**Selected Answer: C**

<https://cloud.google.com/dlp/docs/pseudonymization>

upvoted 1 times

🗲️ 👤 **cloudprincipal** 6 months, 3 weeks ago

**Selected Answer: C**

Tabayashi is correct.

No format preserving and irreversible are the key requirements

upvoted 2 times

You are setting up a CI/CD pipeline to deploy containerized applications to your production clusters on Google Kubernetes Engine (GKE). You need to prevent containers with known vulnerabilities from being deployed. You have the following requirements for your solution:

Must be cloud-native -

- 
- ☞ Must be cost-efficient
- ☞ Minimize operational overhead

How should you accomplish this? (Choose two.)

- A. Create a Cloud Build pipeline that will monitor changes to your container templates in a Cloud Source Repositories repository. Add a step to analyze Container Analysis results before allowing the build to continue.
- B. Use a Cloud Function triggered by log events in Google Cloud's operations suite to automatically scan your container images in Container Registry.
- C. Use a cron job on a Compute Engine instance to scan your existing repositories for known vulnerabilities and raise an alert if a non-compliant container image is found.
- D. Deploy Jenkins on GKE and configure a CI/CD pipeline to deploy your containers to Container Registry. Add a step to validate your container images before deploying your container to the cluster.
- E. In your CI/CD pipeline, add an attestation on your container image when no vulnerabilities have been found. Use a Binary Authorization policy to block deployments of containers with no attestation in your cluster.

**Correct Answer:** AE

Reference:

<https://cloud.google.com/architecture/prep-kubernetes-engine-for-prod>

| Team       | IAM Role                                              | Project | Permissions                                                                                                                                                            |
|------------|-------------------------------------------------------|---------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Developers | roles/container.developer                             | dev     | Can create Kubernetes resources for the existing clusters within the project, and is not allowed to create or delete clusters.                                         |
| Operations | roles/container.admin                                 | prod    | Full administrative access to the clusters and Kubernetes resources running within the project.                                                                        |
| Security   | roles/container.viewer<br>roles/compute.securityAdmin | prod    | Create, modify, and delete firewall rules and SSL certificates as well as view resources that were created within each cluster including the logs of the running pods. |
| Network    | roles/compute.networkAdmin                            | prod    | Create, modify, and delete networking resources, except for firewall rules and SSL certificates.                                                                       |

Community vote distribution



**mikesp** Highly Voted 6 months, 3 weeks ago

**Selected Answer: AE**

On-demand container analysis can be integrated into a Cloud Build Pipeline:  
<https://cloud.google.com/container-analysis/docs/ods-cloudbuild>  
Also binary attestation is a complementary mechanism "cloud-native".  
upvoted 7 times

**zellck** Most Recent 2 months, 3 weeks ago

**Selected Answer: AE**

AE is the answer.

<https://cloud.google.com/container-analysis/docs/container-analysis>  
Container Analysis is a service that provides vulnerability scanning and metadata storage for containers. The scanning service performs vulnerability scans on images in Container Registry and Artifact Registry, then stores the resulting metadata and makes it available for consumption through an API.

<https://cloud.google.com/binary-authorization/docs/attestations>

After a container image is built, an attestation can be created to affirm that a required activity was performed on the image such as a regression test, vulnerability scan, or other test. The attestation is created by signing the image's unique digest. During deployment, instead of repeating the activities, Binary Authorization verifies the attestations using an attestor. If all of the attestations for an image are verified, Binary Authorization allows the image to be deployed.

upvoted 1 times

  **AzureDP900** 1 month, 2 weeks ago

Agreed

upvoted 1 times

  **szl0144** 7 months ago

AE is the answer, C has too much manual operations

upvoted 1 times

  **ExamQnA** 7 months, 1 week ago

Ans: A,E

[https://cloud.google.com/architecture/binary-auth-with-cloud-build-and-gke#setting\\_the\\_binary\\_authorization\\_policy](https://cloud.google.com/architecture/binary-auth-with-cloud-build-and-gke#setting_the_binary_authorization_policy)

upvoted 2 times

Which type of load balancer should you use to maintain client IP by default while using the standard network tier?

- A. SSL Proxy
- B. TCP Proxy
- C. Internal TCP/UDP
- D. TCP/UDP Network

**Correct Answer:** A

*Community vote distribution*

D (100%)

🗳️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: D**

D. TCP/UDP Network  
upvoted 3 times

🗳️ 👤 **zellck** 2 months, 3 weeks ago

**Selected Answer: D**

D is the answer.

[https://cloud.google.com/load-balancing/docs/load-balancing-overview#choosing\\_a\\_load\\_balancer](https://cloud.google.com/load-balancing/docs/load-balancing-overview#choosing_a_load_balancer)  
upvoted 3 times

🗳️ 👤 **piyush\_1982** 4 months, 3 weeks ago

**Selected Answer: D**

Definitely D  
[https://cloud.google.com/load-balancing/docs/load-balancing-overview#backend\\_region\\_and\\_network](https://cloud.google.com/load-balancing/docs/load-balancing-overview#backend_region_and_network)  
upvoted 3 times

🗳️ 👤 **mikesp** 6 months, 3 weeks ago

**Selected Answer: D**

Internal load balancer (C) is also a non-proxied load balancer but it is supported only in premium-tier networks.  
<https://cloud.google.com/load-balancing/docs/load-balancing-overview>  
upvoted 4 times

🗳️ 👤 **szl0144** 7 months ago

TCP Proxy Load Balancing terminates TCP connections from the client and creates new connections to the backends. By default, the original client IP address and port information is not preserved.  
Answer is D  
upvoted 1 times

🗳️ 👤 **AzureDP900** 1 month, 2 weeks ago

Yes, D is right  
upvoted 1 times

🗳️ 👤 **ExamQnA** 7 months ago

**Selected Answer: D**

Ans: D (though it should have been "External TCP/UDP Network load balancers")  
Cant be (C), as they are not supported on standard tier:  
<https://cloud.google.com/load-balancing/docs/load-balancing-overview>  
upvoted 4 times

🗳️ 👤 **Tabayashi** 7 months, 4 weeks ago

Answer is (C).

Use Internal TCP/UDP Load Balancing in the following circumstances:  
You need to forward the original packets unproxied. For example, if you need the client source IP address to be preserved.  
[https://cloud.google.com/load-balancing/docs/internal#use\\_cases](https://cloud.google.com/load-balancing/docs/internal#use_cases)  
upvoted 2 times

🗳️ 👤 **Arturo\_Cloud** 3 months, 2 weeks ago

I disagree with you, both C and D can keep the Client IP, however only TCP/UDP Network is for standard network.

<https://cloud.google.com/load-balancing/docs/network>  
<https://cloud.google.com/load-balancing/docs/load-balancing-overview>

upvoted 1 times



You want to prevent users from accidentally deleting a Shared VPC host project. Which organization-level policy constraint should you enable?

- A. compute.restrictSharedVpcHostProjects
- B. compute.restrictXpnProjectLienRemoval
- C. compute.restrictSharedVpcSubnetworks
- D. compute.sharedReservationsOwnerProjects

**Correct Answer:** B

Reference:

<https://cloud.google.com/vpc/docs/provisioning-shared-vpc>

Console

gcloud

API

Terraform

1. Go to the VM instances page in the Google Cloud Console.

Go to the VM instances page

2. Click **Create**.
3. Specify a **Name** for the instance.
4. Click **Management, security, disks, networking, sole tenancy**.
5. Click **Networking**.
6. Under **Network interfaces**, click the `default` network.
7. Click the **Networks shared with me** radio button.
8. Select the **Shared subnet** where you want to create the instance.
9. Specify any other necessary parameters for the instance.
10. Click **Create**.

Community vote distribution

B (100%)

 **Tabayashi** Highly Voted 7 months, 4 weeks ago

Answer is (B).

This boolean constraint restricts the set of users that can remove a Shared VPC project lien without organization-level permission where this constraint is set to True.

<https://cloud.google.com/resource-manager/docs/organization-policy/org-policy-constraints>

upvoted 5 times

 **zelick** Most Recent 2 months, 3 weeks ago

**Selected Answer: B**

B is the answer.

<https://cloud.google.com/resource-manager/docs/organization-policy/org-policy-constraints#constraints-for-specific-services>

- constraints/compute.restrictXpnProjectLienRemoval

- Restrict shared VPC project lien removal

This boolean constraint restricts the set of users that can remove a Shared VPC host project lien without organization-level permission where this constraint is set to True.

By default, any user with the permission to update liens can remove a Shared VPC host project lien. Enforcing this constraint requires that permission be granted at the organization level.

upvoted 4 times

 **AzureDP900** 1 month, 2 weeks ago

Agree with your explanation and Thank you for sharing the link

upvoted 1 times

 **mikesp** 6 months, 3 weeks ago

**Selected Answer: B**

<https://cloud.google.com/resource-manager/docs/organization-policy/org-policy-constraints>  
upvoted 1 times

## Question #128

Topic 1

Users are reporting an outage on your public-facing application that is hosted on Compute Engine. You suspect that a recent change to your firewall rules is responsible. You need to test whether your firewall rules are working properly. What should you do?

- A. Enable Firewall Rules Logging on the latest rules that were changed. Use Logs Explorer to analyze whether the rules are working correctly.
- B. Connect to a bastion host in your VPC. Use a network traffic analyzer to determine at which point your requests are being blocked.
- C. In a pre-production environment, disable all firewall rules individually to determine which one is blocking user traffic.
- D. Enable VPC Flow Logs in your VPC. Use Logs Explorer to analyze whether the rules are working correctly.

**Correct Answer: A**

Reference:

<https://cloud.google.com/vpc/docs/using-flow-logs>

*Community vote distribution*

A (100%)

 **mikesp** **Highly Voted** 6 months, 3 weeks ago

**Selected Answer: A**

<https://cloud.google.com/vpc/docs/firewall-rules-logging>  
upvoted 5 times

 **AzureDP900** **Most Recent** 1 month, 2 weeks ago

A is right  
upvoted 1 times

 **ExamQnA** 7 months, 1 week ago

Ans:A  
<https://cloud.google.com/vpc/docs/firewall-rules-logging>  
upvoted 4 times



You are a security administrator at your company. Per Google-recommended best practices, you implemented the domain restricted sharing organization policy to allow only required domains to access your projects. An engineering team is now reporting that users at an external partner outside your organization domain cannot be granted access to the resources in a project. How should you make an exception for your partner's domain while following the stated best practices?

- A. Turn off the domain restriction sharing organization policy. Set the policy value to "Allow All."
- B. Turn off the domain restricted sharing organization policy. Provide the external partners with the required permissions using Google's Identity and Access Management (IAM) service.
- C. Turn off the domain restricted sharing organization policy. Add each partner's Google Workspace customer ID to a Google group, add the Google group as an exception under the organization policy, and then turn the policy back on.
- D. Turn off the domain restricted sharing organization policy. Set the policy value to "Custom." Add each external partner's Cloud Identity or Google Workspace customer ID as an exception under the organization policy, and then turn the policy back on.

**Correct Answer:** D

Community vote distribution

D (82%)

C (18%)

  **mikesp** Highly Voted 6 months, 3 weeks ago

**Selected Answer: D**

The question is that is necessary to add identities from another Domain to cloud identity. The only way to do that is by adding the Customer Ids as exception. The procedure does not support adding groups, etc...

The groups and the corresponding users can be added later on with Cloud Identity once that the domain of their organization is allowed: The allowed\_values are Google Workspace customer IDs, such as C03xgje4y. Only identities belonging to a Google Workspace domain from the list of allowed\_values will be allowed on IAM policies once this organization policy has been applied. Google Workspace human users and groups must be part of that Google Workspace domain, and IAM service accounts must be children of an organization resource associated with the given Google Workspace domain

upvoted 8 times

  **AzureDP900** 1 month, 2 weeks ago

Agreed with your explanaiton

upvoted 1 times

  **AwesomeGCP** Most Recent 2 months, 2 weeks ago

**Selected Answer: D**

D. Turn off the domain restricted sharing organization policy. Set the policy value to "Custom." Add each external partner's Cloud Identity or Google Workspace customer ID as an exception under the organization policy, and then turn the policy back on.

upvoted 2 times

  **zellck** 2 months, 3 weeks ago

**Selected Answer: D**

D is the answer.

[https://cloud.google.com/resource-manager/docs/organization-policy/restricting-domains#setting\\_the\\_organization\\_policy](https://cloud.google.com/resource-manager/docs/organization-policy/restricting-domains#setting_the_organization_policy)

The domain restriction constraint is a type of list constraint. Google Workspace customer IDs can be added and removed from the allowed\_values list of a domain restriction constraint. The domain restriction constraint does not support denying values, and an organization policy can't be saved with IDs in the denied\_values list.

All domains associated with a Google Workspace account listed in the allowed\_values will be allowed by the organization policy. All other domains will be denied by the organization policy.

upvoted 2 times

  **AzureDP900** 1 month, 2 weeks ago

Thank you for detailed explanation

upvoted 1 times

  **sumundada** 5 months, 1 week ago

**Selected Answer: D**

[https://cloud.google.com/resource-manager/docs/organization-policy/restricting-domains#setting\\_the\\_organization\\_policy](https://cloud.google.com/resource-manager/docs/organization-policy/restricting-domains#setting_the_organization_policy)

upvoted 2 times

  **Medofree** 7 months ago

The right answer is D.

Because we add the "Customer ID" as exception and not Google group.

[https://cloud.google.com/resource-manager/docs/organization-policy/restricting-domains#setting\\_the\\_organization\\_policy](https://cloud.google.com/resource-manager/docs/organization-policy/restricting-domains#setting_the_organization_policy)  
upvoted 1 times

  **bartlomiejwaw** 7 months, 2 weeks ago

**Selected Answer: C**

Policy should be turned on at the end. Adding the whole group as an exception is far more reasonable than adding all identities.  
upvoted 3 times

  **mT3** 7 months ago

I agree

Ref: [https://cloud.google.com/resource-manager/docs/organization-policy/restricting-domains#setting\\_the\\_organization\\_policy](https://cloud.google.com/resource-manager/docs/organization-policy/restricting-domains#setting_the_organization_policy)  
upvoted 1 times

You plan to use a Google Cloud Armor policy to prevent common attacks such as cross-site scripting (XSS) and SQL injection (SQLi) from reaching your web application's backend. What are two requirements for using Google Cloud Armor security policies? (Choose two.)

- A. The load balancer must be an external SSL proxy load balancer.
- B. Google Cloud Armor Policy rules can only match on Layer 7 (L7) attributes.
- C. The load balancer must use the Premium Network Service Tier.
- D. The backend service's load balancing scheme must be EXTERNAL.
- E. The load balancer must be an external HTTP(S) load balancer.

**Correct Answer:** *BE*

Reference:

<https://cloud.google.com/armor>

*Community vote distribution*

DE (100%)

 **AzureDP900** 1 month, 2 weeks ago

D,E is most appropriate in this case

D. The backend service's load balancing scheme must be EXTERNAL.

E. The load balancer must be an external HTTP(S) load balancer.

upvoted 2 times

 **soltium** 2 months, 1 week ago

**Selected Answer: DE**

DE.

Well technically you can use EXTERNAL\_MANAGED scheme too.

upvoted 1 times

 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: DE**

D. The backend service's load balancing scheme must be EXTERNAL.

E. The load balancer must be an external HTTP(S) load balancer.

upvoted 1 times

 **Jeanphi72** 4 months ago

**Selected Answer: DE**

<https://cloud.google.com/armor/docs/security-policy-overview#requirements> says:

The backend service's load balancing scheme must be EXTERNAL, or EXTERNAL\_MANAGED \*\*\* if you are using global external HTTP(S) load balancer \*\*\*.

Thus D and E fit (A could fit if a suggestion like The backend service's load balancing scheme must \*\* NOT \*\* be EXTERNAL

upvoted 1 times

 **piyush\_1982** 4 months, 3 weeks ago

I am not sure if there is some mistake in the question or in the options given.

<https://cloud.google.com/armor/docs/security-policy-overview#requirements>

As per the link above, below are the requirements for using Google Cloud Armor security policies:

1. The load balancer must be a global external HTTP(S) load balancer, global external HTTP(S) load balancer (classic), external TCP proxy load balancer, or external SSL proxy load balancer.
2. The backend service's load balancing scheme must be EXTERNAL, or EXTERNAL\_MANAGED if you are using a global external HTTP(S) load balancer.
3. The backend service's protocol must be one of HTTP, HTTPS, HTTP/2, TCP, or SSL.

The correct answer seems to be A D and E.

A. The load balancer must be an external SSL proxy load balancer. (external SSL proxy load balancer is one of the load balancing options listed in the link)

D. The backend service's load balancing scheme must be EXTERNAL. (or EXTERNAL\_MANAGED)

E. The load balancer must be an external HTTP(S) load balancer. (Also one of the options listed)

upvoted 1 times

 **zellck** 2 months, 3 weeks ago

Security policy for A does not block XSS and SQLi which is at layer 7.  
<https://cloud.google.com/armor/docs/security-policy-overview#policy-types>  
upvoted 1 times

🗨️ 👤 **nacying** 6 months, 2 weeks ago

**Selected Answer: DE**

These are the requirements for using Google Cloud Armor security policies:

The load balancer must be an external HTTP(S) load balancer, TCP proxy load balancer, or SSL proxy load balancer.  
The backend service's load balancing scheme must be EXTERNAL.  
The backend service's protocol must be one of HTTP, HTTPS, HTTP/2, TCP, or SSL.  
<https://cloud.google.com/armor/docs/security-policy-overview>  
upvoted 1 times

🗨️ 👤 **cloudprincipal** 6 months, 3 weeks ago

**Selected Answer: DE**

DE

Requirements

These are the requirements for using Google Cloud Armor security policies:

- \* The load balancer must be an external HTTP(S) load balancer, TCP proxy load balancer, or SSL proxy load balancer.
- \* The backend service's load balancing scheme must be EXTERNAL.
- \* The backend service's protocol must be one of HTTP, HTTPS, HTTP/2, TCP, or SSL.

See <https://cloud.google.com/armor/docs/security-policy-overview#requirements>  
upvoted 2 times

🗨️ 👤 **szl0144** 7 months ago

Google Cloud Armor security policies are sets of rules that match on attributes from Layer 3 to Layer 7 to protect externally facing applications or services. Each rule is evaluated with respect to incoming traffic.  
I choose DE  
upvoted 1 times

🗨️ 👤 **ExamQnA** 7 months, 1 week ago

Ans:D,E

<https://cloud.google.com/armor/docs/security-policy-overview>

Relevant extracts:

1. Google Cloud Armor security policies enable you to rate-limit or redirect requests to your HTTP(S) Load Balancing, TCP Proxy Load Balancing, or SSL Proxy Load Balancing ...
2. Google Cloud Armor security policies are sets of rules that match on attributes from Layer 3 to Layer 7 to protect externally facing applications or services...
3. The load balancer can be in Premium Tier or Standard Tier.

upvoted 3 times

🗨️ 👤 **Tabayashi** 7 months, 4 weeks ago

Answer is (B), (D).

Google Cloud Armor security policies protect your application by providing Layer 7 filtering and by scrubbing incoming requests for common web attacks or other Layer 7 attributes to potentially block traffic before it reaches your load balanced backend services or backend buckets.  
Google Cloud Armor security policies are available only for backend services behind an external HTTP(S) load balancer, TCP proxy load balancer, or an SSL proxy load balancer. The load balancer can be in Premium Tier or Standard Tier.  
<https://cloud.google.com/armor/docs/security-policy-overview>  
upvoted 3 times

🗨️ 👤 **Taliesyn** 7 months, 2 weeks ago

I'd say D, E, even though it's redundant.  
Definitely not B, as you can filter on IP addresses, ie layer 3.  
upvoted 4 times

You perform a security assessment on a customer architecture and discover that multiple VMs have public IP addresses. After providing a recommendation to remove the public IP addresses, you are told those VMs need to communicate to external sites as part of the customer's typical operations. What should you recommend to reduce the need for public IP addresses in your customer's VMs?

- A. Google Cloud Armor
- B. Cloud NAT
- C. Cloud Router
- D. Cloud VPN

**Correct Answer:** *B*

*Community vote distribution*

B (100%)

🗉 👤 **AzureDP900** 1 month, 2 weeks ago

B Cloud NAT  
upvoted 2 times

🗉 👤 **Random\_Mane** 3 months, 2 weeks ago

**Selected Answer: B**

B. <https://cloud.google.com/nat/docs/overview>  
upvoted 4 times

You are tasked with exporting and auditing security logs for login activity events for Google Cloud console and API calls that modify configurations to Google

Cloud resources. Your export must meet the following requirements:

- ☞ Export related logs for all projects in the Google Cloud organization.
- ☞ Export logs in near real-time to an external SIEM.

What should you do? (Choose two.)

- A. Create a Log Sink at the organization level with a Pub/Sub destination.
- B. Create a Log Sink at the organization level with the includeChildren parameter, and set the destination to a Pub/Sub topic.
- C. Enable Data Access audit logs at the organization level to apply to all projects.
- D. Enable Google Workspace audit logs to be shared with Google Cloud in the Admin Console.
- E. Ensure that the SIEM processes the AuthenticationInfo field in the audit log entry to gather identity information.

**Correct Answer:** AE

Reference:

<https://www.datadoghq.com/blog/monitoring-gcp-audit-logs/>

Community vote distribution

BD (100%)

🗳️ 👤 **cloudprincipal** Highly Voted 6 months, 3 weeks ago

**Selected Answer: BD**

B  
because for all projects

D  
"Google Workspace Login Audit: Login Audit logs track user sign-ins to your domain. These logs only record the login event. They don't record which system was used to perform the login action."  
<https://cloud.google.com/logging/docs/audit/gsuite-audit-logging#services>  
upvoted 9 times

🗳️ 👤 **Medofree** Highly Voted 7 months ago

Correct answers are : B,D

B : to respond to the "logs for all projects" requirement and " near real-time "requirement  
D: to be able de log "login activities" we need to export audit logs from Google Workspace to Google Cloud.  
upvoted 6 times

🗳️ 👤 **AzureDP900** Most Recent 1 month, 2 weeks ago

B,D is right  
upvoted 1 times

🗳️ 👤 **ExamQnA** 7 months, 1 week ago

Ans:B,C  
[https://cloud.google.com/logging/docs/export/aggregated\\_sinks](https://cloud.google.com/logging/docs/export/aggregated_sinks): To use aggregated sinks, you create a sink in a Google Cloud organization or folder, and set the sink's includeChildren parameter to True. That sink can then route log entries from the organization or folder, plus (recursively) from any contained folders, billing accounts, or Cloud projects.  
<https://cloud.google.com/logging/docs/audit#data-access>  
Data Access audit logs-- except for BigQuery Data Access audit logs-- are disabled by default because audit logs can be quite large. If you want Data Access audit logs to be written for Google Cloud services other than BigQuery, you must explicitly enable them  
upvoted 5 times



Your company's Chief Information Security Officer (CISO) creates a requirement that business data must be stored in specific locations due to regulatory requirements that affect the company's global expansion plans. After working on the details to implement this requirement, you determine the following:


- ☞ The services in scope are included in the Google Cloud Data Residency Terms.
- ☞ The business data remains within specific locations under the same organization.
- ☞ The folder structure can contain multiple data residency locations.

You plan to use the Resource Location Restriction organization policy constraint. At which level in the resource hierarchy should you set the constraint?

- A. Folder
- B. Resource
- C. Project
- D. Organization

**Correct Answer: B**

*Community vote distribution*



🗳️ 👤 **mouchu** Highly Voted 7 months, 1 week ago

Answer = C

"The folder structure can contain multiple data residency locations" suggest that restriction should be applied on projects level  
upvoted 12 times

🗳️ 👤 **AzureDP900** 1 month, 1 week ago

Yes, It is C. This is very tricky question and we need to read very carefully. In general Folders will used but in this case Project is right  
upvoted 1 times

🗳️ 👤 **AzureDP900** 1 month, 1 week ago

Q 137 is same  
upvoted 1 times

🗳️ 👤 **piyush\_1982** 4 months, 3 weeks ago

why not D?  
upvoted 1 times

🗳️ 👤 **Taliesyn** Highly Voted 7 months, 2 weeks ago

**Selected Answer: A**

Org policies can't be applied on resources ...  
upvoted 5 times

🗳️ 👤 **adelynllllllllll** Most Recent 3 weeks, 5 days ago

the answer should be B  
<https://cloud.google.com/resource-manager/docs/organization-policy/defining-locations>  
upvoted 1 times

🗳️ 👤 **Rightsaidfred** 1 month ago

**Selected Answer: C**

Different Locations therefore needs to be applied at Project Level.  
upvoted 1 times

🗳️ 👤 **TonytheTiger** 1 month ago

To set an organization policy including a resource locations constraint:  
<https://cloud.google.com/resource-manager/docs/organization-policy/defining-locations>  
upvoted 1 times

🗳️ 👤 **AzureDP900** 1 month, 1 week ago

C is right  
upvoted 1 times

🗳️ 👤 **AzureDP900** 1 month, 1 week ago

137 and 133 questions are similar except ☞ The projects are aligned to specific locations.  
upvoted 1 times

🗨️ 👤 **AzureDP900** 1 month, 1 week ago  
C is right in my poinion.  
upvoted 1 times

🗨️ 👤 **rotorclear** 2 months ago

**Selected Answer: C**

I initially thought it will be folder but there are multiple data residencies in a folder hence it will conflict  
upvoted 1 times

🗨️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: A**

A. Folder  
upvoted 1 times

🗨️ 👤 **cjjatta** 3 months, 1 week ago

**Selected Answer: C**

The folder structure can contain multiple data residency locations  
upvoted 1 times

🗨️ 👤 **GHOST1985** 3 months, 2 weeks ago

**Selected Answer: A**

Answer A,  
the requirement says : "The business data remains within specific locations under the same organization" that's mean you need to set up the restriction at higher level so folder is the correct answer => Answer A  
upvoted 3 times

🗨️ 👤 **piyush\_1982** 4 months, 3 weeks ago

**Selected Answer: D**

I think D is the best solution.

<https://cloud.google.com/assured-workloads/docs/concept-data-residency>

Organizations with data residency requirements can set up a Resource Locations policy that constrains the location of new in-scope resources for their whole organization or for individual projects.

As per the link above it says for such a requirement you can set up the resource location constraint at Org or Individual Project level. Setting up at the individual project level will be exhausting hence setting up at Org level would be the best way to move forward.  
upvoted 2 times

🗨️ 👤 **mikesp** 6 months, 3 weeks ago

**Selected Answer: A**

Since it is possible to set org policy constraints at project and folder level and "The folder structure can contain multiple data residency locations", i would set the policy at folder level.

<https://cloud.google.com/resource-manager/docs/organization-policy/using-constraints>

"You can also apply the organization policy to a folder or a project with the --folder or the --project flags, and the folder ID and project ID, respectively."

upvoted 1 times

🗨️ 👤 **szl0144** 7 months ago

**Selected Answer: C**

C is the answer, already have the organization policy to set all data in the same region but need to change the policy at the folder level  
upvoted 4 times



You need to set up a Cloud interconnect connection between your company's on-premises data center and VPC host network. You want to make sure that on-premises applications can only access Google APIs over the Cloud Interconnect and not through the public internet. You are required to only use APIs that are supported by VPC Service Controls to mitigate against exfiltration risk to non-supported APIs. How should you configure the network?

- A. Enable Private Google Access on the regional subnets and global dynamic routing mode.
- B. Set up a Private Service Connect endpoint IP address with the API bundle of "all-apis", which is advertised as a route over the Cloud interconnect connection.
- C. Use private.googleapis.com to access Google APIs using a set of IP addresses only routable from within Google Cloud, which are advertised as routes over the connection.
- D. Use restricted.googleapis.com to access Google APIs using a set of IP addresses only routable from within Google Cloud, which are advertised as routes over the Cloud Interconnect connection.

**Correct Answer:** B

Reference:

<https://cloud.google.com/network-connectivity/docs/interconnect/concepts/overview>

Community vote distribution

D (100%)

  **Nicky1402** Highly Voted 7 months, 2 weeks ago

I think the correct answer is D.

It is mentioned in the question: "You are required to only use APIs that are supported by VPC Service Controls", from which we can understand that we cannot use private.googleapis.com. Hence, option A & C can be eliminated.

API bundle with all-apis is mentioned in option B which is wrong as we want to use only those APIs supported by VPC service controls. Hence, option B can be eliminated.

Option D has all the solutions we need.

<https://cloud.google.com/vpc/docs/private-service-connect>

An API bundle:

All APIs (all-apis): most Google APIs

(same as private.googleapis.com).

VPC-SC (vpc-sc): APIs that VPC Service Controls supports

(same as restricted.googleapis.com).

VMs in the same VPC network as the endpoint (all regions)

On-premises systems that are connected to the VPC network that contains the endpoint

upvoted 9 times

  **AzureDP900** 1 month, 2 weeks ago

Yes, It is D

upvoted 1 times

  **AzureDP900** 1 month, 2 weeks ago

D. Use restricted.googleapis.com to access Google APIs using a set of IP addresses only routable from within Google Cloud, which are advertised as routes over the Cloud Interconnect connection.

upvoted 1 times

  **samuelmorher** Most Recent 3 days, 19 hours ago

**Selected Answer: D**

it's D

upvoted 1 times

  **marmar11111** 1 month, 1 week ago

**Selected Answer: D**

<https://cloud.google.com/vpc/docs/configure-private-google-access-hybrid>

Choose restricted.googleapis.com when you only need access to Google APIs and services that are supported by VPC Service Controls.

upvoted 1 times

  **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: D**

D. Use restricted.googleapis.com to access Google APIs using a set of IP addresses only routable from within Google Cloud, which are advertised as routes over the Cloud Interconnect connection.

upvoted 2 times

  **zellck** 2 months, 3 weeks ago

**Selected Answer: D**

D is the answer.

<https://cloud.google.com/vpc/docs/configure-private-google-access-hybrid#config-choose-domain>

If you need to restrict users to just the Google APIs and services that support VPC Service Controls, use `restricted.googleapis.com`. Although VPC Service Controls are enforced for compatible and configured services, regardless of the domain you use, `restricted.googleapis.com` provides additional risk mitigation for data exfiltration. Using `restricted.googleapis.com` denies access to Google APIs and services that are not supported by VPC Service Controls.

upvoted 1 times

  **bnikunj** 3 months, 2 weeks ago

D is answer, <https://cloud.google.com/vpc/docs/configure-private-service-connect-apis#supported-apis>

The all-apis bundle provides access to the same APIs as `private.googleapis.com`

Choose `vpc-sc` when you only need access to Google APIs and services that are supported by VPC Service Controls. The `vpc-sc` bundle does not permit access to Google APIs and services that do not support VPC Service Controls. 1

upvoted 1 times

  **cloudprincipal** 6 months, 3 weeks ago

**Selected Answer: D**

Will agree with the others

upvoted 2 times

  **cloudprincipal** 6 months, 1 week ago

This is actually specified in the documentation: <https://cloud.google.com/vpc/docs/configure-private-google-access-hybrid#config-choose-domain>

upvoted 3 times

  **ExamQnA** 7 months, 1 week ago

Ans: D

Note: If you need to restrict users to just the Google APIs and services that support VPC Service Controls, use `restricted.googleapis.com`.

<https://cloud.google.com/vpc/docs/configure-private-google-access-hybrid>

upvoted 3 times

You need to implement an encryption-at-rest strategy that protects sensitive data and reduces key management complexity for non-sensitive data. Your solution has the following requirements:


- ☞ Schedule key rotation for sensitive data.
- ☞ Control which region the encryption keys for sensitive data are stored in.
- ☞ Minimize the latency to access encryption keys for both sensitive and non-sensitive data.

What should you do?

- A. Encrypt non-sensitive data and sensitive data with Cloud External Key Manager.
- B. Encrypt non-sensitive data and sensitive data with Cloud Key Management Service.
- C. Encrypt non-sensitive data with Google default encryption, and encrypt sensitive data with Cloud External Key Manager.
- D. Encrypt non-sensitive data with Google default encryption, and encrypt sensitive data with Cloud Key Management Service.

**Correct Answer: B**

*Community vote distribution*



🗳️ 👤 **GHOST1985** Highly Voted 👍 3 months, 2 weeks ago

**Selected Answer: D**

Answer D  
because "Minimize the latency to access encryption keys"  
upvoted 6 times

🗳️ 👤 **GHOST1985** 2 months, 2 weeks ago

Sorry answer is B  
upvoted 2 times

🗳️ 👤 **TonytheTiger** Most Recent ⌚ 1 month ago

Answer D  
<https://cloud.google.com/docs/security/encryption/default-encryption>  
upvoted 3 times

🗳️ 👤 **marmar11111** 1 month, 1 week ago

**Selected Answer: D**

The default already has low latency! "Because of the high volume of keys at Google, and the need for low latency and high availability, DEKs are stored near the data that they encrypt. DEKs are encrypted with (wrapped by) a key encryption key (KEK), using a technique known as envelope encryption. These KEKs are not specific to customers; instead, one or more KEKs exist for each service."

We need less complexity and low latency so use default on non-sensitive data!  
upvoted 2 times

🗳️ 👤 **AzureDP900** 1 month, 2 weeks ago

B. Encrypt non-sensitive data and sensitive data with Cloud Key Management Service.  
upvoted 1 times

🗳️ 👤 **coco10k** 1 month, 3 weeks ago

**Selected Answer: D**

keeps complexity low  
upvoted 2 times

🗳️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: B**

B. Encrypt non-sensitive data and sensitive data with Cloud Key Management Service.  
upvoted 1 times

🗳️ 👤 **GHOST1985** 2 months, 2 weeks ago

**Selected Answer: B**

☞ Schedule key rotation for sensitive data. :  
=> Cloud KMS allows you to set a rotation schedule for symmetric keys to automatically generate a new key version at a fixed time interval. Multiple versions of a symmetric key can be active at any time for decryption, with only one primary key version used for encrypting new data. With EKM, create an externally managed key directly from the Cloud KSM console.

☞ Control which region the encryption keys for sensitive data are stored in.  
=> If using Cloud KMS, your cryptographic keys will be stored in the region where you deploy the resource. You also have the option of storing

those keys inside a physical Hardware Security Module located in the region you choose with Cloud HSM.

☞ Minimize the latency to access encryption keys for both sensitive and non-sensitive data :

=> Cloud KMS is available in several global locations and across multi-regions, allowing you to place your service where you want for low latency and high availability.

<https://cloud.google.com/security-key-management>

upvoted 2 times


Your security team uses encryption keys to ensure confidentiality of user data. You want to establish a process to reduce the impact of a potentially compromised symmetric encryption key in Cloud Key Management Service (Cloud KMS).

Which steps should your team take before an incident occurs? (Choose two.)

- A. Disable and revoke access to compromised keys.
- B. Enable automatic key version rotation on a regular schedule.
- C. Manually rotate key versions on an ad hoc schedule.
- D. Limit the number of messages encrypted with each key version.
- E. Disable the Cloud KMS API.

**Correct Answer:** AB

Community vote distribution



**zelick** Highly Voted 2 months, 3 weeks ago

**Selected Answer: BD**

BD is the answer. The steps need to be done BEFORE an incident occurs.

upvoted 6 times

**AzureDP900** 1 month, 2 weeks ago

Yes, B and D

upvoted 1 times

**parasthakur** Highly Voted 3 months ago

**Selected Answer: BD**

Should be BD. A is wrong because there is no comprise happened as the question states "before an incident".

As per document "Limiting the number of messages encrypted with the same key version helps prevent attacks enabled by cryptanalysis."  
<https://cloud.google.com/kms/docs/key-rotation>

upvoted 5 times

**PST21** Most Recent 2 days, 19 hours ago

You want to reduce the impact - which will be post the issue has occurred so has to be AB.

If asked for preventive steps then B & D.

upvoted 1 times

**spiritix821** 6 days, 10 hours ago

<https://cloud.google.com/kms/docs/key-rotation> -> "If you suspect that a key version is compromised, disable it and revoke access to it as soon as possible" so A could be correct. do you agree?

upvoted 1 times

**AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: BD**

B. Enable automatic key version rotation on a regular schedule.

D. Limit the number of messages encrypted with each key version.

upvoted 2 times

**GHOST1985** 2 months, 2 weeks ago

**Selected Answer: BD**

Answers BD

upvoted 1 times

**parasthakur** 3 months ago

Should be BD. A is wrong because there is no comprise happened as the question states "before an incident".

As per document "Limiting the number of messages encrypted with the same key version helps prevent attacks enabled by cryptanalysis."  
<https://cloud.google.com/kms/docs/key-rotation>

upvoted 1 times

**satyajeet1999** 3 months, 1 week ago

**Selected Answer: AB**

should be AB.

upvoted 1 times





Your company's chief information security officer (CISO) is requiring business data to be stored in specific locations due to regulatory requirements that affect the company's global expansion plans. After working on a plan to implement this requirement, you determine the following:


- ↻ The services in scope are included in the Google Cloud data residency requirements.
- ↻ The business data remains within specific locations under the same organization.
- ↻ The folder structure can contain multiple data residency locations.
- ↻ The projects are aligned to specific locations.

You plan to use the Resource Location Restriction organization policy constraint with very granular control. At which level in the hierarchy should you set the constraint?

- A. Organization
- B. Resource
- C. Project
- D. Folder

Correct Answer: A

Community vote distribution



- 🗨️

TonytheTiger

1 week, 6 days ago

on the exam

upvoted 1 times
- 🗨️

Littleivy

1 month, 1 week ago

Selected Answer: C

Need to be in project level to have required granularity

upvoted 2 times
- 🗨️

AzureDP900

1 month, 1 week ago

D should be right , This is same as question 133

upvoted 3 times
- 🗨️

AzureDP900

1 month, 1 week ago

sorry it is C

upvoted 2 times
- 🗨️

AwesomeGCP

1 month, 2 weeks ago

Selected Answer: C

C. Project

upvoted 2 times
- 🗨️

coco10k

1 month, 3 weeks ago

Selected Answer: C

most granular

upvoted 1 times
- 🗨️

soltium

2 months, 1 week ago

I think its C.

A and D will inherits the org policy which make it easier to manage, but the opposite of granular.

For B, org policy cannot be applied to resource.

upvoted 1 times
- 🗨️

TheBuckler

2 months, 1 week ago

Answer is C. The key word here is "very granular\* control". Most granular choice here is Project, as you cannot apply policy constraints to resources.

upvoted 3 times
- 🗨️

GHOST1985

3 months, 2 weeks ago

Selected Answer: D

i Woul say D,

Same question then 133, with new requirement for project is aligned for a specific location

i think it is better to set up the restriction at higher level "Organization" so all the childrens (folders, projects) inherit the residency location restriction

upvoted 2 times

  **GHOST1985** 3 months, 2 weeks ago  
sorry i mean Answer A  
upvoted 2 times

  **AzureDP900** 1 month, 1 week ago  
C is right  
upvoted 1 times

  **AzureDP900** 1 month, 1 week ago  
D is right  
upvoted 1 times

  **AzureDP900** 1 month, 1 week ago  
Sorry it is C  
upvoted 1 times

Question #138

Topic 1

A database administrator notices malicious activities within their Cloud SQL instance. The database administrator wants to monitor the API calls that read the configuration or metadata of resources. Which logs should the database administrator review?

- A. Admin Activity
- B. System Event
- C. Access Transparency
- D. Data Access

**Correct Answer:** C

Community vote distribution



  **GHOST1985** **Highly Voted**  3 months, 2 weeks ago  
**Selected Answer: D**  
answer D  
Data Access audit logs contain API calls that read the configuration or metadata of resources, as well as user-driven API calls that create, modify, or read user-provided resource data.  
upvoted 5 times

  **AzureDP900** 1 month, 2 weeks ago  
D. Data Access  
upvoted 1 times

  **AwesomeGCP** **Most Recent**  2 months, 2 weeks ago  
**Selected Answer: D**  
D. Data Access  
upvoted 1 times

  **Random\_Mane** 3 months, 2 weeks ago  
**Selected Answer: D**  
D. <https://cloud.google.com/logging/docs/audit/#data-access>  
"Data Access audit logs contain API calls that read the configuration or metadata of resources, as well as user-driven API calls that create, modify, or read user-provided resource data."  
upvoted 2 times

  **Baburao** 3 months, 3 weeks ago  
Should be D  
<https://cloud.google.com/logging/docs/audit/#data-access>  
upvoted 2 times



You are backing up application logs to a shared Cloud Storage bucket that is accessible to both the administrator and analysts. Analysts should not have access to logs that contain any personally identifiable information (PII). Log files containing PII should be stored in another bucket that is only accessible to the administrator. What should you do?

- A. Upload the logs to both the shared bucket and the bucket with PII that is only accessible to the administrator. Use the Cloud Data Loss Prevention API to create a job trigger. Configure the trigger to delete any files that contain PII from the shared bucket.
- B. On the shared bucket, configure Object Lifecycle Management to delete objects that contain PII.
- C. On the shared bucket, configure a Cloud Storage trigger that is only triggered when PII is uploaded. Use Cloud Functions to capture the trigger and delete the files that contain PII.
- D. Use Pub/Sub and Cloud Functions to trigger a Cloud Data Loss Prevention scan every time a file is uploaded to the administrator's bucket. If the scan does not detect PII, have the function move the objects into the shared Cloud Storage bucket.

**Correct Answer:** *D*

  **jitu028** Highly Voted 2 months, 3 weeks ago

Answer is D

upvoted 5 times

  **AzureDP900** Most Recent 1 month, 2 weeks ago

D. Use Pub/Sub and Cloud Functions to trigger a Cloud Data Loss Prevention scan every time a file is uploaded to the administrator's bucket. If the scan does not detect PII, have the function move the objects into the shared Cloud Storage bucket

upvoted 4 times

You work for an organization in a regulated industry that has strict data protection requirements. The organization backs up their data in the cloud. To comply with data privacy regulations, this data can only be stored for a specific length of time and must be deleted after this specific period.

You want to automate the compliance with this regulation while minimizing storage costs. What should you do?

- A. Store the data in a persistent disk, and delete the disk at expiration time.
- B. Store the data in a Cloud Bigtable table, and set an expiration time on the column families.
- C. Store the data in a BigQuery table, and set the table's expiration time.
- D. Store the data in a Cloud Storage bucket, and configure the bucket's Object Lifecycle Management feature.

**Correct Answer:** C

*Community vote distribution*

D (100%)

🗲️ 👤 **Baburao** Highly Voted 👍 3 months, 3 weeks ago

should be D.

To minimize costs, it's always GCS even though BQ comes as a close 2nd. But, since the question did not specify what kind of data it is (raw files vs tabular data), it is safe to assume GCS is the preferred option with Lifecycle enablement.

upvoted 6 times

🗲️ 👤 **TonytheTiger** Most Recent ⌚ 1 month ago

D is the answer.

<https://cloud.google.com/storage/docs/lifecycle>

upvoted 2 times

🗲️ 👤 **AzureDP900** 1 month, 2 weeks ago

D. Store the data in a Cloud Storage bucket, and configure the bucket's Object Lifecycle Management feature.

upvoted 1 times

🗲️ 👤 **zellck** 2 months, 3 weeks ago

**Selected Answer: D**

D is the answer.

upvoted 1 times

🗲️ 👤 **GHOST1985** 3 months ago

**Selected Answer: D**

GCS is the preferred option with Lifecycle enablement.

upvoted 1 times

You have been tasked with configuring Security Command Center for your organization's Google Cloud environment. Your security team needs to receive alerts of potential crypto mining in the organization's compute environment and alerts for common Google Cloud misconfigurations that impact security. Which Security Command Center features should you use to configure these alerts? (Choose two.)

- A. Event Threat Detection
- B. Container Threat Detection
- C. Security Health Analytics
- D. Cloud Data Loss Prevention
- E. Google Cloud Armor

Correct Answer: AE

Community vote distribution



**zellick** Highly Voted 2 months, 3 weeks ago

**Selected Answer: AC**

AC is the answer.

<https://cloud.google.com/security-command-center/docs/concepts-event-threat-detection-overview>

Event Threat Detection is a built-in service for the Security Command Center Premium tier that continuously monitors your organization and identifies threats within your systems in near-real time.

<https://cloud.google.com/security-command-center/docs/concepts-security-sources#security-health-analytics>

Security Health Analytics managed vulnerability assessment scanning for Google Cloud can automatically detect common vulnerabilities and misconfigurations across:

upvoted 5 times

**TonytheTiger** Most Recent 1 week, 6 days ago

on the exam

upvoted 1 times

**AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: AC**

- A. Event Threat Detection
  - C. Security Health Analytics
- upvoted 1 times

**waikiki** 2 months, 4 weeks ago

Security Command Center and Google Cloud Armor are separate services. The question is asking about the functionality of the Security Command Center.

upvoted 1 times

**Random\_Mane** 3 months, 2 weeks ago

**Selected Answer: AC**

A,C

<https://cloud.google.com/security-command-center/docs/concepts-security-command-center-overview>

upvoted 1 times

You have noticed an increased number of phishing attacks across your enterprise user accounts. You want to implement the Google 2-Step Verification (2SV) option that uses a cryptographic signature to authenticate a user and verify the URL of the login page. Which Google 2SV option should you use?

- A. Titan Security Keys
- B. Google prompt
- C. Google Authenticator app
- D. Cloud HSM keys

**Correct Answer:** C

Community vote distribution

A (100%)

🗳️ **TonytheTiger** 1 week, 6 days ago

A. Titan Security Keys  
on the exam  
upvoted 2 times

🗳️ **AzureDP900** 1 month, 2 weeks ago

A.  
[https://store.google.com/us/product/titan\\_security\\_key?pli=1&hl=en-US](https://store.google.com/us/product/titan_security_key?pli=1&hl=en-US)  
Provides phishing-resistant 2nd factor of authentication for high-value users. Works with many devices, browsers & services. Supports FIDO standards.  
upvoted 1 times

🗳️ **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: A**

A. Titan Security Keys  
upvoted 2 times

🗳️ **zelck** 2 months, 3 weeks ago

**Selected Answer: A**

A is the answer.

<https://cloud.google.com/titan-security-key>

Security keys use public key cryptography to verify a user's identity and URL of the login page ensuring attackers can't access your account even if you are tricked into providing your username and password.

upvoted 3 times

🗳️ **GHOST1985** 3 months ago

**Selected Answer: A**

Titan Security Key: Help prevent account takeovers from phishing attacks.  
upvoted 1 times

🗳️ **satyajeet1999** 3 months, 1 week ago

**Selected Answer: A**

agreed  
upvoted 2 times

🗳️ **Random\_Mane** 3 months, 2 weeks ago

**Selected Answer: A**

A. "Security keys use public key cryptography to verify a user's identity and URL of the login page ensuring attackers can't access your account even if you are tricked into providing your username and password."

<https://cloud.google.com/titan-security-key>

<https://qwiklabs.medium.com/two-factor-authentication-annoying-but-important-5fdb9e731868>

upvoted 2 times

🗳️ **Arturo\_Cloud** 3 months, 2 weeks ago

I totally agree.  
upvoted 3 times

Your organization hosts a financial services application running on Compute Engine instances for a third-party company. The third-party company's servers that will consume the application also run on Compute Engine in a separate Google Cloud organization. You need to configure a secure network connection between the Compute Engine instances. You have the following requirements:


- ☞ The network connection must be encrypted.
- ☞ The communication between servers must be over private IP addresses.

What should you do?

- A. Configure a Cloud VPN connection between your organization's VPC network and the third party's that is controlled by VPC firewall rules.
- B. Configure a VPC peering connection between your organization's VPC network and the third party's that is controlled by VPC firewall rules.
- C. Configure a VPC Service Controls perimeter around your Compute Engine instances, and provide access to the third party via an access level.
- D. Configure an Apigee proxy that exposes your Compute Engine-hosted application as an API, and is encrypted with TLS which allows access only to the third party.

**Correct Answer: A**

*Community vote distribution*



🗳️ **TonytheTiger** 1 month ago

B:  
<https://cloud.google.com/vpc/docs/vpc-peering>  
 upvoted 2 times

🗳️ **TonytheTiger** 1 month ago

Sorry - Ans C - Key point "separate Google Cloud Organization"  
 Private Service Connect allows private consumption of services across VPC networks that belong to different groups, teams, projects, or organizations.  
<https://cloud.google.com/vpc/docs/private-service-connect>  
 upvoted 1 times

🗳️ **Littleivy** 1 month, 1 week ago

**Selected Answer: B**

B

VPC Network Peering gives you several advantages over using external IP addresses or VPNs to connect networks

<https://cloud.google.com/vpc/docs/vpc-peering>  
 upvoted 1 times

🗳️ **AzureDP900** 1 month, 2 weeks ago

B. Configure a VPC peering connection between your organization's VPC network and the third party's that is controlled by VPC firewall rules.  
 upvoted 1 times

🗳️ **soltium** 2 months, 1 week ago

A and B is correct, Cloud VPN are encrypted, VPC Peering might be unencrypted but this docs said it's encrypted.  
[https://cloud.google.com/docs/security/encryption-in-transit#virtual\\_machine\\_to\\_virtual\\_machine](https://cloud.google.com/docs/security/encryption-in-transit#virtual_machine_to_virtual_machine)  
 upvoted 1 times

🗳️ **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: B**

B. Configure a VPC peering connection between your organization's VPC network and the third party's that is controlled by VPC firewall rules.  
 upvoted 1 times

🗳️ **zelck** 2 months, 3 weeks ago

**Selected Answer: B**

B is the answer.  
 upvoted 1 times

🗳️ **satyajeet1999** 2 months, 4 weeks ago

**Selected Answer: B**

final B  
 upvoted 1 times

🗨️ 👤 **GHOST1985** 3 months ago

**Selected Answer: B**

Google encrypts and authenticates data in transit at one or more network layers when data moves outside physical boundaries not controlled by Google or on behalf of Google. All VM-to-VM traffic within a VPC network and peered VPC networks is encrypted.

[https://cloud.google.com/docs/security/encryption-in-transit#cio-level\\_summary](https://cloud.google.com/docs/security/encryption-in-transit#cio-level_summary)

upvoted 4 times

🗨️ 👤 **satyajeet1999** 3 months, 1 week ago

**Selected Answer: A**

sry A

upvoted 1 times

🗨️ 👤 **satyajeet1999** 3 months, 2 weeks ago

**Selected Answer: B**

<https://cloud.google.com/vpc/docs/vpc-peering>

upvoted 1 times

🗨️ 👤 **Arturo\_Cloud** 3 months, 2 weeks ago

The requirement is that the connection is encrypted, so A) is the best solution.

<https://cloud.google.com/network-connectivity/docs/vpn/concepts/overview>

upvoted 1 times

🗨️ 👤 **satyajeet1999** 3 months, 2 weeks ago

The ques states that connection must be over internal IP but vpn uses public IP. So how could it be the right one?

upvoted 1 times

🗨️ 👤 **Arturo\_Cloud** 2 months, 1 week ago

In a VPN the communication between the machines is by Private IP, to establish the communication between VPN with VPN it is with a Public IP.

- Therefore, the communication between machines will be with private IPs.
- The communication between the VPNs is by Public IP.
- A VPN connection is encrypted from end to end.

You have never deployed a VPN...?

upvoted 1 times




Your company's new CEO recently sold two of the company's divisions. Your Director asks you to help migrate the Google Cloud projects associated with those divisions to a new organization node. Which preparation steps are necessary before this migration occurs? (Choose two.)

- A. Remove all project-level custom Identity and Access Management (IAM) roles.
- B. Disallow inheritance of organization policies.
- C. Identify inherited Identity and Access Management (IAM) roles on projects to be migrated.
- D. Create a new folder for all projects to be migrated.
- E. Remove the specific migration projects from any VPC Service Controls perimeters and bridges.

**Correct Answer:** *BE*

*Community vote distribution*



**Don10** Highly Voted 3 months ago

**Selected Answer:** DE

D. [https://cloud.google.com/resource-manager/docs/project-migration#import\\_export\\_folders](https://cloud.google.com/resource-manager/docs/project-migration#import_export_folders)

E. [https://cloud.google.com/resource-manager/docs/project-migration#vpcsc\\_security\\_perimeters](https://cloud.google.com/resource-manager/docs/project-migration#vpcsc_security_perimeters)  
upvoted 6 times

**meh009** Most Recent 3 weeks, 2 days ago

**Selected Answer:** CE

Let's go through each:

[https://cloud.google.com/resource-manager/docs/project-migration#import\\_export\\_folders](https://cloud.google.com/resource-manager/docs/project-migration#import_export_folders)

A - Document refers to organization-level custom IAM role. Since this mentions Project-level, we can eliminate this.

B - As per document, It is not NECESSARY to do this

C - Document doesn't explicitly mention this is a MUST, but mentions that IF you do not do this, you are likely to run into problems.

D - The document does mention this as a 'Consideration' to make the migration smoother

E - The document clearly says you MUST do this otherwise migration will fail.

So going with this above. C and E are the two necessary things to do before migration IMHO  
upvoted 1 times

**csrazdan** 1 week, 5 days ago

E will break connectivity for the projects that have to be migrated. The question is asking for the steps required for the preparation of migration.  
I would go with CD  
upvoted 1 times

**marmar11111** 1 month, 1 week ago

**Selected Answer:** CD

[https://cloud.google.com/resource-manager/docs/project-migration#plan\\_policy](https://cloud.google.com/resource-manager/docs/project-migration#plan_policy)

When you migrate your project, it will no longer inherit the policies from its current place in the resource hierarchy, and will be subject to the effective policy evaluation at its destination. We recommend making sure that the effective policies at the project's destination match as much as possible the policies that the project had in its source location. [https://cloud.google.com/resource-manager/docs/project-migration#import\\_export\\_folders](https://cloud.google.com/resource-manager/docs/project-migration#import_export_folders)

Policy inheritance can cause unintended effects when you are migrating a project, both in the source and destination organization resources. You can mitigate this risk by creating specific folders to hold only projects for export and import, and ensuring that the same policies are inherited by the folders in both organization resources. You can also set permissions on these folders that will be inherited to the projects moved within them, helping to accelerate the project migration process.  
upvoted 4 times

**Littleivy** 1 month, 1 week ago

**Selected Answer:** CE

Answer is CE since D is not a must-have action  
upvoted 3 times

**csrazdan** 1 week, 5 days ago

D is not required from an actual migration standpoint but can be done in the preparation of the migration plan.  
upvoted 1 times

**dat987** 1 month, 1 week ago

agree with you

upvoted 1 times

  **bhunias** 1 month, 3 weeks ago

**Selected Answer: DE**

D

[https://cloud.google.com/resource-manager/docs/project-migration#import\\_export\\_folders](https://cloud.google.com/resource-manager/docs/project-migration#import_export_folders)

Policy inheritance can cause unintended effects when you are migrating a project, both in the source and destination organization resources.

[https://cloud.google.com/resource-manager/docs/project-migration#vpcsc\\_security\\_perimeters](https://cloud.google.com/resource-manager/docs/project-migration#vpcsc_security_perimeters)

E is needed because projects can't be moved if its attached to perimeters and this process is not instant, so we need to do it before the migration.

upvoted 2 times

  **Underverse** 1 month, 3 weeks ago

**Selected Answer: CE**

C) To identify access problems for users that could access projects due inherited rights and might face "closed doors" after migration.

E) Is definitely a requirement listed on the links given.

D) Is NOT needed. There is an official recommendation (aka best practice) that folders can be used to organize the migration process better, but by no way is this a requirement. You can clearly migrate projects without it.

upvoted 3 times

  **soltium** 2 months, 1 week ago

**Selected Answer: CE**

I'm leaning towards CE.

C if inherited IAM is not identified first, after migration some service might not function properly.

E is needed because projects can't be moved if its attached to perimeters and this process is not instant, so we need to do it before the migration.

D will help speeds up the process but this is not mandatory and this steps depends on C too.

upvoted 2 times

  **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: DE**

D. Create a new folder for all projects to be migrated.

E. Remove the specific migration projects from any VPC Service Controls perimeters and bridges.

upvoted 2 times

  **GHOST1985** 2 months, 2 weeks ago

**Selected Answer: CD**

the requirement says "which preparation steps are necessary BEFORE this migration occurs" thats means before starting the migration which action you can do without changeing the actual configuration

C: [https://cloud.google.com/resource-manager/docs/project-migration#custom\\_roles](https://cloud.google.com/resource-manager/docs/project-migration#custom_roles)

D: [https://cloud.google.com/resource-manager/docs/project-migration#import\\_export\\_folders](https://cloud.google.com/resource-manager/docs/project-migration#import_export_folders)

upvoted 2 times

  **GHOST1985** 2 months, 2 weeks ago

Remove all project-level custom IAM are in the Assign permissions not in the migration plan => [https://cloud.google.com/resource-manager/docs/project-migration#assign\\_permissions](https://cloud.google.com/resource-manager/docs/project-migration#assign_permissions)

VPC Services control perimeters are in the handle special cases section not in the migration plan => [https://cloud.google.com/resource-manager/docs/project-migration#special\\_cases](https://cloud.google.com/resource-manager/docs/project-migration#special_cases)

upvoted 1 times

  **zellck** 2 months, 3 weeks ago

**Selected Answer: CD**

CD is the answer.

[https://cloud.google.com/resource-manager/docs/project-migration#plan\\_policy](https://cloud.google.com/resource-manager/docs/project-migration#plan_policy)

When you migrate your project, it will no longer inherit the policies from its current place in the resource hierarchy, and will be subject to the effective policy evaluation at its destination. We recommend making sure that the effective policies at the project's destination match as much as possible the policies that the project had in its source location.

[https://cloud.google.com/resource-manager/docs/project-migration#import\\_export\\_folders](https://cloud.google.com/resource-manager/docs/project-migration#import_export_folders)

Policy inheritance can cause unintended effects when you are migrating a project, both in the source and destination organization resources. You can mitigate this risk by creating specific folders to hold only projects for export and import, and ensuring that the same policies are inherited by the folders in both organization resources. You can also set permissions on these folders that will be inherited to the projects moved within them, helping to accelerate the project migration process.

upvoted 3 times

  **Random\_Mane** 3 months ago

**Selected Answer: AE**

A. [https://cloud.google.com/resource-manager/docs/project-migration#custom\\_roles](https://cloud.google.com/resource-manager/docs/project-migration#custom_roles)

E. [https://cloud.google.com/resource-manager/docs/project-migration#vpcsc\\_security\\_perimeters](https://cloud.google.com/resource-manager/docs/project-migration#vpcsc_security_perimeters)

upvoted 1 times



  **rrvv** 3 months, 1 week ago  
<https://cloud.google.com/resource-manager/docs/project-migration>

AS per above article D. Create a new folder for all projects to be migrated. is part of best practices.  
C. Identify inherited Identity and Access Management (IAM) roles on projects to be migrated.- is also a good candidate as inherited roles won't move and only policies/ roles set at project level will be valid after migration.

E. Remove the specific migration projects from any VPC Service Controls perimeters and bridges.- is also a good candidate as VPC-SC policy conflict can create troubles while project migration.  
upvoted 2 times

  **GHOST1985** 3 months, 2 weeks ago  
any idea about the answers ?  
upvoted 1 times

## Question #145

Topic 1

You are a consultant for an organization that is considering migrating their data from its private cloud to Google Cloud. The organization's compliance team is not familiar with Google Cloud and needs guidance on how compliance requirements will be met on Google Cloud. One specific compliance requirement is for customer data at rest to reside within specific geographic boundaries. Which option should you recommend for the organization to meet their data residency requirements on Google Cloud?

- A. Organization Policy Service constraints
- B. Shielded VM instances
- C. Access control lists
- D. Geolocation access controls
- E. Google Cloud Armor

**Correct Answer: A**

Community vote distribution

A (100%)

  **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: A**

A. Organization Policy Service constraints  
upvoted 2 times

  **rrvv** 3 months, 1 week ago

A. Organization Policy Service constraints to add org policy for Resource Location Restriction

<https://cloud.google.com/resource-manager/docs/organization-policy/using-constraints#list-constraint>  
upvoted 1 times

  **AzureDP900** 1 month, 2 weeks ago

yes A. is right. Organization Policy Service constraints  
upvoted 1 times

  **satyajeet1999** 3 months, 2 weeks ago

**Selected Answer: A**

<https://cloud.google.com/blog/products/identity-security/meet-data-residency-requirements-with-google-cloud>  
upvoted 2 times

Your security team wants to reduce the risk of user-managed keys being mismanaged and compromised. To achieve this, you need to prevent developers from creating user-managed service account keys for projects in their organization. How should you enforce this?

- A. Configure Secret Manager to manage service account keys.
- B. Enable an organization policy to disable service accounts from being created.
- C. Enable an organization policy to prevent service account keys from being created.
- D. Remove the `iam.serviceAccounts.getAccessToken` permission from users.

**Correct Answer:** C

*Community vote distribution*

C (100%)

🗲️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: C**

C. Enable an organization policy to prevent service account keys from being created.  
upvoted 2 times

🗲️ 👤 **Random\_Mane** 3 months, 2 weeks ago

**Selected Answer: C**

C. <https://cloud.google.com/iam/docs/best-practices-for-managing-service-account-keys>  
"To prevent unnecessary usage of service account keys, use organization policy constraints:

At the root of your organization's resource hierarchy, apply the Disable service account key creation and Disable service account key upload constraints to establish a default where service account keys are disallowed.

When needed, override one of the constraints for selected projects to re-enable service account key creation or upload."

upvoted 3 times

🗲️ 👤 **AzureDP900** 1 month, 2 weeks ago

Yes, You are right

Enable an organization policy to prevent service account keys from being created.

upvoted 1 times

🗲️ 👤 **Baburao** 3 months, 3 weeks ago

C seems to be a correct option but there must be an exclusion for CI/CD pipelines or SuperAdmins/OrgAdmins. Otherwise, nobody will be able to create ServiceAccount Keys.

upvoted 2 times

You are responsible for managing your company's identities in Google Cloud. Your company enforces 2-Step Verification (2SV) for all users. You need to reset a user's access, but the user lost their second factor for 2SV. You want to minimize risk. What should you do?

- A. On the Google Admin console, select the appropriate user account, and generate a backup code to allow the user to sign in. Ask the user to update their second factor.
- B. On the Google Admin console, temporarily disable the 2SV requirements for all users. Ask the user to log in and add their new second factor to their account. Re-enable the 2SV requirement for all users.
- C. On the Google Admin console, select the appropriate user account, and temporarily disable 2SV for this account. Ask the user to update their second factor, and then re-enable 2SV for this account.
- D. On the Google Admin console, use a super administrator account to reset the user account's credentials. Ask the user to update their credentials after their first login.

**Correct Answer:** B

*Community vote distribution*

A (100%)

🗲️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: A**

A. On the Google Admin console, select the appropriate user account, and generate a backup code to allow the user to sign in. Ask the user to update their second factor.

upvoted 2 times

🗲️ 👤 **zelck** 2 months, 3 weeks ago

**Selected Answer: A**

A is the answer.

<https://support.google.com/a/answer/9176734>

Use backup codes for account recovery

If you need to recover an account, use backup codes. Accounts are still protected by 2-Step Verification, and backup codes are easy to generate.

upvoted 3 times

🗲️ 👤 **AzureDP900** 1 month, 2 weeks ago

.Agreed, On the Google Admin console, select the appropriate user account, and generate a backup code to allow the user to sign in. Ask the user to update their second factor.

upvoted 1 times

🗲️ 👤 **Random\_Mane** 3 months, 2 weeks ago

**Selected Answer: A**

A. <https://support.google.com/a/answer/9176734?hl=en>

upvoted 3 times

Which Google Cloud service should you use to enforce access control policies for applications and resources?

- A. Identity-Aware Proxy
- B. Cloud NAT
- C. Google Cloud Armor
- D. Shielded VMs

**Correct Answer: A**

*Community vote distribution*



  **AzureDP900** 1 month, 2 weeks ago

A. Identity-Aware Proxy  
upvoted 2 times

  **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: A**

A. Identity-Aware Proxy  
upvoted 1 times

  **Random\_Mane** 3 months, 2 weeks ago

**Selected Answer: A**

A. <https://cloud.google.com/iap/docs/concepts-overview>  
"Use IAP when you want to enforce access control policies for applications and resources."  
upvoted 3 times

You want to update your existing VPC Service Controls perimeter with a new access level. You need to avoid breaking the existing perimeter with this change, and ensure the least disruptions to users while minimizing overhead. What should you do?

- A. Create an exact replica of your existing perimeter. Add your new access level to the replica. Update the original perimeter after the access level has been vetted.
- B. Update your perimeter with a new access level that never matches. Update the new access level to match your desired state one condition at a time to avoid being overly permissive.
- C. Enable the dry run mode on your perimeter. Add your new access level to the perimeter configuration. Update the perimeter configuration after the access level has been vetted.
- D. Enable the dry run mode on your perimeter. Add your new access level to the perimeter dry run configuration. Update the perimeter configuration after the access level has been vetted.

**Correct Answer:** D

*Community vote distribution*

D (100%)

🗲️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: D**

D. Enable the dry run mode on your perimeter. Add your new access level to the perimeter dry run configuration. Update the perimeter configuration after the access level has been vetted.

upvoted 2 times

🗲️ 👤 **zelck** 2 months, 3 weeks ago

**Selected Answer: D**

D is the answer.

<https://cloud.google.com/vpc-service-controls/docs/dry-run-mode>

When using VPC Service Controls, it can be difficult to determine the impact to your environment when a service perimeter is created or modified. With dry run mode, you can better understand the impact of enabling VPC Service Controls and changes to perimeters in existing environments.

upvoted 3 times

🗲️ 👤 **AzureDP900** 1 month, 2 weeks ago

D. Enable the dry run mode on your perimeter. Add your new access level to the perimeter dry run configuration. Update the perimeter configuration after the access level has been vetted.

upvoted 1 times

🗲️ 👤 **Baburao** 3 months, 3 weeks ago

D seems to be correct.

[https://cloud.google.com/vpc-service-controls/docs/manage-dry-run-configurations#updating\\_a\\_dry\\_run\\_configuration](https://cloud.google.com/vpc-service-controls/docs/manage-dry-run-configurations#updating_a_dry_run_configuration)


upvoted 3 times

Your organization's Google Cloud VMs are deployed via an instance template that configures them with a public IP address in order to host web services for external users. The VMs reside in a service project that is attached to a host (VPC) project containing one custom Shared VPC for the VMs. You have been asked to reduce the exposure of the VMs to the internet while continuing to service external users. You have already recreated the instance template without a public IP address configuration to launch the managed instance group (MIG). What should you do?

- A. Deploy a Cloud NAT Gateway in the service project for the MIG.
- B. Deploy a Cloud NAT Gateway in the host (VPC) project for the MIG.
- C. Deploy an external HTTP(S) load balancer in the service project with the MIG as a backend.
- D. Deploy an external HTTP(S) load balancer in the host (VPC) project with the MIG as a backend.

**Correct Answer: C**

Community vote distribution



**Littleivy** Highly Voted 1 month, 1 week ago

**Selected Answer: C**

Answer is C

NAT is for egress. To serve customers, need to have LB in the same project  
upvoted 5 times

**GHOST1985** Most Recent 1 month, 1 week ago

**Selected Answer: C**

No doubt the answer is C, this is the Two-tier web service model , below the example from google cloud documentation  
[https://cloud.google.com/vpc/docs/shared-vpc#two-tier\\_web\\_service](https://cloud.google.com/vpc/docs/shared-vpc#two-tier_web_service)  
upvoted 3 times

**AzureDP900** 1 month, 2 weeks ago

B. Deploy a Cloud NAT Gateway in the host (VPC) project for the MIG.  
upvoted 1 times

**GHOST1985** 1 month, 1 week ago

How Cloud NAT could be able to expose internal IP to the public users !! please refers to the documentation before ansewring !  
<https://cloud.google.com/nat/docs/overview>  
upvoted 1 times

**AzureDP900** 1 month, 1 week ago

Thank you for sharing link, I am changing it to C  
upvoted 1 times

**coco10k** 1 month, 3 weeks ago

**Selected Answer: C**

recently support for host project LBs was introduced but usually the LB stays with the backend services in the service project.  
so answer C  
upvoted 2 times

**asdf12345678** 1 month, 2 weeks ago

the official doc still does not support frontend / backend of global https LB in different projects. so +1 to C ([https://cloud.google.com/load-balancing/docs/features#network\\_topologies](https://cloud.google.com/load-balancing/docs/features#network_topologies))  
upvoted 1 times

**Table2022** 1 month, 4 weeks ago

Answer is C, The first example creates all of the load balancer components and backends in the service project.  
<https://cloud.google.com/load-balancing/docs/https/setting-up-reg-ext-shared-vpc>  
upvoted 1 times

**crisyeb** 2 months ago

**Selected Answer: C**

For me C is the answer.

Cloud NAT is for outbound traffic and LB is to handle external customers' request to web services, so it is a LB.

Between C and D:  
In this documentation



<https://cloud.google.com/load-balancing/docs/https#shared-vpc>

it says that "The global external IP address, the forwarding rule, the target HTTP(S) proxy, and the associated URL map must be defined in the same service project as the backends." and in the statement it says that the MIG are in the service project, so in my opinion the LB components must be in the service project.

upvoted 3 times

  **rotorclear** 2 months, 1 week ago

**Selected Answer: D**

NAT is for outbound while the requirement is to serve external customers who will consume web service. Hence the choice is a LB not NAT

upvoted 1 times

  **soltium** 2 months, 1 week ago

C is the answer.

A B Cloud NAT only handle outbound connection from the VM to internet.

D I'm pretty sure you can't select the service project's MIG as backend when creating LB on the host.

upvoted 1 times

  **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: B**

B. Deploy a Cloud NAT Gateway in the host (VPC) project for the MIG.

upvoted 1 times

  **zelck** 2 months, 3 weeks ago

**Selected Answer: D**

D is the answer.

<https://cloud.google.com/load-balancing/docs/https#shared-vpc>

While you can create all the load balancing components and backends in the Shared VPC host project, this model does not separate network administration and service development responsibilities.

upvoted 2 times

  **rrvv** 3 months, 1 week ago

In shared VPC design, it is possible to create a separate NAT gateway in the service project however as per the best practices, a regional NAT gateway should be created in the host project for each regional subnet/network which is being extended to the attached service projects. Hence I will opt for option B

upvoted 1 times

  **GHOST1985** 3 months ago

the requirement says : "while continuing to service external users" , Cloud NAT does not expose service to external users, Cloud NAT is only used for internet outbound

so Answer C is the best Answer

upvoted 3 times

  **GHOST1985** 3 months, 2 weeks ago

**Selected Answer: D**

you cannot use Cloud NAT to attack another project or VPC network, Cloud NAT is only for internet outbond  
the answer is D because MIG backend are hoster in the services project not the hos project

upvoted 2 times

  **tangac** 3 months, 2 weeks ago

**Selected Answer: B**

as it is mention you're right it's B

upvoted 2 times

  **satyajeet1999** 3 months, 2 weeks ago

**Selected Answer: B**

B is the Right One.

upvoted 3 times

  **GHOST1985** 3 months ago

the requirement says : "while continuing to service external users" , Cloud NAT does not expose service to external users, Cloud NAT is only used for internet outbound

so Answer C is the best Answer

upvoted 1 times

Your privacy team uses crypto-shredding (deleting encryption keys) as a strategy to delete personally identifiable information (PII). You need to implement this practice on Google Cloud while still utilizing the majority of the platform's services and minimizing operational overhead. What should you do?

- A. Use client-side encryption before sending data to Google Cloud, and delete encryption keys on-premises.
- B. Use Cloud External Key Manager to delete specific encryption keys.
- C. Use customer-managed encryption keys to delete specific encryption keys.
- D. Use Google default encryption to delete specific encryption keys.

**Correct Answer:** C

Community vote distribution

C (100%)

🗲️ 👤 **Random\_Mane** Highly Voted 👍 3 months, 1 week ago

**Selected Answer: C**

C. <https://cloud.google.com/sql/docs/mysql/cmek>

"You might have situations where you want to permanently destroy data encrypted with CMEK. To do this, you destroy the customer-managed encryption key version. You can't destroy the keyring or key, but you can destroy key versions of the key."

upvoted 5 times

🗲️ 👤 **AzureDP900** Most Recent ⌚ 1 month, 2 weeks ago

C is right

upvoted 1 times

🗲️ 👤 **rotorclear** 2 months, 1 week ago

**Selected Answer: C**

CMEK allows users to manage their keys on google without operation overhead of managing keys externally

upvoted 3 times

🗲️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: C**

C. Use customer-managed encryption keys to delete specific encryption keys.

upvoted 1 times

🗲️ 👤 **zelck** 2 months, 3 weeks ago

**Selected Answer: C**

C is the answer to minimise operational overhead.

upvoted 2 times



You need to centralize your team's logs for production projects. You want your team to be able to search and analyze the logs using Logs Explorer. What should you do?

- A. Enable Cloud Monitoring workspace, and add the production projects to be monitored.
- B. Use Logs Explorer at the organization level and filter for production project logs.
- C. Create an aggregate org sink at the parent folder of the production projects, and set the destination to a Cloud Storage bucket.
- D. Create an aggregate org sink at the parent folder of the production projects, and set the destination to a logs bucket.

**Correct Answer:** D

*Community vote distribution*

D (100%)

🗲️ 👤 **soltium** 2 months, 1 week ago

D because in C we can't use logs explorer to read data from a bucket.  
upvoted 3 times

🗲️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: D**

D. Create an aggregate org sink at the parent folder of the production projects, and set the destination to a logs bucket.  
upvoted 4 times

🗲️ 👤 **Bill1000** 2 months, 3 weeks ago

C is the answer .  
upvoted 1 times

🗲️ 👤 **zellck** 2 months, 3 weeks ago

**Selected Answer: D**

D is the answer.

[https://cloud.google.com/logging/docs/export/aggregated\\_sinks#supported-destinations](https://cloud.google.com/logging/docs/export/aggregated_sinks#supported-destinations)

You can use aggregated sinks to route logs within or between the same organizations and folders to the following destinations:

- Another Cloud Logging bucket: Log entries held in Cloud Logging log buckets.

upvoted 3 times

🗲️ 👤 **AzureDP900** 1 month, 2 weeks ago

Agree with you, D is right  
upvoted 1 times

🗲️ 👤 **Random\_Mane** 3 months, 2 weeks ago

**Selected Answer: D**

D. <https://cloud.google.com/logging/docs/central-log-storage>  
upvoted 2 times

🗲️ 👤 **GHOST1985** 3 months ago

what this link is for ? ?  
upvoted 1 times

You need to use Cloud External Key Manager to create an encryption key to encrypt specific BigQuery data at rest in Google Cloud. Which steps should you do first?

- A. 1. Create or use an existing key with a unique uniform resource identifier (URI) in your Google Cloud project. 2. Grant your Google Cloud project access to a supported external key management partner system.
- B. 1. Create or use an existing key with a unique uniform resource identifier (URI) in Cloud Key Management Service (Cloud KMS). 2. In Cloud KMS, grant your Google Cloud project access to use the key.
- C. 1. Create or use an existing key with a unique uniform resource identifier (URI) in a supported external key management partner system. 2. In the external key management partner system, grant access for this key to use your Google Cloud project.
- D. 1. Create an external key with a unique uniform resource identifier (URI) in Cloud Key Management Service (Cloud KMS). 2. In Cloud KMS, grant your Google Cloud project access to use the key.

**Correct Answer: C**

*Community vote distribution*

C (100%)

  **zellick** Highly Voted 2 months, 3 weeks ago

**Selected Answer: C**

C is the answer.

[https://cloud.google.com/kms/docs/ekm#how\\_it\\_works](https://cloud.google.com/kms/docs/ekm#how_it_works)

- First, you create or use an existing key in a supported external key management partner system. This key has a unique URI or key path.
- Next, you grant your Google Cloud project access to use the key, in the external key management partner system.
- In your Google Cloud project, you create a Cloud EKM key, using the URI or key path for the externally-managed key.

upvoted 7 times

  **AzureDP900** 1 month, 2 weeks ago

Thank you for detailed explanation, I agree with you

upvoted 1 times

  **AwesomeGCP** Most Recent 2 months, 2 weeks ago

**Selected Answer: C**

C.

- 1. Create or use an existing key with a unique uniform resource identifier (URI) in a supported external key management partner system.
- 2. In the external key management partner system, grant access for this key to use your Google Cloud project.

upvoted 3 times

  **Baburao** 3 months, 3 weeks ago

C seems to be correct option.

[https://cloud.google.com/kms/docs/ekm#how\\_it\\_works](https://cloud.google.com/kms/docs/ekm#how_it_works)

upvoted 3 times

Your company's cloud security policy dictates that VM instances should not have an external IP address. You need to identify the Google Cloud service that will allow VM instances without external IP addresses to connect to the internet to update the VMs. Which service should you use?

- A. Identity Aware-Proxy
- B. Cloud NAT
- C. TCP/UDP Load Balancing
- D. Cloud DNS

**Correct Answer:** *B*

*Community vote distribution*

B (100%)

  **Random\_Mane** Highly Voted  3 months, 2 weeks ago

**Selected Answer: B**

B <https://cloud.google.com/nat/docs/overview>

"Cloud NAT (network address translation) lets certain resources without external IP addresses create outbound connections to the internet."  
upvoted 5 times

  **samuelmorher** Most Recent  3 days, 12 hours ago

**Selected Answer: B**

<https://cloud.google.com/nat/docs/overview>

upvoted 1 times

  **AzureDP900** 1 month, 2 weeks ago

Cloud NAT is right

B

upvoted 1 times

  **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: B**

B. Cloud NAT

upvoted 3 times

You want to make sure that your organization's Cloud Storage buckets cannot have data publicly available to the internet. You want to enforce this across all Cloud Storage buckets. What should you do?

- A. Remove Owner roles from end users, and configure Cloud Data Loss Prevention.
- B. Remove Owner roles from end users, and enforce domain restricted sharing in an organization policy.
- C. Configure uniform bucket-level access, and enforce domain restricted sharing in an organization policy.
- D. Remove \*.setIamPolicy permissions from all roles, and enforce domain restricted sharing in an organization policy.

**Correct Answer:** C

*Community vote distribution*

C (100%)

  **samuelmorher** 3 days, 12 hours ago

**Selected Answer: C**

It's C

upvoted 1 times

  **AzureDP900** 1 month, 2 weeks ago

I agree with C

upvoted 1 times

  **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: C**

C. Configure uniform bucket-level access, and enforce domain restricted sharing in an organization policy.

upvoted 2 times

  **zelck** 2 months, 3 weeks ago

**Selected Answer: C**

C is the answer.

upvoted 3 times

  **GHOST1985** 3 months ago

**Selected Answer: C**

- Uniform bucket-level access: <https://cloud.google.com/storage/docs/uniform-bucket-level-access#should-you-use>

- Domain Restricted Sharing: [https://cloud.google.com/resource-manager/docs/organization-policy/restricting-domains#public\\_data\\_sharing](https://cloud.google.com/resource-manager/docs/organization-policy/restricting-domains#public_data_sharing)


upvoted 3 times

Your company plans to move most of its IT infrastructure to Google Cloud. They want to leverage their existing on-premises Active Directory as an identity provider for Google Cloud. Which two steps should you take to integrate the company's on-premises Active Directory with Google Cloud and configure access management? (Choose two.)

- A. Use Identity Platform to provision users and groups to Google Cloud.
- B. Use Cloud Identity SAML integration to provision users and groups to Google Cloud.
- C. Install Google Cloud Directory Sync and connect it to Active Directory and Cloud Identity.
- D. Create Identity and Access Management (IAM) roles with permissions corresponding to each Active Directory group.
- E. Create Identity and Access Management (IAM) groups with permissions corresponding to each Active Directory group.

**Correct Answer:** BC

Community vote distribution



**rotorclear** 2 months, 1 week ago

**Selected Answer:** BC

Roles should be provided by Auth provider  
upvoted 1 times

**AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer:** CE

C. Install Google Cloud Directory Sync and connect it to Active Directory and Cloud Identity.  
E. Create Identity and Access Management (IAM) groups with permissions corresponding to each Active Directory group.  
upvoted 3 times

**GHOST1985** 3 months ago

**Selected Answer:** CE

<https://cloud.google.com/architecture/identity/federating-gcp-with-active-directory-synchronizing-user-accounts?hl=en>  
  
[https://cloud.google.com/architecture/identity/federating-gcp-with-active-directory-synchronizing-user-accounts?hl=en#deciding\\_where\\_to\\_deploy\\_gcids](https://cloud.google.com/architecture/identity/federating-gcp-with-active-directory-synchronizing-user-accounts?hl=en#deciding_where_to_deploy_gcids)  
upvoted 4 times

**AzureDP900** 1 month, 2 weeks ago

CE sounds good  
upvoted 2 times

**Test114** 2 months, 4 weeks ago

How about BE?  
<https://cloud.google.com/architecture/identity/federating-gcp-with-active-directory-introduction>  
"Single sign-on: Whenever a user needs to authenticate, Google Cloud delegates the authentication to Active Directory by using the Security Assertion Markup Language (SAML) protocol."  
upvoted 1 times

**zellick** 2 months, 3 weeks ago

SAML is used for authentication, not provisioning.  
upvoted 2 times

**GHOST1985** 3 months, 1 week ago

C: Answer C : <https://cloud.google.com/architecture/identity/federating-gcp-with-active-directory-synchronizing-user-accounts?hl=en>  
upvoted 1 times

**AzureDP900** 1 month, 2 weeks ago

CE is right, u need to select 2 answers  
upvoted 1 times

You are in charge of creating a new Google Cloud organization for your company. Which two actions should you take when creating the super administrator accounts? (Choose two.)

- A. Create an access level in the Google Admin console to prevent super admin from logging in to Google Cloud.
- B. Disable any Identity and Access Management (IAM) roles for super admin at the organization level in the Google Cloud Console.
- C. Use a physical token to secure the super admin credentials with multi-factor authentication (MFA).
- D. Use a private connection to create the super admin accounts to avoid sending your credentials over the Internet.
- E. Provide non-privileged identities to the super admin users for their day-to-day activities.

**Correct Answer:** AC

*Community vote distribution*

CE (100%)

🗲️ 👤 **Baburao** Highly Voted 👍 3 months, 3 weeks ago

I think CE makes a better option. See documentation below:  
<https://cloud.google.com/resource-manager/docs/super-admin-best-practices>  
upvoted 7 times

🗲️ 👤 **samuelmorher** Most Recent ⌚ 3 days, 11 hours ago

**Selected Answer: CE**

it's CE  
upvoted 2 times

🗲️ 👤 **AzureDP900** 1 month, 2 weeks ago

CE is good  
upvoted 2 times

🗲️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: CE**

C. Use a physical token to secure the super admin credentials with multi-factor authentication (MFA).  
E. Provide non-privileged identities to the super admin users for their day-to-day activities.  
upvoted 4 times

🗲️ 👤 **zellick** 2 months, 4 weeks ago

**Selected Answer: CE**

CE is the answer.

[https://cloud.google.com/resource-manager/docs/super-admin-best-practices#discourage\\_super\\_admin\\_account\\_usage](https://cloud.google.com/resource-manager/docs/super-admin-best-practices#discourage_super_admin_account_usage)  
- Use a security key or other physical authentication device to enforce two-step verification  
- Give super admins a separate account that requires a separate login  
upvoted 2 times

🗲️ 👤 **AzureDP900** 1 month, 2 weeks ago

Thanks  
upvoted 1 times



You are deploying a web application hosted on Compute Engine. A business requirement mandates that application logs are preserved for 12 years and data is kept within European boundaries. You want to implement a storage solution that minimizes overhead and is cost-effective. What should you do?

- A. Create a Cloud Storage bucket to store your logs in the EUROPE-WEST1 region. Modify your application code to ship logs directly to your bucket for increased efficiency.
- B. Configure your Compute Engine instances to use the Google Cloud's operations suite Cloud Logging agent to send application logs to a custom log bucket in the EUROPE-WEST1 region with a custom retention of 12 years.
- C. Use a Pub/Sub topic to forward your application logs to a Cloud Storage bucket in the EUROPE-WEST1 region.
- D. Configure a custom retention policy of 12 years on your Google Cloud's operations suite log bucket in the EUROPE-WEST1 region.

**Correct Answer:** B

Community vote distribution

B (100%)

  **GHOST1985** Highly Voted 3 months, 1 week ago

**Selected Answer: B**

A: Google recommend to avoid developping new code while it propose service for that => incorrect  
B: seem to reponse for this needs => correct  
C: Pub/sub is not using for forwarding log, it is an event notification, and no configuration for the retention 12 years is proposed => incorrect  
D: how the application will forward the logs to the bucket ? => incorrect  
upvoted 6 times

  **AwesomeGCP** Most Recent 2 months, 2 weeks ago

**Selected Answer: B**

B. Configure your Compute Engine instances to use the Google Cloud's operations suite Cloud Logging agent to send application logs to a custom log bucket in the EUROPE-WEST1 region with a custom retention of 12 years.  
upvoted 3 times

  **rrvv** 3 months, 1 week ago

B. Configure your Compute Engine instances to use the Google Cloud's operations suite Cloud Logging agent to send application logs to a custom log bucket in the EUROPE-WEST1 region with a custom retention of 12 years.

read more about it here- <https://youtu.be/MI4iG2GIZMA>

upvoted 3 times

  **tangac** 3 months, 2 weeks ago

The A and the C are the two possible (12 years retention is not possible with Cloud Logging...max 3650 days)  
so now the question is...pub/sub or not pub/sub....  
in my opinion when it's said...limit overhead, i should go with the A....but not really sure  
upvoted 4 times

  **meh009** 3 weeks, 2 days ago

Correct. Tested and can verify this. Between A and C. and I would choose A.

upvoted 1 times

You discovered that sensitive personally identifiable information (PII) is being ingested to your Google Cloud environment in the daily ETL process from an on- premises environment to your BigQuery datasets. You need to redact this data to obfuscate the PII, but need to re-identify it for data analytics purposes. Which components should you use in your solution? (Choose two.)

- A. Secret Manager
- B. Cloud Key Management Service
- C. Cloud Data Loss Prevention with cryptographic hashing
- D. Cloud Data Loss Prevention with automatic text redaction
- E. Cloud Data Loss Prevention with deterministic encryption using AES-SIV

**Correct Answer:** *DE*

*Community vote distribution*

BE (100%)

  **GHOST1985** Highly Voted 3 months, 1 week ago

**Selected Answer: BE**

B: you need KMS to store the CryptoKey <https://cloud.google.com/dlp/docs/reference/rest/v2/projects.deidentifyTemplates#crypt>

E: for the de-identity you need to use CryptoReplaceFfxFpeConfig or CryptoDeterministicConfig  
<https://cloud.google.com/dlp/docs/reference/rest/v2/projects.deidentifyTemplates#cryptodeterministicconfig>  
<https://cloud.google.com/dlp/docs/deidentify-sensitive-data>

upvoted 6 times

  **AzureDP900** Most Recent 1 month, 2 weeks ago

B & E is right

upvoted 1 times

  **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: BE**

B. Cloud Key Management Service  
E. Cloud Data Loss Prevention with deterministic encryption using AES-SIV  
upvoted 2 times

  **zelck** 2 months, 3 weeks ago

**Selected Answer: BE**

BE is the answer.

upvoted 3 times

  **waikiki** 2 months, 4 weeks ago

No. As a result of checking the documentation, crypto key = This is a data encryption key (DEK) (as opposed to a key encryption key (KEK) stored by Cloud Key Management Service (Cloud KMS).

upvoted 1 times



You are working with a client that is concerned about control of their encryption keys for sensitive data. The client does not want to store encryption keys at rest in the same cloud service provider (CSP) as the data that the keys are encrypting. Which Google Cloud encryption solutions should you recommend to this client?

(Choose two.)

- A. Customer-supplied encryption keys.
- B. Google default encryption
- C. Secret Manager
- D. Cloud External Key Manager
- E. Customer-managed encryption keys

**Correct Answer:** AD

*Community vote distribution*

AD (100%)

🗲️ 👤 **AzureDP900** 1 month, 2 weeks ago

A,D is perfect  
upvoted 3 times

🗲️ 👤 **soltium** 2 months, 1 week ago

I'm leaning towards D because CSEK is so limited.  
upvoted 1 times

🗲️ 👤 **soltium** 2 months, 1 week ago

whoops didn't read I need to select two, so AD it is.  
upvoted 1 times

🗲️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: AD**

A. Customer-supplied encryption keys.  
D. Cloud External Key Manager  
upvoted 3 times

🗲️ 👤 **DST** 2 months, 3 weeks ago

**Selected Answer: AD**

CSEK & EKM both store keys outside of GCP  
upvoted 3 times

You are implementing data protection by design and in accordance with GDPR requirements. As part of design reviews, you are told that you need to manage the encryption key for a solution that includes workloads for Compute Engine, Google Kubernetes Engine, Cloud Storage, BigQuery, and Pub/Sub. Which option should you choose for this implementation?

- A. Cloud External Key Manager
- B. Customer-managed encryption keys
- C. Customer-supplied encryption keys
- D. Google default encryption

**Correct Answer:** B

Community vote distribution

B (65%)

A (35%)

  **zellick** Highly Voted 2 months, 4 weeks ago

**Selected Answer: B**

B is the answer.

[https://cloud.google.com/kms/docs/using-other-products#cmek\\_integrations](https://cloud.google.com/kms/docs/using-other-products#cmek_integrations)

[https://cloud.google.com/kms/docs/using-other-products#cmek\\_integrations](https://cloud.google.com/kms/docs/using-other-products#cmek_integrations)

CMEK is supported for all the listed google services.

upvoted 7 times

  **marmar11111** Most Recent 1 month ago

**Selected Answer: B**

"Customer-supplied encryption keys (CSEK) are a feature in Google Cloud Storage and Google Compute Engine. If you supply your own encryption keys, Google uses your key to protect the Google-generated keys used to encrypt and decrypt your data." so it can't be A because A doesn't support all these services.

The answer is B as this still allows you to manage your keys!

upvoted 2 times

  **GHOST1985** 1 month, 1 week ago

**Selected Answer: A**

<https://cloud.google.com/blog/products/compliance/how-google-cloud-helps-customers-stay-current-with-gdpr>

upvoted 1 times

  **Littleivy** 1 month, 1 week ago

**Selected Answer: A**

Obviously A is the better answer. Based on the GCP blog [1], you can utilize Cloud External Key Manager (Cloud EKM) to manage customer key easily and fulfill the compliance requirements as Key Access Justifications is already GA.

Also, Cloud EKM supports all the services listed in the questions per the reference [2]

[1] <https://cloud.google.com/blog/products/compliance/how-google-cloud-helps-customers-stay-current-with-gdpr>

[2] [https://cloud.google.com/kms/docs/ekm#supported\\_services](https://cloud.google.com/kms/docs/ekm#supported_services)

upvoted 1 times

  **Table2022** 1 month, 4 weeks ago

Answer is A, Customers can also require detailed justification and approval each time a key is requested to decrypt data using External Key Manager, and deny Google the ability to decrypt their data for any reason using Key Access Justifications, which is now in General Availability.

<https://cloud.google.com/blog/products/compliance/how-google-cloud-helps-customers-stay-current-with-gdpr>

upvoted 1 times

  **rotorclear** 2 months, 1 week ago

**Selected Answer: B**

External key and CSEK can only be used with CE and CS not all the mentioned services

upvoted 2 times

  **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: A**

A. Cloud External Key Manager

upvoted 1 times

  **DST** 2 months, 3 weeks ago

**Selected Answer: A**

The GDPR requirements makes me lean towards EKM which supports all the services detailed.  
upvoted 1 times

  **GHOST1985** 3 months, 1 week ago

**Selected Answer: A**

A: is the best answer because you the requirement is "you manage your key" and "the solution workload includes GKE, cloud storage, BigQuery, Pub/sub etc..."  
[https://cloud.google.com/kms/docs/ekm#supported\\_services](https://cloud.google.com/kms/docs/ekm#supported_services)  
B & C: both keys are for encryption only cloud storage : <https://cloud.google.com/storage/docs/encryption/customer-managed-keys>  
D: incorrect  
upvoted 2 times

  **zellick** 2 months, 4 weeks ago

CMEK supports more than just cloud storage.  
[https://cloud.google.com/kms/docs/using-other-products#cmek\\_integrations](https://cloud.google.com/kms/docs/using-other-products#cmek_integrations)  
upvoted 3 times

Question #162

Topic 1

Which Identity-Aware Proxy role should you grant to an Identity and Access Management (IAM) user to access HTTPS resources?

- A. Security Reviewer
- B. IAP-Secured Tunnel User
- C. IAP-Secured Web App User
- D. Service Broker Operator

**Correct Answer: B**

*Community vote distribution*

C (100%)

  **AzureDP900** 1 month, 2 weeks ago

C is right  
IAP-secured Web App User  
(roles/iap.httpsResourceAccessor)

Provides permission to access HTTPS resources which use Identity-Aware Proxy.  
upvoted 2 times

  **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: C**

C. IAP-Secured Web App User  
upvoted 1 times

  **GHOST1985** 3 months ago

**Selected Answer: C**

Answer C  
IAP-Secured Tunnel User: Grants access to tunnel resources that use IAP.  
IAP-Secured Web App User: Access HTTPS resources which use Identity-Aware Proxy, Grants access to App Engine, Cloud Run, and Compute Engine resources.  
upvoted 3 times

  **Random\_Mane** 3 months, 2 weeks ago

**Selected Answer: C**

C, <https://cloud.google.com/iap/docs/managing-access>  
"IAP-Secured Web App User: Grants access to the app and other HTTPS resources that use IAP."  
upvoted 3 times

  **Baburao** 3 months, 3 weeks ago

Should be C. It is clearly mentioned here in Documentation:  
<https://cloud.google.com/iap/docs/managing-access#roles>  
IAP-Secured Web App User (roles/iap.httpsResourceAccessor)  
upvoted 3 times

You need to audit the network segmentation for your Google Cloud footprint. You currently operate Production and Non-Production infrastructure-as-a-service (IaaS) environments. All your VM instances are deployed without any service account customization. After observing the traffic in your custom network, you notice that all instances can communicate freely `` despite tag-based VPC firewall rules in place to segment traffic properly `` with a priority of 1000. What are the most likely reasons for this behavior?

- A. All VM instances are missing the respective network tags.
- B. All VM instances are residing in the same network subnet.
- C. All VM instances are configured with the same network route.
- D. A VPC firewall rule is allowing traffic between source/targets based on the same service account with priority 999. E . A VPC firewall rule is allowing traffic between source/targets based on the same service account with priority 1001.

Correct Answer: AC

Community vote distribution



- adelynllllllllll

1 month ago

D: a 999 will overwrite 1000

upvoted 1 times
- Littleivy

1 month, 1 week ago

Selected Answer: D

The answer is D

upvoted 1 times
- rotorclear

2 months, 1 week ago

Selected Answer: AD

1001 is lower priority

upvoted 2 times
- soltium

2 months, 1 week ago

D. priority 999 is a higher priority than 1000, so if 999 has allow all policy then any deny policy with lower priority will not be applied.

upvoted 3 times
- AwesomeGCP

2 months, 2 weeks ago

Selected Answer: C

C. All VM instances are configured with the same network route.

upvoted 2 times
- dat987

2 months, 1 week ago

Do you have any documents for this? Thanks

upvoted 1 times
- redgoose6810

2 months, 3 weeks ago

maybe A . any idea please.

upvoted 3 times
- maxth3mad

2 months, 1 week ago

maybe B too ... same subnet ...

upvoted 1 times
- maxth3mad

2 months, 1 week ago

but if a firewall rule is in place, probably A

upvoted 1 times

You are creating a new infrastructure CI/CD pipeline to deploy hundreds of ephemeral projects in your Google Cloud organization to enable your users to interact with Google Cloud. You want to restrict the use of the default networks in your organization while following Google-recommended best practices. What should you do?

- A. Enable the constraints/compute.skipDefaultNetworkCreation organization policy constraint at the organization level.
- B. Create a cron job to trigger a daily Cloud Function to automatically delete all default networks for each project.
- C. Grant your users the IAM Owner role at the organization level. Create a VPC Service Controls perimeter around the project that restricts the compute.googleapis.com API.
- D. Only allow your users to use your CI/CD pipeline with a predefined set of infrastructure templates they can deploy to skip the creation of the default networks.

**Correct Answer: A**

*Community vote distribution*

A (100%)

🗲️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: A**

A. Enable the constraints/compute.skipDefaultNetworkCreation organization policy constraint at the organization level.  
upvoted 3 times

🗲️ 👤 **zellck** 2 months, 3 weeks ago

**Selected Answer: A**

A is the answer.

<https://cloud.google.com/resource-manager/docs/organization-policy/org-policy-constraints>  
- constraints/compute.skipDefaultNetworkCreation

This boolean constraint skips the creation of the default network and related resources during Google Cloud Platform Project resource creation where this constraint is set to True. By default, a default network and supporting resources are automatically created when creating a Project resource.

upvoted 3 times

🗲️ 👤 **AzureDP900** 1 month, 2 weeks ago

Agreed

upvoted 1 times

🗲️ 👤 **Random\_Mane** 3 months, 2 weeks ago

**Selected Answer: A**

A. <https://cloud.google.com/resource-manager/docs/organization-policy/org-policy-constraints>

"This boolean constraint skips the creation of the default network and related resources during Google Cloud Platform Project resource creation where this constraint is set to True. By default, a default network and supporting resources are automatically created when creating a Project resource.

constraints/compute.skipDefaultNetworkCreation"


upvoted 2 times

You are a security administrator at your company and are responsible for managing access controls (identification, authentication, and authorization) on Google Cloud. Which Google-recommended best practices should you follow when configuring authentication and authorization? (Choose two.)

- A. Use Google default encryption.
- B. Manually add users to Google Cloud.
- C. Provision users with basic roles using Google's Identity and Access Management (IAM) service.
- D. Use SSO/SAML integration with Cloud Identity for user authentication and user lifecycle management.
- E. Provide granular access with predefined roles.

**Correct Answer:** DE

Community vote distribution



**Littleivy** 1 month, 1 week ago

**Selected Answer: DE**

Answer is DE of course  
upvoted 3 times

**AzureDP900** 1 month, 2 weeks ago

DE is perfect  
upvoted 2 times

**AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: DE**

D. Use SSO/SAML integration with Cloud Identity for user authentication and user lifecycle management.  
E. Provide granular access with predefined roles.  
upvoted 3 times

**zelck** 2 months, 4 weeks ago

**Selected Answer: DE**

DE is the answer.  
upvoted 4 times

**zelck** 2 months, 3 weeks ago

[https://cloud.google.com/iam/docs/using-iam-securely#least\\_privilege](https://cloud.google.com/iam/docs/using-iam-securely#least_privilege)

Basic roles include thousands of permissions across all Google Cloud services. In production environments, do not grant basic roles unless there is no alternative. Instead, grant the most limited predefined roles or custom roles that meet your needs.

upvoted 1 times

**GHOST1985** 3 months ago

**Selected Answer: D**

Answer : DE  
upvoted 2 times



You have been tasked with inspecting IP packet data for invalid or malicious content. What should you do?

- A. Use Packet Mirroring to mirror traffic to and from particular VM instances. Perform inspection using security software that analyzes the mirrored traffic.
- B. Enable VPC Flow Logs for all subnets in the VPC. Perform inspection on the Flow Logs data using Cloud Logging.
- C. Configure the Fluentd agent on each VM Instance within the VPC. Perform inspection on the log data using Cloud Logging.
- D. Configure Google Cloud Armor access logs to perform inspection on the log data.

**Correct Answer:** B

*Community vote distribution*

A (100%)

🗲️ 👤 **AzureDP900** 1 month, 2 weeks ago

A is right  
upvoted 2 times

🗲️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: A**

A. Use Packet Mirroring to mirror traffic to and from particular VM instances. Perform inspection using security software that analyzes the mirrored traffic.  
upvoted 3 times

🗲️ 👤 **zelick** 2 months, 3 weeks ago

**Selected Answer: A**

A is the answer.

<https://cloud.google.com/vpc/docs/packet-mirroring>

Packet Mirroring clones the traffic of specified instances in your Virtual Private Cloud (VPC) network and forwards it for examination. Packet Mirroring captures all traffic and packet data, including payloads and headers.

upvoted 4 times

🗲️ 👤 **Random\_Mane** 3 months, 2 weeks ago

**Selected Answer: A**

A. [https://cloud.google.com/vpc/docs/packet-mirroring#enterprise\\_security](https://cloud.google.com/vpc/docs/packet-mirroring#enterprise_security)

"Packet Mirroring clones the traffic of specified instances in your Virtual Private Cloud (VPC) network and forwards it for examination. Packet Mirroring captures all traffic and packet data, including payloads and headers."

upvoted 3 times

🗲️ 👤 **Baburao** 3 months, 3 weeks ago


Sorry, it should be A, not B.  
upvoted 3 times

🗲️ 👤 **Baburao** 3 months, 3 weeks ago

Should be B. VPC FLOW logs cannot capture packet information.  
<https://cloud.google.com/vpc/docs/using-packet-mirroring>  
upvoted 1 times

You have the following resource hierarchy. There is an organization policy at each node in the hierarchy as shown. Which load balancer types are denied in VPC


A?



- A. All load balancer types are denied in accordance with the global node's policy.
- B. INTERNAL\_TCP\_UDP, INTERNAL\_HTTP\_HTTPS is denied in accordance with the folder's policy.
- C. EXTERNAL\_TCP\_PROXY, EXTERNAL\_SSL\_PROXY are denied in accordance with the project's policy.
- D. EXTERNAL\_TCP\_PROXY, EXTERNAL\_SSL\_PROXY, INTERNAL\_TCP\_UDP, and INTERNAL\_HTTP\_HTTPS are denied in accordance with the folder and project's policies.

**Correct Answer: D**

Community vote distribution



**tangac** Highly Voted 3 months, 3 weeks ago

**Selected Answer: A**

the good answer is A as indicated here : <https://cloud.google.com/load-balancing/docs/org-policy-constraints#gcloud>  
upvoted 6 times

**AzureDP900** 1 month, 2 weeks ago

yes, It is A  
upvoted 1 times

**madhu81321** Most Recent 3 weeks, 6 days ago

**Selected Answer: D**

There are restrictions at folder level too.  
upvoted 1 times

**TheBuckler** 2 months, 1 week ago

NVM - the answer actually is A. The Org has it's own restrictions too!  
upvoted 3 times

**Table2022** 1 month, 4 weeks ago

Agreed with A, good one!  
upvoted 2 times



  **TheBuckler** 2 months, 1 week ago

The answer is D. We also need to consider the Load Balancer types that are restricted at the Folder level as well as the Project level.  
upvoted 1 times

  **satyajeet1999** 3 months, 2 weeks ago

**Selected Answer: A**

It's A.  
upvoted 2 times

## Question #168

Topic 1

Your security team wants to implement a defense-in-depth approach to protect sensitive data stored in a Cloud Storage bucket. Your team has the following requirements:

- ↔ The Cloud Storage bucket in Project A can only be readable from Project B.
- ↔ The Cloud Storage bucket in Project A cannot be accessed from outside the network.
- ↔ Data in the Cloud Storage bucket cannot be copied to an external Cloud Storage bucket.

What should the security team do?

- A. Enable domain restricted sharing in an organization policy, and enable uniform bucket-level access on the Cloud Storage bucket.
- B. Enable VPC Service Controls, create a perimeter around Projects A and B, and include the Cloud Storage API in the Service Perimeter configuration.
- C. Enable Private Access in both Project A and B's networks with strict firewall rules that allow communication between the networks.
- D. Enable VPC Peering between Project A and B's networks with strict firewall rules that allow communication between the networks.

**Correct Answer: D**

Community vote distribution

B (100%)

  **Baburao** **Highly Voted**  3 months, 3 weeks ago

Should be B.  
VPC Peering is between organizations not between Projects in an organization. That is Shared VPC. In this case, both projects are in same organization so having VPC Service Controls around both projects with necessary rules should be fine.  
upvoted 6 times

  **GHOST1985** 3 months, 1 week ago

Answer is B, but you can ave vpc peering between two projects in the same organization, nothing prevents that if you have only two prjects to communicates vpc peering i better than shared vpc ;)  
upvoted 2 times

  **TonytheTiger** **Most Recent**  1 month ago

B:  
<https://cloud.google.com/vpc-service-controls/docs/overview>  
upvoted 2 times

  **AzureDP900** 1 month, 2 weeks ago

B is right  
upvoted 2 times

  **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: B**

B. Enable VPC Service Controls, create a perimeter around Projects A and B, and include the Cloud Storage API in the Service Perimeter configuration.  
upvoted 3 times

  **tangac** 3 months, 3 weeks ago

**Selected Answer: B**

<https://www.examttopics.com/discussions/google/view/33958-exam-professional-cloud-security-engineer-topic-1-question/>  
upvoted 3 times

You need to create a VPC that enables your security team to control network resources such as firewall rules. How should you configure the network to allow for separation of duties for network resources?

- A. Set up multiple VPC networks, and set up multi-NIC virtual appliances to connect the networks.
- B. Set up VPC Network Peering, and allow developers to peer their network with a Shared VPC.
- C. Set up a VPC in a project. Assign the Compute Network Admin role to the security team, and assign the Compute Admin role to the developers.
- D. Set up a Shared VPC where the security team manages the firewall rules, and share the network with developers via service projects.

**Correct Answer:** *D*

*Community vote distribution*

D (100%)

🗲️ 👤 **AzureDP900** 1 month, 2 weeks ago

D. Set up a Shared VPC where the security team manages the firewall rules, and share the network with developers via service projects.  
upvoted 3 times

🗲️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: D**

D. Set up a Shared VPC where the security team manages the firewall rules, and share the network with developers via service projects.  
upvoted 2 times

🗲️ 👤 **zellck** 2 months, 2 weeks ago

**Selected Answer: D**

D is the answer.  
upvoted 1 times

🗲️ 👤 **jitu028** 2 months, 3 weeks ago

Answer is D  
upvoted 1 times

You are onboarding new users into Cloud Identity and discover that some users have created consumer user accounts using the corporate domain name. How should you manage these consumer user accounts with Cloud Identity?

- A. Use Google Cloud Directory Sync to convert the unmanaged user accounts.
- B. Create a new managed user account for each consumer user account.
- C. Use the transfer tool for unmanaged user accounts.
- D. Configure single sign-on using a customer's third-party provider.

**Correct Answer: C**

*Community vote distribution*

C (100%)

🗨️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: C**

C. Use the transfer tool for unmanaged user accounts.  
upvoted 2 times

🗨️ 👤 **zellck** 2 months, 4 weeks ago

**Selected Answer: C**

C is the answer.

<https://support.google.com/a/answer/6178640?hl=en>

The transfer tool enables you to see what unmanaged users exist, and then invite those unmanaged users to the domain.  
upvoted 2 times

🗨️ 👤 **AzureDP900** 1 month, 2 weeks ago

C is right  
upvoted 1 times

🗨️ 👤 **GHOST1985** 3 months, 1 week ago

**Selected Answer: C**

[https://cloud.google.com/architecture/identity/migrating-consumer-accounts#finding\\_unmanaged\\_user\\_accounts](https://cloud.google.com/architecture/identity/migrating-consumer-accounts#finding_unmanaged_user_accounts)  
upvoted 4 times

🗨️ 👤 **Random\_Mane** 3 months, 2 weeks ago

**Selected Answer: C**

C.  
<https://cloud.google.com/architecture/identity/migrating-consumer-accounts>  
"In addition to showing you all unmanaged accounts, the transfer tool for unmanaged users lets you initiate an account transfer by sending an account transfer request. Initially, an account is listed as Not yet invited, indicating that no transfer request has been sent."  
upvoted 1 times

🗨️ 👤 **Baburao** 3 months, 3 weeks ago

C seems to be correct option in these situation.  
<https://support.google.com/cloudidentity/answer/7062710?hl=en>  
upvoted 2 times


You have created an OS image that is hardened per your organization's security standards and is being stored in a project managed by the security team. As a

Google Cloud administrator, you need to make sure all VMs in your Google Cloud organization can only use that specific OS image while minimizing operational overhead. What should you do? (Choose two.)

- A. Grant users the compute.imageUser role in their own projects.
- B. Grant users the compute.imageUser role in the OS image project.
- C. Store the image in every project that is spun up in your organization.
- D. Set up an image access organization policy constraint, and list the security team managed project in the project's allow list.
- E. Remove VM instance creation permission from users of the projects, and only allow you and your team to create VM instances.

**Correct Answer: AD**

Community vote distribution



**Littleivy** 1 month, 1 week ago

**Selected Answer: BD**

Need to grant permission of project owned the image  
upvoted 1 times

**rrvv** 2 months, 1 week ago

Answer should be B and D  
review the example listed here to grant the IAM policy to a service account

[https://cloud.google.com/deployment-manager/docs/configuration/using-images-from-other-projects-for-vm-instances#granting\\_access\\_to\\_images](https://cloud.google.com/deployment-manager/docs/configuration/using-images-from-other-projects-for-vm-instances#granting_access_to_images)  
upvoted 1 times

**Littleivy** 1 month, 1 week ago

Need to grant permission of project owned the image  
upvoted 1 times

**AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: BD**

B. Grant users the compute.imageUser role in the OS image project.  
D. Set up an image access organization policy constraint, and list the security team managed project in the project's allow list.  
upvoted 2 times

**zellck** 2 months, 3 weeks ago

**Selected Answer: BD**

BD is the answer.

<https://cloud.google.com/resource-manager/docs/organization-policy/org-policy-constraints>  
- constraints/compute.trustedImageProjects

This list constraint defines the set of projects that can be used for image storage and disk instantiation for Compute Engine. If this constraint is active, only images from trusted projects will be allowed as the source for boot disks for new instances.  
upvoted 4 times

**AzureDP900** 1 month, 2 weeks ago

Thank you for sharing link, BD correct  
upvoted 1 times

**GHOST1985** 3 months, 1 week ago

**Selected Answer: AD**

the compute.imageUser is a Permission to list and read images without having other permissions on the image. Granting this role at the project level gives users the ability to list all images in the project and create resources, such as instances and persistent disks, based on images in the project.  
<https://cloud.google.com/compute/docs/access/iam#compute.imageUser>  
upvoted 2 times

**GHOST1985** 3 months ago

Sorry Answer BD  
upvoted 2 times

  **Baburao** 3 months, 3 weeks ago

I think it should be BD instead of AD.

Users should have access to the project where the secured image is stored which is "Security Team's project". Users will obviously need permission to create VM in their own project but to use image from another project, they need "imageUser" permission on that project.

upvoted 3 times

## Question #172

Topic 1

You're developing the incident response plan for your company. You need to define the access strategy that your DevOps team will use when reviewing and investigating a deployment issue in your Google Cloud environment. There are two main requirements:

- ☞ Least-privilege access must be enforced at all times.
- ☞ The DevOps team must be able to access the required resources only during the deployment issue.

How should you grant access while following Google-recommended best practices?

- A. Assign the Project Viewer Identity and Access Management (IAM) role to the DevOps team.
- B. Create a custom IAM role with limited list/view permissions, and assign it to the DevOps team.
- C. Create a service account, and grant it the Project Owner IAM role. Give the Service Account User Role on this service account to the DevOps team.
- D. Create a service account, and grant it limited list/view permissions. Give the Service Account User Role on this service account to the DevOps team.

**Correct Answer: B**

*Community vote distribution*

D (100%)

  **Baburao** **Highly Voted**  3 months, 3 weeks ago

I think the answer should D. Option B gives them "Always On" permissions but the question asks for "Just in time" permissions. So, this is possible only with a Service Account. Once the incident response team resolves the issue, the service account key can be disabled.

upvoted 7 times

  **AzureDP900** **Most Recent**  1 month, 2 weeks ago

D is right

upvoted 2 times

  **soltium** 2 months, 1 week ago

I think its B, option D lacks Service Account Token Creator Role so the DevOps team can't impersonate the SA.

upvoted 2 times

  **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: D**

D. Create a service account, and grant it limited list/view permissions. Give the Service Account User Role on this service account to the DevOps team.

upvoted 2 times

  **GHOST1985** 3 months, 1 week ago

**Selected Answer: D**

answer should be D

upvoted 2 times

You are working with a client who plans to migrate their data to Google Cloud. You are responsible for recommending an encryption service to manage their encrypted keys. You have the following requirements:

- ☞ The master key must be rotated at least once every 45 days.
- ☞ The solution that stores the master key must be FIPS 140-2 Level 3 validated.
- ☞ The master key must be stored in multiple regions within the US for redundancy.

Which solution meets these requirements?

- A. Customer-managed encryption keys with Cloud Key Management Service
- B. Customer-managed encryption keys with Cloud HSM
- C. Customer-supplied encryption keys
- D. Google-managed encryption keys

**Correct Answer: D**

*Community vote distribution*

B (100%)

🗳️ **TonytheTiger** 3 weeks, 6 days ago

Answer: B

<https://cloud.google.com/docs/security/cloud-hsm-architecture>

upvoted 1 times

🗳️ **Littleivy** 1 month, 1 week ago

**Selected Answer: B**

Cloud HSM is right answer

upvoted 1 times

🗳️ **AzureDP900** 1 month, 2 weeks ago

Cloud HSM is right answer is B

upvoted 1 times

🗳️ **soltium** 2 months, 1 week ago

B.Cloud HSM can be rotated automatically(same front end as KMS), FIPS 140-2 level 3 validated, support multi-region.

upvoted 1 times

🗳️ **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: B**

B. Customer-managed encryption keys with Cloud HSM

upvoted 4 times

🗳️ **zellck** 2 months, 4 weeks ago

**Selected Answer: B**

B is the answer.

upvoted 3 times

🗳️ **Sav94** 3 months, 2 weeks ago

Both A and B. But question ask for redundancy. So I think it's A.

upvoted 1 times

🗳️ **Random\_Mane** 3 months, 2 weeks ago

**Selected Answer: B**

B. <https://cloud.google.com/docs/security/key-management-deep-dive>

<https://cloud.google.com/kms/docs/faq>

"Keys generated with protection level HSM, and the cryptographic operations performed with them, comply with FIPS 140-2 Level 3."

upvoted 3 times

🗳️ **Baburao** 3 months, 3 weeks ago

This should be definitely A. Only Cloud KMS supports FIPS 140-2 levels1, 2 and 3.

<https://cloud.google.com/kms/docs/faq#standards>

upvoted 1 times

🗳️ **Arturo\_Cloud** 3 months, 2 weeks ago



I disagree with you, you are being asked only for FIPS 140-2 Level 3 and multiple availability, so B) is the best answer.  
Here is the much more detailed evidence.  
<https://cloud.google.com/docs/security/cloud-hsm-architecture>  
upvoted 4 times

Question #174

Topic 1

You manage your organization's Security Operations Center (SOC). You currently monitor and detect network traffic anomalies in your VPCs based on network logs. However, you want to explore your environment using network payloads and headers. Which Google Cloud product should you use?

- A. Cloud IDS
- B. VPC Service Controls logs
- C. VPC Flow Logs
- D. Google Cloud Armor
- E. Packet Mirroring

Correct Answer: E

Community vote distribution

E (100%)

- kalyan\_krishna742020** 2 weeks, 3 days ago

It should be A..  
Cloud IDS inspects not only the IP header of the packet, but also the payload.  
<https://cloud.google.com/blog/products/identity-security/how-google-cloud-ids-helps-detect-advanced-network-threats>  
upvoted 2 times
- AzureDP900** 1 month, 2 weeks ago

E  
Packet Mirroring captures all traffic and packet data, including payloads and headers. The capture can be configured for both egress and ingress traffic, only ingress traffic, or only egress traffic.  
upvoted 2 times
- hello\_gcp\_devops** 1 month, 3 weeks ago

Packet Mirroring clones the traffic of specified instances in your Virtual Private Cloud (VPC) network and forwards it for examination. Packet Mirroring captures all traffic and packet data, including payloads and headers. The capture can be configured for both egress and ingress traffic, only ingress traffic, or only egress traffic.  
upvoted 1 times
- hello\_gcp\_devops** 1 month, 3 weeks ago

E is the answer  
upvoted 2 times
- zelck** 2 months, 4 weeks ago

Selected Answer: E

E is the answer.

<https://cloud.google.com/vpc/docs/packet-mirroring>  
Packet Mirroring clones the traffic of specified instances in your Virtual Private Cloud (VPC) network and forwards it for examination. Packet Mirroring captures all traffic and packet data, including payloads and headers.  
upvoted 4 times
- Random\_Mane** 3 months, 2 weeks ago

Selected Answer: E

<https://cloud.google.com/vpc/docs/packet-mirroring>  
upvoted 2 times


You are consulting with a client that requires end-to-end encryption of application data (including data in transit, data in use, and data at rest) within Google Cloud.

Which options should you utilize to accomplish this? (Choose two.)

- A. External Key Manager
- B. Customer-supplied encryption keys
- C. Hardware Security Module
- D. Confidential Computing and Istio
- E. Client-side encryption

**Correct Answer:** AB

Community vote distribution



  **Baburao** Highly Voted  3 months, 3 weeks ago

I feel this should be DE.

Confidential Computing enables encryption for "data-in-use"

Client Side encryption enables security for "data in transit" from Customer site to GCP

Once data is at rest, use Google's default encryption for "data at rest"

upvoted 5 times

  **pmrizzo** Most Recent  6 days, 10 hours ago

[https://cloud.google.com/compute/confidential-vm/docs/about-cvm#end-to-end\\_encryption](https://cloud.google.com/compute/confidential-vm/docs/about-cvm#end-to-end_encryption)

upvoted 1 times

  **Littleivy** 1 month, 1 week ago

Selected Answer: DE

Google Cloud customers with additional requirements for encryption of data over WAN can choose to implement further protections for data as it moves from a user to an application, or virtual machine to virtual machine. These protections include IPSec tunnels, Gmail S/MIME, managed SSL certificates, and Istio.

<https://cloud.google.com/docs/security/encryption-in-transit>

upvoted 1 times

  **AwesomeGCP** 2 months, 2 weeks ago

Selected Answer: DE

D. Confidential Computing and Istio

E. Client-side encryption

upvoted 2 times

  **zelick** 2 months, 3 weeks ago

Selected Answer: AE

AE is my answer.

upvoted 1 times

  **GHOST1985** 2 months, 3 weeks ago

Selected Answer: DE

Confidential Computing enables encryption for "data-in-use"

Client Side encryption enables security for "data in transit" from Customer site to GCP

Once data is at rest, use Google's default encryption for "data at rest"

upvoted 4 times



You need to enforce a security policy in your Google Cloud organization that prevents users from exposing objects in their buckets externally. There are currently no buckets in your organization. Which solution should you implement proactively to achieve this goal with the least operational overhead?

- A. Create an hourly cron job to run a Cloud Function that finds public buckets and makes them private.
- B. Enable the constraints/storage.publicAccessPrevention constraint at the organization level.
- C. Enable the constraints/storage.uniformBucketLevelAccess constraint at the organization level.
- D. Create a VPC Service Controls perimeter that protects the storage.googleapis.com service in your projects that contains buckets. Add any new project that contains a bucket to the perimeter.

**Correct Answer: B**

*Community vote distribution*

B (100%)

🗳️ **TonytheTiger** 2 weeks, 2 days ago

Exam Question Dec 2022

upvoted 2 times

🗳️ **AzureDP900** 1 month, 2 weeks ago

B is right

upvoted 1 times

🗳️ **AzureDP900** 1 month, 2 weeks ago

Public access prevention protects Cloud Storage buckets and objects from being accidentally exposed to the public. When you enforce public access prevention, no one can make data in applicable buckets public through IAM policies or ACLs. There are two ways to enforce public access prevention:

You can enforce public access prevention on individual buckets.

If your bucket is contained within an organization, you can enforce public access prevention by using the organization policy constraint storage.publicAccessPrevention at the project, folder, or organization level.

upvoted 1 times

🗳️ **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: B**

B. Enable the constraints/storage.publicAccessPrevention constraint at the organization level.

upvoted 1 times

🗳️ **zellick** 2 months, 4 weeks ago

**Selected Answer: B**

B is the answer.

<https://cloud.google.com/storage/docs/public-access-prevention>

Public access prevention protects Cloud Storage buckets and objects from being accidentally exposed to the public.

If your bucket is contained within an organization, you can enforce public access prevention by using the organization policy constraint storage.publicAccessPrevention at the project, folder, or organization level.

upvoted 3 times

🗳️ **Random\_Mane** 3 months, 2 weeks ago

**Selected Answer: B**

B. <https://cloud.google.com/storage/docs/org-policy-constraints>

"When you apply the publicAccessPrevention constraint on a resource, public access is restricted for all buckets and objects, both new and existing, under that resource."

upvoted 2 times

Your company requires the security and network engineering teams to identify all network anomalies and be able to capture payloads within VPCs. Which method should you use?

- A. Define an organization policy constraint.
- B. Configure packet mirroring policies.
- C. Enable VPC Flow Logs on the subnet.
- D. Monitor and analyze Cloud Audit Logs.

**Correct Answer:** B

*Community vote distribution*

B (100%)

🗲️ 👤 **AwesomeGCP** 2 months, 2 weeks ago

**Selected Answer: B**

B. Configure packet mirroring policies.  
upvoted 2 times

🗲️ 👤 **zellck** 2 months, 4 weeks ago

**Selected Answer: B**

B is the answer.

<https://cloud.google.com/vpc/docs/packet-mirroring>

Packet Mirroring clones the traffic of specified instances in your Virtual Private Cloud (VPC) network and forwards it for examination. Packet Mirroring captures all traffic and packet data, including payloads and headers.

upvoted 4 times

🗲️ 👤 **AzureDP900** 1 month, 2 weeks ago

B is right .  
upvoted 2 times

🗲️ 👤 **AzureDP900** 1 month, 2 weeks ago

Packet Mirroring is useful when you need to monitor and analyze your security status. It exports all traffic, not only the traffic between sampling periods. For example, you can use security software that analyzes mirrored traffic to detect all threats or anomalies. Additionally, you can inspect the full traffic flow to detect application performance issues.

upvoted 2 times

🗲️ 👤 **Random\_Mane** 3 months, 2 weeks ago

**Selected Answer: B**

<https://cloud.google.com/vpc/docs/packet-mirroring>

upvoted 2 times